

Evaluation Choices Determine Apparent Quantum Kernel Robustness under Distribution Shift

Roberto Fernández-Barrios^[ORCID]^{1*}, Iker Pastor-López^[ORCID]¹,
Asier González-Santocildes^[ORCID]¹, Pablo García Bringas^[ORCID]¹

^{1*}Faculty of Engineering, University of Deusto, Avda. de las
Universidades, 24., Bilbao, 48007, Spain.

*Corresponding author(s). E-mail(s): roberto.fernandez.b@deusto.es;
Contributing authors: iker.pastor@deusto.es; gonzalez.asier@deusto.es;
pablo.garcia.bringas@deusto.es;

Abstract

Quantum-kernel robustness claims depend on how candidates are tuned and selected. We audit this dependence with a controlled kernel-swap benchmark, a ten-specification curve, and a prospectively frozen external validation on three real-world domain shifts. In security benchmarks, fixed regularization, weak references and same-test oracle selection produce apparent gains up to $+0.037$, whereas train-only tuning, target-label-free selection and equal budgets yield quantum-minus-classical effects of -0.005 for support-vector classification and -0.002 for Gaussian-process classification. Externally, the controlled support-vector effect is -0.0119 (conditional 95% interval $[-0.0205, -0.0033]$), while the weak-oracle effect is $+0.0089$ and protocol contraction is $+0.0208$. Contraction also remains positive for the probabilistic classifier, although its controlled effect crosses zero. Geometry is associated with accuracy but shared by both families, and finite-shot estimation distorts exact-statevector rank. The apparent robustness advantage is therefore an evaluation-dependent result in the low-qubit regimes studied.

Keywords: quantum machine learning; quantum kernels; distribution shift;
out-of-distribution robustness; kernel methods; kernel-target alignment; Gaussian
processes; malware detection; network intrusion detection

1 Introduction

Quantum kernel methods embed classical data into quantum states and use state overlaps as similarities for otherwise classical learners [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#); [Schuld \(2021\)](#). This separation makes them a clean setting for asking whether quantum-induced feature geometry provides useful inductive bias. It also makes comparison unusually sensitive to evaluation design: the downstream learner can be held fixed, but conclusions still depend on the encoding, bandwidth, regularization, classical reference set, and rule used to select a candidate [Huang et al. \(2021\)](#); [Schuld et al. \(2021\)](#); [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#).

Recent benchmarks have accordingly shifted the question from whether an isolated quantum model can perform well to whether a claimed gain survives strong and symmetric controls. Broad QML studies find competitive classical baselines once pipelines are aligned [Bowles et al. \(2024\)](#); quantum-kernel benchmarks show substantial sensitivity to feature map and hyperparameter choice [Schnabel and Roth \(2025\)](#); and a concurrent tabular study using nested cross-validation, spectral analysis, and hardware validation finds no significant quantum win across its tested pairs [Kakavand et al. \(2026\)](#). These studies evaluate predominantly identically distributed train–test settings. They do not address a further deployment question: what happens when the target distribution changes and its labels cannot be used to choose the deployed kernel?

Distribution shift makes model selection part of the estimand. Work on domain generalization and real-world shift has shown that robustness gains can disappear when search budgets, hyperparameters, and selection criteria are standardized [Gulrajani and Lopez-Paz \(2021\)](#); [Koh et al. \(2021\)](#); [Yao et al. \(2022\)](#); [Zhou et al. \(2023\)](#). A selector that inspects the shifted test labels estimates an oracle upper bound, not deployable performance; a family with more candidates has more opportunities to win that selection. Fair quantum–classical evaluation under shift must therefore control

not only the classifier and data split, but also regularization, candidate budget, and label access.

Security detection provides a useful stress test because shift is intrinsic to the application. Malware evolves over time, and experimental bias can overstate detector performance [Anderson and Roth \(2018\)](#); [Pendlebury et al. \(2019\)](#); [Arp et al. \(2022\)](#); network traffic changes with infrastructure and attack campaigns [Moustafa and Slay \(2015\)](#); [Alsaedi et al. \(2020\)](#). Quantum classifiers have already been evaluated on intrusion-detection benchmarks, sometimes with reported gains over small classical reference sets on i.i.d. splits [Abreu et al. \(2024, 2025\)](#); [Bellante et al. \(2025\)](#). We use these datasets not to propose a security system, but to expose kernel families to reproducible constructed shifts and to a held-out attack-campaign shift.

The closest robustness-specific concurrent work studies additive input corruption and proposes a spectral phase encoding [Herrero Gómez et al. \(2026\)](#). Our question is complementary: rather than optimize a new encoding, we audit how evaluation choices change a family-level conclusion under constructed and held-out-campaign shifts. The distinguishing axes are no-OOD-label selection, equal candidate budgets, per-configuration train-only regularization, fixed-case inference, and a geometry analysis tied specifically to behaviour after the distribution moves.

We implement a controlled *kernel-swap* benchmark in which preprocessing, splits, classifier, and family-internal selection logic are fixed while only the kernel changes. Four security scenarios from EMBER, UNSW-NB15, and ToN-IoT yield eight scenario-groups across constructed and held-out-campaign shifts. Support-vector and Laplace Gaussian-process classifiers consume the same precomputed Gram matrices. The quantum family contains four fidelity feature-map families with angle-scale sweeps; the classical family extends linear and RBF references with polynomial, Laplacian, and Matérn kernels and symmetric bandwidth freedom. Each configuration’s regularization is tuned from training data alone. The primary deployed candidate is

selected on a disjoint ID-validation split, and the two extended families receive equal candidate budgets. Fifteen pipeline realizations per setting quantify conditional variability; the benchmarks are treated as fixed case studies rather than as random draws from an undefined population.

The comparison is coupled to a mechanistic audit. We measure effective rank, centered kernel-target alignment, concentration, and geometric difference [Roy and Vetterli \(2007\)](#); [Cristianini et al. \(2001\)](#); [Cortes et al. \(2012\)](#); [Huang et al. \(2021\)](#); use partial association, leave-one-dataset-out prediction, and effective-rank matching to restrict interpretation; and perturb fidelity estimates with finite shots to test whether exact statevector geometry survives measurement.

The evaluation choices change the conclusion. With same-test oracle selection, fixed regularization, and customary linear+RBF baselines, fidelity kernels appear to improve OOD balanced accuracy by as much as +0.037. With train-only regularization, no-OOD-label selection, and equal candidate budgets, the dataset-equal quantum-minus-classical delta is -0.005 for SVC and -0.002 for the GP classifier. The primary SVC comparison is classical-favoured or conditionally tied in seven of eight scenario-groups. Effective rank and OOD alignment remain associated with accuracy, but the association is regime-dependent and shared by both families; at matched effective rank, the classical configuration is at least as accurate in every group. Finite-shot estimation inflates effective rank and destabilizes candidate selection.

The contribution is therefore an evaluation result rather than an advantage claim: it identifies which protocol choices create the optimistic quantum-kernel robustness result, supplies a deployable equal-budget comparison under shift, and shows that the observed geometry does not uniquely identify quantum provenance. The scope is deliberately limited to the low-qubit fidelity kernels, datasets, and shift constructions studied here.




Fig. 1 The controlled kernel-swap protocol. Benchmark pools feed reproducible shift constructions; all kernels operate on a shared train-only embedding, receive per-configuration regularization tuned on the training split, and are consumed by the same precomputed-kernel classifiers. Family-level conclusions are drawn by honest, equal-budget selection on an in-distribution validation split (P1'), summarized as fixed case studies with conditional intervals; the Gram matrices feed the geometry-association analysis and the finite-shot model.

2 Results

2.1 Overview

We report results under the controlled kernel-swap protocol described in Section 4. All comparisons share the same classifier, preprocessing pipeline, and ID/OOD split definitions within a scenario-group; the only varying component is the kernel, and every number below aggregates fifteen pipeline realizations per setting. The argument proceeds in three movements. First, under test-peeking oracle selection and fixed regularization, the fidelity kernels show an apparent OOD advantage over linear and RBF baselines (Section 2.2). Second, we tune each configuration on training data alone, select on an in-distribution validation split disjoint from both test splits, and match the candidate budgets of the extended families. Under this controlled comparison no robust family-level advantage remains; the primary SVC comparison is classical-favoured or tied in seven of eight scenario-groups (Section 2.4). Third, a geometry analysis shows that effective rank and OOD alignment are associated with—but do

not causally determine—performance, that the association is regime-dependent, and that it favours neither family (Section 2.6); a finite-shot analysis then shows the exact geometry does not transfer unchanged to finite measurement (Section 2.7). Throughout, we treat the eight scenario-groups as fixed case studies and report effect sizes with conditional intervals rather than a population p -value, because the settings within a benchmark share data pools and are not exchangeable (Section 4).

2.2 The apparent advantage under oracle selection

We begin with the comparison that dominates the literature: fidelity-based quantum kernels against linear and RBF baselines, with each family’s representative chosen by Best-by-OOD selection—the configuration with the highest balanced accuracy on the OOD test split itself—and the classical kernels left at the default regularization $C = 1$. Under these conventions, on EMBER, the quantum family shows an OOD advantage of +0.037 under SVC and +0.028 under the GP classifier against linear+RBF, shrinking to +0.004/+0.002 once the classical family is enlarged with heavier-tailed kernels. Supplementary Table 2 gives the complete oracle results.

Two features favour this apparent quantum result. Best-by-OOD is an *oracle*: it uses the labels it then reports. Fixing $C = 1$ also leaves classical candidates at a potentially unsuitable regularization point. These values are therefore descriptive upper bounds, not deployable effects.

2.3 A specification curve locates the reversal

To separate the contribution of each evaluation choice from the choice of a preferred endpoint, we enumerated all ten logically valid, already-computed specifications fixed before this extension (Figure 2). The specifications vary four axes in a prespecified order: same-test OOD-oracle versus ID-based selection, fixed versus train-cross-validated SVC regularization, customary linear/RBF versus extended classical




Fig. 2 A frozen specification curve shows that evaluation choices reverse the apparent family effect. Grey points are the eight security scenario-groups; diamonds are dataset-equal means. Specifications are shown in their prespecified S1–S10 order, not sorted by effect. They vary target-label access, SVC regularization, classical reference strength, and candidate budget. Positive values favour the fidelity-kernel family. GPC has no C -regularization axis. No specification-level hypothesis test is performed.

references, and native versus equal candidate budgets. Every specification retains all eight security scenario-groups.

The family-level conclusion spans both sides of zero. For SVC, the dataset-equal quantum-minus-classical mean ranges from $+0.047$ under legacy ID selection with fixed $C = 1$ and the customary reference (S2) to -0.006 under ID-validation selection, train-only regularization, the extended reference, and equal budgets (S10). For the GP classifier, for which the regularization axis is not applicable, the corresponding range is $+0.042$ to -0.003 . Expanding the classical reference contracts the positive effect before honest selection or budget matching is imposed; train-only SVC regularization and ID-validation selection then move the controlled estimates to zero or slightly negative. The path is not monotone in every intermediate specification, which is precisely why reporting only an oracle and a final controlled endpoint would be insufficient. The curve shows that the apparent advantage is conditional on a bundle of evaluation decisions rather than a stable property of quantum provenance.

323 2.4 A fair comparison removes the advantage

324
325 A deployable comparison must tune each configuration’s regularization without look-
326 ing at the test data, select the deployed configuration without OOD labels, and
327 compare families at the same candidate budget. We implement all three (Section 4).
328
329 For every kernel candidate, the SVC regularization C is chosen by five-fold cross-
330 validation on the training Gram matrix alone. The deployed configuration is then the
331 one with the best balanced accuracy on an in-distribution *validation* split, held out
332 from the ID test split (protocol P1’); no OOD label is ever consulted in selection.
333
334 The primary comparison uses the same candidate budget for both extended fami-
335 lies (60 per family in full-coverage groups and 20 in reduced-coverage groups), with
336 the full classical pool retained as a sensitivity. We treat the eight scenario-groups
337 as fixed case studies and report, per group, the quantum-minus-classical OOD delta
338 with a conditional interval over the five split realizations, and no population p -value
339 (Section 4).
340
341

342 Table 1 and Figure 3 give the result: no robust advantage remains against either
343 reference. Against the equal-budget extended classical family, the P1’ delta is classical-
344 favoured or within the conditional uncertainty of zero in seven of the eight scenario-
345 groups under the primary SVC comparison, with a dataset-equal-weighted mean of
346 -0.005 (leave-one-dataset-out range $[-0.006, -0.004]$); under the GP classifier the
347 mean is -0.002 ($[-0.003, +0.000]$). Against the customary linear+RBF baseline—the
348 comparison that produced the $+0.037$ oracle result of Section 2.2—the no-OOD-label
349 delta is -0.004 under SVC and $+0.004$ under the GP classifier, with heterogeneous
350 signs across scenario-groups. The single regime that retains a small quantum lean
351 under both equal-budget comparisons is UNSW-NB15 reconnaissance under held-out
352 attack-campaign shift ($+0.005$ for both classifiers); both conditional intervals include
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368

Table 1 No-OOD-label selection results (eight scenario-groups, fifteen pipeline realizations per setting). Quantum-minus-classical OOD balanced-accuracy delta under P1' (select on ID-validation, tune per-configuration C on train, evaluate on OOD test). The extended-family comparison is the frozen equal-budget analysis: 60 candidates per family in full-coverage groups and 20 in reduced-coverage groups. Its interval is conditional pipeline-realization uncertainty over the five split seeds. The customary linear+RBF comparison is contextual rather than an equal-budget endpoint; the full-pool extended-family sensitivity remains in the machine-readable artifact. No population p -value is reported.

Scenario	Clf.	vs lin./RBF	vs ext. (equal B)	conditional interval
EMBER $m1$	SVC	-0.0095	-0.0058	$[-0.0340, +0.0223]$
	GPC	-0.0050	-0.0009	$[-0.0101, +0.0083]$
EMBER $m2$	SVC	+0.0139	-0.0064	$[-0.0137, +0.0009]$
	GPC	+0.0149	+0.0027	$[-0.0027, +0.0081]$
UNSW-DoS (campaign)	SVC	-0.0043	-0.0033	$[-0.0139, +0.0073]$
	GPC	+0.0041	-0.0057	$[-0.0120, +0.0005]$
UNSW-DoS (constructed)	SVC	-0.0085	-0.0044	$[-0.0119, +0.0031]$
	GPC	-0.0056	-0.0108	$[-0.0177, -0.0039]$
UNSW-Recon (campaign)	SVC	+0.0004	+0.0050	$[-0.0058, +0.0157]$
	GPC	+0.0074	+0.0052	$[-0.0019, +0.0123]$
UNSW-Recon (constructed)	SVC	-0.0076	-0.0089	$[-0.0146, -0.0031]$
	GPC	-0.0083	-0.0075	$[-0.0208, +0.0059]$
ToN-IoT (campaign)	SVC	-0.0172	-0.0148	$[-0.0268, -0.0029]$
	GPC	-0.0046	-0.0090	$[-0.0190, +0.0010]$
ToN-IoT (constructed)	SVC	-0.0016	-0.0011	$[-0.0049, +0.0028]$
	GPC	+0.0313	+0.0109	$[+0.0017, +0.0202]$
Dataset-equal mean, SVC		-0.0043	-0.0050	$[-0.0060, -0.0040]$
Dataset-equal mean, GPC		+0.0043	-0.0019	$[-0.0028, +0.0003]$

zero. The full-pool extended-family sensitivity is similar in direction and scale (dataset-equal means -0.004 for both classifiers), so the conclusion does not depend on giving the classical family its larger pool.

Rank matching does not reveal a residual quantum advantage.

We pair each quantum configuration with the classical configuration of nearest training effective rank, within the same run and dimension, and compare their OOD accuracy. Figure 4 reports the outcome: in all eight scenario-groups the median rank-matched delta is negative (quantum minus classical, between -0.001 and -0.009), and the quantum configuration is the more accurate of the pair in only 22–37% of matches. Within




Fig. 3 No robust quantum advantage under no-OOD-label selection. Per-scenario-group P1' OOD deltas (quantum minus classical), for SVC (left) and the GP classifier (right). Filled circles compare against the equal-budget extended classical family with conditional pipeline-realization intervals; open diamonds compare against the customary linear+RBF reference. All eight scenario-groups are shown.

the overlap covered by this observational matching, classical kernels are therefore at least as accurate at comparable effective rank.

2.5 Protocol sensitivity replicates on real-world domain shifts

We next applied the prospectively frozen protocol to College Scorecard, diabetes readmission, and ACS income, preserving the published TableShift domain splits and withholding all OOD labels from the controlled selector. The complete external grid contains 30 task-size-seed units and 37,800 split-level results; every prespecified unit passed the independent audit. Figure 5 shows all task-size effects and conditional seed-cluster intervals.

The primary SVC result satisfies the frozen criterion for *replicated protocol sensitivity*. The task-equal controlled effect is -0.0119 (conditional 95% interval $[-0.0205, -0.0033]$), whereas the weak-reference, fixed- C , same-test-oracle effect is $+0.0089$ ($[+0.0047, +0.0137]$). Their contraction is $+0.0208$ ($[+0.0113, +0.0307]$), with a leave-one-task-out range of $[+0.0116, +0.0260]$. No task has a positive controlled




Fig. 4 At matched effective rank, classical kernels are at least as accurate. For each quantum configuration we find the classical configuration of nearest training effective rank (within run and dimension) and take the OOD balanced-accuracy difference; bars are per-scenario-group medians, annotated with the fraction of matched pairs in which the quantum configuration is the more accurate.

SVC effect whose interval excludes zero in both sizes. Size-equal controlled effects are -0.0160 for College Scorecard, -0.0066 for diabetes readmission, and -0.0131 for ACS income.

The secondary GP classifier shows a larger contraction of $+0.0324$ ($[+0.0223, +0.0427]$), but its task-equal controlled effect is $+0.0077$ with an interval spanning zero ($[-0.0015, +0.0183]$). College Scorecard leans quantum under this classifier, diabetes leans classical, and ACS is slightly positive. This classifier dependence limits the conclusion appropriately: the external study replicates the sensitivity of the family verdict to evaluation design, not universal classical dominance.




Fig. 5 External domain shifts separate optimistic and deployable conclusions. Effects are quantum minus classical OOD balanced accuracy for the weak-reference fixed-regularization same-test oracle, the equal-budget train-tuned same-test oracle, and the primary equal-budget train-tuned validation selector. Rows show both nested sizes for all three fixed tasks plus the task-equal mean. Intervals are conditional on the five seed clusters; no population p -value is reported.

2.6 Geometry is associated with, but does not determine, robustness

The two families are close because both reach similar kernel geometry, and geometry—not quantumness—is what carries information about behaviour under shift. That information, however, is partial: it is a regime-dependent association, not the single governing law that optimistic accounts (our own earlier report included) proposed.

The geometry–OOD link is a regime-dependent association, not a law.

Two descriptors of the training and OOD Gram matrices carry information about which kernels generalize: the effective rank of the training kernel and the centered kernel-target alignment on the OOD split. But their association with OOD accuracy is regime-dependent, not a single mechanism. Within scenario-groups (Figure 6 and Supplementary Table 3), the median Spearman correlation between effective rank and OOD accuracy ranges from +0.80 (EMBER $m2$) down to -0.15 (UNSW-NB15 reconnaissance under the constructed shift); partial correlations that control for family




Fig. 6 The geometry–OOD association is regime-dependent. Median within-unit Spearman correlation between each geometry descriptor and OOD balanced accuracy, per scenario-group (SVC). Effective rank orders accuracy strongly in some regimes (EMBER $m2$, ToN-IoT $m2$) and weakly or negatively in others (the UNSW constructed-shift scenarios); OOD alignment is more consistently positive but never uniform. No single descriptor governs all regimes.

label and length scale track the raw values closely. OOD alignment is more consistently positive (+0.22 to +0.64) but likewise varies by regime. A geometry model fit on three datasets and used to rank the held-out dataset’s kernels achieves only Spearman 0.23–0.40. We therefore describe a genuine but partial association, not a causal mechanism or deployable predictor.

2.7 Finite-shot estimation: the geometry is a statevector idealization

All quantum kernels above are computed by exact statevector simulation, which is the right idealization for a geometry study but not what a device would return. To probe how the picture degrades away from that idealization we apply a *finite-shot fidelity-estimation perturbation model*: each fidelity entry is replaced by a binomial estimate from s shots, train blocks are symmetrized, unit-diagonal, clipped to $[0, 1]$, and projected to the positive-semidefinite cone (Section 4). This is not a hardware simulation—no device noise, mitigation, or transpilation is modelled—but it isolates the statistical cost of estimating a fidelity from finite measurements. Table 2 reports the effect on the $q1000$ subset. OOD alignment and OOD accuracy are robust: median

Table 2 Finite-shot fidelity-estimation perturbation (EMBER+network $q1000$ subset, SVC). “OOD acc. dev.” is the median absolute deviation of OOD balanced accuracy from the statevector-exact value; “KTA dev.” the same for OOD alignment; “rank ratio” the median ratio of perturbed to exact effective rank; “sel. agree” the fraction of runs whose ID-validation-selected configuration is unchanged; “regret” the mean OOD accuracy lost by deploying the perturbed selection.

Shots	OOD acc. dev.	KTA dev.	rank ratio	sel. agree	regret
128	0.0327	0.0004	2.12	7%	+0.0007
512	0.0273	0.0001	1.56	10%	+0.0017
2048	0.0240	0.0001	1.29	19%	−0.0025
8192	0.0181	0.0000	1.15	21%	−0.0001

KTA_{OOD} deviates by at most 0.0004 from the exact value even at 128 shots, and OOD balanced accuracy by 0.018–0.033. Effective rank is not: finite sampling inflates it by a median factor of 2.1 at 128 shots, falling to 1.15 only at 8192 shots, because sampling noise adds spurious spectral mass. The within-family configuration selected by ID-validation changes identity under perturbation (agreement 7–21%), but the accuracy cost of that instability is negligible (mean regret ≤ 0.003), so the flatness of the family optimum—already visible in the exact results—persists. The practical reading: the alignment description we can measure transfers to finite shots, but the effective-rank description, and any story that leans on it, is a statevector-exact quantity that would need many shots to reproduce.

2.8 Predictive uncertainty under shift

The Laplace GP probabilities provide a secondary uncertainty check (Supplementary Figure 1). For the honestly selected candidate, predictive entropy rises from ID to OOD for both families in every scenario-group, while expected calibration error changes moderately and comparably. Because neither family has an accuracy advantage, the fidelity kernels show no calibration benefit unavailable to classical kernels. We do not interpret rising entropy alone as evidence of calibration.

3 Discussion

The study’s central interpretive move is to replace the question “do quantum kernels beat classical kernels under shift?” with “which properties of a kernel are associated with its behaviour under shift, and does either family have exclusive access to them?”. Read through that lens, the results are clear and negative. The apparent advantage over linear and RBF baselines does not survive the combined correction of oracle selection, fixed regularization, and unequal candidate families. It does not reappear against the equal-budget extended family or at matched effective rank. The geometry descriptors carry information as regime-dependent associations rather than a governing law, and they are shared: a well-tuned heavy-tailed classical kernel occupies the same observed geometric regime as the fidelity maps. On present evidence there is no robust out-of-distribution advantage for fidelity kernels in the low-qubit settings studied, and the only residual—a modest in-distribution separability edge under the probabilistic classifier—is small, classifier-dependent, and itself shrinks under no-OOD-label selection.

From the perspective of the distribution-shift literature, this reinforces a general lesson: conclusions about robustness are only as credible as the protocol used to obtain them [Quiñonero-Candela et al. \(2008\)](#); [Gulrajani and Lopez-Paz \(2021\)](#); [Koh et al. \(2021\)](#); [Yao et al. \(2022\)](#). Our results add a QML-specific corollary. The strength of the classical baseline, whether its regularization is tuned, whether model selection may consult the target labels, and the size of each family’s candidate budget are not implementation details: each can move the quantum-versus-classical verdict, and together they turn an apparent +0.037 advantage into a null. A demonstration that fixes any of them in the quantum family’s favour, reported alone, would mislead—and much of the optimistic quantum-kernel literature fixes several at once.

The external validation shows that this is not only a peculiarity of constructed security shifts. Under a protocol frozen before result inspection, three education, health,

691 and socioeconomic tasks reproduce a positive protocol contraction and a negative con-
692 trolled SVC effect. The result is stable to omitting any one task, yet heterogeneous
693 by task and classifier: the GP endpoint retains a small positive task-equal lean whose
694 interval crosses zero. The strongest defensible conclusion is therefore about evaluation
695 sensitivity, not the universal ordering of kernel families.

699 In-distribution separability is the one axis on which the fidelity kernels retain any
700 edge, and even it must be stated carefully. Under the oracle that selects and evaluates
701 on the same in-distribution split, the quantum family separates the classes slightly bet-
702 ter than linear+RBF in every EMBER setting; but under honest selection—choosing
703 on a validation split and reporting on a disjoint test split—the edge falls to +0.003
704 under SVC (mixed in sign across groups) and +0.012 under the probabilistic classifier,
705 and it does not translate into any out-of-distribution benefit. This ID–OOD dissocia-
706 tion is real but modest, and it cuts against, not for, an operational recommendation:
707 a family that is marginally more separable in-distribution yet no more robust under
708 shift offers nothing to a deployment whose whole difficulty is the shift. The pattern is
709 consistent with the “accuracy on the line” literature [Miller et al. \(2021\)](#), in which ID
710 and OOD accuracy usually move together; here they decouple weakly, and geometry—
711 shared by both families—rather than quantumness marks the axis along which they
712 do.

722 The geometry association connects several strands of literature. That length scale
723 is decisive for kernel behaviour is known for both families [Shaydulin and Wild \(2022\)](#);
724 [Canatar et al. \(2023\)](#); our contribution is the comparative consequence—once length
725 scale is tuned symmetrically and regularization is tuned at all, the classical fam-
726 ily reaches the same geometric regime, so asymmetric tuning does not merely bias
727 quantum-versus-classical comparisons, it manufactures the apparent advantage. The
728 geometric difference of Huang et al. [Huang et al. \(2021\)](#) does not predict which fam-
729 ily wins a scenario, consistent with its definition as a bound on the *potential* for a
730

separation rather than its sign; a large geometric difference is compatible with equal or worse OOD accuracy, as we observe. That effective rank and OOD alignment are associated with performance is a statement about kernels in general, not about quantum kernels—which is precisely why the fidelity maps hold no privileged position: they are one costly route into a geometric regime that a short-length-scale classical kernel reaches directly.

The high-rank finding must also be reconciled with two spectral results that superficially point the other way. Kübler et al. argue that quantum kernels generalize in-distribution only when the relevant spectral mass is concentrated in few components Kübler et al. (2021), and Thanasilp et al. show that fidelity kernels concentrate exponentially—toward effectively rank-one, uninformative Gram matrices—as qubit counts grow Thanasilp et al. (2024). There is no contradiction, but there is a genuine tension that our results locate precisely. Both results concern the asymptotic or i.i.d. regime; our effective ranks (roughly 10–150 on training sets of 10^3 – 4×10^3 samples at 4–12 qubits) sit far below the concentration regime, so no contradiction arises. Our finite-shot analysis, however, adds an empirical caveat these theoretical results anticipate: effective rank measured from finitely estimated fidelities inflates substantially (Section 2.7), so the exact effective ranks we report are a statevector idealization, and any account that leans on effective rank as the operative quantity would need many shots to reproduce even at the modest qubit counts studied here.

We emphasize what the geometry analysis does not deliver. It is a post-hoc association, not a deployable selection rule: choosing a configuration by training-set geometry alone does not reliably identify the better kernel, and OOD alignment requires OOD labels. Because the two families are statistically indistinguishable under honest selection, the more useful open problem is no longer “which family wins” but whether any label-free, train-only signal can select a robust kernel at all—a question that concerns classical and quantum kernels equally.

783 From a practical perspective, the paper suggests three takeaways for security ML.
784
785 The first is substantive: on the evidence here, fidelity-based quantum kernels do not
786 offer a robustness advantage under distribution shift over well-tuned classical ker-
787 nels, while their pairwise overlap estimation introduces circuit and shot costs absent
788 from direct evaluation of the classical baselines. A median-heuristic Laplacian kernel,
789 especially one with a shortened length scale and its regularization tuned, is a strong
790 and inexpensive robustness baseline, and it matches the fidelity maps under no-OOD-
791 label selection on both modalities. Practitioners weighing a quantum-kernel pipeline
792 for OOD-sensitive security detection should, on this evidence, prefer a tuned classical
793 kernel unless a specific regime demonstrably favours the quantum construction after
794 the same controls.
795
796
797
798
799
800

801 The second is methodological. Independently of the preferred kernel family, the
802 protocol used here provides a reusable blueprint for robust evaluation in security
803 settings:
804
805

- 806 1. keep the classifier and preprocessing fixed, and swap only the kernel;
- 807 2. define explicit OOD splits tied to interpretable shift mechanisms, including at least
808 one non-constructed shift mechanism where the data permit;
- 809 3. give every family the same freedom—the same candidate budget, symmetric length-
810 scale tuning, and per-configuration regularization tuned on training data alone;
- 811 4. select the deployed configuration without out-of-distribution labels, and keep an
812 in-distribution test split disjoint from the validation split used for selection;
- 813 5. repeat the pipeline, report variability as conditional intervals, and treat correlated
814 benchmarks as fixed case studies rather than pooling them into a single population
815 p -value.
816
817
818
819
820
821
822
823
824
825
826
827
828

Security ML frequently suffers from comparisons in which model class, tuning effort, and evaluation conditions change simultaneously; the way the verdict of Sections 2.2–2.4 moves from a large apparent advantage to a null as the protocol is made fair is a concrete demonstration of how much the conclusion depends on getting this right.

The third concerns diagnostics: the geometry descriptors are cheap to compute from Gram matrices that the pipeline already builds, and the alignment-survival analysis can be run post hoc on any deployed kernel system with a labeled drift sample. Practitioners should nonetheless resist reading these results as a blanket recommendation to replace classical kernels; the responsible recommendation is to test the candidate families under the expected shift regime, using exactly the kind of protocol described here.

Several limitations define the scope of our conclusions.

Our shift operationalization covers three reproducible mechanisms—within-class sparsity-tail extremes, train-geometry-tail extremes, and a held-out attack-campaign shift on network traffic (an unseen attack campaign plus a capture-partition change, not a timestamped temporal split)—but not the full space of real-world drift processes. Explicit temporal splits on timestamped malware, family-based drift, and adversarially induced shift remain outside the study’s scope, and our conclusions should be re-tested under them rather than assumed.

Kernel and encoding coverage also remains finite. The classical family now spans seven kernels plus a bandwidth sweep, and both families receive symmetric tuning freedom, but neither candidate set is exhaustive: the quantum side is restricted to four commonly used fidelity feature maps at low qubit counts, and trained or task-adapted encodings, projected kernels, and deeper maps are not evaluated. The mechanism gives a concrete prediction for such extensions—what should matter is their Gram-matrix geometry, not their provenance—which future work can falsify. At substantially larger qubit counts, exponential kernel concentration [Thanasilp et al. \(2024\)](#) would erode

875 the alignment-preserving rank the mechanism rewards, so our conclusions should not
876 be extrapolated beyond the low-qubit regime studied here.

878 Selection is necessarily over finite candidate sets. P1' (ID-validation selection),
879 P2 (cross-realization OOD-supervised selection), and P3 (same-test OOD oracle) are
880 selection procedures over finite candidate sets. The primary comparison matches can-
881 didate budgets, and the full-pool and budget curves provide sensitivities, but residual
882 dependence on candidate definitions and search granularity remains. Fully nested
883 selection over additional independent benchmark pools would strengthen the analysis
884 further.

889 The quantum kernels are simulated rather than executed on hardware. The study
890 therefore characterizes the geometry of fidelity kernels, not the behavior of noisy
891 devices: finite sampling of kernel entries, hardware noise, and error mitigation would all
892 perturb the Gram matrices [Wang et al. \(2021\)](#), and exponential concentration implies
893 that with polynomially many shots the very spectral richness our mechanism rewards
894 becomes expensive to resolve at larger scales [Thanasilp et al. \(2024\)](#); the effect of either
895 on alignment survival is unknown and worth measuring. We make no claim about
896 hardware feasibility, resource advantage, or quantum speedup—the only rigorously
897 established quantum-classical learning separation concerns a specifically constructed
898 problem, not natural data [Liu et al. \(2021\)](#). At the dimensions studied, the fidelity
899 kernels are also classically simulable [Schreiber et al. \(2023\)](#); [Jerbi et al. \(2024\)](#), which
900 is precisely what makes them available as practical kernel constructions today.

901 Dataset scope and feature caveats provide a final boundary. Four benchmark
902 scenarios across two modalities is a substantial widening relative to single-dataset
903 studies, but security is broader still: dynamic malware analysis and encrypted traf-
904 fic remain untested. Within the network scenarios, the feature sets inherited from
905 the public benchmarks retain port numbers and protocol indicators, which are legiti-
906 mate but known shortcut-prone features in NIDS benchmarks [Engelen et al. \(2021\)](#);
907
908
909
910
911
912
913
914
915
916
917
918
919
920

because every kernel receives identical features, this affects absolute accuracies, not the family comparisons. Finally, the mechanism’s effective-rank component is itself regime-dependent (strong on EMBER and ToN-IoT, weak on UNSW-Reconnaissance), a gradation we report rather than smooth over.

In light of these limitations, the appropriate interpretation of the paper is deliberately conservative. We do not claim universal superiority, complexity-theoretic quantum advantage, or hardware relevance—nor do we claim to have refuted quantum kernels in general. We claim three narrower things, supported by conditional effect estimates rather than a population p -value: that the apparent out-of-distribution advantage of fidelity kernels over customary baselines disappears when same-test oracle selection and untuned regularization are removed; that under equal candidate budgets and no-OOD-label selection no robust family-level advantage remains against either the customary or the extended classical reference, across eight fixed scenario-groups and two classifiers; and that the geometric descriptors associated with out-of-distribution behaviour occur in both families and do not identify kernel provenance as the source of robustness. This view is consistent with recent calls for careful benchmarking in QML [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#); [Kakavand et al. \(2026\)](#) and is more informative than either an unqualified advantage claim or an unqualified dismissal: it identifies the evaluation choices that generate the optimistic result and shows how the conclusion changes when they are controlled.

This paper asked why kernel families behave differently under distribution shift, using the quantum-versus-classical comparison as both motivation and stress test. A controlled kernel-swap protocol—classifier, preprocessing, splits, selection logic, and tuning freedom held fixed while only the kernel varies—was instantiated on four security scenarios spanning two modalities and three shift mechanisms, two classifier families including a probabilistic one, and 72 principal settings with fifteen

967 pipeline realizations. Effects are reported as fixed-case summaries with conditional
968
969 pipeline-realization intervals, not as population significance tests.

970 The empirical arc has three movements. Under same-test oracle selection and fixed
971
972 regularization, fidelity-based quantum kernels appear to beat the customary linear and
973
974 RBF baselines under shift. That result does not survive the controlled comparison:
975 once each configuration is tuned on training data, selection uses no OOD labels, and
976
977 both extended families search an equal budget of kernel geometries, the dataset-equal-
978
979 weighted OOD difference is -0.005 under the primary SVC analysis and -0.002 under
980
981 the GP classifier. The primary comparison is classical-favoured or tied in seven of eight
982
983 scenario-groups. At matched effective rank the classical kernels are at least as accurate
984
985 everywhere. Effective rank and OOD alignment are associated with performance in a
986
987 regime-dependent way, are shared by both families, and, a finite-shot analysis shows,
988
989 do not transfer unchanged from exact statevectors to finite measurement.

990 The contribution is therefore threefold. Empirically, the study provides a con-
991
992 trolled, variability-aware comparison of quantum and classical kernels under shift
993
994 whose central finding is a clean negative obtained under fair conditions: no robust
995
996 family-level advantage survives budget matching, regularization tuning, and honest
997
998 selection. Methodologically, it isolates the specific evaluation choices—oracle selection,
999
1000 fixed regularization, mismatched budgets, and pooled significance testing over corre-
1001
1002 lated settings—that manufacture apparent quantum-kernel advantages, and it leaves
1003
1004 behind a reproducible protocol that removes each. Analytically, it shows that the
1005
1006 geometric quantities often invoked to explain such advantages describe both families
1007
1008 equally and do not transfer unchanged to finite-shot estimation.

1009 The most useful open problem this leaves is not “which family wins” but whether
1010
1011 any label-free, train-only signal can select a robust kernel under shift at all; the security
1012
1013 community’s label-free drift detectors [Yang et al. \(2021\)](#) and the theory of task-model
1014
1015 alignment [Bordelon et al. \(2020\)](#); [Canatar et al. \(2021\)](#); [Ma et al. \(2023\)](#) are natural

starting points. It is a question about kernels in general, and answering it would benefit classical and quantum methods alike—which is, in the end, the point: the productive question is not whether one family dominates, but which similarity geometry aligns with the shift regime one actually faces, and how to find it without peeking.

4 Methods

4.1 Problem setting

We study binary security classification—malware versus benign software, and attack versus benign traffic—under distribution shift. Let D_{ID} denote the in-distribution data-generating process and D_{OOD} an out-of-distribution process induced by a constructed stress test or a held-out attack campaign. Given training data $T \sim D_{\text{ID}}$, we evaluate on an ID test set S_{ID} and an OOD test set S_{OOD} .

To isolate the kernel, classifier, preprocessing, optimization protocol, and split design are fixed. Classical and quantum Gram matrices are therefore consumed by the same downstream learner on the same samples.

4.2 Datasets and split construction

The four scenarios span static malware (EMBER) [Anderson and Roth \(2018\)](#) and network flows: UNSW-NB15 DoS, UNSW-NB15 Reconnaissance [Moustafa and Slay \(2015\)](#), and ToN-IoT Scanning [Alsaedi et al. \(2020\)](#). Every setting has shared T , S_{ID} , and S_{OOD} samples with identical class balance across kernel families.

EMBER uses two label-conditional tail stress tests. The $m1$ split holds out within-class extremes of feature-block sparsity; $m2$ holds out within-class extremes of distance to train-fitted class centroids. Network scenarios use an analogous train-centroid construction and a held-out attack-campaign regime. In the latter, training and ID pools contain benign traffic plus non-target attacks, whereas OOD contains benign traffic plus the target campaign; UNSW pools additionally differ in capture partition. These

are reproducible stress tests, not timestamped temporal splits or certified covariate shifts. Exact algorithms, feature exclusions, leakage checks, and pool definitions are in the Supplementary Note “Dataset and shift construction.”

4.3 Controlled kernel-swap protocol

Figure 1 summarizes the controlled pipeline. Both SVC and a binary Gaussian-process classifier with a logistic-link Laplace approximation consume unit-diagonal precomputed Gram matrices [Cortes and Vapnik \(1995\)](#); [Rasmussen and Williams \(2006\)](#). Comparisons are always within a classifier. The GP classifier supplies approximate probabilities for uncertainty diagnostics; no post-hoc calibration is applied.

Training-fitted MaxAbs scaling and TruncatedSVD produce shared representations at $d \in \{4, 6, 8, 10, 12\}$. Quantum inputs are then mapped linearly to $[0, \pi]$. Dimension is an explicit candidate axis.

4.4 Kernel families

The customary classical reference contains linear and RBF kernels. The extended reference adds degree-2 and degree-3 polynomial, Laplacian, and Matérn-3/2 and 5/2 kernels. RBF, Laplacian, and Matérn bandwidths sweep factors $\{0.1, 0.3, 1, 3, 10\}$ around training-only scale or median-distance heuristics. The quantum family uses fidelity

$$K_Q(x, x') = |\langle \phi_Q(x), \phi_Q(x') \rangle|^2$$

for two ZZ depths, PauliXZ, and Z feature maps [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#). Every map sweeps angle scales $\{0.5, 1, 2\}$. Classical bandwidth and quantum angle-scale freedom are symmetric in purpose [Shaydulin and Wild \(2022\)](#); [Canatar et al. \(2023\)](#). Exact formulas and grids appear in the Supplementary Note “Controlled kernel-swap protocol.”

4.5 Selection, candidate budgets, and inference

For each geometry, SVC regularization $C \in \{0.01, 0.1, 1, 10, 100\}$ is selected by five-fold stratified cross-validation on the training Gram matrix. The primary P1' protocol divides the ID holdout into disjoint validation and test halves. Within each family, the candidate with the highest ID-validation balanced accuracy is deployed on ID-test and OOD-test; OOD labels never enter this choice. Same-test OOD-oracle and cross-realization OOD-supervised protocols are descriptive sensitivities.

Best-of-family selection favours a larger pool. The primary extended-family endpoint therefore uses 60 candidates per family where coverage is complete and 20 where it is reduced, generated by 5,000 classical-pool subsamples under three schemes. The full 115-candidate classical pool is secondary.

Every configuration has fifteen pipeline realizations (five split seeds by three model seeds). Two EMBER and six network scenario-groups are crossed with three master seeds and $(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$ regimes (1000, 500, 500), (2000, 1000, 1000), and (4000, 1800, 1800). Because scenario-groups share source pools and are not exchangeable, they are fixed case studies. Per group we average over the five split-seed clusters and report a conditional interval; across groups we use dataset-equal means and leave-one-dataset-out ranges. We attach no population p -value to eight fixed, correlated cases. The complete grid is in Supplementary Table 1 and resampling details in the Supplementary Note "Controlled kernel-swap protocol."

4.6 Specification curve

The specification curve is a descriptive multiverse over every logically valid endpoint for which the frozen legacy and v4 artifacts contain complete paired data. Its ten specifications cross, where available, legacy or v4 generation; same-test OOD-oracle, legacy ID-based, or disjoint ID-validation selection; fixed $C = 1$ or five-fold train-CV regularization for SVC; customary linear/RBF or extended classical references; and

native or equal candidate budgets. The order S1–S10 was fixed before aggregation and is never sorted by effect size. Within each specification, pipeline realizations are first averaged within setting and then within scenario-group; groups are averaged within source dataset before the three datasets receive equal weight. All eight group effects are displayed, and no specification-level hypothesis test is performed. Exact definitions and machine-readable group- and dataset-level values are provided in the Supplementary Information and versioned result artifact.

4.7 Prospectively frozen external validation

Before acquiring or inspecting any external quantum–classical result, we froze the tasks, sampling, preprocessing, candidate pools, estimands, uncertainty procedure, and interpretation rule in a version-controlled protocol. The external source is TableShift, a benchmark of real-world distribution shifts in tabular binary classification [Gardner et al. \(2023\)](#). We fixed three public tasks from distinct domains and with distinct published split variables: College Scorecard (institution type), diabetes readmission (admission source), and ACS income (geographic region). TableShift’s predefined `train`, `validation`, `id_test`, and `ood_test` roles are preserved; its `ood_validation` split is excluded from preprocessing, tuning, selection, and reporting.

Each task has two nested sample-size strata: $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{OOD}}) = (1000, 250, 250, 500)$ and $(2000, 500, 500, 1000)$. Five label-blind deterministic subsampling seeds form the conditional resampling clusters. Training-only column filtering, numeric median and categorical mode imputation, one-hot encoding, MaxAbs scaling, TruncatedSVD, standardization, and range mapping produce the same dimensions and kernel pools as in the security study. The primary deployed configuration is selected on `validation`; `ood_test` is touched only for final evaluation. The same-test OOD oracle is retained as a non-deployable diagnostic.

For each task, size, seed, and classifier, the 60-candidate quantum pool is compared with 5,000 equal-size, whole-kernel-block subsamples of the 115-candidate classical pool. The external primary effect is the selected quantum OOD accuracy minus the expected selected classical OOD accuracy over these budget-matched subsamples. The protocol-contraction diagnostic subtracts this controlled effect from the weak-reference, fixed- C , same-test-oracle effect. We average seeds within task and size, give the two nested sizes equal weight within task, and give the three fixed tasks equal weight. Conditional 95% intervals use 9,999 resamples of seed clusters within task; no population p -value is attached to the three deliberately selected tasks. Full acquisition provenance, integrity gates, and the frozen outcome-classification rule are reported in the Supplementary Information.

4.8 Metrics and geometry analysis

OOD balanced accuracy is primary. Family effects are quantum minus classical after selection. ID balanced accuracy and ID-to-OOD degradation are secondary. For the GP classifier we also record log loss, Brier score, 15-bin expected calibration error [Guo et al. \(2017\)](#), and predictive entropy.

From the classifier Gram matrices we compute effective rank [Roy and Vetterli \(2007\)](#), centered kernel-target alignment [Cristianini et al. \(2001\)](#); [Cortes et al. \(2012\)](#), entry concentration, cross-split similarity, and regularized geometric difference [Huang et al. \(2021\)](#). We relate geometry to OOD accuracy within scenario-groups, use partial association to control family and length scale, test portability by leaving one dataset out, and pair quantum candidates with nearest-effective-rank classical candidates within run and dimension. These analyses are associational. Formulas and diagnostic definitions appear in the Supplementary Note “Metrics and geometry descriptors.”

1243 4.9 Finite-shot perturbation

1244

1245 Exact fidelity entries are replaced by binomial estimates at $s \in \{128, 512, 2048, 8192\}$
1246 shots. Training blocks are symmetrized, clipped, set to unit diagonal, and projected
1247 to the positive-semidefinite cone. The model isolates sampling error, not device noise.
1248 Projection diagnostics and interpretation are in the Supplementary Note “Finite-shot
1249 fidelity-estimation model.”

1250

1251

1252 4.10 Use of generative AI

1253

1254 OpenAI Codex was used during manuscript revision to assist with language edit-
1255 ing, code inspection, and consistency checking. The authors reviewed and take
1256 responsibility for all analyses, claims, citations, and text; the system is not an author.

1257

1258 Declarations

1259

1260

1261 **Data availability.** This study uses the public EMBER [Anderson and Roth \(2018\)](#),
1262 UNSW-NB15 [Moustafa and Slay \(2015\)](#), ToN-IoT [Alsaedi et al. \(2020\)](#), and TableShift
1263 [Gardner et al. \(2023\)](#) benchmarks, which remain subject to their original access
1264 conditions. Raw TableShift records are not redistributed. Derived split definitions,
1265 deterministic row-position hashes, shift constructions, experiment manifests, aggre-
1266 gated results, and summary tables are available in the Zenodo repository at [https://](https://doi.org/10.5281/zenodo.19147649)
1267 doi.org/10.5281/zenodo.19147649. Before submission, this DOI will be verified against
1268 the exact archived release used to build the manuscript.

1269

1270 **Code availability.** Code for constructing the splits, running the kernel experiments,
1271 reproducing the confirmatory analysis, and regenerating every table and figure is avail-
1272 able at https://github.com/roberto-fernandez-barrios/kernel_shift_framework and in
1273 the versioned Zenodo archive above. The platform-independent Python 3.12 environ-
1274 ment is pinned in `environment.lock.yml`; the code is released under BSD-3-Clause
1275 with no restriction on non-academic use.

Competing interests.	The authors declare that they have no competing interests.	1289
		1290
Authors' contributions.	RFB: conceptualization, methodology, software, data	1291
	curation, formal analysis, visualization, investigation, writing—original draft,	1292
	writing—review and editing. IPL: supervision, validation, writing—review and editing.	1293
		1294
	AGS: validation, project administration, writing—review and editing. PGB: supervi-	1295
	sion, resources, writing—review and editing. All authors read and approved the final	1296
	manuscript.	1297
		1298
		1299
		1300
Acknowledgements.	The authors received no specific funding for the submitted	1301
	work.	1302
		1303
		1304
		1305
References		1306
		1307
		1308
Abreu, D., Esteve Rothenberg, C., Abelém, A.: QML-IDS: Quantum machine learning		1309
intrusion detection system. In: 2024 IEEE Symposium on Computers and Commu-		1310
nications (ISCC), pp. 1–6. IEEE, ??? (2024). https://doi.org/10.1109/ISCC61673.		1311
		1312
		1313
		1314
		1315
		1316
Abreu, D., Moura, D., Esteve Rothenberg, C., Abelém, A.: QuantumNetSec: Quan-		1317
tum machine learning for network security. International Journal of Network		1318
Management 35 (4), 70018 (2025) https://doi.org/10.1002/nem.70018		1319
		1320
		1321
		1322
Alsaedi, A., Moustafa, N., Tari, Z., Mahmood, A., Anwar, A.: TON_IoT teleme-		1323
try dataset: A new generation dataset of IoT and IIoT for data-driven intrusion		1324
detection systems. IEEE Access 8 , 165130–165150 (2020) https://doi.org/10.1109/		1325
		1326
		1327
		1328
		1329
Arp, D., Quiring, E., Pendlebury, F., Warnecke, A., Pierazzi, F., Wressnegger, C.,		1330
Cavallaro, L., Rieck, K.: Dos and don'ts of machine learning in computer security.		1331
In: 31st USENIX Security Symposium, pp. 3971–3988 (2022)		1332
		1333
		1334

1335 Anderson, H.S., Roth, P.: Ember: An open dataset for training static PE malware
 1336 machine learning models. arXiv **abs/1804.04637** (2018)
 1337
 1338
 1339 Bowles, J., Ahmed, S., Schuld, M.: Better than classical? the subtle art of bench-
 1340 marking quantum machine learning models. arXiv **abs/2403.07059** (2024) [https:](https://doi.org/10.48550/arXiv.2403.07059)
 1342 [//doi.org/10.48550/arXiv.2403.07059](https://doi.org/10.48550/arXiv.2403.07059)
 1343
 1344
 1345 Bordelon, B., Canatar, A., Pehlevan, C.: Spectrum dependent learning curves in ker-
 1346 nel regression and wide neural networks. In: Proceedings of the 37th International
 1347 Conference on Machine Learning (ICML). PMLR, vol. 119, pp. 1024–1034 (2020)
 1348
 1349
 1350
 1351 Bellante, A., Fioravanti, T., Carminati, M., Zanero, S., Luongo, A.: Evaluating the
 1352 potential of quantum machine learning in cybersecurity: A case-study on PCA-
 1353 based intrusion detection systems. Computers & Security **154**, 104341 (2025) [https:](https://doi.org/10.1016/j.cose.2025.104341)
 1355 [//doi.org/10.1016/j.cose.2025.104341](https://doi.org/10.1016/j.cose.2025.104341)
 1356
 1357
 1358 Canatar, A., Bordelon, B., Pehlevan, C.: Spectral bias and task-model align-
 1359 ment explain generalization in kernel regression and infinitely wide neural
 1360 networks. Nature Communications **12**, 2914 (2021) [https://doi.org/10.1038/](https://doi.org/10.1038/s41467-021-23103-1)
 1362 [s41467-021-23103-1](https://doi.org/10.1038/s41467-021-23103-1)
 1363
 1364
 1365
 1366 Cortes, C., Mohri, M., Rostamizadeh, A.: Algorithms for learning kernels based on
 1367 centered alignment. Journal of Machine Learning Research **13**(28), 795–828 (2012)
 1368
 1369
 1370 Canatar, A., Peters, E., Pehlevan, C., Wild, S.M., Shaydulin, R.: Bandwidth enables
 1371 generalization in quantum kernel models. Transactions on Machine Learning
 1372 Research (2023)
 1373
 1374
 1375
 1376 Cristianini, N., Shawe-Taylor, J., Elisseeff, A., Kandola, J.: On kernel-target align-
 1377 ment. In: Advances in Neural Information Processing Systems (NeurIPS), vol. 14
 1378 (2001)
 1379
 1380

Cortes, C., Vapnik, V.: Support-vector networks. Machine Learning 20 (3), 273–297 (1995) https://doi.org/10.1007/BF00994018	1381 1382 1383 1384
Engelen, G., Rimmer, V., Joosen, W.: Troubleshooting an intrusion detection dataset: The CICIDS2017 case study. In: 2021 IEEE Security and Privacy Workshops (SPW), pp. 7–12. IEEE, ??? (2021). https://doi.org/10.1109/SPW53761.2021.00009	1385 1386 1387 1388 1389 1390
Gulrajani, I., Lopez-Paz, D.: In search of lost domain generalization. In: International Conference on Learning Representations (2021). https://openreview.net/forum?id=IQdXeXDoWtl	1391 1392 1393 1394 1395 1396
Gardner, J., Popovic, Z., Schmidt, L.: Benchmarking distribution shift in tabular data with TableShift. In: Advances in Neural Information Processing Systems, vol. 36 (2023). https://doi.org/10.52202/075280-2324	1397 1398 1399 1400 1401 1402
Guo, C., Pleiss, G., Sun, Y., Weinberger, K.Q.: On calibration of modern neural networks. In: Proceedings of the 34th International Conference on Machine Learning (ICML). PMLR, vol. 70, pp. 1321–1330 (2017)	1403 1404 1405 1406 1407 1408
Huang, H.-Y., Broughton, M., Mohseni, M., Babbush, R., Boixo, S., Neven, H., McClean, J.R.: Power of data in quantum machine learning. Nature Communications 12 , 2631 (2021) https://doi.org/10.1038/s41467-021-22539-9	1409 1410 1411 1412 1413 1414
Havlíček, V., Córcoles, A.D., Temme, K., Harrow, A.W., Kandala, A., Chow, J.M., Gambetta, J.M.: Supervised learning with quantum-enhanced feature spaces. Nature 567 , 209–212 (2019) https://doi.org/10.1038/s41586-019-0980-2	1415 1416 1417 1418 1419 1420
Herrero Gómez, P., Jimeno Morenilla, A., Muñoz-Hernández, D., Mora Mora, H.: Spectral Phase Encoding for Quantum Kernel Methods. arXiv:2602.19644 (2026)	1421 1422 1423 1424 1425 1426

1427 Jerbi, S., Gyurik, C., Marshall, S.C., Molteni, R., Dunjko, V.: Shadows of quan-
 1428 tum machine learning. *Nature Communications* **15**, 5676 (2024) [https://doi.org/10.](https://doi.org/10.1038/s41467-024-49877-8)
 1429 [1038/s41467-024-49877-8](https://doi.org/10.1038/s41467-024-49877-8)
 1430
 1431
 1432 Kübler, J.M., Buchholz, S., Schölkopf, B.: The inductive bias of quantum kernels.
 1433 In: *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp.
 1434 12661–12673 (2021)
 1435
 1436
 1437
 1438 Koh, P.W., Sagawa, S., Marklund, H., Xie, S.M., Zhang, M., Balsubramani, A., Hu,
 1439 W., Yasunaga, M., Phillips, R.L., Gao, I., Lee, T., David, E., Stavness, I., Guo,
 1440 W., Earnshaw, B., Haque, I., Beery, S.M., Leskovec, J., Kundaje, A., Pierson, E.,
 1441 Levine, S., Finn, C., Liang, P.: Wilds: A benchmark of in-the-wild distribution
 1442 shifts. In: *Proceedings of the 38th International Conference on Machine Learning.*
 1443 *Proceedings of Machine Learning Research*, vol. 139, pp. 5637–5664. PMLR, ???
 1444 (2021). <https://proceedings.mlr.press/v139/koh21a.html>
 1445
 1446
 1447
 1448
 1449
 1450
 1451 Kakavand, S., Strohmeier, C., Schlotter, M.: Benchmarking quantum kernel support
 1452 vector machines against classical baselines on tabular data: A rigorous empirical
 1453 study with hardware validation. *arXiv abs/2604.18837* (2026) [https://doi.org/10.](https://doi.org/10.48550/arXiv.2604.18837)
 1454 [48550/arXiv.2604.18837](https://doi.org/10.48550/arXiv.2604.18837)
 1455
 1456
 1457
 1458 Liu, Y., Arunachalam, S., Temme, K.: A rigorous and robust quantum speed-up in
 1459 supervised machine learning. *Nature Physics* **17**, 1013–1017 (2021) [https://doi.org/](https://doi.org/10.1038/s41567-021-01287-z)
 1460 [10.1038/s41567-021-01287-z](https://doi.org/10.1038/s41567-021-01287-z)
 1461
 1462
 1463
 1464 Ma, C., Pathak, R., Wainwright, M.J.: Optimally tackling covariate shift in RKHS-
 1465 based nonparametric regression. *The Annals of Statistics* **51**(2), 738–761 (2023)
 1466 <https://doi.org/10.1214/23-AOS2268>
 1467
 1468
 1469
 1470 Moustafa, N., Slay, J.: UNSW-NB15: A comprehensive data set for network intrusion
 1471
 1472

detection systems (UNSW-NB15 network data set),. In: 2015 Military Commu-	1473
nications and Information Systems Conference (MilCIS), pp. 1–6 (2015). https:	1474
//doi.org/10.1109/MilCIS.2015.7348942	1475
	1476
	1477
	1478
Miller, J.P., Taori, R., Raghunathan, A., Sagawa, S., Koh, P.W., Shankar, V., Liang,	1479
P., Carmon, Y., Schmidt, L.: Accuracy on the line: On the strong correlation between	1480
out-of-distribution and in-distribution generalization. In: Proceedings of the 38th	1481
International Conference on Machine Learning (ICML). PMLR, vol. 139, pp. 7721–	1482
7735 (2021)	1483
	1484
	1485
	1486
	1487
Pendlebury, F., Pierazzi, F., Jordaney, R., Kinder, J., Cavallaro, L.: TESSERACT:	1488
Eliminating experimental bias in malware classification across space and time. In:	1489
28th USENIX Security Symposium, pp. 729–746 (2019)	1490
	1491
	1492
	1493
Quiñonero-Candela, J., Sugiyama, M., Schwaighofer, A., Lawrence, N.D. (eds.):	1494
Dataset Shift in Machine Learning. MIT Press, Cambridge, MA (2008). https:	1495
//doi.org/10.7551/mitpress/9780262170055.001.0001	1496
	1497
	1498
	1499
Roy, O., Vetterli, M.: The effective rank: A measure of effective dimensionality. In:	1500
2007 15th European Signal Processing Conference (EUSIPCO), pp. 606–610 (2007)	1501
	1502
	1503
Rasmussen, C.E., Williams, C.K.I.: Gaussian Processes for Machine Learning. MIT	1504
Press, Cambridge, MA (2006). https://doi.org/10.7551/mitpress/3206.001.0001	1505
	1506
	1507
Schuld, M.: Supervised quantum machine learning models are kernel methods. arXiv	1508
abs/2101.11020 (2021) https://doi.org/10.48550/arXiv.2101.11020	1509
	1510
	1511
Schreiber, F.J., Eisert, J., Meyer, J.J.: Classical surrogates for quantum learn-	1512
ing models. Physical Review Letters 131 , 100803 (2023) https://doi.org/10.1103/	1513
PhysRevLett.131.100803	1514
	1515
	1516
	1517
	1518

1519 Schuld, M., Killoran, N.: Quantum machine learning in feature hilbert spaces. Physical
 1520 Review Letters **122**(4), 040504 (2019) [https://doi.org/10.1103/PhysRevLett.122.](https://doi.org/10.1103/PhysRevLett.122.040504)
 1521 [040504](https://doi.org/10.1103/PhysRevLett.122.040504)
 1522
 1523
 1524 Schnabel, J., Roth, M.: Quantum kernel methods under scrutiny: A benchmark-
 1525 ing study. Quantum Machine Intelligence **7**, 58 (2025) [https://doi.org/10.1007/](https://doi.org/10.1007/s42484-025-00273-5)
 1526 [s42484-025-00273-5](https://doi.org/10.1007/s42484-025-00273-5)
 1527
 1528
 1529 Schuld, M., Sweke, R., Meyer, J.J.: Effect of data encoding on the expressive power of
 1530 variational quantum-machine-learning models. Physical Review A **103**(3), 032430
 1531 (2021) <https://doi.org/10.1103/PhysRevA.103.032430>
 1532
 1533
 1534 Shaydulin, R., Wild, S.M.: Importance of kernel bandwidth in quantum machine learn-
 1535 ing. Physical Review A **106**(4), 042407 (2022) [https://doi.org/10.1103/PhysRevA.](https://doi.org/10.1103/PhysRevA.106.042407)
 1536 [106.042407](https://doi.org/10.1103/PhysRevA.106.042407)
 1537
 1538
 1539 Thanasilp, S., Wang, S., Cerezo, M., Holmes, Z.: Exponential concentration in
 1540 quantum kernel methods. Nature Communications **15**, 5200 (2024)
 1541
 1542 Wang, X., Du, Y., Luo, Y., Tao, D.: Towards understanding the power of quan-
 1543 tum kernels in the NISQ era. Quantum **5**, 531 (2021) [https://doi.org/10.22331/](https://doi.org/10.22331/q-2021-08-30-531)
 1544 [q-2021-08-30-531](https://doi.org/10.22331/q-2021-08-30-531)
 1545
 1546 Yao, H., Choi, C., Cao, B., Lee, Y., Koh, P.W., Finn, C.: Wild-time: A benchmark of in-
 1547 the-wild distribution shift over time. In: Advances in Neural Information Processing
 1548 Systems (2022). Datasets and Benchmarks Track
 1549
 1550 Yang, L., Guo, W., Hao, Q., Ciptadi, A., Ahmadzadeh, A., Xing, X., Wang, G.: CADE:
 1551 Detecting and explaining concept drift samples for security applications. In: 30th
 1552 USENIX Security Symposium, pp. 2327–2344 (2021)
 1553
 1554
 1555
 1556
 1557
 1558
 1559
 1560
 1561
 1562
 1563
 1564

Zhou, K., Liu, Z., Qiao, Y., Xiang, T., Loy, C.C.: Domain generalization: A survey.	1565
IEEE Transactions on Pattern Analysis and Machine Intelligence 45 (4), 4396–4415	1566
(2023) https://doi.org/10.1109/TPAMI.2022.3195549	1567
	1568
	1569
	1570
	1571
	1572
	1573
	1574
	1575
	1576
	1577
	1578
	1579
	1580
	1581
	1582
	1583
	1584
	1585
	1586
	1587
	1588
	1589
	1590
	1591
	1592
	1593
	1594
	1595
	1596
	1597
	1598
	1599
	1600
	1601
	1602
	1603
	1604
	1605
	1606
	1607
	1608
	1609
	1610