

Sharp Target-Domain Certificates for Quantum-Kernel Advantage under Distribution Shift

Roberto Fernández-Barrios^{[ORCID]¹*}, Iker Pastor-López^{[ORCID]¹},
Asier González-Santocildes^{[ORCID]¹}, Pablo García Bringas^{[ORCID]¹}

¹*Faculty of Engineering, University of Deusto, Avda. de las
Universidades, 24., Bilbao, 48007, Spain.

*Corresponding author(s). E-mail(s): roberto.fernandez.b@deusto.es;
Contributing authors: iker.pastor@deusto.es; gonzalez.asier@deusto.es;
pablo.garcia.bringas@deusto.es;

Abstract

Claims of quantum predictive advantage under distribution shift are usually evaluated only after target labels are known. We derive the assumption-free sharp identified set for a fixed model’s finite-batch advantage over the best member of a prespecified family under any bounded loss; zero–one accuracy has a closed form and exact monotone partial-label updates. Across eight security shifts, zero-label upper endpoints against 115 classical kernels span 0.002–0.088. A retrospective audit reduces every endpoint to 0.010 with 0–33 of 500 labels, including entangling-ZZ models. In an internally locked prospective replication, four eligible task–classifier medians require 0–3 labels (overall median 0; maximum 76), and all 20 realized effects are negative; a third task fails its frozen feature gate. Protocol controls reverse an apparent advantage, while finite-shot noise can increase predictive distinctness without useful advantage. The framework separates protocol validity, target predictive indispensability, and quantum relevance.

Keywords: quantum machine learning; quantum kernels; quantum advantage; partial identification; active testing; distribution shift; classical surrogates; kernel methods

047 1 Introduction

048
049 Quantum kernel methods embed classical data into quantum states and use state
050 overlaps as similarities for otherwise classical learners [1–3]. This separation makes
051 them a clean setting for studying quantum inductive bias, but a predictive-advantage
052 claim must pass three distinct tests. The comparison protocol must be valid; the target
053 predictions must resist reproduction by an appropriate classical reference family; and
054 the surviving behaviour must depend on a non-trivial quantum construction whose
055 resources are counted. A model can pass one test and fail another.

056
057 Recent benchmarks emphasize the first test. Broad QML comparisons find compet-
058 itive classical baselines once pipelines are aligned [4]; quantum-kernel results depend
059 strongly on feature maps and hyperparameters [5]; and controlled tabular studies
060 often find no stable quantum win. A concurrent finance preprint likewise shows an
061 apparent quantum-kernel advantage vanishing under point-in-time, kernel-swap, and
062 budget controls [6]. Under distribution shift the selection rule becomes part of the esti-
063 mand: same-test selection is an oracle, train-only regularization can change rankings,
064 and a larger candidate family has more opportunities to win [7–10]. Fair equal-budget
065 benchmarking answers who wins under symmetric search, but not whether any clas-
066 sical witness in a declared family can reproduce the selected quantum model on the
067 deployment inputs.

068
069 Geometric and surrogate analyses approach that second question from complemen-
070 tary directions. Huang et al. use a label-independent geometric difference to screen
071 whether a quantum kernel can have a prediction advantage [11]; classical-surrogate
072 constructions reproduce input–output relations of trained quantum models [12]; and
073 recent theory studies when quantum learners depart from minimum-norm classical
074 solutions [13]. Those objects characterize geometric potential, global approximation,
075 or learning-rule separation. Our target-batch question is different. A train-only Gram
076 matrix does not determine how a kernel transports predictions to a shifted target
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092

because deployment also depends on the target–train block, while different geometries can yield identical decisions on the observed target inputs. We therefore audit predictive indispensability directly on target covariates while keeping their labels protected.

That protection creates an identification problem. Disagreement itself is established: co-validation relates independently trained classifiers’ disagreement to error and variance [14], discordant-pair evaluation concentrates adjudication where two classifiers differ [15], Active Testing constructs low-variance risk estimates from actively sampled test labels [16], and limited-label model selection aims to choose the best pretrained classifier [17]. Weak-supervision work instead partially identifies individual performance metrics from fitted label-model marginals [18]. We do not claim any of those ingredients. We identify the *relative* finite-batch advantage of one fixed model against the best member of a prespecified family, exactly over every completion of the unaudited labels, without estimating those labels, postulating their marginals, or assuming a target sampling model.

We make four contributions. First, for any additive bounded loss on a finite label space, we derive the sharp partial-label identified set for advantage over the best member of a finite reference family and show that no assumption-free rule using the same information can return a uniformly smaller valid interval. Zero–one accuracy has an unusually simple closed form, including after any adaptive partial audit. Second, nested classical reference families and target-label audits define a reference-breadth–target-supervision evidence frontier. Third, we give an exact prevalence-conditional balanced-accuracy formulation and an explicit PSD construction showing why train-only geometry cannot determine target predictions. Fourth, we embed these objects in a controlled quantum-kernel benchmark and internally lock a prospective Gate 2 replication before opening target labels; entanglement strata and repeated finite-shot perturbations separately test when predictive distinctness is quantum-relevant.

139 The benchmark uses four security scenarios from EMBER, UNSW-NB15, and ToN-
140 IoT, producing eight fixed constructed or held-out-campaign shifts. Support-vector
141 and Laplace Gaussian-process classifiers consume identical precomputed Gram matri-
142 ces. The quantum family contains entangling-ZZ and separable-product fidelity maps
143 at 4–12 qubits; the classical family contains 115 linear, RBF, polynomial, Laplacian,
144 and Matérn candidates. Models are tuned on training data and selected on a disjoint
145 ID-validation split. Equal 60-candidate pools support fair performance comparison,
146 while the complete prespecified classical family supports the separate indispensability
147 audit. One prospectively frozen three-task TableShift study corroborates protocol sen-
148 sitivity; a second, internally locked experiment tests transfer of Gate 2 on previously
149 unused task identifiers.

157 Against the full classical family, the zero-label upper endpoint of accuracy advan-
158 tage ranges from 0.002 to 0.050 for SVC and 0.006 to 0.088 for GPC. In a retrospective
159 fixed-batch analysis, an adaptive bottleneck-coverage audit reduces it to at most
160 0.010 using 0–33 of 500 labels (median 13 over 16 fixed case-classifier cells). Dynamic
161 random-disagreement and non-adaptive initial-coverage controls both have median
162 12.5, so the large gain over prediction-independent ordering is attributable to exclud-
163 ing non-actionable queries rather than to the proposed coverage heuristic. The result
164 persists for selected entangling-ZZ models. In the prospective Gate-2 replication, two
165 technically eligible tasks produce four task-classifier medians of 0–3 labels (over-
166 all seed-level median 0; maximum 76), satisfying the frozen strong-transfer rule; the
167 third task fails its prespecified minimum-feature gate before model execution. All 20
168 prospective realized accuracy effects are negative. The certificate remains conditional
169 on the finite classical family and target batch, not evidence of classical simulability or
170 complexity-theoretic separation. It converts the earlier benchmark conclusion—that
171
172
173
174
175
176
177
178
179
180
181
182
183
184




Fig. 1 The controlled kernel-swap protocol. Benchmark pools feed reproducible shift constructions; all kernels operate on a shared train-only embedding, receive per-configuration regularization tuned on the training split, and are consumed by the same precomputed-kernel classifiers. Family-level conclusions are drawn by target-label-free, equal-budget selection on an in-distribution validation split (P1'). Each scenario-group interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the fixed benchmark pools and design. Gram matrices feed the associational geometry audit and repeated finite-shot sensitivity.

controls remove an apparent gain—into a quantitative question: how much advantage remains possible, and how much classical search and target supervision suffice to falsify a material claim?

2 Results

2.1 Overview

The results are organized as three gates. Gate 1 asks whether the comparison protocol itself is valid: regularization is fitted on training data, deployment candidates are selected without OOD labels, and family search budgets are distinguished from unrestricted adversarial reference breadth. Gate 2 asks whether the selected quantum model is predictively indispensable on the observed target inputs. This is the new target-batch certificate and evidence frontier. Gate 3 asks whether any surviving behaviour is quantum-relevant, separating entangling- ZZ maps from product maps and

accounting for finite-shot estimation. The eight security groups are fixed cases; conditional intervals over q-split clusters describe pipeline variation given the benchmark pools and design, not uncertainty over a population of datasets.

2.2 Sharp target-batch identification of predictive advantage

Consider a finite target batch, a fixed quantum prediction, and M fixed classical predictions. For a finite label space $\mathcal{Y}$ and an additive bounded loss, let $\ell_i^Q(y)$ and $\ell_{ij}^C(y)$ be their losses on item i if its label were y , and define the classical-minus-quantum contrast $a_{ij}(y) = \ell_{ij}^C(y) - \ell_i^Q(y)$. Positive advantage favours the quantum model:

$$\Delta_\ell(y) = \min_{j \leq M} \frac{1}{n} \sum_{i=1}^n a_{ij}(y_i).$$

If labels have been audited on L and $S_j(L) = \sum_{i \in L} a_{ij}(y_i)$, every assignment on L^c remains admissible.

Theorem 1 (Sharp bounded-loss region and information optimality) *The smallest interval containing $\Delta_\ell(y)$ for every completion of the unaudited labels is $[L_\ell(L), U_\ell(L)]$, where*

$$L_\ell(L) = \frac{1}{n} \min_j \left\{ S_j(L) + \sum_{i \notin L} \min_{y \in \mathcal{Y}} a_{ij}(y) \right\},$$

$$U_\ell(L) = \frac{1}{n} \max_{y_{L^c} \in \mathcal{Y}^{|L^c|}} \min_j \left\{ S_j(L) + \sum_{i \notin L} a_{ij}(y_i) \right\}.$$

Both endpoints are attained. Consequently, any assumption-free interval computed from the same fixed losses and audited labels and valid for every admissible completion must contain both endpoints.

The lower endpoint separates over labels after choosing a witness. The upper endpoint is an exact finite max-min problem, expressible as a mixed-integer linear program with one-hot label variables. Auditing any additional label removes

completions, so the lower endpoint cannot decrease and the upper endpoint cannot increase. This general statement is not a claim that loss-based partial identification or active evaluation is new; the contribution is the exact relative-family estimand and its information-optimal finite-batch certificate. Zero-one classification has further structure: assigning an unaudited label to the quantum hard prediction simultaneously maximizes its contrast with every disagreeing classical witness. This yields the following closed forms used in all experiments.

Fix the hard target predictions of a selected quantum classifier h_Q and a pre-specified classical family $\mathcal{C}_B = \{h_{C_1}, \dots, h_{C_M}\}$ on n target inputs. For an otherwise unrestricted target labelling y , define

$$\Delta_B(y) = \text{Acc}(h_Q; y) - \max_{j \leq M} \text{Acc}(h_{C_j}; y), \quad d_j = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{h_Q(x_i) \neq h_{C_j}(x_i)\}.$$

Theorem 2 (Sharp zero-label region) *The smallest interval containing $\Delta_B(y)$ for every possible target labelling is*

$$\mathcal{I}_B^{(0)} = \left[-\max_{j \leq M} d_j, \min_{j \leq M} d_j \right].$$

Both endpoints are attainable, including for multiclass hard predictions.

The upper endpoint is not a probabilistic bound or an estimate: it is the exact largest advantage compatible with the observed prediction matrix when every target label remains unknown. It is attained by setting $y = h_Q$. The lower endpoint is attained by assigning the predictions of the classical model farthest from h_Q . Hence no smaller label-free interval is valid without extra assumptions. In particular, if $\min_j d_j < \tau$, any accuracy advantage of at least τ over this classical family is falsified on the batch, even before a target label is revealed. Whenever $\min_j d_j > 0$, the same target covariates and prediction matrix remain compatible with both signs of advantage; disagreement alone cannot identify which family wins.

The zero-label certificate is already restrictive in several cases (Figure 2a). Against the complete 115-candidate classical family, the upper endpoint ranges from 0.002 to 0.050 for SVC and from 0.006 to 0.088 for GPC (Table 1). A small endpoint says that at least one classical witness nearly reproduces the selected quantum decisions; it does not say that the quantum kernel is efficiently simulable or globally reproducible.

Theorem 3 (Sharp arbitrary-partial-label region) *Let L be any audited subset. For classical witness j , let*

$$S_j(L) = \sum_{i \in L} [\mathbf{1}\{h_Q(x_i) = y_i\} - \mathbf{1}\{h_{C_j}(x_i) = y_i\}]$$

and let $R_j(L)$ count unaudited points on which h_Q and h_{C_j} disagree. Then the sharp identified interval over all completions of the unaudited labels is

$$\mathcal{I}_B(L) = \left[\min_j \frac{S_j(L) - R_j(L)}{n}, \min_j \frac{S_j(L) + R_j(L)}{n} \right].$$

The upper endpoint is non-increasing as labels are audited and collapses to the realized family advantage when L contains the full batch.

We use this result as an active falsification audit, not as a new model-selection rule. At each step the models remain fixed; one tested rule queries the point covering the largest number of witnesses currently tying at the smallest reducible upper bound (Methods). In this retrospective fixed-batch analysis, the rule reduces the 0.010 endpoint using 0–33 of 500 labels, with median 13 across the 16 case–classifier cells (Figure 2b). Stronger controls remove any claim that the coverage heuristic itself is unusually efficient: dynamic random sampling restricted to current actionable disagreements and a fixed initial-coverage ordering both have median 12.5 labels. Against random-active sampling, the adaptive rule wins, ties, and loses in 5, 5, and 6 cells; against initial coverage the counts are 6, 4, and 6. An exact ex-post label oracle needs median 6.5 labels, and the adaptive rule’s median excess is 4.5. The main acquisition result is therefore that prediction-aware restriction makes a small audit possible,




Fig. 2 Sharp target-batch certificates against the full 115-candidate classical family. **a**, Zero-label upper endpoints for the SVC and GPC quantum models fixed by P1'. Colour separates selected entangling-ZZ and product-map models; the dashed line marks 0.010 accuracy. **b**, Labels needed to reduce the endpoint to at most 0.010 (symmetric-log scale). Horizontal ticks are the exact retrospective label oracle; filled markers are the fixed adaptive coverage rule; open diamonds and triangles are medians for dynamic random-active and non-adaptive initial-coverage controls, with bars spanning 2.5–97.5% across 200 SHA-256 tie orders. All certificates use the canonical q1000 realization (master, q-split, and model seeds 42) and a 500-point target batch. Comparator ranges are tie-order sensitivities, not confidence intervals.

Table 1 Sharp accuracy certificates for the eight fixed target batches against the full 115-candidate classical family. U_0 is the zero-label upper endpoint; $n_{.01}$ is the fixed adaptive-coverage count for $U \leq 0.010$; Δ_{full} is the point-identified quantum advantage after all 500 labels are unlocked. The canonical run, quantum model, and every classical candidate were fixed before this audit; target labels are never used to retune or reselect them.

Target case	SVC			GPC		
	U_0	$n_{.01}$	Δ_{full}	U_0	$n_{.01}$	Δ_{full}
EMBER $m1$	.050	12	-.024	.080	29	-.030
EMBER $m2$	.054	15	-.036	.074	33	-.016
ToN-IoT constructed	.002	0	-.004	.016	3	-.014
ToN-IoT campaign	.002	0	-.018	.006	0	-.012
UNSW DoS constructed	.010	0	-.016	.066	22	-.052
UNSW DoS campaign	.034	11	-.026	.028	12	-.032
UNSW Recon constructed	.046	14	-.032	.088	22	-.066
UNSW Recon campaign	.034	28	-.026	.030	14	+.002

not that one heuristic dominates. Prediction-independent ordering remains a weak diagnostic (median 245.5) and is reported only in Supplementary Information.

2.3 A reference-breadth–target-supervision evidence frontier

Nested classical reference families and partial labels expose two different costs of falsification. Increasing the classical family can only decrease the certificate’s upper

415 endpoint; auditing informative target labels can only decrease it further. We therefore
 416 report the surface $(B, |L|) \mapsto U_B(L)$ and, for a materiality threshold τ , the smallest
 417 reference breadth and target supervision required to establish $U_B(L) \leq \tau$. Here B
 418 indexes a prespecified *reference breadth*, not runtime or computational complexity.

421 The frozen block-ordering analysis already provides three nested breadths: 30, 60,
 422 and 115 candidates. For each cell we take the median zero-label endpoint over 5,000
 423 prespecified whole-block orderings, then summarize the 16 case-classifier cells. The
 424 cross-cell median contracts from 0.042 at $B = 30$ to 0.034 at $B = 60$ and remains
 425 0.034 at $B = 115$; the number of cells already at or below 0.010 changes from 3
 426 to 4 to 4 (Supplementary reference-breadth table). Thus most aggregate contraction
 427 occurs by 60 candidates, but the full family remains necessary for the deliberately
 428 adversarial question and for individual cases. These ordering distributions quantify
 429 dependence on reference breadth; they are not confidence intervals or computational-
 430 cost measurements.

431 The prediction-aware evidence curves contract rapidly because only disagreements
 432 capable of lowering a current witness endpoint need adjudication (corresponding curve
 433 display in Supplementary Information). The adaptive rule is not consistently better
 434 than random-active sampling, although recomputing the active set avoids the 286.5-
 435 label median outlier produced by a frozen initial-coverage ordering. The result persists
 436 when the selected quantum model uses an entangling-ZZ map: across the 12 entangling
 437 case-classifier cells, the adaptive rule needs 0–33 labels (median 14) for the 0.010
 438 endpoint, while random-active cell medians are 0–34 (median 12.5). Thus the finding
 439 is not an artefact of including classically factorizable product maps.

440 Reference breadth matters separately from fair benchmarking. For ToN-IoT scan-
 441 ning under GPC, the customary 30-candidate linear/RBF family yields a realized
 442 quantum accuracy advantage of +0.028, so no number of labels can falsify a one-
 443 point claim against that restricted reference. Against the full 115-candidate family the
 444

realized advantage is -0.014 , the zero-label upper endpoint is 0.016 , and three targeted labels reduce it below 0.010 . Equal 60-versus-60 budgets remain the appropriate performance comparison; the full family asks the different, deliberately adversarial question of predictive indispensability.

2.4 Gate 2 transfers prospectively in a technically limited replication

The security acquisition curves above are retrospective. We therefore froze a separate Gate-2 replication before acquiring its three TableShift tasks, computing a model prediction, or opening a target label. The task identifiers—BRFSS Diabetes, ACS Food Stamps, and NHANES Lead—were inherited from an unused fallback list in the earlier external protocol. The immutable specification fixed five label-blind sampling seeds, the q1000 split sizes, P1' selection, both classifiers, the 60-quantum and 115-classical pools, physical prediction-label separation, three thresholds, 200 SHA-256 orders for each strong control, and numerical outcome categories. The aggregate prediction lock was hashed before the one-time label audit (Methods).

All five NHANES seeds produced only 11 non-constant post-encoding features, below the frozen 12-dimensional grid requirement. The task therefore failed a pre-specified minimum-feature gate before any model candidate was executed; it was not replaced or repaired by reducing the grid. The replication is consequently technically limited to ten task-seed units, or 20 fixed task-seed-classifier cells, rather than evidence from three complete tasks.

The available result meets the frozen strong-transfer category (Figure 3). Against the full classical family, zero-label endpoints range from 0 to 0.168 (median 0.010), and 11/20 cells already satisfy $U \leq 0.010$ without target labels. The fixed adaptive audit reduces every remaining endpoint to that threshold with 0–76 labels (overall median 0). The four task-classifier medians are 0 for ACS SVC, 3 for ACS GPC, 3 for BRFSS




Fig. 3 Prospectively locked Gate-2 replication against the full 115-candidate classical family. **a**, Zero-label upper endpoints for five deterministic sampling seeds in each technically eligible task-classifier cell; black ticks mark medians and the dashed line marks 0.010 accuracy. **b**, Labels required to reduce the endpoint to at most 0.010. Coloured points are the five fixed adaptive counts, black ticks are exact retrospective label-oracle counts, and open markers with bars are medians and 2.5–97.5% quantiles for dynamic random-active and non-adaptive initial-coverage policies across 1,000 seed-order combinations (five seeds times 200 SHA-256 orders). The 50-label line is the frozen task-classifier median criterion; the dotted 100-label line is the frozen seed-level ceiling. Comparator ranges are algorithmic sensitivities, not confidence intervals. NHANES Lead failed the prespecified 12-feature gate before model execution and is absent.

SVC, and 0 for BRFSS GPC; all are below the 50-label criterion, and the worst seed remains below the frozen 100-label ceiling. Every realized accuracy effect is negative, ranging from -0.156 to -0.002 .

The prospective controls reproduce the earlier negative algorithmic conclusion. Against both dynamic random-active and non-adaptive initial-coverage medians, adaptive coverage wins, ties, and loses in 3, 15, and 2 of the 20 cells, respectively; the median paired difference is zero. Prediction-independent ordering is far less efficient in the non-trivial cells. Thus the transferred result concerns the sharp certificate and restriction to actionable disagreements, not superiority of the coverage heuristic. All prospectively selected quantum winners are product maps, so this replication does not independently validate the entangling-ZZ stratum; that stratum remains supported only by the retrospective security cases.

2.5 The apparent advantage under oracle selection

We begin with a deliberately optimistic reference comparison: fidelity-based quantum kernels against linear and RBF baselines, with each family’s representative chosen by Best-by-OOD selection—the configuration with the highest balanced accuracy on the OOD test split itself—and the classical kernels left at the default regularization $C = 1$. Related intrusion-detection studies motivate the small-reference context [19–21], but we do not claim that this exact protocol dominates QML practice. Under these conventions, on EMBER, the quantum family shows an OOD advantage of +0.037 under SVC and +0.028 under the GP classifier against linear+RBF, shrinking to +0.004/+0.002 once the classical family is enlarged with heavier-tailed kernels. Supplementary Table 5 gives the complete oracle results.

Two features favour this apparent quantum result. Best-by-OOD is an *oracle*: it uses the labels it then reports. Fixing $C = 1$ also leaves classical candidates at a potentially unsuitable regularization point. These values are therefore descriptive upper bounds, not deployable effects.

2.6 The controlled comparison does not preserve the apparent advantage

A deployable comparison must tune each configuration’s regularization without looking at the test data, select the deployed configuration without OOD labels, and compare families at the same candidate budget. We implement all three (Section 4). For every kernel candidate, the SVC regularization C is chosen by five-fold cross-validation on the training Gram matrix alone. The deployed configuration is then the one with the best balanced accuracy on an in-distribution *validation* split, held out from the ID test split (protocol P1’); no OOD label is ever consulted in selection. The primary comparison uses 60 candidates per extended family in every scenario-group and classifier, with the full 115-candidate classical pool retained as a sensitivity.




Fig. 4 Controlled target-label-free family effects in eight fixed scenario-groups. P1' OOD differences (quantum minus classical) are shown for SVC (left) and GPC (right). Filled circles compare against the equal-budget extended classical family; error bars are pointwise 95% conditional cluster- t intervals over five q-split clusters ($n = 5$), conditional on the fixed benchmark pools and design. Open diamonds compare against the customary linear+RBF reference. The intervals are not simultaneous and do not imply population inference or equivalence.

We treat the eight scenario-groups as fixed case studies and report, per group, the quantum-minus-classical OOD delta with a conditional interval over the five split realizations, and no population p -value (Section 4).

Table 2 and Figure 4 show no consistent evidence of a quantum advantage under the controlled endpoint. Against the equal-budget extended classical family, the source-dataset-equal difference is -0.00565 for SVC (leave-one-source-dataset-out range $[-0.00702, -0.00450]$) and -0.00094 for GPC ($[-0.00190, +0.00094]$). Against the customary linear+RBF reference, the corresponding differences are -0.00407 and $+0.00591$, with heterogeneous group signs. UNSW-NB15 reconnaissance under held-out-campaign shift has a small positive mean ($+0.005$ for both classifiers), with pointwise conditional intervals spanning zero. The full-pool extended-family sensitivity is similar in direction and scale (-0.00378 for SVC and -0.00293 for GPC).

Table 2 Target-label-free selection results for eight fixed scenario-groups. Values are quantum-minus-classical OOD balanced-accuracy differences under $P1'$ (select on ID-validation, tune per-configuration C on train, evaluate on OOD test). The frozen extended-family analysis uses 60 candidates per family in every group-classifier cell. Each scenario interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the benchmark pools and design; the fifteen pipeline realizations per setting are nested within these clusters. An interval containing zero is not an equivalence result. Source-dataset-equal rows average groups within three source datasets and report leave-one-source-dataset-out ranges, not confidence intervals. The customary linear+RBF comparison is contextual rather than equal-budget.

Scenario	Clf.	vs lin./RBF	vs ext. (equal B)	interval / LODO range
EMBER $m1$	SVC	−0.0095	−0.0058	[−0.0340, +0.0223]
	GPC	−0.0050	−0.0009	[−0.0101, +0.0083]
EMBER $m2$	SVC	+0.0139	−0.0064	[−0.0137, +0.0009]
	GPC	+0.0149	+0.0027	[−0.0027, +0.0081]
UNSW-DoS (campaign)	SVC	−0.0043	−0.0033	[−0.0139, +0.0073]
	GPC	+0.0041	−0.0057	[−0.0120, +0.0005]
UNSW-DoS (constructed)	SVC	−0.0085	−0.0044	[−0.0119, +0.0031]
	GPC	−0.0056	−0.0108	[−0.0177, −0.0039]
UNSW-Recon (campaign)	SVC	+0.0004	+0.0050	[−0.0058, +0.0157]
	GPC	+0.0074	+0.0052	[−0.0019, +0.0123]
UNSW-Recon (constructed)	SVC	−0.0076	−0.0089	[−0.0146, −0.0031]
	GPC	−0.0083	−0.0075	[−0.0208, +0.0059]
ToN-IoT (campaign)	SVC	−0.0172	−0.0148	[−0.0268, −0.0029]
	GPC	−0.0046	−0.0090	[−0.0190, +0.0010]
ToN-IoT (constructed)	SVC	−0.0016	−0.0011	[−0.0049, +0.0028]
	GPC	+0.0313	+0.0109	[+0.0017, +0.0202]
Source-dataset-equal mean, SVC		−0.0041	−0.0056	[−0.0070, −0.0045]
Source-dataset-equal mean, GPC		+0.0059	−0.0009	[−0.0019, +0.0009]

The candidate-budget result is insensitive to the sampling scheme.

The primary `kernel.blocked` scheme samples 12 whole classical kernel-shape blocks crossed with all five dimensions, matching both the quantum count and its feature-map-by-dimension structure. Uniform and kernel-stratified samples match the count but not this block structure. The source-dataset-equal resampling estimates remain close under all three schemes: −0.00576, −0.00555, and −0.00548 for SVC, and −0.00185, −0.00232, and −0.00253 for GPC, respectively. The largest scheme range in any group-classifier cell is 0.00353. Supplementary Table 2 reports all 48 cells. Candidate-count equality is therefore distinct from search-space isomorphism, but this distinction does not drive the conclusion here.

Rank-matched differences remain negative within the predefined overlap.

We pair each quantum configuration with the classical configuration of nearest training effective rank within the same run and dimension. This observational matching reuses classical candidates and applies the prespecified caliper $\max(r_Q, r_C)/\min(r_Q, r_C) \leq 1.25$. It retains 48,708 of 64,800 pairs (75.2%; 69.3–83.3% by group). Supplementary Figure 3 reports negative median quantum-minus-classical differences in all eight groups (-0.0013 to -0.0094); the quantum member is more accurate in 22–37% of retained pairs. The pooled median is -0.0056 . Minimum-cost one-to-one matching retains 24,675 pairs at the same caliper and yields a pooled median of -0.0053 , again negative in every group. Results remain negative over calipers 1.10, 1.50, 2.00, and without a caliper (Supplementary Table 3). These pairings describe the predefined region of rank overlap; they do not identify a causal effect of kernel family.

2.7 The quantum pool mixes entangling and product kernels

Circuit decomposition changes the interpretation of the aggregate family. The two ZZFeatureMap variants contain all-to-all two-qubit Pauli interactions, whereas the Pauli-X/Z and ZFeatureMap variants contain only one-qubit Pauli strings. For the latter pair, the prepared state and its fidelity kernel factorize over coordinates (Methods), permitting direct classical evaluation as a product of one-qubit overlaps. Thus two of the four map families cannot supply evidence that entanglement or hard quantum kernel estimation is responsible for their behaviour.

Table 3 reports logical, unrouted resource counts. At 12 qubits, one ZZFeatureMap state preparation requires 132 CX gates for one repetition and 264 for two; the corresponding compute–uncompute fidelity templates require 264 and 528 CX gates. Both product-map templates require zero CX gates at every studied dimension. Under P1' selection, product maps supply 594/1,080 (55.0%) of the selected quantum SVC configurations and 348/1,080 (32.2%) of the GP selections. The composition alone does

Table 3 Logical circuit audit for the four quantum feature maps. Depth and CX columns give the values at 4 and 12 qubits after Qiskit 2.3.1 decomposition to **rz**, **sx**, **x**, and **cx**, optimization level 0, without device routing. “Fidelity CX” counts the compute–uncompute template for one kernel entry. Complete values at all five dimensions and one-qubit gate counts are provided in Supplementary Table 13 and the machine-readable artifact. These logical counts are not hardware-execution estimates.

Map	Pauli strings	Stratum	Product factorization	map depth (4/12)	map CX (4/12)	fidelity CX (4/12)
ZZ, 1 rep.	Z, ZZ	entangling ZZ	no	19/67	12/132	24/264
ZZ, 2 reps.	Z, ZZ	entangling ZZ	no	35/107	24/264	48/528
Pauli-X/Z, 1 rep.	X, Z	separable product	yes	11/11	0/0	0/0
Z, 2 reps.	Z	separable product	yes	8/8	0/0	0/0

not compare predictive quality; the matched stratum analysis below does so at 30 candidates per stratum.

At matched 30-candidate budgets, the source-dataset-equal SVC effect is -0.00928 for the entangling-ZZ stratum and -0.00547 for the product-map stratum; the constituent source-dataset means range from -0.01158 to -0.00580 and from -0.01167 to $+0.00021$, respectively. The entangling-ZZ point estimate is negative in every fixed scenario-group (-0.01960 to -0.00132); product-map estimates range from -0.02245 to $+0.00571$. The full 60-candidate pool remains at -0.00566 . Thus the controlled result is not produced by pooling product maps with entangling circuits, and restricting the quantum family to ZZ interactions does not reveal a hidden positive family effect. Supplementary Table 8 reports every pointwise conditional interval; these post-review strata are fixed-case sensitivities, not new population tests.

2.8 Within-v4 sensitivities separate the evaluation bundle

The within-v4 factorial removes the generation boundary from the historical specification map. In the q1000 SVC stratum, the fixed- C , same-test-oracle, customary-reference, native-budget endpoint is $+0.0131$, whereas the train-CV, ID-validation, extended-reference, equal-count endpoint is -0.0047 , a descriptive contraction of 0.0178 . The 16 three-source-dataset-equal endpoints range from -0.0084 to $+0.0131$ (Figure 5; Supplementary Table 10). Thus the sign change is present within one result generation and does not require the legacy-to-v4 boundary.

Averaging paired changes over the other factorial axes gives -0.0084 for fixed $C = 1$ to train-CV, -0.0079 for customary to extended reference, -0.0037 for same-test oracle to ID validation, and -0.0002 for native to equal-count budget. These finite-design averages are not independent or causal contributions. In particular, the mean regularization-by-reference difference-in-differences is $+0.0139$, whereas the other five mean pairwise interactions have absolute magnitude at most 0.0034 (Supplementary Table 11). Regularization and reference strength must therefore be interpreted as part of an interacting evaluation bundle rather than as an additive attribution.

Removing the frozen port-, protocol-, and service-related fields also changes the relative comparison, confirming that common feature access alone does not make short-cut use family-invariant. Across the two network source datasets, the controlled effect moves from -0.0031 to -0.0245 (change -0.0215 ; source-dataset changes -0.0447 for ToN-IoT and $+0.0018$ for UNSW-NB15). The six group changes range from -0.0794 to $+0.0099$, and every pointwise change interval includes zero (Supplementary Table 12). The large negative change is concentrated in the ToN-IoT campaign shift, where removing direct traffic descriptors reduces both families’ absolute accuracy but the selected fidelity kernel more strongly. This sensitivity neither restores a positive aggregate quantum effect nor establishes that all shortcuts have been removed.

2.9 Protocol sensitivity is prospectively corroborated on external domain shifts

We next applied the prospectively frozen protocol to College Scorecard, diabetes read-mission, and ACS income, preserving the published TableShift domain splits and withholding all OOD labels from the controlled selector. The complete external grid contains 30 task-size-seed units and 37,800 split-level results; every prespecified unit passed the automated author-provided audit. Figure 6 shows all task-size effects and conditional seed-cluster intervals.




Fig. 5 Circuit-aware and within-v4 fixed-case sensitivities for SVC. **a**, Quantum-minus-classical OOD balanced-accuracy effects for all quantum maps ($B = 60$), entangling-ZZ maps ($B = 30$), and separable-product maps ($B = 30$), each against a blocked equal-count extended-classical reference. Coloured points are the eight scenario-groups and black diamonds are three-source-dataset-equal means. **b**, Original and ablated controlled effects after removing the frozen port-, protocol-, and service-related fields in the six q1000 network groups. Connecting lines show paired group changes. **c**, Three-source-dataset-equal effects for the q1000 within-v4 $2 \times 2 \times 2 \times 2$ factorial. Positive values favour the fidelity-kernel family. All quantities are descriptive summaries conditional on the fixed datasets and finite candidate pools; panels do not show population uncertainty or causal attribution.

The primary SVC result satisfies the machine-readable outcome category frozen as `replicated_protocol_sensitivity`; narratively, protocol sensitivity is prospectively corroborated across three fixed external tasks. The task-equal controlled effect is -0.0119 (95% conditional seed-cluster interval $[-0.0205, -0.0033]$, $n = 5$




Fig. 6 External domain shifts separate optimistic and deployable conclusions. Effects are quantum minus classical OOD balanced accuracy for the weak-reference fixed-regularization same-test oracle, the equal-budget train-tuned same-test oracle, and the primary equal-budget train-tuned validation selector. Rows show both nested sizes for all three fixed tasks plus the task-equal mean. Error bars are 95% conditional seed-cluster intervals over five deterministic subsampling seeds per task ($n = 5$), preserving the nested-size relation. They quantify pipeline variation conditional on the fixed tasks and design, not uncertainty over a population of tasks.

deterministic subsampling seeds per task), whereas the weak-reference, fixed- C , same-test-oracle effect is $+0.0089$ ($[+0.0047, +0.0137]$). Their contraction is $+0.0208$ ($[+0.0113, +0.0307]$), with a leave-one-task-out range of $[+0.0116, +0.0260]$. Size-equal controlled point estimates are -0.0160 for College Scorecard, -0.0066 for diabetes readmission, and -0.0131 for ACS income; their pointwise intervals and both nested sizes are shown in Figure 6.

The secondary GP classifier shows a larger contraction of $+0.0324$ ($[+0.0223, +0.0427]$), but its task-equal controlled effect is $+0.0077$ with an interval spanning zero ($[-0.0015, +0.0183]$). College Scorecard leans quantum under this classifier, diabetes leans classical, and ACS is slightly positive. This classifier dependence limits the conclusion appropriately: the external study corroborates the sensitivity of the family verdict to evaluation design, not universal classical dominance.

2.10 Geometry is associated with robustness only partially

The results are consistent with both families reaching similar informative geometries. Geometry carries partial information about behaviour under shift, but neither the analysis nor the matching identifies it as a causal mediator: the association is regime-dependent, not the single governing law that optimistic accounts (our own earlier report included) proposed.

The geometry- OOD link is a regime-dependent association, not a law.

Two descriptors of the training and OOD Gram matrices carry information about which kernels generalize: the effective rank of the training kernel and the centered kernel-target alignment on the OOD split. But their association with OOD accuracy is regime-dependent, not a single mechanism. Within scenario-groups (Supplementary Figure 4 and Supplementary Table 4), the median Spearman correlation between effective rank and OOD accuracy ranges from $+0.80$ (EMBER $m2$) down to -0.15 (UNSW-NB15 reconnaissance under the constructed shift); partial correlations that control for family label and length scale track the raw values closely. OOD alignment is more consistently positive ($+0.22$ to $+0.64$) but likewise varies by regime. Because OOD KTA uses the labels of the same OOD split on which accuracy is measured, it is a post hoc, label-reusing diagnostic rather than an independently validated mechanism. A geometry model trained on two source datasets and used to rank kernels in the third achieves only Spearman 0.22 – 0.40 (EMBER 0.28 , UNSW-NB15 0.22 , ToN-IoT 0.40). We therefore describe a genuine but partial association, not a causal mechanism or deployable predictor.

2.11 Finite shots can create distinctness without useful advantage

All primary quantum kernels use exact statevectors. To test whether finite estimation can create apparent non-emulability, we apply the same zero-label certificate to 30 repeated binomial fidelity estimates at 128, 512, 2,048, and 8,192 shots for one fixed SVC quantum model per case. Each replicate reselects C from training data under unprojected, independently projected, and coherent Nyström-PSD conditions (Methods).

At 128 shots the perturbed model disagrees with its exact-statevector counterpart on a median 7.2–7.4% of target decisions, depending on PSD handling. Yet the across-case median increase in the full-family zero-label endpoint is only +0.003 for Nyström, +0.005 unprojected, and +0.009 for independent projection; individual cases range from decreases of 0.014 to increases of 0.040. The corresponding median realized balanced-accuracy changes are -0.0010 , -0.0044 , and -0.0005 , respectively (Figure 7). As shots increase, median certificate changes generally approach zero, while the case ranges remain heterogeneous.

Finite-shot noise can therefore move a quantum predictor away from every exact classical witness without making it better. This is a specifically important insufficiency result for the certificate: predictive non-emulability is necessary for a material advantage against the reference family, but it is not sufficient, and measurement noise can manufacture distinctness. Effective-rank inflation and the original performance sensitivities remain in the Supplementary Information. One case–replicate requires 1,624,250 distinct fidelity estimates; the full four-shot-level sensitivity projects 4.24 trillion circuit shots before routing and device overhead, halved if the four selected product-map cases are evaluated by direct factorization instead.




Fig. 7 Finite-shot perturbations can increase apparent non-emulability without increasing performance. Lines are medians across the eight fixed case-level Monte Carlo medians and ribbons span their minima and maxima. **a**, Change from the exact-statevector zero-label accuracy upper endpoint against the full classical family. **b**, Change in realized OOD balanced-accuracy advantage. Each case has 30 measurement replicates at each shot level and PSD condition. Ribbons are fixed-case ranges, not confidence intervals.

3 Discussion

The central result is an exact answer to a deliberately narrower question than quantum advantage in the complexity-theoretic sense: given fixed target inputs, fixed predictions, a bounded loss, and a prespecified classical reference family, how much predictive advantage remains compatible with the labels not yet inspected? The finite max–min gives the sharp region after any partial audit without a sampling, calibration, prevalence, or shift assumption. No interval based only on those observables can be uniformly smaller. For zero–one accuracy, the disagreement matrix collapses that optimization to a closed form. This turns “a classical model looks similar” into an information-optimal finite-batch falsification statement with an explicit materiality threshold.

This scope distinguishes the result from its nearest ingredients. Co-validation uses disagreement between retrained classifiers to study error and variance [14]; Active Testing targets unbiased, low-variance risk estimation [16]; limited-label model selection seeks the best model to deploy [17]; and weak-supervision partial identification combines fitted label marginals with Fréchet bounds for individual metrics [18]. Our models remain fixed, the estimand is advantage against the best member of a declared

1059 family, and every unaudited label completion remains admissible. The acquisition rule
1060 is therefore secondary: it chooses which ambiguity to remove, whereas the certifi-
1061 cate remains exact for any chosen subset. In QML, this target-decision object also
1062 complements Huang’s train-sample geometric sieve rather than replacing it [11].
1063

1064 The empirical frontier is much sharper than the benchmark effect alone. Against
1065 115 classical kernels, at most 33 of 500 labels under the fixed adaptive rule are required
1066 to rule out an accuracy advantage above one point in every retrospective SVC and
1067 GPC case. Strong prediction-aware controls are at least as efficient on average, with
1068 median 12.5 labels versus 13, whereas an exact retrospective oracle needs median 6.5.
1069 This negative algorithmic result matters: most of the gain comes from adjudicating
1070 actionable disagreements, not from a privileged acquisition heuristic. The same certifi-
1071 cate contraction holds for the selected entangling-ZZ subset. In the internally locked
1072 prospective replication, all four eligible task–classifier medians require at most three
1073 labels and the overall seed-level median is zero; the controls again show no acquisition-
1074 heuristic advantage. It does not establish universal classical emulation: it identifies the
1075 quantum model’s decisions relative to a finite reference family on these particular tar-
1076 get covariates. Its value is that every remaining ambiguity is visible in the certificate
1077 rather than hidden inside an average performance comparison.
1078

1079 The three-gate interpretation prevents stronger conclusions from being conflated.
1080 Protocol validity asks whether the reported comparison could be deployed; predictive
1081 indispensability asks whether an appropriate classical witness reproduces the fixed
1082 target decisions; quantum relevance asks whether any surviving behaviour depends on
1083 a non-factorizable circuit and remains meaningful after accounting for estimation cost.
1084 Passing all three gates creates a candidate for deeper advantage analysis, not a proof
1085 of speedup. Here Gate 1 reverses the optimistic sign, Gate 2 rules out a material batch
1086 advantage with limited auditing and transfers prospectively on two new eligible tasks,
1087 and Gate 3 finds no rescue from the retrospective entangling stratum while finite-shot
1088

noise sometimes increases distinctness without performance. Because all prospective winners are product maps, the prospective experiment strengthens Gate 2 but adds no independent evidence about Gate 3.

From the perspective of distribution-shift research, this reinforces that robustness conclusions are conditional on the protocol [7–9, 22]. Within v4, moving from the fixed- C , same-test-oracle, customary-reference, native-budget corner to the fully controlled corner changes the q1000 estimate from +0.0131 to −0.0047. Train-CV regularization and the extended reference have the largest mean paired changes, but their sizeable interaction precludes an additive or causal reading. The legacy/v4 display still cannot isolate regularization, and the larger historical +0.037 endpoint remains interpretable only as a complete evaluation bundle.

The external analysis provides prospective corroboration beyond the constructed security shifts. Under a protocol frozen before result inspection, three education, health, and socioeconomic tasks corroborate a positive protocol contraction and a negative controlled SVC effect. The result is stable to omitting any one task, yet heterogeneous by task and classifier: the GP endpoint retains a small positive task-equal lean whose interval crosses zero. The strongest defensible conclusion is therefore about evaluation sensitivity, not the universal ordering of kernel families.

Length scale is decisive for both classical and quantum kernel behaviour [23, 24]. With symmetric scale freedom and tuned regularization, the families occupy overlapping effective-rank and alignment regimes. These quantities are kernel descriptors, not unique signatures of quantum provenance; OOD alignment additionally reuses target labels and remains strictly post hoc.

The high-rank finding must also be reconciled with two spectral results that superficially point the other way. Kübler et al. argue that quantum kernels generalize in-distribution only when the relevant spectral mass is concentrated in few components [25], and Thanasilp et al. show that fidelity kernels concentrate exponentially—toward

effectively rank-one, uninformative Gram matrices—as qubit counts grow [26]. There is no contradiction, but there is a genuine tension that our results locate precisely. Both results concern the asymptotic or i.i.d. regime; our effective ranks (roughly 10–150 on training sets of 10^3 – 4×10^3 samples at 4–12 qubits) sit far below the concentration regime. The repeated finite-shot sensitivity adds an empirical caveat these theoretical results anticipate: effective rank from sampled fidelities is inflated by a median factor of 1.93 at 128 shots and 1.13 at 8,192 shots (Section 2.11). Exact effective rank is therefore a statevector idealization even at these modest qubit counts.

From a practical perspective, no consistent fidelity-kernel advantage is observed over the tuned classical family, while overlap estimation introduces circuit and shot costs absent from direct classical-kernel evaluation. A shortened-length-scale Laplacian is therefore an important inexpensive reference, not a universally superior choice. More generally, comparisons should keep representation and learner fixed, tune all regularization on training data, match search budgets, select without target labels, include a non-constructed shift where possible, and report correlated benchmarks as fixed cases. Geometry descriptors remain useful diagnostics, but not deployable selectors unless validated without target-label access.

Several limitations define the scope of our conclusions.

Our shift operationalization covers three reproducible mechanisms—within-class sparsity-tail extremes, train-geometry-tail extremes, and a held-out attack-campaign shift on network traffic (an unseen attack campaign plus a capture-partition change, not a timestamped temporal split)—but not the full space of real-world drift processes. Explicit temporal splits on timestamped malware, family-based drift, and adversarially induced shift remain outside the study’s scope, and our conclusions should be re-tested under them rather than assumed.

Kernel and encoding coverage also remains finite. The classical family spans seven kernels plus a bandwidth sweep, and both families receive symmetric tuning freedom,

but neither candidate set is exhaustive: the quantum side is restricted to four com- 1197
monly used fidelity feature maps at low qubit counts, and trained or task-adapted 1198
encodings, projected kernels, and deeper maps are not evaluated. The observed asso- 1199
ciation suggests that future extensions should report their Gram-matrix geometry as 1200
well as provenance, but it does not predict their outcome. At substantially larger 1201
qubit counts, exponential kernel concentration [26] creates an additional regime not 1202
represented here, so our conclusions should not be extrapolated beyond the low-qubit 1203
study. 1204
1205
1206
1207
1208
1209

The certificate is relative to a finite, prespecified classical family and a fixed 1210
observed target batch. It does not rule out a quantum advantage against a differ- 1211
ent classical family, establish a surrogate outside the observed covariates, or convert 1212
candidate count into computational complexity. Likewise, the batch interval is not 1213
a confidence interval for a future deployment population. The optional population 1214
result requires i.i.d. target sampling and models fixed before observing that sample, 1215
assumptions that do not describe the constructed tail batches. 1216
1217
1218
1219
1220
1221

The partial-label result is mathematically valid for any audit subset, but its secu- 1222
rity acquisition curve was formulated after the eight target label files were unlocked, 1223
and the stronger comparators were added after the first retrospective curves were 1224
known. The separate Gate-2 experiment addresses that chronology internally: tasks, 1225
prediction locks, policies, thresholds, and success categories were fixed and hashed 1226
before label opening, and it meets the strong-transfer rule on every technically eli- 1227
gible task-classifier cell. This was not a public preregistration or an independently 1228
timestamped commitment. Its scope remains narrow. Only two tasks completed the 1229
frozen feature grid because NHANES failed a prespecified gate, the seeds are deter- 1230
ministic subsamples rather than independent task replicates, and every prospective 1231
quantum winner is a product map. The result supports transfer of the certificate 1232
and actionable-disagreement principle to these fixed tasks, not a population claim, 1233
1234
1235
1236
1237
1238
1239
1240
1241
1242

1243 entangling-kernel replication, or superiority of bottleneck coverage. Balanced-accuracy
1244 certification additionally requires a prespecified prevalence or prevalence range; using
1245 the retrospectively observed prevalence is only a post-unlock sensitivity.
1246

1247 The quantum kernels are simulated rather than executed on hardware. The study
1248 therefore characterizes fidelity-kernel geometry, not noisy-device behaviour. Repeated
1249 binomial sampling quantifies one component of measurement uncertainty, but hard-
1250 ware noise, transpilation, readout, mitigation, and circuit-dependent correlations
1251 remain unmodelled [27]; exponential concentration may also make informative spectral
1252 structure expensive to resolve at larger scales [26]. We make no claim about hardware
1253 feasibility, resource advantage, or quantum speedup—the only rigorously established
1254 quantum-classical learning separation concerns a specifically constructed problem, not
1255 natural data [28]. At the dimensions studied, the fidelity kernels are also classically
1256 simulable [12, 29], which is precisely what makes them available as practical kernel
1257 constructions today.

1266 Dataset scope and feature caveats provide a final boundary. Dynamic malware anal-
1267 ysis and encrypted traffic remain untested. The network exports retain port-derived
1268 counters and, for ToN-IoT, direct port, protocol, and service indicators that can sup-
1269 port shortcuts in NIDS benchmarks [30]. Sharing features controls the inputs but
1270 does not imply that distinct kernels exploit them equally. In the frozen q1000 abla-
1271 tion, removing these fields shifts the two-source-dataset-equal effect from -0.0031 to
1272 -0.0245 , but the change is highly heterogeneous and concentrated in ToN-IoT. This
1273 is evidence that feature access can affect the relative comparison, not proof that every
1274 shortcut has been removed or that the ablated tasks define a preferred deployment
1275 benchmark.

1283 We do not claim universal classical superiority, predictive equivalence, complexity-
1284 theoretic advantage, hardware relevance, or a refutation of quantum kernels. The
1285 proposed audit instead makes a candidate advantage progressively falsifiable: first by
1286
1287
1288

protocol controls, then by a prespecified classical witness family on target covariates, and finally by limited adjudication of the disagreements that keep a material advantage possible. Future work should broaden the prospective replication to more tasks and entangling winners, and pair candidate breadth with measured computational resources. This interpretation is consistent with calls for controlled QML benchmarking [4, 5, 31].

4 Methods

4.1 Problem setting and estimands

We study binary security classification—malware versus benign software and attack versus benign traffic—under distribution shift. Let D_{ID} denote the in-distribution process and D_{OOD} a constructed stress-test or held-out-campaign process. Each run contains a training set T , an ID holdout that is divided into validation V_{ID} and test S_{ID} , and an OOD test set S_{OOD} . Preprocessing, samples, classifier implementation, tuning rule, and family-selection rule are shared; only the kernel changes. Comparisons are always within a classifier.

The primary endpoint is OOD balanced accuracy after P1' deployment selection. A family effect is

$$\Delta_{\text{OOD}} = \text{BAcc}_{\text{OOD}}^{(Q)} - \text{BAcc}_{\text{OOD}}^{(C)},$$

so positive values favour the quantum family. ID balanced accuracy and ID-to-OOD degradation are secondary. P2 selects from OOD information in other realizations, and P3 selects and evaluates on the same OOD test labels; both are explicitly non-deployable descriptive oracles.

1335 4.2 Sharp finite-target identification

1336

1337 For Theorem 1, fix the prediction-dependent loss tables before the audit and write

1338

1339 $a_{ij}(y) = \ell_{ij}^C(y) - \ell_i^Q(y)$. For any completion of the labels outside L ,

1340

1341

1342

1343

1344

1345

1346 For the lower endpoint, minima commute and the label choices separate by item:

1347

1348

1349

1350

1351

1352

1353 For the upper endpoint introduce one-hot variables $z_{iy} \in \{0, 1\}$ with $\sum_{y \in \mathcal{Y}} z_{iy} = 1$

1354

1355

1356

1357

1358

1359

1360

$$nt \leq S_j(L) + \sum_{i \notin L} \sum_{y \in \mathcal{Y}} a_{ij}(y) z_{iy} \quad \text{for every } j.$$

1361 This mixed-integer linear program enumerates the finite max–min implicitly. Finite-

1362

1363 ness of $\mathcal{Y}^{|L^c|}$ guarantees that both extrema are attained. Any interval based on the

1364

1365

1366

1367

1368

1369

1370

1371

1372

1373 The empirical accuracy certificate changes the estimand from a selected family

1374

1375

1376

1377

1378

1379

1380

mean to a fixed-model comparison. Let $q_i = h_Q(x_i)$ and $c_{ij} = h_{C_j}(x_i)$ be hard predictions on a target batch of size n . The quantum model is the P1' winner fixed using ID validation; the classical candidates, their train-only hyperparameters, and their target predictions are also fixed before any target-label audit. Every displayed certificate uses

the same inherited canonical realization per security group: $n_{\text{train}} = 1000$, $n_{\text{ID}} = 500$,
 $n_{\text{OOD}} = 500$, master seed 42, q-split seed 42, and model seed 42. The exact scenario
run prefixes and selected SVC/GPC configurations appear in the canonical-run table
in Supplementary Information. For a candidate labelling y ,

$$g_j(y) = \text{Acc}(q; y) - \text{Acc}(c_j; y), \quad \Delta_B(y) = \min_j g_j(y).$$

Only disagreements contribute to g_j . If $d_j = n^{-1} \sum_i \mathbf{1}\{q_i \neq c_{ij}\}$, then $g_j(y) \in$
 $[-d_j, d_j]$. The labelling $y_i = q_i$ attains $g_j = d_j$ for every j simultaneously, so
 $\max_y \Delta_B(y) = \min_j d_j$. If k maximizes d_j , the labelling $y_i = c_{ik}$ attains $\Delta_B(y) \leq$
 $g_k(y) = -d_k$; no g_j can be below $-d_j \geq -d_k$, so $\min_y \Delta_B(y) = -\max_j d_j$. This proves
Theorem 2 and its sharpness.

For an arbitrary audited subset L , define the observed signed contribution $S_j(L)$
and remaining disagreement count $R_j(L)$ as in Theorem 3. Every completion satisfies

$$\frac{S_j - R_j}{n} \leq g_j(y) \leq \frac{S_j + R_j}{n}.$$

Assigning every unaudited label to q_i attains all witness upper endpoints simultane-
ously. Assigning them to the witness that minimizes $S_j - R_j$ attains the lower endpoint
of Δ_B . This proves Theorem 3. In binary classification, if D_j is the original disagree-
ment count and the audit has revealed a_j disagreements favouring q and b_j favouring
 c_j , the interval simplifies to

$$\left[\min_j \frac{-D_j + 2a_j}{n}, \min_j \frac{D_j - 2b_j}{n} \right].$$

Thus the upper endpoint cannot increase under arbitrary adaptive querying. A
classical-favouring label on a binary disagreement reduces that witness endpoint by

exactly $2/n$; a quantum-favouring label leaves its upper endpoint unchanged. These statements are deterministic for the observed batch and require neither i.i.d. sampling nor a model for the unaudited labels.

For completeness, an optional population form is available when the target inputs are i.i.d. and every model is fixed before observing them. A union bound and Hoeffding’s inequality give, with probability at least $1 - \delta$,

$$\sup_y \Delta_B^{P_T}(y) \leq \min_{j \leq M} \hat{d}_j + \sqrt{\frac{\log(2M/\delta)}{2n}}.$$

We do not apply this result to the constructed tail batches, for which the finite-batch certificate is the appropriate estimand.

4.3 Reference tiers and active target-label audit

The partial-label audits use two nested operational tiers. The customary tier contains the 30 linear and RBF candidates; the full tier contains all 115 prespecified classical candidates. Separately, the frozen zero-label breadth sensitivity evaluates 30, 60, and 115 candidates using 5,000 nested whole-block orderings rooted at `ksf-v9-frontier-20260731`; a 60-candidate draw contains 12 of 23 five-dimension blocks and the full tier contains all blocks. Tier size measures reference breadth only. It is not a computational-resource claim because training time, memory, and model-construction cost are not equalized by candidate count.

At every audit step, we calculate the witness-specific numerator $S_j + R_j$. Among witnesses with unaudited disagreements, the adaptive coverage rule identifies those tying at the smallest reducible numerator and queries an unaudited input that covers the largest number of these bottleneck witnesses. Ties follow a fixed SHA-256

order. The rule uses target covariates, predictions, and previously revealed out-
comes, but never an unrevealed label. It is evaluated retrospectively for thresholds
 $\tau \in \{0.005, 0.010, 0.020\}$.

We compare it with three stronger references. Dynamic random-active sampling
selects by SHA-256 order among all unaudited points disagreeing with at least one
current reducible bottleneck witness. Non-adaptive initial coverage ranks points once
by how many zero-label bottleneck witnesses they cover, without label feedback. Both
use 200 frozen SHA-256 tie orders. Finally, an undeployable label oracle gives the exact
ex-post minimum number of queries over every possible audit subset. If D_j is witness
 j 's total disagreement count, crossing threshold τ requires

$$r_j(\tau) = \max\left\{0, \left\lceil \frac{D_j - n\tau}{2} \right\rceil\right\}$$

revealed disagreements whose realized label favours witness j ; the oracle minimum is
the smallest feasible $r_j(\tau)$ over witnesses. This is exact because the family endpoint
crosses when any one witness endpoint crosses. A prediction-independent complete
ordering is retained only as a weak diagnostic. Comparator quantiles describe frozen
algorithmic orders, not sampling uncertainty.

The zero-label SVC definition and go/no-go criterion were frozen before pre-
diction outcomes were inspected. The arbitrary-partial-label acquisition extension
was declared only after the target labels had been unlocked; the strong-comparator
amendment was declared after the original acquisition curves were known. Both are
exploratory on these eight cases. GPC repeats the mathematics and policies but is
not an independent prospective validation. Target labels are used only to evaluate
certificate contraction, the oracle, and the final realized effect; no kernel, classifier,
regularization, threshold, or P1' winner is refitted or reselected. Prediction archives

1519 and the separately stored label file are joined only after SHA-256, row-index, model,
1520 and target-length gates pass.

1522

1523 4.4 Prospectively frozen Gate-2 replication 1524

1525 The Gate-2 replication specification was hashed on 1 August 2026 before any of its
1526 three tasks was acquired, any model prediction was computed, or any target label
1527 was inspected. The source is the same TableShift implementation pinned at commit
1528 `fca9429814703a07e3902d005d46563a207b7f0a` [32]. BRFSS Diabetes (race shift),
1529 ACS Food Stamps (Census-division shift), and NHANES Lead (poverty shift) were
1530 exactly the unused fallback identifiers recorded before the earlier external campaign;
1531 no project quantum result or target sample for them existed at freeze time. All three
1532 were attempted and none could be replaced based on availability or outcome.

1533 For each task, seeds $\{42, 123, 999, 7, 2024\}$ define deterministic label-blind sam-
1534 ples with $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{target}}) = (1000, 250, 250, 500)$. The v5 train-only semantic
1535 encoding, scaling, TruncatedSVD dimensions $\{4, 6, 8, 10, 12\}$, and angular representa-
1536 tion are reused verbatim. The 60 quantum candidates and 115 classical candidates,
1537 P1' validation selection, candidate-specific train-CV for SVC, fixed Laplace GPC, and
1538 decision thresholds are unchanged. The customary 30-candidate tier is secondary; the
1539 full tier and $\tau = 0.010$ define the prospective outcome. Secondary thresholds are 0.005
1540 and 0.020.

1541 The pipeline is fail-closed. A prediction-lock stage copies target labels to a separate
1542 sealed tree, removes them from the model state, and writes hard target predictions,
1543 row-order digests, selected validation-only quantum configurations, and SHA-256
1544 hashes without calculating or printing target performance, prevalence, disagreement,
1545 or an acquisition curve. Candidate checkpoints contain only target predictions and
1546 validation-selection quantities. An aggregate manifest is accepted only when every
1547 available task-seed unit has all 175 kernel configurations for both classifiers or a
1548

permitted pre-model technical-failure record. The audit stage verifies that aggregate hash, records the one-time label opening and analysis-code hash, then joins labels by source-position digest. No model is refitted after the join.

For each classifier and the two classical tiers, the fixed policies are adaptive bottle-neck coverage, 200 dynamic random-active SHA-256 orders, 200 non-adaptive initial-coverage orders, the exact ex-post label oracle, and a weak prediction-independent order. The hash root appends task, seed, model, tier, policy, and draw. The primary summaries are four or six task-classifier medians over five seeds, depending on technical availability. Strong transfer requires every available task-classifier median to be at most 50 labels, the overall seed-level median to be at most 25, and no seed to exceed 100. With exactly two available tasks, all four cells are required and the result remains explicitly technically limited. Partial transfer requires all four medians to be at most 50 and an overall median at most 50; all other complete two-task outcomes fail to transfer. Fewer than two available tasks is technically incomplete. These categories were evaluated exactly once.

NHANES generated 11 non-constant encoded features in all five seeds, fewer than required for the frozen 12-dimensional candidate grid. Minimum-feature failure was an explicitly permitted technical-unavailability gate. No NHANES model candidate was executed, the task was not replaced, and dimensions were not changed. BRFSS Diabetes and ACS Food Stamps passed every gate, leaving ten task-seed units for the technically limited prospective outcome.

4.5 Balanced-accuracy identification

Balanced accuracy is not label-free unless target prevalence is known or restricted. At a fixed positive count n_+ , we group inputs by their joint prediction signature $z = (q_i, c_{i1}, \dots, c_{iM})$. Let n_z be the size of group z and $k_z \in \{0, \dots, n_z\}$ its unknown number of positive labels, with $\sum_z k_z = n_+$. For classifier h with group-level prediction

1611 $h_z \in \{0, 1\}$,

$$1612 \quad \text{BAcc}_h(k) = \frac{1}{2} \left[\frac{\sum_z k_z h_z}{n_+} + \frac{\sum_z (n_z - k_z)(1 - h_z)}{n - n_+} \right].$$

1613
1614 The sharp upper endpoint is the mixed-integer linear program that maximizes t subject
1615 to

$$1616 \quad t \leq \text{BAcc}_Q(k) - \text{BAcc}_{C_j}(k) \quad \text{for every } j,$$

1617
1618 together with the integer and prevalence constraints above. The lower endpoint is the
1619 minimum over witnesses of their corresponding linear minima. Relaxing k_z to a con-
1620 tinuous interval gives the sharp distributional or fractional-label LP. When prevalence
1621 is known only to lie in an interval, we maximize over its admissible integer counts.
1622 We report this as a sensitivity because the natural target prevalence is unavailable
1623 without labels; accuracy remains the main assumption-free certificate.
1624
1625
1626
1627
1628
1629
1630

1631 4.6 Why train-only geometry cannot determine target 1632 prediction 1633 1634

1635 The geometric difference and effective rank summarize training Gram structure,
1636 whereas deployment also depends on the target–train block. An explicit normalized
1637 PSD construction makes the distinction exact. Set $K_{TT} = I_n$. One valid joint Gram
1638 extension gives a unit-norm target feature orthogonal to all training features, hence
1639 $K_{UT} = 0$. A second gives the target the same feature as the first training point, hence
1640 $K_{UT} = e_1^\top$. Both joint matrices are PSD, have unit diagonal, and share identical K_{TT} .
1641 For kernel ridge regularization $\lambda > 0$, however, their target prediction operators are
1642
1643
1644
1645
1646
1647

$$1648 \quad 0 \quad \text{and} \quad e_1^\top (I_n + n\lambda I_n)^{-1} = \frac{e_1^\top}{1 + n\lambda},$$

1649
1650
1651
1652
1653
1654
1655
1656

whose spectral-norm difference is $1/(1 + n\lambda)$. Train-only geometry can therefore screen possible separation but cannot characterize target transport. We use the hard-prediction certificate for SVC and GPC rather than treating this KRR construction as an endpoint result.

4.7 Security datasets and shared split sizes

The source datasets are EMBER [33], UNSW-NB15 [34], and ToN-IoT [35]. Four binary scenarios—EMBER malware, UNSW-NB15 DoS, UNSW-NB15 Reconnaissance, and ToN-IoT Scanning—produce two shift mechanisms each and hence eight scenario-groups. Classical and quantum kernels always receive identical row indices and class balance.

The 72 principal settings cross eight groups with master seeds $\{42, 123, 999\}$ and $(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$ equal to $(1000, 500, 500)$, $(2000, 1000, 1000)$, or $(4000, 1800, 1800)$. Each setting is instantiated with q-split seeds $\{42, 123, 999, 7, 2024\}$ and model seeds $\{42, 123, 999\}$, yielding fifteen pipeline realizations arranged in five q-split clusters. Indices are stored explicitly; automated gates verify split-role disjointness, required cardinalities, both classes, identical family inputs, and absence of label-derived features.

4.8 EMBER shift construction

EMBER uses the histogram and byte-entropy feature blocks and two label-conditional tail stress tests. They are not described as certified covariate shifts because conditioning the split on class does not establish invariance of $P(Y | X)$. For $m1$, the sparsity score is the number of non-zero entries in the selected feature block. Within each class, low- and high-score extremes supply the exact-size OOD set; training and ID samples are stratified draws from the remainder.

For $m2$, the final training set is fixed first. A MaxAbs scaler and TruncatedSVD representation are fitted on training data only, and a centroid is computed for each

class. Every non-training sample is scored by Euclidean distance to its own-class centroid. Within-class low- and high-distance extremes form S_{OOD} , and the ID holdout is sampled from the remaining pool. This construction probes unusual geometry relative to the fixed training set without using test information to fit the representation.

4.9 Network-flow shift construction

Numeric preparation removes identifiers and explicit label derivatives, leaving 39 UNSW-NB15 and 88 ToN-IoT features. The UNSW-NB15 export excludes direct categorical protocol, service, and state fields but retains three port-derived counters; ToN-IoT retains two port fields plus protocol and service indicators. A leakage audit records every exclusion and feature-label association; the largest observed absolute correlation is 0.62 and belongs to a retained traffic feature. Each scenario has reference and current pools and uses two mechanisms.

The *m2-centroid* mechanism is the network analogue of EMBER *m2*: train-only MaxAbs+TruncatedSVD class centroids define within-class distance tails, which supply the OOD set. The *attack-campaign* mechanism performs no geometric selection. Training and ID pools contain benign traffic and all non-target attack categories; OOD contains benign traffic and the held-out target campaign. UNSW-NB15 reference and OOD pools additionally use different official capture partitions. This is an unseen-campaign plus capture-partition shift, not a timestamped temporal split.

4.10 Shared preprocessing and representations

For every run and dimension $d \in \{4, 6, 8, 10, 12\}$, a MaxAbs scaler and TruncatedSVD are fitted on T ; the reduced coordinates are standardized and mapped linearly with training-fitted parameters so that the *training* coordinates lie in $[0, \pi]$. Held-out coordinates receive the same affine map and may extrapolate outside that interval under shift. The same embedded samples are supplied to both families, making the comparison a kernel swap rather than a representation swap. Dimension is a candidate axis.

No validation, ID-test, or OOD-test row fits a transformation. The SVD random state equals the model seed.

4.11 Classifiers and train-only regularization

The primary learner is an SVC with a precomputed Gram matrix and balanced class weights [36]. For every kernel geometry, $C \in \{0.01, 0.1, 1, 10, 100\}$ is selected by five-fold stratified cross-validation on the training Gram matrix; ties favour the smaller C . When $n_{\text{train}} > 2000$, this selection step uses a class-stratified random-state-20260717 subsample capped at 2,000 rows, while the final model is always refitted on the complete training block. Because C is resolved before family selection, its grid points are not counted as separate family candidates.

The secondary learner is a binary Gaussian-process classifier with a logistic link and Laplace approximation [37]. It consumes the same precomputed kernel, uses diagonal jitter 10^{-8} , at most 100 Newton iterations, and convergence tolerance 10^{-6} . The logistic-probit approximation supplies predictive probabilities; no post-hoc calibration is applied.

4.12 Classical kernel family

The customary reference contains a cosine-normalized linear kernel and RBF kernel. The extended family adds cosine-normalized degree-2 and degree-3 polynomial kernels with $\gamma = 1/d$ and intercept 1, the Laplacian kernel $\exp(-\|x-x'\|_1/\ell_1)$, and Matérn-3/2 and Matérn-5/2 kernels. The base ℓ_1 and ℓ_2 scales are median pairwise Manhattan and Euclidean distances, respectively, calculated from at most 2,000 training rows with the model seed; the RBF base scale is fitted from training variance. RBF, Laplacian, and both Matérn kernels use factors $\{0.1, 0.3, 1, 3, 10\}$ around these training-only bases. Deduplication yields 23 kernel-shape/scale blocks crossed with five dimensions, or 115 extended-classical candidates in full-coverage groups.

1795 4.13 Quantum fidelity-kernel family

1796

1797 The quantum family uses exact statevector fidelities

1798

1799

1800

1801

1802

1803

1804 [1, 2]. Qiskit 2.3.1 implements four map blocks [38]: fully connected ZZ maps with one

1805 and two repetitions; a one-repetition Pauli map with single-qubit strings **X** and **Z**; and

1807 a two-repetition **Z** map. Only the ZZ maps contain two-qubit interactions. The latter

1808

1809 two prepare product states,

1810

1811

1812

1813

1814

1815

1816 so the constructor argument `entanglement="full"` stored for the Pauli-X/Z map has

1817

1818 no effect without a multiqubit Pauli string. The product can be evaluated classically

1819 in $O(d)$ one-qubit-overlap calculations. We call these the entangling-ZZ and separable-

1820

1821 product strata. Each map uses angle scales $\{0.5, 1, 2\}$ and all five dimensions, giving

1822

1823

1824 60 candidates (30 per stratum) at 4–12 simulated qubits. Classical bandwidth and

1825

1826

1827 4.14 Circuit and sampling-resource audit

1828

1829 Circuit counts use Qiskit 2.3.1 and the logical basis $\{\mathbf{rz}, \mathbf{sx}, \mathbf{x}, \mathbf{cx}\}$ at transpiler

1830

1831 optimization level 0, without backend layout or routing. We report each feature-map

1832

1833 template and the compute–uncompute fidelity template formed by composing the map

1834 at x with the inverse map at x' after assigning distinct numeric inputs. Counts are

1835

1836 compiler- and basis-conditional logical resources, not device schedules.

1837

1838

1839

1840

For the repeated q1000 finite-shot sensitivity, each case and measurement replicate estimates the off-diagonal training triangle, the combined 500-row ID-holdout-to-training rectangle (partitioned downstream into validation and test roles), the OOD-test-to-training rectangle, and the off-diagonal OOD square triangle; diagonals are fixed to one. Thus one kernel configuration requires

$$\binom{1000}{2} + 500(1000) + 500(1000) + \binom{500}{2} = 1,624,250$$

distinct fidelity estimates per shot level and replicate. Multiplying this count by shots gives the idealized circuit-shot total. The three PSD conditions reuse the same sampled entries and add only classical post-processing. Routing, calibration, mitigation, queueing, and device noise are excluded.

4.15 Target-label-free family selection

The original ID holdout is divided deterministically with class stratification into V_{ID} and S_{ID} . Within each class, global row indices are ordered by SHA-256 of the salt `ksf-v4-idsplit::` and index, then assigned alternately to validation and test. The halves are disjoint and approximately class-balanced. P1' deploys, within each family, the candidate with highest V_{ID} balanced accuracy and reports it once on S_{ID} and S_{OOD} . Ties follow lexicographic candidate order. OOD labels never enter regularization, preprocessing, or P1' selection.

4.16 Candidate-budget matching

Best-of-family selection rewards a larger pool. The confirmatory coverage audit found 115 extended-classical and 60 quantum candidates in every one of the eight scenario-groups and both classifiers. The primary comparison therefore draws 60 candidates from each family in all 16 group-classifier cells. For each cell, 5,000 extended-classical subsamples are generated with seed 20260715. The primary `kernel_blocked` scheme

1887 samples 12 of the 23 complete kernel-shape/scale blocks, each crossed with all five
 1888 dimensions, to mirror the 12 quantum feature-map/angle-scale blocks crossed with the
 1889 same dimensions. This matches both candidate count and block structure, although it
 1890 does not make the two search spaces identical. `uniform` samples individual candidates,
 1891 while `kernel_stratified` distributes the draw across kernel-shape/scale strata. Each
 1892 subsample selects its P1' winner; the classical endpoint is the resampling expectation.
 1893 The full 115-candidate pool is secondary. Resampling percentiles describe sensitivity
 1894 to the sampled candidate set, not population uncertainty.

1901 4.17 Entangling and separable quantum strata

1902 The reviewer-motivated v0.8 sensitivity was committed before its performance con-
 1903 trasts were calculated. It reuses all v4 summaries and compares three SVC pools under
 1904 train-CV and P1': the full 60-candidate quantum family, the 30 entangling-ZZ candi-
 1905 dates, and the 30 separable-product candidates. The 60-candidate pool is compared
 1906 with 12 complete classical blocks and each 30-candidate stratum with six complete
 1907 classical blocks, always crossed with all five dimensions. Each classical endpoint is the
 1908 expectation over 5,000 `kernel_blocked` draws using SHA-256-derived seeds rooted at
 1909 20260729. This tests dependence on map stratum, not computational advantage.

1918 4.18 Within-v4 factorial and shortcut sensitivity

1919 The within-generation SVC factorial uses all 360 q1000 runs and crosses: same-test
 1920 OOD versus disjoint ID-validation selection; fixed $C = 1$ versus train-CV C ; cus-
 1921 tomary versus extended classical reference; and native versus equal-count budgets.
 1922 Fixed- C predictions are recomputed with the deterministic v4 preprocessing, kernel
 1923 constructors, samples, and split roles; an integrity gate requires exact candidate-split
 1924 key identity with every archived v4 SVC summary and the original ID-split hash salt.
 1925 With the customary 30-candidate reference, equal count samples six of twelve quan-
 1926 tum blocks; with the extended reference, it samples 12 of 23 classical blocks. The

same 5,000 frozen block draws are paired across selection and regularization cells. Sixteen fixed-case endpoints and their paired axis contrasts are descriptive; they do not identify causal effects.

The q1000 shortcut sensitivity uses the same design on all 270 network runs after removing features before the train-only embedding. UNSW-NB15 excludes `ct_src_dport_ltm`, `ct_dst_sport_ltm`, and `is_sm_ips_ports`; ToN-IoT excludes source/destination ports, all `proto_*`, `service`, and `service_*` fields. This removes 3 of 39 and 14 of 88 fields, respectively. The manuscript-facing contrast uses train-CV, P1', and 60-candidate block matching with paired draws. It is a port/protocol/service-field ablation, not a guarantee that every shortcut has been removed.

4.19 Security fixed-case aggregation and conditional intervals

The effective replication axis within a scenario-group is q-split seed. Master seed is a design stratum rather than a replicate, some EMBER OOD pools recur across master seeds, and training sizes are nested. Audited OOD overlap is at most 1.2% for EMBER and 7.8% for network groups. Within each q-split cluster, model seeds are averaged within master seed, master seeds within size, and the three sizes with equal weight. The five resulting cluster means T_j give the group effect $\bar{T}$ and 95% conditional cluster- t interval

$$\bar{T} \pm t_{4,0.975} \frac{s_T}{\sqrt{5}}.$$

This few-cluster t construction follows the cluster-mean approach of Ibragimov and Müller [39]. The interval ($n = 5$ q-split clusters) is conditional on the fixed benchmark pools and experimental design; inclusion of zero is not evidence of equivalence. A 9,999-draw BCa bootstrap over cluster means and leave-one-q-split calculations are secondary diagnostics, reported for all 16 primary cells in Supplementary Table 6. The five cluster values behind every interval are exposed in Supplementary Table 7.

Across groups, EMBER $m1/m2$ are averaged within source dataset, the four UNSW-NB15 groups within UNSW-NB15, and the two ToN-IoT groups within ToN-IoT. These three source-dataset means receive equal weight. Leave-one-source-dataset-out ranges omit one of the three at a time. The eight scenario-groups and three datasets were not sampled from defined populations; therefore no population p -value or dataset-population confidence interval is reported.

4.20 Specification curve

The S1–S10 display is retained as a historical cross-generation sensitivity map over complete legacy and v4 endpoints. S1–S4 use fixed- C legacy generation and S5–S10 train-CV v4 generation, so their boundary cannot isolate regularization. The rows remain in their prespecified order and are never sorted by effect. Individual-axis interpretation instead uses the within-v4 factorial above. All summaries average realizations within setting and group, then groups within the three source datasets, and finally source datasets with equal weight; no specification-level hypothesis test is performed.

4.21 Prospectively frozen external corroboration

Before acquiring TableShift data or inspecting any external result, we committed the tasks, sampling, preprocessing, candidate pools, estimands, uncertainty procedure, and outcome-category rule. The source implementation is pinned to commit `fca9429814703a07e3902d005d46563a207b7f0a` [32]. Three public tasks span distinct domains and published shift variables: College Scorecard (institution type), diabetes readmission (admission source), and ACS income (geographic region). TableShift’s `train`, `validation`, `id_test`, and `ood_test` roles are preserved; `ood_validation` is never loaded. College Scorecard is read from its public CC0 archive because the pinned downloader invokes authenticated Kaggle access; the pinned parser, schema, and domain split are unchanged. Source-file SHA-256 hashes and the container digest are recorded.

Each task has nested strata $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{OOD}}) = (1000, 250, 250, 500)$ and
 (2000, 500, 500, 1000). For seeds $\{42, 123, 999, 7, 2024\}$, source rows are ordered by
 SHA-256 of task, split, seed, and original position, giving deterministic label-blind
 samples. Gates verify source-role disjointness, size nesting, both classes, source-schema
 preservation, and at least 12 non-constant encoded features. All 30 task-size-seed
 units passed.

All transformations fit **train** only. Constant columns are removed; numeric vari-
 ables receive median imputation; Boolean, categorical, and string variables receive
 most-frequent imputation and one-hot encoding with unknown categories ignored.
 MaxAbs scaling, TruncatedSVD at the five study dimensions, standardization, and
 a training-fitted affine map that places training coordinates in $[0, \pi]$ follow; held-out
 coordinates may extrapolate beyond that range. The SVD random state is 20260729.
 The security kernel and classifier pools are then reused.

For each task, size, seed, and classifier, the validation-selected 60-candidate quan-
 tum endpoint is compared with 5,000 **kernel_blocked** 60-candidate subsamples of the
 115-candidate classical pool using seed 20260729. The primary effect is quantum minus
 expected classical OOD balanced accuracy. The equal-budget OOD oracle selects on
ood_test; the weak-fixed oracle uses linear/RBF, fixes SVC at $C = 1$, and also selects
 on the reported OOD labels. Protocol contraction is weak-fixed-oracle effect minus
 controlled P1' effect.

Seeds are averaged within task and size, nested sizes receive equal weight within
 task, and the three fixed tasks receive equal weight. Conditional 95% intervals use 9,999
 resamples of the five seed clusters within task while preserving nested sizes. Leave-
 one-task-out ranges are descriptive. The intervals quantify deterministic-subsampling
 variability conditional on three fixed tasks; they do not represent a population of tasks.

2071 4.22 Performance and uncertainty metrics

2072

2073 Balanced accuracy on S_{OOD} is primary. Accuracy, macro and positive-class F1,
 2074 ROC AUC, and precision–recall AUC are recorded on every split. The GP analysis
 2075 additionally records log loss, Brier score, expected calibration error with 15 equal-
 2076 width probability bins [40], and mean predictive entropy. No threshold or calibration
 2077 parameter is fitted on OOD labels.

2081

2082 4.23 Geometry descriptors and association analyses

2083

2084 For positive-semidefinite K with eigenvalues λ_i and $p_i = \lambda_i / \sum_j \lambda_j$, effective rank is

2085

2086

2087

2088

2089

2090

2091

2092 [41]. For $y \in \{-1, +1\}^n$ and doubly centered K_c , centered kernel-target alignment is

2093

2094

2095

2096

2097

2098

2099

2100

2101

2102

2103

2104

2105

2106

2107

2108

2109

2110

2111

2112

2113

2114

2115

2116

$$r_{\text{eff}} = \exp\left(-\sum_i p_i \log p_i\right)$$

$$\text{KTA}(K, y) = \frac{\langle K_c, yy^\top \rangle_F}{n \|K_c\|_F}$$

2099 [42, 43]. KTA is computed separately on training, ID, and OOD square blocks, along-
 2100 side off-diagonal and cross-split similarities. OOD KTA reuses the same OOD labels as
 2101 accuracy and is therefore a post hoc diagnostic, not a label-independent mechanism.
 2102 Within each run–dimension cell, Spearman correlations relate effective rank, OOD
 2103 KTA, and KTA survival (OOD minus ID) to SVC OOD balanced accuracy for all,
 2104 classical-only, and quantum-only candidates. Partial Spearman association residualizes
 2105 ranked variables against family and log length/angle scale. For portability, a descrip-
 2106 tive model bins log effective rank into ten training-dataset quantiles, learns mean
 2107 OOD accuracy per bin from two source datasets, and ranks cells in the third; held-out
 2108 Spearman correlation is reported once per source dataset.

Rank matching is observational. Within each run and dimension, every quantum candidate is first paired to the classical candidate with smallest absolute effective-rank difference; classical candidates may be reused. Rank ratio $\max(r_Q, r_C) / \min(r_Q, r_C)$ is filtered at prespecified calipers 1.10, 1.25, 1.50, and 2.00, plus unfiltered reporting. A sensitivity uses minimum-cost one-to-one assignment with absolute log-rank difference, leaving excess candidates unmatched. Counts, retained fractions, rank discrepancies, mean and median OOD differences, and the fraction favouring quantum are reported for every caliper. None of these analyses identifies a causal geometry or family effect.

4.24 Repeated finite-shot perturbation

The original v0.6.0 finite-shot analysis and the reviewer-motivated v0.7.0 out-of-sample PSD extension were each frozen before their outcomes were computed. They fix one $q1000$, master-seed-42, q-split-42, model-seed-42 run per security group and, within each, the exact-statevector quantum SVC candidate selected by $P1'$ on V_{ID} . Lexicographic candidate order resolves ties. Input `summary_v4.csv` SHA-256 hashes, chosen kernels, dimensions, and exact C values are stored in the frozen specifications.

For shots $s \in \{128, 512, 2048, 8192\}$ and measurement replicates $r = 0, \dots, 29$, each exact fidelity p is sampled independently as $\text{Binomial}(s, p)/s$. A separate NumPy generator is used for train, ID-validation, ID-test, OOD-test, and OOD-square blocks. Its unsigned 64-bit seed is the first eight SHA-256 bytes of

```
ksf-v6-shots::::<kernel>::::<shots>::::<block>.
```

Python's process-randomized hash is not used.

Square train and OOD blocks are sampled on the upper triangle, mirrored, clipped to $[0, 1]$, and assigned unit diagonal; rectangular blocks receive binomial sampling and clipping. Three prespecified conditions use the same sampled blocks. The unprojected

condition makes no PSD correction. The independent-square heuristic clips negative eigenvalues of the train and OOD square matrices separately while leaving every rectangular evaluation–train block unchanged; it is an algorithmic intervention for a precomputed SVC, not a coherent correction of one out-of-sample kernel.

The coherent Nyström condition eigendecomposes the sampled training matrix as $K_T = U \text{diag}(\lambda)U^\top$ and retains $\lambda_i > \tau$, where $\tau = \max\{10^{-12}, \epsilon n \max(1, \max_i |\lambda_i|)\}$. For retained U_+ and λ_+ ,

$$K_T^+ = U_+ \text{diag}(\lambda_+)U_+^\top, \quad \Phi_E = K_{E,T}U_+ \text{diag}(\lambda_+)^{-1/2}.$$

The corrected evaluation–train and evaluation-square blocks are $\Phi_E \Phi_T^\top$ and $\Phi_E \Phi_E^\top$, respectively, with $\Phi_T = U_+ \text{diag}(\lambda_+)^{1/2}$. Thus the OOD square used for KTA is derived from the sampled OOD–train block rather than independently projected. In all three conditions, C is reselected by the same five-fold train-only procedure and the final SVC is refitted. Outputs include signed and absolute OOD change from the fixed exact endpoint, effective-rank ratio, OOD-KTA change, retained rank, eigenvalue tolerance, negative-eigenvalue fraction, and Frobenius changes for every block. Within each fixed Gram matrix, medians and central 2.5–97.5% Monte Carlo intervals summarize 30 replicates. Across the eight fixed matrices, medians and ranges are descriptive.

For the v0.9 extension we export the hard target prediction of every perturbed SVC and compare it with the unchanged exact-statevector classical witnesses. For each replicate and classical tier we record self-disagreement from the exact quantum predictor, the zero-label interval, and realized accuracy and balanced-accuracy advantages after the integrity-gated label join. This deliberately asks whether quantum measurement perturbation creates behaviour absent from the existing classical reference; it

does not give the classical family an additional noise or retuning budget. The model
contains no device or correlated-noise process.

4.25 Statistics and reproducibility

No null-hypothesis test is used for the eight security groups, three source datasets,
or three external tasks. Every uncertainty object names its effective n , resampling
unit, and conditioning set: security intervals use five q-split clusters; external intervals
use five deterministic seed clusters per task; finite-shot intervals use 30 measurement
replicates conditional on one fixed Gram matrix. The fifteen security pipeline realiza-
tions are never treated as independent. Security intervals use the t_4 reference and are
pointwise, not simultaneous; the five underlying cluster means are published. Inclusion
of zero is never interpreted as equivalence. No multiplicity adjustment was prespec-
ified. Because multiple intervals and diagnostics are displayed, they are interpreted
descriptively rather than as a familywise error-controlled decision procedure; the BCa
analysis with five clusters is retained only as a fragile sensitivity.

Sample sizes were design- and compute-constrained and fixed before outcome anal-
ysis; no power calculation was used. No completed group, task, specification, budget
scheme, caliper, or finite-shot replicate was removed based on its result. Randomized
operations use recorded seeds; external sampling, ID division, and audit tie-breaking
use deterministic SHA-256 ordering. The test suite validates source-dataset mapping,
endpoint equality across analyses, candidate-budget coverage, matching uniqueness,
seed stability, exact-reference hashes, prediction/label separation, theorem endpoints
by exhaustive enumeration, MILP consistency, monotone partial-label contraction,
and complete Monte Carlo cells.

2255 **4.26 Software environment**

2256

2257 Experiments use Python 3.11/3.12, NumPy 2.3.5, pandas 2.2.3, scikit-learn 1.8.0,
2258 Matplotlib 3.11.0, and Qiskit 2.3.1. The Windows-specific Conda snapshot and cross-
2260 platform package specifications serve different archival purposes; neither is an exact
2261 machine image of every full-scale run.

2262

2263

2265 **4.27 Use of generative AI**

2266

2267 OpenAI Codex was used during manuscript revision for language editing, code
2268 inspection, analysis implementation, and consistency checking. It did not determine
2269 authorship, study accountability, or acceptance of any scientific claim. The authors
2270 inspected the generated changes, verified all reported analyses and citations, and take
2272 full responsibility for the manuscript.

2273

2274

2277 **Declarations**

2278

2279 **Data Availability.** This study uses the public EMBER [33], UNSW-NB15 [34],
2280 ToN-IoT [35], and TableShift [32] benchmarks, which remain subject to their original
2281 access conditions. Raw source records are not redistributed. Derived split definitions,
2282 deterministic row-position hashes, shift constructions, experiment manifests, complete
2283 aggregated results, target-prediction matrices, physically separated audit labels, evi-
2284 dence frontiers, technical-failure records, and integrity manifests are archived in the
2285 immutable v1.1.3 release at <https://doi.org/10.5281/zenodo.21751137>. The all-version
2286 concept DOI is <https://doi.org/10.5281/zenodo.19147649>; the exact v0.8.0 predecessor
2287 is <https://doi.org/10.5281/zenodo.21717074> [44].

2294

2295 **Code Availability.** Code for constructing the splits, running the kernel experi-
2296 ments, reproducing every analysis, and regenerating every table and figure is available
2297 at https://github.com/roberto-fernandez-barrios/kernel_shift_framework and in the
2298
2299
2300

versioned Zenodo archive above. A Windows-specific Conda snapshot and platform-neutral package specifications document the software dependencies; neither is claimed as an exact machine image of every full-scale run. The code is released under BSD-3-Clause with no restriction on non-academic use.

Competing interests. All authors declare no financial or non-financial competing interests.

Ethics approval and consent to participate. This computational study uses only public benchmark releases and involved no participant recruitment, intervention, or access by the authors to direct identifiers. No new human data were collected, and no new participant consent or ethics approval was required for this secondary benchmark analysis. The original data providers' access conditions and governance remain applicable.

Authors' contributions. R.F.-B.: conceptualization, methodology, software, data curation, formal analysis, visualization, investigation, writing—original draft, and writing—review and editing. I.P.-L.: supervision, validation, and writing—review and editing. A.G.-S.: validation, project administration, and writing—review and editing. P.G.B.: supervision, resources, and writing—review and editing. All authors read and approved the final manuscript.

Acknowledgements. This study received no funding.

References

- [1] Havlíček, V. *et al.* Supervised learning with quantum-enhanced feature spaces. *Nature* **567**, 209–212 (2019).
- [2] Schuld, M. & Killoran, N. Quantum machine learning in feature hilbert spaces. *Physical Review Letters* **122**, 040504 (2019).

- 2347 [3] Schuld, M. Supervised quantum machine learning models are kernel methods.
2348 *arXiv* **abs/2101.11020** (2021). URL <https://arxiv.org/abs/2101.11020>.
2349
2350
- 2351 [4] Bowles, J., Ahmed, S. & Schuld, M. Better than classical? the subtle art of bench-
2352 marking quantum machine learning models. *arXiv* **abs/2403.07059** (2024).
2353 URL <https://arxiv.org/abs/2403.07059>.
2354
2355
- 2356 [5] Schnabel, J. & Roth, M. Quantum kernel methods under scrutiny: A benchmark-
2357 ing study. *Quantum Machine Intelligence* **7**, 58 (2025).
2358
2359
- 2360 [6] Shen, J. Quantum kernels and the cross-section of stock returns: Anatomy of a
2361 vanishing advantage (2026). URL <https://arxiv.org/abs/2607.20168>. Preprint,
2362 arXiv:2607.20168.
2363
2364
- 2365 [7] Gulrajani, I. & Lopez-Paz, D. *In search of lost domain generalization*. *Inter-*
2366 *national Conference on Learning Representations* (International Conference
2367 on Learning Representations, 2021). URL <https://openreview.net/forum?id=lQdXeXDoWtI>.
2368
2369
- 2370 [8] Koh, P. W. *et al.* *Wilds: A benchmark of in-the-wild distribution shifts*. *Pro-*
2371 *ceedings of the 38th International Conference on Machine Learning*, Vol. 139
2372 of *Proceedings of Machine Learning Research*, 5637–5664. International Con-
2373 ference on Machine Learning (PMLR, Cambridge, MA, USA, 2021). URL
2374 <https://proceedings.mlr.press/v139/koh21a.html>.
2375
2376
- 2377 [9] Yao, H. *et al.* *Wild-time: A benchmark of in-the-wild distribution shift over*
2378 *time*. *Advances in Neural Information Processing Systems* (Neural Information
2379 Processing Systems Foundation, 2022). Datasets and Benchmarks Track.
2380
2381
- 2382 [10] Zhou, K., Liu, Z., Qiao, Y., Xiang, T. & Loy, C. C. Domain generalization: A
2383 survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **45**,
2384
2385
2386
2387
2388
2389
2390
2391
2392

4396–4415 (2023).	2393
	2394
[11] Huang, H.-Y. <i>et al.</i> Power of data in quantum machine learning. <i>Nature Communications</i> 12 , 2631 (2021).	2395
	2396
	2397
	2398
	2399
[12] Schreiber, F. J., Eisert, J. & Meyer, J. J. Classical surrogates for quantum learning models. <i>Physical Review Letters</i> 131 , 100803 (2023).	2400
	2401
	2402
	2403
[13] Thabet, S., Monbroussou, L., Mamon, E. Z. & Landman, J. When quantum and classical models disagree: Learning beyond minimum norm least square. <i>npj Quantum Information</i> 12 , 81 (2026).	2404
	2405
	2406
	2407
	2408
	2409
[14] Madani, O., Pennock, D. M. & Flake, G. W. <i>Co-validation: Using model disagreement on unlabeled data to validate classification algorithms. Advances in Neural Information Processing Systems</i> , Vol. 17, 873–880 (2004). URL https://proceedings.neurips.cc/paper_files/paper/2004/hash/92f54963fc39a9d87c2253186808ea61-Abstract.html .	2410
	2411
	2412
	2413
	2414
	2415
	2416
	2417
	2418
[15] Musgrove, D., Radtke, A. & Haddad, T. Discordant pair analysis for sample efficient model evaluation. <i>Scientific Reports</i> 13 , 20987 (2023).	2419
	2420
	2421
	2422
[16] Kossen, J., Farquhar, S., Gal, Y. & Rainforth, T. <i>Active testing: Sample-efficient model evaluation. Proceedings of the 38th International Conference on Machine Learning</i> , Vol. 139 of <i>Proceedings of Machine Learning Research</i> , 5753–5763 (PMLR, 2021). URL https://proceedings.mlr.press/v139/kossen21a.html .	2423
	2424
	2425
	2426
	2427
	2428
	2429
	2430
[17] Okanovic, P. <i>et al.</i> All models are wrong, some are useful: Model selection with limited labels. <i>Proceedings of The 28th International Conference on Artificial Intelligence and Statistics</i> , Vol. 258 of <i>Proceedings of Machine Learning Research</i> , 2035–2043 (PMLR, 2025). URL https://proceedings.mlr.press/v258/okanovic25a.html .	2431
	2432
	2433
	2434
	2435
	2436
	2437
	2438

- 2439 [18] Polo, F. M., Maity, S., Yurochkin, M., Banerjee, M. & Sun, Y. *Weak supervision*
2440 *performance evaluation via partial identification. Advances in Neural Information*
2441 *Processing Systems*, Vol. 37 (2024).
2442
2443
2444
- 2445 [19] Abreu, D., Esteve Rothenberg, C. & Abelém, A. *QML-IDS: Quantum machine*
2446 *learning intrusion detection system. 2024 IEEE Symposium on Computers and*
2447 *Communications (ISCC)*, 1–6. IEEE Computer Society (IEEE, Paris, France,
2448 2024).
2449
2450
2451
- 2452 [20] Abreu, D., Moura, D., Esteve Rothenberg, C. & Abelém, A. QuantumNetSec:
2453 Quantum machine learning for network security. *International Journal of Network*
2454 *Management* **35**, e70018 (2025).
2455
2456
2457
- 2458 [21] Bellante, A., Fioravanti, T., Carminati, M., Zanero, S. & Luongo, A. Evaluat-
2459 ing the potential of quantum machine learning in cybersecurity: A case-study
2460 on PCA-based intrusion detection systems. *Computers & Security* **154**, 104341
2461 (2025).
2462
2463
2464
- 2465 [22] Quiñonero-Candela, J., Sugiyama, M., Schwaighofer, A. & Lawrence, N. D. (eds)
2466 *Dataset Shift in Machine Learning* (MIT Press, Cambridge, MA, 2008).
2467
2468
2469
- 2470 [23] Shaydulin, R. & Wild, S. M. Importance of kernel bandwidth in quantum machine
2471 learning. *Physical Review A* **106**, 042407 (2022).
2472
2473
- 2474 [24] Canatar, A., Peters, E., Pehlevan, C., Wild, S. M. & Shaydulin, R. Band-
2475 width enables generalization in quantum kernel models. *Transactions on Machine*
2476 *Learning Research* (2023).
2477
2478
2479
- 2480 [25] Kübler, J. M., Buchholz, S. & Schölkopf, B. *The inductive bias of quantum ker-*
2481 *nels. Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34,
2482 12661–12673 (Neural Information Processing Systems Foundation, 2021).
2483
2484

- [26] Thanasilp, S., Wang, S., Cerezo, M. & Holmes, Z. Exponential concentration in quantum kernel methods. *Nature Communications* **15**, 5200 (2024).
- [27] Wang, X., Du, Y., Luo, Y. & Tao, D. Towards understanding the power of quantum kernels in the NISQ era. *Quantum* **5**, 531 (2021).
- [28] Liu, Y., Arunachalam, S. & Temme, K. A rigorous and robust quantum speed-up in supervised machine learning. *Nature Physics* **17**, 1013–1017 (2021).
- [29] Jerbi, S., Gyurik, C., Marshall, S. C., Molteni, R. & Dunjko, V. Shadows of quantum machine learning. *Nature Communications* **15**, 5676 (2024).
- [30] Engelen, G., Rimmer, V. & Joosen, W. *Troubleshooting an intrusion detection dataset: The CICIDS2017 case study. 2021 IEEE Security and Privacy Workshops (SPW)*, 7–12. IEEE Computer Society (IEEE, San Francisco, CA, USA, 2021).
- [31] Kakavand, S., Strohmeyer, C. & Schlotter, M. Benchmarking quantum kernel support vector machines against classical baselines on tabular data: A rigorous empirical study with hardware validation. *arXiv* **abs/2604.18837** (2026). URL <https://arxiv.org/abs/2604.18837>.
- [32] Gardner, J., Popovic, Z. & Schmidt, L. *Benchmarking distribution shift in tabular data with TableShift. Advances in Neural Information Processing Systems*, Vol. 36 (Neural Information Processing Systems Foundation, 2023).
- [33] Anderson, H. S. & Roth, P. Ember: An open dataset for training static PE malware machine learning models. *arXiv* **abs/1804.04637** (2018). URL <https://arxiv.org/abs/1804.04637>.

- 2531 [34] Moustafa, N. & Slay, J. *UNSW-NB15: A comprehensive data set for network*
2532 *intrusion detection systems (UNSW-NB15 network data set),. 2015 Military*
2533 *Communications and Information Systems Conference (MilCIS)*, 1–6 (IEEE,
2534 2015).
2535
2536
2537
2538 [35] Alsaedi, A., Moustafa, N., Tari, Z., Mahmood, A. & Anwar, A. TON_IoT teleme-
2539 try dataset: A new generation dataset of IoT and IIoT for data-driven intrusion
2540 detection systems. *IEEE Access* **8**, 165130–165150 (2020).
2541
2542
2543
2544 [36] Cortes, C. & Vapnik, V. Support-vector networks. *Machine Learning* **20**, 273–297
2545 (1995).
2546
2547
2548 [37] Rasmussen, C. E. & Williams, C. K. I. *Gaussian Processes for Machine Learning*
2549 (MIT Press, Cambridge, MA, 2006).
2550
2551
2552 [38] Javadi-Abhari, A. *et al.* Quantum computing with qiskit. *arXiv* **abs/2405.08810**
2553 (2024). URL <https://arxiv.org/abs/2405.08810>.
2554
2555
2556 [39] Ibragimov, R. & Müller, U. K. Inference with few heterogeneous clusters. *The*
2557 *Review of Economics and Statistics* **98**, 83–96 (2016).
2558
2559
2560 [40] Guo, C., Pleiss, G., Sun, Y. & Weinberger, K. Q. *On calibration of modern neural*
2561 *networks. Proceedings of the 34th International Conference on Machine Learn-*
2562 *ing (ICML)*, Vol. 70 of *PMLR*, 1321–1330 (International Conference on Machine
2563 Learning, 2017).
2564
2565
2566
2567 [41] Roy, O. & Vetterli, M. *The effective rank: A measure of effective dimension-*
2568 *ality. 2007 15th European Signal Processing Conference (EUSIPCO)*, 606–610
2569 (EURASIP, 2007).
2570
2571
2572
2573
2574
2575
2576

[42] Cristianini, N., Shawe-Taylor, J., Elisseeff, A. & Kandola, J. <i>On kernel-target alignment. Advances in Neural Information Processing Systems (NeurIPS)</i> , Vol. 14 (Neural Information Processing Systems Foundation, 2001).	2577 2578 2579 2580 2581 2582
[43] Cortes, C., Mohri, M. & Rostamizadeh, A. Algorithms for learning kernels based on centered alignment. <i>Journal of Machine Learning Research</i> 13 , 795–828 (2012).	2583 2584 2585 2586 2587 2588
[44] Fernández-Barrios, R., Pastor-López, I., González-Santocildes, A. & García Bringas, P. Evaluation choices shape apparent quantum kernel robustness under distribution shift (2026). URL https://doi.org/10.5281/zenodo.21717074 .	2589 2590 2591 2592 2593 2594 2595 2596 2597 2598 2599 2600 2601 2602 2603 2604 2605 2606 2607 2608 2609 2610 2611 2612 2613 2614 2615 2616 2617 2618 2619 2620 2621 2622