

Sharp Target-Domain Certificates for Quantum-Kernel Advantage under Distribution Shift

Roberto Fernández-Barrios^{[ORCID]1*}, Iker Pastor-López^{[ORCID]1},
Asier González-Santocildes^{[ORCID]1}, Pablo García Bringas^{[ORCID]1}

^{1*}Faculty of Engineering, University of Deusto, Avda. de las
Universidades, 24., Bilbao, 48007, Spain.

*Corresponding author(s). E-mail(s): roberto.fernandez.b@deusto.es;
Contributing authors: iker.pastor@deusto.es; gonzalez.asier@deusto.es;
pablo.garcia.bringas@deusto.es;

Abstract

Quantum predictive advantage under shift usually requires known target labels. We derive the assumption-free sharp identified interval for the finite-batch advantage of a fixed candidate over the best member of a prespecified fixed classical-kernel family under any bounded loss and unrestricted completions of the unaudited labels; zero-one accuracy has exact closed-form partial-label updates. Across eight security shifts, zero-label upper endpoints against 115 classical kernels span 0.002–0.088. Retrospective auditing reduces every endpoint to 0.010 with 0–33 of 500 labels, including entangling-ZZ models. Prospective corroboration on two eligible tasks met predefined criteria: four task-classifier medians require 0–3 labels (overall median 0; maximum 76), and all 20 realized effects are negative; a third task fails its feature gate. Protocol controls reverse an apparent advantage, while finite-shot noise can increase predictive distinctness without useful advantage. The framework separates protocol validity, target predictive indispensability, and quantum relevance.

Keywords: quantum machine learning; quantum kernels; quantum advantage; partial identification; active testing; distribution shift; classical surrogates; kernel methods

047 1 Introduction

048
049 Quantum kernel methods embed classical data into quantum states and use state
050 overlaps as similarities for otherwise classical learners [1–3]. This separation makes
051 them a clean setting for studying quantum inductive bias, but a predictive-advantage
052 claim must pass three distinct tests. The comparison protocol must be valid; the target
053 predictions must resist reproduction by an appropriate classical reference family; and
054 the surviving behaviour must depend on a non-trivial quantum construction whose
055 resources are counted. A model can pass one test and fail another.

056
057 Recent benchmarks emphasize the first test. Broad QML comparisons find compet-
058 itive classical baselines once pipelines are aligned [4]; Slattery et al. show numerically
059 that general-purpose hyperparameter tuning can make fidelity kernels well approxi-
060 mated by classical kernels on classical data [5]; and quantum-kernel results depend
061 strongly on feature maps and hyperparameters [6]. Those studies concern benchmark
062 performance or global kernel approximation. Our contribution instead gives an exact
063 finite-target interval conditional on fixed predictions and a prespecified classical-kernel
064 reference family, including under partial target-label auditing. A concurrent finance
065 preprint likewise shows an apparent quantum-kernel advantage vanishing under point-
066 in-time, kernel-swap, and budget controls [7]. Under distribution shift the selection
067 rule becomes part of the estimand: same-test selection is an oracle, train-only regu-
068 larization can change rankings, and a larger candidate family has more opportunities
069 to win [8–11]. Fair equal-budget benchmarking answers who wins under symmetric
070 search, but not whether any classical witness in a declared family can reproduce the
071 selected quantum model on the deployment inputs.

072
073 Geometric and surrogate analyses approach that second question from complemen-
074 tary directions. Huang et al. use a label-independent geometric difference to screen
075 whether a quantum kernel can have a prediction advantage [12]; classical-surrogate
076 constructions reproduce input–output relations of trained quantum models [13]; and
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092

recent theory studies when quantum learners depart from minimum-norm classical solutions [14]. Those objects characterize geometric potential, global approximation, or learning-rule separation. Our target-batch question is different. A train-only Gram matrix does not determine how a kernel transports predictions to a shifted target because deployment also depends on the target–train block, while different geometries can yield identical decisions on the observed target inputs. We therefore audit predictive indispensability directly on target covariates while keeping their labels protected.

That protection creates an identification problem in the tradition of partial identification: observed information and maintained restrictions can determine a set of compatible values without point-identifying the target [15, 16]. Here the observed object is unusually concrete—one finite batch and fixed prediction matrices—and the maintained label restriction is deliberately empty. The estimand is the relative functional “fixed candidate versus best member of a prespecified fixed family.” We derive its sharp interval over unrestricted completions of the unaudited labels, rather than estimate a population parameter or add assumptions until point identification is recovered. Weak-supervision work also uses partial identification, but combines fitted label-model marginals with Fréchet restrictions for individual metrics [17]; those marginals are neither available nor required here.

Label-efficient evaluation pursues a different objective. Active evaluation has used adaptive stratified sampling [18], importance-weighted Active Testing [19], learned loss surrogates [20], and prediction-powered control variates [21] to estimate risk accurately with fewer labels. Limited-label model selection instead chooses which pretrained classifier to deploy [22], while semi-supervised multi-model evaluation borrows information across classifier scores and unlabeled data under a mixture model [23]. Under distribution shift, ATC predicts target accuracy from confidence [24]; disagreement-discrepancy methods provide assumption-dependent error bounds

139 [25, 26]. Our certificate estimates neither labels nor risk: it retains every comple-
140 tion and reports the exact range of a fixed relative functional. Co-validation and
141 discordant-pair evaluation remain complementary uses of disagreement [27, 28].
142

143 The distinction from minimax transductive learning is equally important. Balsub-
144 ramani and Freund formulate ensemble prediction on an unlabeled test set as a game,
145 using prior classifier-label correlation constraints to choose a new predictor that min-
146 imizes worst-case error [29]. We do not optimize or combine predictions and impose
147 no such label constraints. The candidate and reference predictions remain fixed; the
148 adversarial completion is used only to identify the sharp range of their realized finite-
149 batch advantage. Thus transductive minimax aggregation is a decision rule under
150 supplied performance information, whereas our object is an assumption-free identifi-
151 cation certificate conditional on a declared prediction matrix. The distinction concerns
152 the target, admissible information, and action, not only the computational form.
153

154 We make four contributions. First, for any additive bounded loss on a finite label
155 space, we derive the interval hull of the partial-label identified set for advantage over
156 the best member of a prespecified classical-kernel reference family and show that
157 no assumption-free rule using the same information can return a uniformly smaller
158 valid interval. Zero-one accuracy has an unusually simple closed form and is pathwise
159 exact after any realized audit subset, including subsets selected adaptively. Second,
160 nested classical reference families and target-label audits define a reference-breadth-
161 target-supervision evidence frontier. Third, we give an exact prevalence-conditional
162 balanced-accuracy formulation and an explicit PSD construction showing why train-
163 only geometry cannot determine target predictions. Fourth, we embed these objects
164 in a controlled quantum-kernel benchmark and prospectively specify a Gate 2 repli-
165 cation before opening target labels; entanglement strata and repeated finite-shot
166 perturbations separately test when predictive distinctness is quantum-relevant.
167

168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184

The benchmark uses four security scenarios from EMBER, UNSW-NB15, and ToN-IoT, producing eight fixed constructed or held-out-campaign shifts. Support-vector and Laplace Gaussian-process classifiers consume identical precomputed Gram matrices. The quantum family contains entangling-ZZ and separable-product fidelity maps at 4–12 qubits; the classical family contains 115 linear, RBF, polynomial, Laplacian, and Matérn candidates. Models are tuned on training data and selected on a disjoint ID-validation split. Equal 60-candidate pools support fair performance comparison, while the complete prespecified classical-kernel reference family supports the separate Gate 2 indispensability audit. One prospectively specified three-task TableShift study corroborates protocol sensitivity; a second tests Gate 2 prospectively on previously unused task identifiers.

Against the full prespecified classical-kernel reference family, the zero-label upper endpoint of accuracy advantage ranges from 0.002 to 0.050 for SVC and 0.006 to 0.088 for GPC. In a retrospective fixed-batch analysis, an adaptive bottleneck-coverage audit reduces it to at most 0.010 using 0–33 of 500 labels (median 13 over 16 fixed case-classifier cells). Dynamic random-disagreement and non-adaptive initial-coverage controls both have median 12.5, so the large gain over prediction-independent ordering is attributable to excluding non-actionable queries rather than to the proposed coverage heuristic. The result persists for selected entangling-ZZ models. Prospective corroboration on two technically eligible tasks met the predefined criteria: four task-classifier medians were 0–3 labels, the overall seed-level median was 0, and the maximum was 76. All 20 prospective realized accuracy effects were negative. The third task failed its prespecified minimum-feature gate before model execution. The certificate remains conditional on the prespecified classical-kernel reference family and target batch, not evidence of classical simulability or complexity-theoretic separation. It converts the earlier benchmark conclusion—that controls remove an apparent gain—into




Fig. 1 The controlled kernel-swap protocol. Benchmark pools feed reproducible shift constructions; all kernels operate on a shared train-only embedding, receive per-configuration regularization tuned on the training split, and are consumed by the same precomputed-kernel classifiers. Family-level conclusions are drawn by target-label-free, equal-budget selection on an in-distribution validation split (P1'). Each scenario-group interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the fixed benchmark pools and design. Gram matrices feed the associational geometry audit and repeated finite-shot sensitivity.

a quantitative question: how much advantage remains possible, and how much classical search and target supervision suffice to falsify a material claim?

2 Results

2.1 Overview

The results are organized as three gates. Gate 1 asks whether the comparison protocol itself is valid: regularization is fitted on training data, deployment candidates are selected without OOD labels, and family search budgets are distinguished from unrestricted adversarial reference breadth. Gate 2 asks whether the selected quantum model is predictively indispensable relative to the prespecified classical-kernel reference family on the observed target inputs. This is the new target-batch certificate and evidence frontier. Gate 3 asks whether any surviving behaviour is quantum-relevant, separating entangling- ZZ maps from product maps and accounting for finite-shot

estimation. The eight security groups are fixed cases; conditional intervals over q-
split clusters describe pipeline variation given the benchmark pools and design, not
uncertainty over a population of datasets.

2.2 Sharp target-batch identification of predictive advantage

Consider a finite target batch, a fixed quantum prediction, and M fixed classical
predictions. For a finite label space $\mathcal{Y}$ and an additive bounded loss, let $\ell_i^Q(y)$ and $\ell_{ij}^C(y)$
be their losses on item i if its label were y , and define the classical-minus-quantum
contrast $a_{ij}(y) = \ell_{ij}^C(y) - \ell_i^Q(y)$. Positive advantage favours the quantum model:

$$\Delta_\ell(y) = \min_{j \leq M} \frac{1}{n} \sum_{i=1}^n a_{ij}(y_i).$$

If labels have been audited on L and $S_j(L) = \sum_{i \in L} a_{ij}(y_i)$, every assignment on L^c
remains admissible.

Theorem 1 (Sharp bounded-loss interval hull: sharpness and minimality) *The smallest
interval containing $\Delta_\ell(y)$ for every completion of the unaudited labels is $[L_\ell(L), U_\ell(L)]$, where*

$$L_\ell(L) = \frac{1}{n} \min_j \left\{ S_j(L) + \sum_{i \notin L} \min_{y \in \mathcal{Y}} a_{ij}(y) \right\},$$

$$U_\ell(L) = \frac{1}{n} \max_{y_{L^c} \in \mathcal{Y}^{|L^c|}} \min_j \left\{ S_j(L) + \sum_{i \notin L} a_{ij}(y_i) \right\}.$$

*Both endpoints are attained. Consequently, any assumption-free interval computed from the
same fixed losses and audited labels and valid for every admissible completion must contain
both endpoints.*

The lower endpoint separates over labels after choosing a witness. The upper
endpoint is an exact finite max-min problem, expressible as a mixed-integer linear

program with one-hot label variables. Auditing any additional label removes completions, so the lower endpoint cannot decrease and the upper endpoint cannot increase. This general statement is not a claim that loss-based partial identification or active evaluation is new; the contribution is the exact relative-family estimand and its sharp, minimal finite-batch certificate. Zero-one classification has further structure: assigning an unaudited label to the quantum hard prediction simultaneously maximizes its contrast with every disagreeing classical witness. This yields the following closed forms used in all experiments.

Fix the hard target predictions of a selected quantum classifier h_Q and a pre-specified classical family $\mathcal{C}_B = \{h_{C_1}, \dots, h_{C_M}\}$ on n target inputs. For an otherwise unrestricted target labelling y , define

$$\Delta_B(y) = \text{Acc}(h_Q; y) - \max_{j \leq M} \text{Acc}(h_{C_j}; y), \quad d_j = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{h_Q(x_i) \neq h_{C_j}(x_i)\}.$$

Theorem 2 (Sharp zero-label region) *The smallest interval containing $\Delta_B(y)$ for every possible target labelling is*

$$\mathcal{I}_B^{(0)} = \left[-\max_{j \leq M} d_j, \min_{j \leq M} d_j \right].$$

Both endpoints are attainable, including for multiclass hard predictions.

The upper endpoint is not a probabilistic bound or an estimate: it is the exact largest advantage compatible with the observed prediction matrix when every target label remains unknown. It is attained by setting $y = h_Q$. The lower endpoint is attained by assigning the predictions of the classical model farthest from h_Q . Hence no smaller label-free interval is valid without extra assumptions. In particular, if $\min_j d_j < \tau$, any accuracy advantage of at least τ over this classical family is falsified on the batch, even before a target label is revealed. Whenever $\min_j d_j > 0$, the same target covariates and prediction matrix remain compatible with both signs of advantage; disagreement alone cannot identify which family wins.

The zero-label certificate is already restrictive in several cases (Figure 2a). Against the complete 115-candidate classical family, the upper endpoint ranges from 0.002 to 0.050 for SVC and from 0.006 to 0.088 for GPC (Table 1). A small endpoint says that at least one classical witness nearly reproduces the selected quantum decisions; it does not say that the quantum kernel is efficiently simulable or globally reproducible.

Theorem 3 (Sharp arbitrary-partial-label region) *Let L be any audited subset. For classical witness j , let*

$$S_j(L) = \sum_{i \in L} [\mathbf{1}\{h_Q(x_i) = y_i\} - \mathbf{1}\{h_{C_j}(x_i) = y_i\}]$$

and let $R_j(L)$ count unaudited points on which h_Q and h_{C_j} disagree. Then the sharp identified interval over all completions of the unaudited labels is

$$\mathcal{I}_B(L) = \left[\min_j \frac{S_j(L) - R_j(L)}{n}, \min_j \frac{S_j(L) + R_j(L)}{n} \right].$$

The upper endpoint is non-increasing as labels are audited and collapses to the realized family advantage when L contains the full batch.

We use this result as an active falsification audit, not as a new model-selection rule. At each step the models remain fixed; one tested rule queries the point covering the largest number of witnesses currently tying at the smallest reducible upper bound (Methods). In this retrospective fixed-batch analysis, the rule reduces the 0.010 endpoint using 0–33 of 500 labels, with median 13 across the 16 case–classifier cells (Figure 2b). Stronger controls remove any claim that the coverage heuristic itself is unusually efficient: dynamic random sampling restricted to current actionable disagreements and a fixed initial-coverage ordering both have median 12.5 labels. Against random-active sampling, the adaptive rule wins, ties, and loses in 5, 5, and 6 cells; against initial coverage the counts are 6, 4, and 6. An exact ex-post label oracle needs median 6.5 labels, and the adaptive rule’s median excess is 4.5. The main acquisition result is therefore that prediction-aware restriction makes a small audit possible,




Fig. 2 Sharp target-batch certificates against the full 115-candidate classical family. **a**, Zero-label upper endpoints for the SVC and GPC quantum models fixed by $P1'$. Colour separates selected entangling-ZZ and product-map models; the dashed line marks 0.010 accuracy. **b**, Labels needed to reduce the endpoint to at most 0.010 (symmetric-log scale). Horizontal ticks are the exact retrospective label oracle; filled markers are the fixed adaptive coverage rule; open diamonds and triangles are medians for dynamic random-active and non-adaptive initial-coverage controls, with bars spanning 2.5–97.5% across 200 SHA-256 tie orders. All certificates use the canonical q1000 realization (master, q-split, and model seeds 42) and a 500-point target batch. Comparator ranges are tie-order sensitivities, not confidence intervals.

Table 1 Sharp accuracy certificates for the eight fixed target batches against the full 115-candidate classical family. U_0 is the zero-label upper endpoint; $n_{.01}$ is the fixed adaptive-coverage count for $U \leq 0.010$; Δ_{full} is the point-identified quantum advantage after all 500 labels are unlocked. The canonical run, quantum model, and every classical candidate were fixed before this audit; target labels are never used to retune or reselect them.

Target case	SVC			GPC		
	U_0	$n_{.01}$	Δ_{full}	U_0	$n_{.01}$	Δ_{full}
EMBER $m1$	.050	12	-.024	.080	29	-.030
EMBER $m2$	.054	15	-.036	.074	33	-.016
ToN-IoT constructed	.002	0	-.004	.016	3	-.014
ToN-IoT campaign	.002	0	-.018	.006	0	-.012
UNSW DoS constructed	.010	0	-.016	.066	22	-.052
UNSW DoS campaign	.034	11	-.026	.028	12	-.032
UNSW Recon constructed	.046	14	-.032	.088	22	-.066
UNSW Recon campaign	.034	28	-.026	.030	14	+.002

not that one heuristic dominates. Prediction-independent ordering remains a weak diagnostic (median 245.5) and is reported only in Supplementary Information.

2.3 A reference-breadth–target-supervision evidence frontier

Nested classical reference families and partial labels expose two different costs of falsification. Increasing the classical family can only decrease the certificate’s upper

endpoint; auditing informative target labels can only decrease it further. We therefore report the surface $(B, |L|) \mapsto U_B(L)$ and, for a materiality threshold τ , the smallest reference breadth and target supervision required to establish $U_B(L) \leq \tau$. Here B indexes a prespecified *reference breadth*, not runtime or computational complexity.

The frozen block-ordering analysis already provides three nested breadths: 30, 60, and 115 candidates. For each cell we take the median zero-label endpoint over 5,000 prespecified whole-block orderings, then summarize the 16 case-classifier cells. The cross-cell median contracts from 0.042 at $B = 30$ to 0.034 at $B = 60$ and remains 0.034 at $B = 115$; the number of cells already at or below 0.010 changes from 3 to 4 to 4 (Supplementary reference-breadth table). Thus most aggregate contraction occurs by 60 candidates, but the full family remains necessary for the deliberately adversarial question and for individual cases. These ordering distributions quantify dependence on reference breadth; they are not confidence intervals or computational-cost measurements.

The prediction-aware evidence curves contract rapidly because only disagreements capable of lowering a current witness endpoint need adjudication (corresponding curve display in Supplementary Information). The adaptive rule is not consistently better than random-active sampling, although recomputing the active set avoids the 286.5-label median outlier produced by a frozen initial-coverage ordering. The result persists when the selected quantum model uses an entangling-ZZ map: across the 12 entangling case-classifier cells, the adaptive rule needs 0–33 labels (median 14) for the 0.010 endpoint, while random-active cell medians are 0–34 (median 12.5). Thus the finding is not an artefact of including classically factorizable product maps.

Reference breadth matters separately from fair benchmarking. For ToN-IoT scanning under GPC, the customary 30-candidate linear/RBF family yields a realized quantum accuracy advantage of +0.028, so no number of labels can falsify a one-point claim against that restricted reference. Against the full 115-candidate family the

507 realized advantage is -0.014 , the zero-label upper endpoint is 0.016 , and three tar-
508 getted labels reduce it below 0.010 . Equal 60-versus-60 budgets remain the appropriate
509 performance comparison; the full family asks the different, deliberately adversarial
510 question of predictive indispensability.
511

512 2.4 Prospective Gate-2 corroboration against the prespecified 513 classical-kernel reference family 514

515 The security acquisition curves above are retrospective. We therefore specified a sepa-
516 rate Gate-2 replication against the prespecified classical-kernel reference family before
517 acquiring its three TableShift tasks, computing a model prediction, or opening a target
518 label. The task identifiers were BRFSS Diabetes, ACS Food Stamps, and NHANES
519 Lead. The immutable specification fixed five label-blind sampling seeds, the q1000 split
520 sizes, $P1'$ selection, both classifiers, the 60-quantum and 115-classical pools, physical
521 prediction-label separation, three thresholds, 200 SHA-256 orders for each strong con-
522 trol, and numerical outcome criteria. The aggregate prediction lock was hashed before
523 the one-time label audit (Methods).
524

525 All five NHANES seeds produced only 11 non-constant post-encoding features,
526 below the frozen 12-dimensional grid requirement. The task therefore failed a pre-
527 specified minimum-feature gate before any model candidate was executed; it was not
528 replaced or repaired by reducing the grid. The replication is consequently techni-
529 cally limited to ten task-seed units, or 20 fixed task-seed-classifier cells, rather than
530 evidence from three complete tasks.
531

532 The two technically eligible tasks met the predefined numerical criteria (Figure 3).
533 Against the full prespecified classical-kernel reference family, zero-label endpoints
534 range from 0 to 0.168 (median 0.010), and $11/20$ cells already satisfy $U \leq 0.010$ with-
535 out target labels. The fixed adaptive audit reduces every remaining endpoint to that
536 threshold with 0–76 labels (overall median 0). The four task-classifier medians are 0
537
538




Fig. 3 Prospectively specified Gate-2 replication against the full 115-candidate prespecified classical-kernel reference family. **a**, Zero-label upper endpoints for five deterministic sampling seeds in each technically eligible task-classifier cell; black ticks mark medians and the dashed line marks 0.010 accuracy. **b**, Labels required to reduce the endpoint to at most 0.010. Coloured points are the five fixed adaptive counts, black ticks are exact retrospective label-oracle counts, and open markers with bars are medians and 2.5–97.5% quantiles for dynamic random-active and non-adaptive initial-coverage policies across 1,000 seed-order combinations (five seeds times 200 SHA-256 orders). The 50-label line is the predefined task-classifier median criterion; the dotted 100-label line is the predefined seed-level ceiling. Comparator ranges are algorithmic sensitivities, not confidence intervals. NHANES Lead failed the prespecified 12-feature gate before model execution and is absent.

for ACS SVC, 3 for ACS GPC, 3 for BRFSS SVC, and 0 for BRFSS GPC; all are below the 50-label criterion, and the worst seed remains below the predefined 100-label ceiling. Every realized accuracy effect is negative, ranging from -0.156 to -0.002 .

The prospective controls reproduce the earlier negative algorithmic conclusion. Against both dynamic random-active and non-adaptive initial-coverage medians, adaptive coverage wins, ties, and loses in 3, 15, and 2 of the 20 cells, respectively; the median paired difference is zero. Prediction-independent ordering is far less efficient in the non-trivial cells. Thus the transferred result concerns the sharp certificate and restriction to actionable disagreements, not superiority of the coverage heuristic. All prospectively selected quantum winners are product maps, so this replication does not independently validate the entangling-ZZ stratum; that stratum remains supported only by the retrospective security cases.

599 2.5 The apparent advantage under oracle selection

600
601 We begin with a deliberately optimistic reference comparison: fidelity-based quantum
602 kernels against linear and RBF baselines, with each family’s representative chosen
603 by Best-by-OOD selection—the configuration with the highest balanced accuracy on
604 the OOD test split itself—and the classical kernels left at the default regularization
605 $C = 1$. Related intrusion-detection studies motivate the small-reference context [30–
606 32], but we do not claim that this exact protocol dominates QML practice. Under
607 these conventions, on EMBER, the quantum family shows an OOD advantage of
608 $+0.037$ under SVC and $+0.028$ under the GP classifier against linear+RBF, shrinking
609 to $+0.004/+0.002$ once the classical family is enlarged with heavier-tailed kernels.
610 Supplementary Table 5 gives the complete oracle results.

611
612 Two features favour this apparent quantum result. Best-by-OOD is an *oracle*:
613 it uses the labels it then reports. Fixing $C = 1$ also leaves classical candidates at
614 a potentially unsuitable regularization point. These values are therefore descriptive
615 upper bounds, not deployable effects.

624 2.6 The controlled comparison does not preserve the apparent 625 advantage

626
627 A deployable comparison must tune each configuration’s regularization without look-
628 ing at the test data, select the deployed configuration without OOD labels, and
629 compare families at the same candidate budget. We implement all three (Section 4).
630 For every kernel candidate, the SVC regularization C is chosen by five-fold cross-
631 validation on the training Gram matrix alone. The deployed configuration is then the
632 one with the best balanced accuracy on an in-distribution *validation* split, held out
633 from the ID test split (protocol P1’); no OOD label is ever consulted in selection.
634 The primary comparison uses 60 candidates per extended family in every scenario-
635 group and classifier, with the full 115-candidate classical pool retained as a sensitivity.
636
637
638
639
640
641
642
643
644




Fig. 4 Descriptive target-label-free family contrasts in eight fixed scenario-groups. P1' OOD differences (quantum minus classical) are shown for SVC (left) and GPC (right). Filled circles compare against the equal-budget extended classical family; error bars are pointwise 95% conditional cluster- t intervals over five q-split clusters ($n = 5$), conditional on the fixed benchmark pools and design. Open diamonds compare against the customary linear+RBF reference. The intervals are not simultaneous and do not imply population inference or equivalence.

We treat the eight scenario-groups as fixed case studies and report, per group, the quantum-minus-classical OOD delta with a conditional interval over the five split realizations, and no population p -value (Section 4).

Table 2 and Figure 4 show no consistent evidence of a quantum advantage under the controlled endpoint. Against the equal-budget extended classical family, the source-dataset-equal difference is -0.00565 for SVC (leave-one-source-dataset-out range $[-0.00702, -0.00450]$) and -0.00094 for GPC ($[-0.00190, +0.00094]$). Against the customary linear+RBF reference, the corresponding differences are -0.00407 and $+0.00591$, with heterogeneous group signs. UNSW-NB15 reconnaissance under held-out-campaign shift has a small positive mean ($+0.005$ for both classifiers), with pointwise conditional intervals spanning zero. The full-pool extended-family sensitivity is similar in direction and scale (-0.00378 for SVC and -0.00293 for GPC).

Two secondary checks do not alter this fixed-case conclusion. Source-dataset-equal estimates remain close across blocked, uniform, and kernel-stratified equal-count

Table 2 Target-label-free selection results for eight fixed scenario-groups. Values are quantum-minus-classical OOD balanced-accuracy differences under $P1'$ (select on ID-validation, tune per-configuration C on train, evaluate on OOD test). The frozen extended-family analysis uses 60 candidates per family in every group-classifier cell. Each scenario interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the benchmark pools and design; the fifteen pipeline realizations per setting are nested within these clusters. An interval containing zero is not an equivalence result. Source-dataset-equal rows average groups within three source datasets and report leave-one-source-dataset-out ranges, not confidence intervals. The customary linear+RBF comparison is contextual rather than equal-budget.

Scenario	Clf.	vs lin./RBF	vs ext. (equal B)	interval / LODO range
EMBER $m1$	SVC	-0.0095	-0.0058	$[-0.0340, +0.0223]$
	GPC	-0.0050	-0.0009	$[-0.0101, +0.0083]$
EMBER $m2$	SVC	+0.0139	-0.0064	$[-0.0137, +0.0009]$
	GPC	+0.0149	+0.0027	$[-0.0027, +0.0081]$
UNSW-DoS (campaign)	SVC	-0.0043	-0.0033	$[-0.0139, +0.0073]$
	GPC	+0.0041	-0.0057	$[-0.0120, +0.0005]$
UNSW-DoS (constructed)	SVC	-0.0085	-0.0044	$[-0.0119, +0.0031]$
	GPC	-0.0056	-0.0108	$[-0.0177, -0.0039]$
UNSW-Recon (campaign)	SVC	+0.0004	+0.0050	$[-0.0058, +0.0157]$
	GPC	+0.0074	+0.0052	$[-0.0019, +0.0123]$
UNSW-Recon (constructed)	SVC	-0.0076	-0.0089	$[-0.0146, -0.0031]$
	GPC	-0.0083	-0.0075	$[-0.0208, +0.0059]$
ToN-IoT (campaign)	SVC	-0.0172	-0.0148	$[-0.0268, -0.0029]$
	GPC	-0.0046	-0.0090	$[-0.0190, +0.0010]$
ToN-IoT (constructed)	SVC	-0.0016	-0.0011	$[-0.0049, +0.0028]$
	GPC	+0.0313	+0.0109	$[+0.0017, +0.0202]$
Source-dataset-equal mean, SVC		-0.0041	-0.0056	$[-0.0070, -0.0045]$
Source-dataset-equal mean, GPC		+0.0059	-0.0009	$[-0.0019, +0.0009]$

sampling (SVC -0.00576 to -0.00548 ; GPC -0.00253 to -0.00185). Observational nearest-rank pairings at the prespecified 1.25 caliper retain 75.2% of pairs and have negative median quantum-minus-classical differences in all eight groups; one-to-one and caliper sensitivities agree. Supplementary Tables 2–4 and Supplementary Figure 3 give the complete schemes, counts, and cautions; none identifies a causal family effect.

2.7 The quantum pool mixes entangling and product kernels

Circuit decomposition changes the interpretation of the aggregate family. The two ZZFeatureMap variants contain all-to-all two-qubit Pauli interactions, whereas the Pauli-X/Z and ZFeatureMap variants contain only one-qubit Pauli strings. For the latter pair, the prepared state and its fidelity kernel factorize over coordinates (Methods),

Table 3 Logical circuit audit for the four quantum feature maps. Depth and CX columns give the values at 4 and 12 qubits after Qiskit 2.3.1 decomposition to **rz**, **sx**, **x**, and **cx**, optimization level 0, without device routing. “Fidelity CX” counts the compute–uncompute template for one kernel entry. Complete values at all five dimensions and one-qubit gate counts are provided in Supplementary Table 13 and the machine-readable artifact. These logical counts are not hardware-execution estimates.

Map	Pauli strings	Stratum	Product factorization	map depth (4/12)	map CX (4/12)	fidelity CX (4/12)
ZZ, 1 rep.	Z, ZZ	entangling ZZ	no	19/67	12/132	24/264
ZZ, 2 reps.	Z, ZZ	entangling ZZ	no	35/107	24/264	48/528
Pauli-X/Z, 1 rep.	X, Z	separable product	yes	11/11	0/0	0/0
Z, 2 reps.	Z	separable product	yes	8/8	0/0	0/0

permitting direct classical evaluation as a product of one-qubit overlaps. Thus two of the four map families cannot supply evidence that entanglement or hard quantum kernel estimation is responsible for their behaviour.

Table 3 reports logical, unrouted resource counts. At 12 qubits, one ZZFeatureMap state preparation requires 132 CX gates for one repetition and 264 for two; the corresponding compute–uncompute fidelity templates require 264 and 528 CX gates. Both product-map templates require zero CX gates at every studied dimension. Under P1’ selection, product maps supply 594/1,080 (55.0%) of the selected quantum SVC configurations and 348/1,080 (32.2%) of the GP selections. The composition alone does not compare predictive quality; the matched stratum analysis below does so at 30 candidates per stratum.

At matched 30-candidate budgets, the source-dataset-equal SVC effect is -0.00928 for the entangling-ZZ stratum and -0.00547 for the product-map stratum; the constituent source-dataset means range from -0.01158 to -0.00580 and from -0.01167 to $+0.00021$, respectively. The entangling-ZZ point estimate is negative in every fixed scenario-group (-0.01960 to -0.00132); product-map estimates range from -0.02245 to $+0.00571$. The full 60-candidate pool remains at -0.00566 . Thus the controlled result is not produced by pooling product maps with entangling circuits, and restricting the quantum family to ZZ interactions does not reveal a hidden positive family effect. Supplementary Table 8 reports every pointwise conditional interval; these strata are fixed-case sensitivities, not new population tests.

783 2.8 Within-v4 sensitivities separate the evaluation bundle

784
785 The within-v4 factorial holds the data generation fixed while varying four evaluation
786 choices. In the q1000 SVC stratum, the fixed- C , same-test-oracle, customary-reference,
787 native-budget endpoint is $+0.0131$, whereas the train-CV, ID-validation, extended-
788 reference, equal-count endpoint is -0.0047 , a descriptive contraction of 0.0178 . The
789 16 three-source-dataset-equal endpoints range from -0.0084 to $+0.0131$ (Figure 5;
790 Supplementary Table 10).

791
792 The largest paired mean changes arise from train-CV regularization (-0.0084)
793 and the extended reference (-0.0079), but a regularization-by-reference interaction
794 of $+0.0139$ prevents additive or causal attribution (Supplementary Table 11). These
795 controls therefore form an interacting evaluation bundle.

796
797 Removing the frozen port-, protocol-, and service-related fields also changes the
798 relative comparison, confirming that common feature access alone does not make short-
799 cut use family-invariant. Across the two network source datasets, the controlled effect
800 moves from -0.0031 to -0.0245 (change -0.0215 ; source-dataset changes -0.0447
801 for ToN-IoT and $+0.0018$ for UNSW-NB15). The six group changes range from
802 -0.0794 to $+0.0099$, and every pointwise change interval includes zero (Supplementary
803 Table 12). The large negative change is concentrated in the ToN-IoT campaign shift,
804 where removing direct traffic descriptors reduces both families' absolute accuracy but
805 the selected fidelity kernel more strongly. This sensitivity neither restores a positive
806 aggregate quantum effect nor establishes that all shortcuts have been removed.

807 2.9 Protocol sensitivity is prospectively corroborated on 808 external domain shifts

809
810 We next applied the prospectively frozen protocol to College Scorecard, diabetes read-
811 mission, and ACS income, preserving the published TableShift domain splits and
812 withholding all OOD labels from the controlled selector. The complete external grid
813
814




Fig. 5 Circuit-aware and within-v4 fixed-case sensitivities for SVC. **a**, Quantum-minus-classical OOD balanced-accuracy effects for all quantum maps ($B = 60$), entangling-ZZ maps ($B = 30$), and separable-product maps ($B = 30$), each against a blocked equal-count extended-classical reference. Coloured points are the eight scenario-groups and black diamonds are three-source-dataset-equal means. **b**, Original and ablated controlled effects after removing the frozen port-, protocol-, and service-related fields in the six q1000 network groups. Connecting lines show paired group changes. **c**, Three-source-dataset-equal effects for the q1000 within-v4 $2 \times 2 \times 2 \times 2$ factorial. Positive values favour the fidelity-kernel family. All quantities are descriptive summaries conditional on the fixed datasets and finite candidate pools; panels do not show population uncertainty or causal attribution.

contains 30 task-size-seed units and 37,800 split-level results; every prespecified unit passed the automated author-provided audit. Figure 6 shows all task-size effects and conditional seed-cluster intervals.

875 The primary SVC result satisfies the machine-readable outcome category frozen
876 as `replicated_protocol_sensitivity`; narratively, protocol sensitivity is prospec-
877 tively corroborated across three fixed external tasks. The task-equal controlled
878 effect is -0.0119 (95% conditional seed-cluster interval $[-0.0205, -0.0033]$, $n = 5$
881 deterministic subsampling seeds per task), whereas the weak-reference, fixed- C ,
882 same-test-oracle effect is $+0.0089$ ($[+0.0047, +0.0137]$). Their contraction is $+0.0208$
883 ($[+0.0113, +0.0307]$), with a leave-one-task-out range of $[+0.0116, +0.0260]$. Size-equal
886 controlled point estimates are -0.0160 for College Scorecard, -0.0066 for diabetes
887 readmission, and -0.0131 for ACS income; their pointwise intervals and both nested
888 sizes are shown in Figure 6.

891 The secondary GP classifier shows a larger contraction of $+0.0324$
892 ($[+0.0223, +0.0427]$), but its task-equal controlled effect is $+0.0077$ with an interval
893 spanning zero ($[-0.0015, +0.0183]$). College Scorecard leans quantum under this clas-
894 sifier, diabetes leans classical, and ACS is slightly positive. This classifier dependence
895 limits the conclusion appropriately: the external study corroborates the sensitivity of
896 the family verdict to evaluation design, not universal classical dominance.

903 2.10 Geometry is associated with robustness only partially

905 Effective rank and OOD kernel-target alignment are associated with OOD accuracy,
906 but the association is regime-dependent. Across scenario-groups, median Spearman
907 correlations span -0.15 to $+0.80$ for effective rank and $+0.22$ to $+0.64$ for OOD align-
908 ment; a leave-one-source-dataset-out ranking model reaches only 0.22 – 0.40 . Because
909 OOD alignment reuses the labels on which accuracy is measured, it is a post hoc diag-
910 nostic. Supplementary Figure 4 and Table 4 report the full analysis; neither descriptor
911 is interpreted as a causal mediator or deployable selector.




Fig. 6 External domain shifts separate optimistic and deployable conclusions. Effects are quantum minus classical OOD balanced accuracy for the weak-reference fixed-regularization same-test oracle, the equal-budget train-tuned same-test oracle, and the primary equal-budget train-tuned validation selector. Rows show both nested sizes for all three fixed tasks plus the task-equal mean. Error bars are 95% conditional seed-cluster intervals over five deterministic subsampling seeds per task ($n = 5$), preserving the nested-size relation. They quantify pipeline variation conditional on the fixed tasks and design, not uncertainty over a population of tasks.

2.11 Finite shots can create distinctness without useful advantage

All primary quantum kernels use exact statevectors. To test whether finite estimation can create apparent non-emulability, we apply the same zero-label certificate to 30 repeated binomial fidelity estimates at 128, 512, 2,048, and 8,192 shots for one fixed SVC quantum model per case. Each replicate reselects C from training data under unprojected, independently projected, and coherent Nyström-PSD conditions (Methods).

At 128 shots the perturbed model disagrees with its exact-statevector counterpart on a median 7.2–7.4% of target decisions, depending on PSD handling. Yet the across-case median increase in the full-family zero-label endpoint is only +0.003 for Nyström, +0.005 unprojected, and +0.009 for independent projection; individual cases range from decreases of 0.014 to increases of 0.040. The corresponding median




Fig. 7 Finite-shot perturbations can increase apparent non-emulability without increasing performance. Lines are medians across the eight fixed case-level Monte Carlo medians and ribbons span their minima and maxima. **a**, Change from the exact-statevector zero-label accuracy upper endpoint against the full classical family. **b**, Change in realized OOD balanced-accuracy advantage. Each case has 30 measurement replicates at each shot level and PSD condition. Ribbons are fixed-case ranges, not confidence intervals.

realized balanced-accuracy changes are -0.0010 , -0.0044 , and -0.0005 , respectively (Figure 7). As shots increase, median certificate changes generally approach zero, while the case ranges remain heterogeneous.

Finite-shot noise can therefore move a quantum predictor away from every exact classical witness without making it better. This is a specifically important insufficiency result for the certificate: predictive non-emulability is necessary for a material advantage against the reference family, but it is not sufficient, and measurement noise can manufacture distinctness. Effective-rank inflation and the original performance sensitivities remain in the Supplementary Information. One case-replicate requires 1,624,250 distinct fidelity estimates; the full four-shot-level sensitivity projects 4.24 trillion circuit shots before routing and device overhead, halved if the four selected product-map cases are evaluated by direct factorization instead.

3 Discussion

The central result is an exact answer to a deliberately narrower question than quantum advantage in the complexity-theoretic sense: given fixed target inputs, fixed predictions, a bounded loss, and a prespecified classical-kernel reference family, how much

predictive advantage remains compatible with the labels not yet inspected? The finite max-min gives the interval hull of the identified set after any partial audit without a sampling, calibration, prevalence, or shift assumption. No interval based only on those observables can be uniformly smaller. For zero-one accuracy, the disagreement matrix collapses that optimization to a closed form. This turns “a classical model looks similar” into a sharp and minimal finite-batch falsification statement with an explicit materiality threshold.

This scope also separates the result from certifiable-surrogate generalization bounds, which control a target model’s population risk through a surrogate with a generalization guarantee and unlabeled disagreement [33]. Our models remain fixed, the estimand is advantage against the best member of a declared family, and every unaudited label completion remains admissible. The acquisition rule is therefore secondary: it chooses which ambiguity to remove, while the certificate is pathwise exact after any realized audit subset. In QML, this target-decision object complements Huang’s train-sample geometric sieve rather than replacing it [12].

The empirical frontier is much sharper than the benchmark effect alone. Against 115 classical kernels, at most 33 of 500 labels under the fixed adaptive rule are required to rule out an accuracy advantage above one point in every retrospective SVC and GPC case. Strong prediction-aware controls are at least as efficient on average, with median 12.5 labels versus 13, whereas an exact retrospective oracle needs median 6.5. This negative algorithmic result matters: most of the gain comes from adjudicating actionable disagreements, not from a privileged acquisition heuristic. The same certificate contraction holds for the selected entangling-ZZ subset. In the prospectively specified replication, all four eligible task-classifier medians require at most three labels and the overall seed-level median is zero; the controls again show no acquisition-heuristic advantage. It does not establish universal classical emulation: it identifies the quantum model’s decisions relative to the prespecified classical-kernel reference family on these

1059 particular target covariates. Its value is that every remaining ambiguity is visible in
 1060 the certificate rather than hidden inside an average performance comparison.
 1061
 1062 The three-gate interpretation prevents stronger conclusions from being conflated.
 1063
 1064 Protocol validity asks whether the reported comparison could be deployed; predictive
 1065 indispensability asks whether a witness in the prespecified classical-kernel reference
 1066 family reproduces the fixed target decisions; quantum relevance asks whether any sur-
 1067 viving behaviour depends on a non-factorizable circuit and remains meaningful after
 1068 accounting for estimation cost. Passing all three gates creates a candidate for deeper
 1069 advantage analysis, not a proof of speedup. Here Gate 1 reverses the optimistic sign,
 1070 Gate 2 rules out a material batch advantage with limited auditing, and prospective cor-
 1071 roboration in two technically eligible tasks meets the predefined criteria. Gate 3 finds
 1072 no rescue from the retrospective entangling stratum while finite-shot noise sometimes
 1073 increases distinctness without performance. Because all prospective winners are prod-
 1074 uct maps, the prospective experiment corroborates Gate 2 but adds no independent
 1075 evidence about Gate 3.
 1076
 1077 From the perspective of distribution-shift research, this reinforces that robust-
 1078 ness conclusions are conditional on the protocol [8–10, 34]. Within v4, moving from
 1079 the fixed- C , same-test-oracle, customary-reference, native-budget corner to the fully
 1080 controlled corner changes the q1000 estimate from +0.0131 to -0.0047 . Train-CV reg-
 1081 ularization and the extended reference have the largest mean paired changes, but their
 1082 sizeable interaction precludes an additive or causal reading. The legacy/v4 display
 1083 still cannot isolate regularization, and the larger historical +0.037 endpoint remains
 1084 interpretable only as a complete evaluation bundle.
 1085
 1086 The external analysis provides prospective corroboration beyond the constructed
 1087 security shifts. Under a protocol frozen before result inspection, three education,
 1088 health, and socioeconomic tasks corroborate a positive protocol contraction and a
 1089
 1090
 1091
 1092
 1093
 1094
 1095
 1096
 1097
 1098
 1099
 1100
 1101
 1102
 1103
 1104

negative controlled SVC effect. The result is stable to omitting any one task, yet heterogeneous by task and classifier: the GP endpoint retains a small positive task-equal lean whose interval crosses zero. The strongest defensible conclusion is therefore about evaluation sensitivity, not the universal ordering of kernel families.

Length scale is decisive for both classical and quantum kernel behaviour [35, 36]. With symmetric scale freedom and tuned regularization, the families occupy overlapping effective-rank and alignment regimes. These quantities are kernel descriptors, not unique signatures of quantum provenance; OOD alignment additionally reuses target labels and remains strictly post hoc.

The high-rank finding must also be reconciled with two spectral results that superficially point the other way. Kübler et al. argue that quantum kernels generalize in-distribution only when the relevant spectral mass is concentrated in few components [37], and Thanasilp et al. show that fidelity kernels concentrate exponentially—toward effectively rank-one, uninformative Gram matrices—as qubit counts grow [38]. There is no contradiction, but there is a genuine tension that our results locate precisely. Both results concern the asymptotic or i.i.d. regime; our effective ranks (roughly 10–150 on training sets of 10^3 – 4×10^3 samples at 4–12 qubits) sit far below the concentration regime. The repeated finite-shot sensitivity adds an empirical caveat these theoretical results anticipate: effective rank from sampled fidelities is inflated by a median factor of 1.93 at 128 shots and 1.13 at 8,192 shots (Section 2.11). Exact effective rank is therefore a statevector idealization even at these modest qubit counts.

From a practical perspective, no consistent fidelity-kernel advantage is observed over the tuned classical family, while overlap estimation introduces circuit and shot costs absent from direct classical-kernel evaluation. A shortened-length-scale Laplacian is therefore an important inexpensive reference, not a universally superior choice. More generally, comparisons should keep representation and learner fixed, tune all regularization on training data, match search budgets, select without target labels, include a

1151 non-constructed shift where possible, and report correlated benchmarks as fixed cases.
 1152 Geometry descriptors remain useful diagnostics, but not deployable selectors unless
 1153 validated without target-label access.
 1154
 1155 Several limitations define the scope of our conclusions.
 1156
 1157 Our shift operationalization covers three reproducible mechanisms—within-class
 1158 sparsity-tail extremes, train-geometry-tail extremes, and a held-out attack-campaign
 1159 shift on network traffic (an unseen attack campaign plus a capture-partition change,
 1160 not a timestamped temporal split)—but not the full space of real-world drift processes.
 1161 Explicit temporal splits on timestamped malware, family-based drift, and adversarially
 1162 induced shift remain outside the study’s scope, and our conclusions should be re-tested
 1163 under them rather than assumed.
 1164
 1165 Kernel and encoding coverage also remains finite. The classical family spans seven
 1166 kernels plus a bandwidth sweep, and both families receive symmetric tuning freedom,
 1167 but neither candidate set is exhaustive: the quantum side is restricted to four com-
 1168 monly used fidelity feature maps at low qubit counts, and trained or task-adapted
 1169 encodings, projected kernels, and deeper maps are not evaluated. The observed asso-
 1170 ciation suggests that future extensions should report their Gram-matrix geometry as
 1171 well as provenance, but it does not predict their outcome. At substantially larger
 1172 qubit counts, exponential kernel concentration [38] creates an additional regime not
 1173 represented here, so our conclusions should not be extrapolated beyond the low-qubit
 1174 study.
 1175
 1176 The certificate is relative to a finite, prespecified classical family and a fixed
 1177 observed target batch. It does not rule out a quantum advantage against a differ-
 1178 ent classical family, establish a surrogate outside the observed covariates, or convert
 1179 candidate count into computational complexity. Likewise, the batch interval is not
 1180 a confidence interval for a future deployment population. The optional population
 1181
 1182
 1183
 1184
 1185
 1186
 1187
 1188
 1189
 1190
 1191
 1192
 1193
 1194
 1195
 1196

result requires i.i.d. target sampling and models fixed before observing that sample,
assumptions that do not describe the constructed tail batches.

The partial-label result is mathematically valid for any audit subset, but the security acquisition curve is retrospective. In the separate prospective Gate-2 experiment against the prespecified classical-kernel reference family, tasks, prediction locks, policies, thresholds, and numerical criteria were fixed and hashed before label opening. The two technically eligible tasks met those predefined criteria. Its scope remains narrow: NHANES failed a prespecified feature gate, the seeds are deterministic subsamples rather than independent task replicates, and every prospective quantum winner is a product map. The result corroborates the certificate and actionable-disagreement principle on these fixed tasks, not a population claim, entangling-kernel replication, or superiority of bottleneck coverage. Balanced-accuracy certification additionally requires a prespecified prevalence or prevalence range; using the retrospectively observed prevalence is only a post-unlock sensitivity.

The quantum kernels are simulated rather than executed on hardware. The study therefore characterizes fidelity-kernel geometry, not noisy-device behaviour. Repeated binomial sampling quantifies one component of measurement uncertainty, but hardware noise, transpilation, readout, mitigation, and circuit-dependent correlations remain unmodelled [39]. Finite-shot cost is not merely an implementation detail: Gentinetta et al. place kernel-estimation uncertainty inside the end-to-end complexity of quantum support-vector machines [40], while exponential concentration can make informative spectral structure expensive to resolve at larger scales [38]. Recent QML evidence is not uniformly negative. Agliardi et al. report a favorable 10-qubit hardware instance on real data (87% accuracy versus 80% for their classical benchmark) and comparable utility-scale results on 156-feature synthetic data (80% versus 83%) [41]. Those study-specific, single-instance comparisons motivate continued hardware evaluation but do not alter our fixed-batch result or validate advantage under our

1243 shift protocol. We make no claim about hardware feasibility, resource advantage, or
1244 quantum speedup—the rigorously established quantum-classical learning separation
1245 concerns a specifically constructed problem, not natural data [42]. At the dimensions
1246 studied, the fidelity kernels are also classically simulable [13, 43], which is precisely
1247 what makes them available as practical kernel constructions today.

1251 Dataset scope and feature caveats provide a final boundary. Dynamic malware anal-
1252 ysis and encrypted traffic remain untested. The network exports retain port-derived
1253 counters and, for ToN-IoT, direct port, protocol, and service indicators that can sup-
1254 port shortcuts in NIDS benchmarks [44]. Sharing features controls the inputs but
1255 does not imply that distinct kernels exploit them equally. In the frozen q1000 abla-
1256 tion, removing these fields shifts the two-source-dataset-equal effect from -0.0031 to
1257 -0.0245 , but the change is highly heterogeneous and concentrated in ToN-IoT. This
1258 is evidence that feature access can affect the relative comparison, not proof that every
1259 shortcut has been removed or that the ablated tasks define a preferred deployment
1260 benchmark.

1268 We do not claim universal classical superiority, predictive equivalence, complexity-
1269 theoretic advantage, hardware relevance, or a refutation of quantum kernels. The
1270 proposed audit instead makes a candidate advantage progressively falsifiable: first by
1271 protocol controls, then by a prespecified classical witness family on target covari-
1272 ates, and finally by limited adjudication of the disagreements that keep a material
1273 advantage possible. Future work should broaden the prospective replication to more
1274 tasks and entangling winners, and pair candidate breadth with measured compu-
1275 tational resources. This interpretation is consistent with calls for controlled QML
1276 benchmarking [4, 6, 45].

1281
1282
1283
1284
1285
1286
1287
1288

4 Methods

4.1 Problem setting and estimands

We study binary security classification—malware versus benign software and attack versus benign traffic—under distribution shift. Let D_{ID} denote the in-distribution process and D_{OOD} a constructed stress-test or held-out-campaign process. Each run contains a training set T , an ID holdout that is divided into validation V_{ID} and test S_{ID} , and an OOD test set S_{OOD} . Preprocessing, samples, classifier implementation, tuning rule, and family-selection rule are shared; only the kernel changes. Comparisons are always within a classifier.

The primary endpoint is OOD balanced accuracy after P1' deployment selection. A family effect is

$$\Delta_{\text{OOD}} = \text{BAcc}_{\text{OOD}}^{(Q)} - \text{BAcc}_{\text{OOD}}^{(C)},$$

so positive values favour the quantum family. ID balanced accuracy and ID-to-OOD degradation are secondary. P2 selects from OOD information in other realizations, and P3 selects and evaluates on the same OOD test labels; both are explicitly non-deployable descriptive oracles.

4.2 Sharp finite-target identification

For Theorem 1, fix the prediction-dependent loss tables before the audit and write $a_{ij}(y) = \ell_{ij}^C(y) - \ell_i^Q(y)$. For any completion of the labels outside L ,

$$\Delta_\ell(y) = \frac{1}{n} \min_j \left\{ S_j(L) + \sum_{i \notin L} a_{ij}(y_i) \right\}.$$

1335 For the lower endpoint, minima commute and the label choices separate by item:

1336

1337

1338

1339

1340

1341

1342 For the upper endpoint introduce one-hot variables $z_{iy} \in \{0, 1\}$ with $\sum_{y \in \mathcal{Y}} z_{iy} = 1$

1343

1344 and maximize t subject to

1345

1346

1347

1348

1349

$$nt \leq S_j(L) + \sum_{i \notin L} \sum_{y \in \mathcal{Y}} a_{ij}(y) z_{iy} \quad \text{for every } j.$$

1350

1351 This mixed-integer linear program enumerates the finite max–min implicitly. Finite-

1352 ness of $\mathcal{Y}^{|L^c|}$ guarantees that both extrema are attained. Any interval based on the

1353

1354 same loss tables and audited labels that excluded either endpoint would fail on the

1355

1356 attaining completion, proving sharpness and minimality. Revealing a label restricts

1357 the completion set, which proves monotone contraction. The implementation is tested

1358

1359 against exhaustive enumeration for binary Brier loss, arbitrary multiclass bounded-loss

1360

1361 tables, arbitrary partial labels, and the zero–one specialization.

1362 The empirical accuracy certificate changes the estimand from a selected family

1363

1364 mean to a fixed-model comparison. Let $q_i = h_Q(x_i)$ and $c_{ij} = h_{C_j}(x_i)$ be hard predic-

1365

1366 tions on a target batch of size n . The quantum model is the P1' winner fixed using ID

1367 validation; the classical candidates, their train-only hyperparameters, and their target

1368

1369 predictions are also fixed before any target-label audit. Every displayed certificate uses

1370 the same inherited canonical realization per security group: $n_{\text{train}} = 1000$, $n_{\text{ID}} = 500$,

1371

1372 $n_{\text{OOD}} = 500$, master seed 42, q-split seed 42, and model seed 42. The exact scenario

1373

1374 run prefixes and selected SVC/GPC configurations appear in the canonical-run table

1375

1376 in Supplementary Information. For a candidate labelling y ,

1377

1378

1379

1380

$$g_j(y) = \text{Acc}(q; y) - \text{Acc}(c_j; y), \quad \Delta_B(y) = \min_j g_j(y).$$

Only disagreements contribute to g_j . If $d_j = n^{-1} \sum_i \mathbf{1}\{q_i \neq c_{ij}\}$, then $g_j(y) \in [-d_j, d_j]$. The labelling $y_i = q_i$ attains $g_j = d_j$ for every j simultaneously, so $\max_y \Delta_B(y) = \max_j d_j$. If k maximizes d_j , the labelling $y_i = c_{ik}$ attains $\Delta_B(y) \leq g_k(y) = -d_k$; no g_j can be below $-d_j \geq -d_k$, so $\min_y \Delta_B(y) = -\max_j d_j$. This proves Theorem 2 and its sharpness.

For an arbitrary audited subset L , define the observed signed contribution $S_j(L)$ and remaining disagreement count $R_j(L)$ as in Theorem 3. Every completion satisfies

$$\frac{S_j - R_j}{n} \leq g_j(y) \leq \frac{S_j + R_j}{n}.$$

Assigning every unaudited label to q_i attains all witness upper endpoints simultaneously. Assigning them to the witness that minimizes $S_j - R_j$ attains the lower endpoint of Δ_B . This proves Theorem 3. In binary classification, if D_j is the original disagreement count and the audit has revealed a_j disagreements favouring q and b_j favouring c_j , the interval simplifies to

$$\left[\min_j \frac{-D_j + 2a_j}{n}, \min_j \frac{D_j - 2b_j}{n} \right].$$

Thus the upper endpoint cannot increase under arbitrary adaptive querying. A classical-favouring label on a binary disagreement reduces that witness endpoint by exactly $2/n$; a quantum-favouring label leaves its upper endpoint unchanged. These statements are deterministic for the observed batch and require neither i.i.d. sampling nor a model for the unaudited labels.

For completeness, an optional population form is available when the target inputs are i.i.d. and every model is fixed before observing them. A union bound and

1427 Hoeffding’s inequality give, with probability at least $1 - \delta$,

1428

1429

1430

1431

1432

1433

1434 We do not apply this result to the constructed tail batches, for which the finite-batch
1435 certificate is the appropriate estimand.

1436

1437

1438 **4.3 Reference tiers and active target-label audit**

1439

1440 The partial-label audits use two nested operational tiers. The customary tier con-
1441 tains the 30 linear and RBF candidates; the full tier contains all 115 prespecified
1442 classical candidates. Separately, the frozen zero-label breadth sensitivity evaluates
1443 30, 60, and 115 candidates using 5,000 nested whole-block orderings rooted at
1444 `ksf-v9-frontier-20260731`; a 60-candidate draw contains 12 of 23 five-dimension
1445 blocks and the full tier contains all blocks. Tier size measures reference breadth
1446 only. It is not a computational-resource claim because training time, memory, and
1447 model-construction cost are not equalized by candidate count.

1448

1449 At every audit step, we calculate the witness-specific numerator $S_j + R_j$. Among
1450 witnesses with unaudited disagreements, the adaptive coverage rule identifies those
1451 tying at the smallest reducible numerator and queries an unaudited input that cov-
1452 ers the largest number of these bottleneck witnesses. Ties follow a fixed SHA-256
1453 order. The rule uses target covariates, predictions, and previously revealed out-
1454 comes, but never an unrevealed label. It is evaluated retrospectively for thresholds
1455 $\tau \in \{0.005, 0.010, 0.020\}$.

1456

1457 We compare it with three stronger references. Dynamic random-active sampling
1458 selects by SHA-256 order among all unaudited points disagreeing with at least one
1459 current reducible bottleneck witness. Non-adaptive initial coverage ranks points once
1460 by how many zero-label bottleneck witnesses they cover, without label feedback. Both

1461

1462

use 200 frozen SHA-256 tie orders. Finally, an undeployable label oracle gives the exact
ex-post minimum number of queries over every possible audit subset. If D_j is witness
 j 's total disagreement count, crossing threshold τ requires

$$r_j(\tau) = \max\left\{0, \left\lceil \frac{D_j - n\tau}{2} \right\rceil\right\}$$

revealed disagreements whose realized label favours witness j ; the oracle minimum is
the smallest feasible $r_j(\tau)$ over witnesses. This is exact because the family endpoint
crosses when any one witness endpoint crosses. A prediction-independent complete
ordering is retained only as a weak diagnostic. Comparator quantiles describe frozen
algorithmic orders, not sampling uncertainty.

The zero-label SVC definition and go/no-go criterion were frozen before pre-
diction outcomes were inspected. The arbitrary-partial-label acquisition extension
was declared only after the target labels had been unlocked; the strong-comparator
amendment was declared after the original acquisition curves were known. Both are
exploratory on these eight cases. GPC repeats the mathematics and policies but is
not an independent prospective validation. Target labels are used only to evaluate
certificate contraction, the oracle, and the final realized effect; no kernel, classifier,
regularization, threshold, or P1' winner is refitted or reselected. Prediction archives
and the separately stored label file are joined only after SHA-256, row-index, model,
and target-length gates pass.

4.4 Prospectively specified Gate-2 replication against the prespecified classical-kernel reference family

The Gate-2 replication against the prespecified classical-kernel reference fam-
ily was specified and hashed on 1 August 2026 before any of its three tasks
was acquired, any model prediction was computed, or any target label was

1519 inspected. The source is the same TableShift implementation pinned at com-
 1520 mit fca9429814703a07e3902d005d46563a207b7f0a [46]. The task identifiers were
 1521 BRFSS Diabetes (race shift), ACS Food Stamps (Census-division shift), and NHANES
 1522 Lead (poverty shift). All three were attempted and none could be replaced based on
 1523 availability or outcome.
 1524
 1525 For each task, seeds $\{42, 123, 999, 7, 2024\}$ define deterministic label-blind sam-
 1526 ples with $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{target}}) = (1000, 250, 250, 500)$. The v5 train-only semantic
 1527 encoding, scaling, TruncatedSVD dimensions $\{4, 6, 8, 10, 12\}$, and angular representa-
 1528 tion are reused verbatim. The 60 quantum candidates and 115 classical candidates,
 1529 P1' validation selection, candidate-specific train-CV for SVC, fixed Laplace GPC, and
 1530 decision thresholds are unchanged. The customary 30-candidate tier is secondary; the
 1531 full tier and $\tau = 0.010$ define the prospective outcome. Secondary thresholds are 0.005
 1532 and 0.020.
 1533
 1534 The pipeline is fail-closed. A prediction-lock stage copies target labels to a separate
 1535 sealed tree, removes them from the model state, and writes hard target predictions,
 1536 row-order digests, selected validation-only quantum configurations, and SHA-256
 1537 hashes without calculating or printing target performance, prevalence, disagreement,
 1538 or an acquisition curve. Candidate checkpoints contain only target predictions and
 1539 validation-selection quantities. An aggregate manifest is accepted only when every
 1540 available task-seed unit has all 175 kernel configurations for both classifiers or a
 1541 permitted pre-model technical-failure record. The audit stage verifies that aggregate
 1542 hash, records the one-time label opening and analysis-code hash, then joins labels by
 1543 source-position digest. No model is refitted after the join.
 1544
 1545 For each classifier and the two classical tiers, the fixed policies are adaptive bottle-
 1546 neck coverage, 200 dynamic random-active SHA-256 orders, 200 non-adaptive initial-
 1547 coverage orders, the exact ex-post label oracle, and a weak prediction-independent
 1548 order. The hash root appends task, seed, model, tier, policy, and draw. The primary
 1549

summaries are four or six task–classifier medians over five seeds, depending on technical availability. The predefined primary criterion requires every available task–classifier median to be at most 50 labels, the overall seed-level median to be at most 25, and no seed to exceed 100. With exactly two available tasks, all four cells are required and the result is explicitly technically limited. A secondary criterion requires all four medians to be at most 50 and an overall median at most 50; fewer than two available tasks is technically incomplete. The criteria were evaluated exactly once.

NHANES generated 11 non-constant encoded features in all five seeds, fewer than required for the frozen 12-dimensional candidate grid. Minimum-feature failure was an explicitly permitted technical-unavailability gate. No NHANES model candidate was executed, the task was not replaced, and dimensions were not changed. BRFSS Diabetes and ACS Food Stamps passed every gate, leaving ten task–seed units for the technically limited prospective outcome.

4.5 Balanced-accuracy identification

Balanced accuracy is not label-free unless target prevalence is known or restricted. At a fixed positive count n_+ , we group inputs by their joint prediction signature $z = (q_i, c_{i1}, \dots, c_{iM})$. Let n_z be the size of group z and $k_z \in \{0, \dots, n_z\}$ its unknown number of positive labels, with $\sum_z k_z = n_+$. For classifier h with group-level prediction $h_z \in \{0, 1\}$,

$$\text{BAcc}_h(k) = \frac{1}{2} \left[\frac{\sum_z k_z h_z}{n_+} + \frac{\sum_z (n_z - k_z)(1 - h_z)}{n - n_+} \right].$$

The sharp upper endpoint is the mixed-integer linear program that maximizes t subject to

$$t \leq \text{BAcc}_Q(k) - \text{BAcc}_{C_j}(k) \quad \text{for every } j,$$

together with the integer and prevalence constraints above. The lower endpoint is the minimum over witnesses of their corresponding linear minima. Relaxing k_z to a continuous interval gives the sharp distributional or fractional-label LP. When prevalence

is known only to lie in an interval, we maximize over its admissible integer counts. We report this as a sensitivity because the natural target prevalence is unavailable without labels; accuracy remains the main assumption-free certificate.

4.6 Why train-only geometry cannot determine target prediction

The geometric difference and effective rank summarize training Gram structure, whereas deployment also depends on the target–train block. An explicit normalized PSD construction makes the distinction exact. Set $K_{TT} = I_n$. One valid joint Gram extension gives a unit-norm target feature orthogonal to all training features, hence $K_{UT} = 0$. A second gives the target the same feature as the first training point, hence $K_{UT} = e_1^\top$. Both joint matrices are PSD, have unit diagonal, and share identical K_{TT} . For kernel ridge regularization $\lambda > 0$, however, their target prediction operators are

$$0 \quad \text{and} \quad e_1^\top (I_n + n\lambda I_n)^{-1} = \frac{e_1^\top}{1 + n\lambda},$$

whose spectral-norm difference is $1/(1 + n\lambda)$. Train-only geometry can therefore screen possible separation but cannot characterize target transport. We use the hard-prediction certificate for SVC and GPC rather than treating this KRR construction as an endpoint result.

4.7 Security datasets and shared split sizes

The source datasets are EMBER [47], UNSW-NB15 [48], and ToN-IoT [49]. Four binary scenarios—EMBER malware, UNSW-NB15 DoS, UNSW-NB15 Reconnaissance, and ToN-IoT Scanning—produce two shift mechanisms each and hence eight scenario-groups. Classical and quantum kernels always receive identical row indices and class balance.

The 72 principal settings cross eight groups with master seeds $\{42, 123, 999\}$ and $(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$ equal to $(1000, 500, 500)$, $(2000, 1000, 1000)$, or $(4000, 1800, 1800)$. Each setting is instantiated with q-split seeds $\{42, 123, 999, 7, 2024\}$ and model seeds $\{42, 123, 999\}$, yielding fifteen pipeline realizations arranged in five q-split clusters. Indices are stored explicitly; automated gates verify split-role disjointness, required cardinalities, both classes, identical family inputs, and absence of label-derived features.

4.8 EMBER shift construction

EMBER uses the histogram and byte-entropy feature blocks and two label-conditional tail stress tests. They are not described as certified covariate shifts because conditioning the split on class does not establish invariance of $P(Y | X)$. For $m1$, the sparsity score is the number of non-zero entries in the selected feature block. Within each class, low- and high-score extremes supply the exact-size OOD set; training and ID samples are stratified draws from the remainder.

For $m2$, the final training set is fixed first. A MaxAbs scaler and TruncatedSVD representation are fitted on training data only, and a centroid is computed for each class. Every non-training sample is scored by Euclidean distance to its own-class centroid. Within-class low- and high-distance extremes form S_{OOD} , and the ID holdout is sampled from the remaining pool. This construction probes unusual geometry relative to the fixed training set without using test information to fit the representation.

4.9 Network-flow shift construction

Numeric preparation removes identifiers and explicit label derivatives, leaving UNSW-NB15 and 88 ToN-IoT features. The UNSW-NB15 export excludes direct categorical protocol, service, and state fields but retains three port-derived counters; ToN-IoT retains two port fields plus protocol and service indicators. A leakage audit records every exclusion and feature-label association; the largest observed absolute

correlation is 0.62 and belongs to a retained traffic feature. Each scenario has reference and current pools and uses two mechanisms.

The *m2-centroid* mechanism is the network analogue of EMBER *m2*: train-only MaxAbs+TruncatedSVD class centroids define within-class distance tails, which sup- ply the OOD set. The *attack-campaign* mechanism performs no geometric selection. Training and ID pools contain benign traffic and all non-target attack categories; OOD contains benign traffic and the held-out target campaign. UNSW-NB15 refer- ence and OOD pools additionally use different official capture partitions. This is an unseen-campaign plus capture-partition shift, not a timestamped temporal split.

4.10 Shared preprocessing and representations

For every run and dimension $d \in \{4, 6, 8, 10, 12\}$, a MaxAbs scaler and TruncatedSVD are fitted on T ; the reduced coordinates are standardized and mapped linearly with training-fitted parameters so that the *training* coordinates lie in $[0, \pi]$. Held-out coordinates receive the same affine map and may extrapolate outside that interval under shift. The same embedded samples are supplied to both families, making the comparison a kernel swap rather than a representation swap. Dimension is a candidate axis. No validation, ID-test, or OOD-test row fits a transformation. The SVD random state equals the model seed.

4.11 Classifiers and train-only regularization

The primary learner is an SVC with a precomputed Gram matrix and balanced class weights [50]. For every kernel geometry, $C \in \{0.01, 0.1, 1, 10, 100\}$ is selected by five-fold stratified cross-validation on the training Gram matrix; ties favour the smaller C . When $n_{\text{train}} > 2000$, this selection step uses a class-stratified random-state-20260717 subsample capped at 2,000 rows, while the final model is always refitted on the complete training block. Because C is resolved before family selection, its grid points are not counted as separate family candidates.

The secondary learner is a binary Gaussian-process classifier with a logistic link and Laplace approximation [51]. It consumes the same precomputed kernel, uses diagonal jitter 10^{-8} , at most 100 Newton iterations, and convergence tolerance 10^{-6} . The logistic-probit approximation supplies predictive probabilities; no post-hoc calibration is applied.

4.12 Classical kernel family

The customary reference contains a cosine-normalized linear kernel and RBF kernel. The extended family adds cosine-normalized degree-2 and degree-3 polynomial kernels with $\gamma = 1/d$ and intercept 1, the Laplacian kernel $\exp(-\|x-x'\|_1/\ell_1)$, and Matérn-3/2 and Matérn-5/2 kernels. The base ℓ_1 and ℓ_2 scales are median pairwise Manhattan and Euclidean distances, respectively, calculated from at most 2,000 training rows with the model seed; the RBF base scale is fitted from training variance. RBF, Laplacian, and both Matérn kernels use factors $\{0.1, 0.3, 1, 3, 10\}$ around these training-only bases. Deduplication yields 23 kernel-shape/scale blocks crossed with five dimensions, or 115 extended-classical candidates in full-coverage groups.

4.13 Quantum fidelity-kernel family

The quantum family uses exact statevector fidelities

$$K_Q(x, x') = |\langle \phi_Q(x) | \phi_Q(x') \rangle|^2$$

[1, 2]. Qiskit 2.3.1 implements four map blocks [52]: fully connected ZZ maps with one and two repetitions; a one-repetition Pauli map with single-qubit strings X and Z; and a two-repetition Z map. Only the ZZ maps contain two-qubit interactions. The latter

1795 two prepare product states,

1796

1797

1798

1799

1800

1801

1802 so the constructor argument `entanglement="full"` stored for the Pauli-X/Z map has

1803

1804 no effect without a multiqubit Pauli string. The product can be evaluated classically

1805 in $O(d)$ one-qubit-overlap calculations. We call these the entangling-ZZ and separable-

1806

1807 product strata. Each map uses angle scales $\{0.5, 1, 2\}$ and all five dimensions, giving

1808

1809 60 candidates (30 per stratum) at 4–12 simulated qubits. Classical bandwidth and

1810 quantum angle-scale freedom both control locality [35, 36]; no circuit is trained.

1811

1812

1813 **4.14 Circuit and sampling-resource audit**

1814

1815 Circuit counts use Qiskit 2.3.1 and the logical basis $\{\mathbf{rz}, \mathbf{sx}, \mathbf{x}, \mathbf{cx}\}$ at transpiler

1816

1817 optimization level 0, without backend layout or routing. We report each feature-map

1818

1819 template and the compute–uncompute fidelity template formed by composing the map

1820

1821 at x with the inverse map at x' after assigning distinct numeric inputs. Counts are

1822 compiler- and basis-conditional logical resources, not device schedules.

1823

1824 For the repeated q1000 finite-shot sensitivity, each case and measurement replicate

1825

1826 estimates the off-diagonal training triangle, the combined 500-row ID-holdout-to-

1827 training rectangle (partitioned downstream into validation and test roles), the

1828

1829 OOD-test-to-training rectangle, and the off-diagonal OOD square triangle; diagonals

1830 are fixed to one. Thus one kernel configuration requires

1831

1832

1833

1834

1835

1836

1837

1838

1839

1840

$$\binom{1000}{2} + 500(1000) + 500(1000) + \binom{500}{2} = 1,624,250$$

distinct fidelity estimates per shot level and replicate. Multiplying this count by shots
gives the idealized circuit-shot total. The three PSD conditions reuse the same sam-
pled entries and add only classical post-processing. Routing, calibration, mitigation,
queueing, and device noise are excluded.

4.15 Target-label-free family selection

The original ID holdout is divided deterministically with class stratification into
 V_{ID} and S_{ID} . Within each class, global row indices are ordered by SHA-256 of the
salt `ksf-v4-idsplit::` and index, then assigned alternately to validation and test.
The halves are disjoint and approximately class-balanced. P1' deploys, within each
family, the candidate with highest V_{ID} balanced accuracy and reports it once on
 S_{ID} and S_{OOD} . Ties follow lexicographic candidate order. OOD labels never enter
regularization, preprocessing, or P1' selection.

4.16 Candidate-budget matching

Best-of-family selection rewards a larger pool. The confirmatory coverage audit found
115 extended-classical and 60 quantum candidates in every one of the eight scenario-
groups and both classifiers. The primary comparison therefore draws 60 candidates
from each family in all 16 group-classifier cells. For each cell, 5,000 extended-classical
subsamples are generated with seed 20260715. The primary `kernel_blocked` scheme
samples 12 of the 23 complete kernel-shape/scale blocks, each crossed with all five
dimensions, to mirror the 12 quantum feature-map/angle-scale blocks crossed with the
same dimensions. This matches both candidate count and block structure, although it
does not make the two search spaces identical. `uniform` samples individual candidates,
while `kernel_stratified` distributes the draw across kernel-shape/scale strata. Each
subsample selects its P1' winner; the classical endpoint is the resampling expectation.
The full 115-candidate pool is secondary. Resampling percentiles describe sensitivity
to the sampled candidate set, not population uncertainty.

1887 4.17 Entangling and separable quantum strata

1888
1889 The circuit-stratum sensitivity reuses all v4 summaries and compares three SVC pools
1890 under train-CV and P1': the full 60-candidate quantum family, the 30 entangling-
1891 ZZ candidates, and the 30 separable-product candidates. The 60-candidate pool is
1892 compared with 12 complete classical blocks and each 30-candidate stratum with six
1893 complete classical blocks, always crossed with all five dimensions. Each classical end-
1894 point is the expectation over 5,000 `kernel_blocked` draws using SHA-256-derived
1895 seeds rooted at 20260729. This tests dependence on map stratum, not computational
1896 advantage.

1902 4.18 Within-v4 factorial and shortcut sensitivity

1903
1904 The within-generation SVC factorial uses all 360 q1000 runs and crosses: same-test
1905 OOD versus disjoint ID-validation selection; fixed $C = 1$ versus train-CV C ; cus-
1906 tomary versus extended classical reference; and native versus equal-count budgets.
1907 Fixed- C predictions are recomputed with the deterministic v4 preprocessing, kernel
1908 constructors, samples, and split roles; an integrity gate requires exact candidate-split
1909 key identity with every archived v4 SVC summary and the original ID-split hash salt.
1910 With the customary 30-candidate reference, equal count samples six of twelve quan-
1911 tum blocks; with the extended reference, it samples 12 of 23 classical blocks. The
1912 same 5,000 frozen block draws are paired across selection and regularization cells. Six-
1913 teen fixed-case endpoints and their paired axis contrasts are descriptive; they do not
1914 identify causal effects.

1924 The q1000 shortcut sensitivity uses the same design on all 270 network runs
1925 after removing features before the train-only embedding. UNSW-NB15 excludes
1926 `ct_src_dport_ltm`, `ct_dst_sport_ltm`, and `is_sm_ips_ports`; ToN-IoT excludes
1927 source/destination ports, all `proto_*`, `service`, and `service_*` fields. This removes 3
1928 of 39 and 14 of 88 fields, respectively. The manuscript-facing contrast uses train-CV,
1929
1930
1931
1932

P1', and 60-candidate block matching with paired draws. It is a port/protocol/service-field ablation, not a guarantee that every shortcut has been removed.

4.19 Security fixed-case aggregation and conditional intervals

The effective replication axis within a scenario-group is q-split seed. Master seed is a design stratum rather than a replicate, some EMBER OOD pools recur across master seeds, and training sizes are nested. Audited OOD overlap is at most 1.2% for EMBER and 7.8% for network groups. Within each q-split cluster, model seeds are averaged within master seed, master seeds within size, and the three sizes with equal weight. The five resulting cluster means T_j give the group effect $\bar{T}$ and 95% conditional cluster- t interval

$$\bar{T} \pm t_{4,0.975} \frac{s_T}{\sqrt{5}}.$$

This few-cluster t construction follows the cluster-mean approach of Ibragimov and Müller [53]. The interval ($n = 5$ q-split clusters) is conditional on the fixed benchmark pools and experimental design; inclusion of zero is not evidence of equivalence. A 9,999-draw BCa bootstrap over cluster means and leave-one-q-split calculations are secondary diagnostics, reported for all 16 primary cells in Supplementary Table 6. The five cluster values behind every interval are exposed in Supplementary Table 7.

Across groups, EMBER $m1/m2$ are averaged within source dataset, the four UNSW-NB15 groups within UNSW-NB15, and the two ToN-IoT groups within ToN-IoT. These three source-dataset means receive equal weight. Leave-one-source-dataset-out ranges omit one of the three at a time. The eight scenario-groups and three datasets were not sampled from defined populations; therefore no population p -value or dataset-population confidence interval is reported.

1979 4.20 Specification curve

1980

1981 The S1–S10 display is retained as a historical cross-generation sensitivity map over
1982 complete legacy and v4 endpoints. S1–S4 use fixed- C legacy generation and S5–S10
1983 train-CV v4 generation, so their boundary cannot isolate regularization. The rows
1984 remain in their prespecified order and are never sorted by effect. Individual-axis inter-
1985 pretation instead uses the within-v4 factorial above. All summaries average realizations
1986 within setting and group, then groups within the three source datasets, and finally
1987 source datasets with equal weight; no specification-level hypothesis test is performed.
1988

1993

1994 4.21 Prospectively frozen external corroboration

1995

1996 Before acquiring TableShift data or inspecting any external result, we committed
1997 the tasks, sampling, preprocessing, candidate pools, estimands, uncertainty proce-
1998 dure, and outcome-category rule. The source implementation is pinned to commit
1999 `fca9429814703a07e3902d005d46563a207b7f0a` [46]. Three public tasks span distinct
2000 domains and published shift variables: College Scorecard (institution type), diabetes
2001 readmission (admission source), and ACS income (geographic region). TableShift’s
2002 `train`, `validation`, `id_test`, and `ood_test` roles are preserved; `ood_validation` is
2003 never loaded. College Scorecard is read from its public CC0 archive because the
2004 pinned downloader invokes authenticated Kaggle access; the pinned parser, schema,
2005 and domain split are unchanged. Source-file SHA-256 hashes and the container digest
2006 are recorded.

2014

2015 Each task has nested strata $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{OOD}}) = (1000, 250, 250, 500)$ and
2016 $(2000, 500, 500, 1000)$. For seeds $\{42, 123, 999, 7, 2024\}$, source rows are ordered by
2017 SHA-256 of task, split, seed, and original position, giving deterministic label-blind
2018 samples. Gates verify source-role disjointness, size nesting, both classes, source-schema
2019 preservation, and at least 12 non-constant encoded features. All 30 task–size–seed
2020 units passed.

2024

All transformations fit `train` only. Constant columns are removed; numeric variables receive median imputation; Boolean, categorical, and string variables receive most-frequent imputation and one-hot encoding with unknown categories ignored. MaxAbs scaling, TruncatedSVD at the five study dimensions, standardization, and a training-fitted affine map that places training coordinates in $[0, \pi]$ follow; held-out coordinates may extrapolate beyond that range. The SVD random state is 20260729. The security kernel and classifier pools are then reused.

For each task, size, seed, and classifier, the validation-selected 60-candidate quantum endpoint is compared with 5,000 `kernel_blocked` 60-candidate subsamples of the 115-candidate classical pool using seed 20260729. The primary effect is quantum minus expected classical OOD balanced accuracy. The equal-budget OOD oracle selects on `ood_test`; the weak-fixed oracle uses linear/RBF, fixes SVC at $C = 1$, and also selects on the reported OOD labels. Protocol contraction is weak-fixed-oracle effect minus controlled $P1'$ effect.

Seeds are averaged within task and size, nested sizes receive equal weight within task, and the three fixed tasks receive equal weight. Conditional 95% intervals use 9,999 resamples of the five seed clusters within task while preserving nested sizes. Leave-one-task-out ranges are descriptive. The intervals quantify deterministic-subsampling variability conditional on three fixed tasks; they do not represent a population of tasks.

4.22 Performance and uncertainty metrics

Balanced accuracy on S_{OOD} is primary. Accuracy, macro and positive-class F1, ROC AUC, and precision-recall AUC are recorded on every split. The GP analysis additionally records log loss, Brier score, expected calibration error with 15 equal-width probability bins [54], and mean predictive entropy. No threshold or calibration parameter is fitted on OOD labels.

2071 4.23 Geometry descriptors and association analyses

2072

2073 For positive-semidefinite K with eigenvalues λ_i and $p_i = \lambda_i / \sum_j \lambda_j$, effective rank is
2074

2075

2076

2077

2078

2079

$$r_{\text{eff}} = \exp\left(-\sum_i p_i \log p_i\right)$$

2080 [55]. For $y \in \{-1, +1\}^n$ and doubly centered K_c , centered kernel-target alignment is
2081

2082

2083

2084

2085

$$\text{KTA}(K, y) = \frac{\langle K_c, yy^\top \rangle_F}{n \|K_c\|_F}$$

2086

2087 [56, 57]. KTA is computed separately on training, ID, and OOD square blocks, along-
2088 side off-diagonal and cross-split similarities. OOD KTA reuses the same OOD labels as
2089 accuracy and is therefore a post hoc diagnostic, not a label-independent mechanism.
2090

2091

2092 Within each run–dimension cell, Spearman correlations relate effective rank, OOD
2093 KTA, and KTA survival (OOD minus ID) to SVC OOD balanced accuracy for all,
2094

2095 classical-only, and quantum-only candidates. Partial Spearman association residualizes
2096

2097 ranked variables against family and log length/angle scale. For portability, a descrip-
2098 tive model bins log effective rank into ten training-dataset quantiles, learns mean
2099

2100 OOD accuracy per bin from two source datasets, and ranks cells in the third; held-out
2101 Spearman correlation is reported once per source dataset.

2103

2104 Rank matching is observational. Within each run and dimension, every quantum
2105 candidate is first paired to the classical candidate with smallest absolute effective-rank
2106 difference; classical candidates may be reused. Rank ratio $\max(r_Q, r_C) / \min(r_Q, r_C)$ is
2107 filtered at prespecified calipers 1.10, 1.25, 1.50, and 2.00, plus unfiltered reporting. A
2108

2109 sensitivity uses minimum-cost one-to-one assignment with absolute log-rank difference,
2110 leaving excess candidates unmatched. Counts, retained fractions, rank discrepancies,
2111

2112 mean and median OOD differences, and the fraction favouring quantum are reported
2113

2114 for every caliper. None of these analyses identifies a causal geometry or family effect.
2115
2116

4.24 Repeated finite-shot perturbation

The finite-shot analysis and its out-of-sample PSD extension were each specified before their outcomes were computed. They fix one $q1000$, master-seed-42, q-split-42, model-seed-42 run per security group and, within each, the exact-statevector quantum SVC candidate selected by $P1'$ on V_{ID} . Lexicographic candidate order resolves ties. Input `summary_v4.csv` SHA-256 hashes, chosen kernels, dimensions, and exact C values are stored in the frozen specifications.

For shots $s \in \{128, 512, 2048, 8192\}$ and measurement replicates $r = 0, \dots, 29$, each exact fidelity p is sampled independently as $\text{Binomial}(s, p)/s$. A separate NumPy generator is used for train, ID-validation, ID-test, OOD-test, and OOD-square blocks. Its unsigned 64-bit seed is the first eight SHA-256 bytes of

```
ksf-v6-shots::::<kernel>::::<shots>::::<block>.
```

Python's process-randomized hash is not used.

Square train and OOD blocks are sampled on the upper triangle, mirrored, clipped to $[0, 1]$, and assigned unit diagonal; rectangular blocks receive binomial sampling and clipping. Three prespecified conditions use the same sampled blocks. The unprojected condition makes no PSD correction. The independent-square heuristic clips negative eigenvalues of the train and OOD square matrices separately while leaving every rectangular evaluation–train block unchanged; it is an algorithmic intervention for a precomputed SVC, not a coherent correction of one out-of-sample kernel.

The coherent Nyström condition eigendecomposes the sampled training matrix as $K_T = U \text{diag}(\lambda) U^\top$ and retains $\lambda_i > \tau$, where $\tau = \max\{10^{-12}, \epsilon n \max(1, \max_i |\lambda_i|)\}$. For retained U_+ and λ_+ ,

$$K_T^+ = U_+ \text{diag}(\lambda_+) U_+^\top, \quad \Phi_E = K_{E,T} U_+ \text{diag}(\lambda_+)^{-1/2}.$$

2163 The corrected evaluation–train and evaluation-square blocks are $\Phi_E \Phi_T^\top$ and $\Phi_E \Phi_E^\top$,
 2164 respectively, with $\Phi_T = U_+ \text{diag}(\lambda_+)^{1/2}$. Thus the OOD square used for KTA is
 2165 derived from the sampled OOD–train block rather than independently projected. In
 2166 all three conditions, C is reselected by the same five-fold train-only procedure and
 2167 the final SVC is refitted. Outputs include signed and absolute OOD change from the
 2171 fixed exact endpoint, effective-rank ratio, OOD-KTA change, retained rank, eigen-
 2172 value tolerance, negative-eigenvalue fraction, and Frobenius changes for every block.
 2173 Within each fixed Gram matrix, medians and central 2.5–97.5% Monte Carlo inter-
 2174 vals summarize 30 replicates. Across the eight fixed matrices, medians and ranges are
 2175 descriptive.

2179 For the v0.9 extension we export the hard target prediction of every perturbed
 2180 SVC and compare it with the unchanged exact-statevector classical witnesses. For each
 2181 replicate and classical tier we record self-disagreement from the exact quantum predic-
 2182 tor, the zero-label interval, and realized accuracy and balanced-accuracy advantages
 2183 after the integrity-gated label join. This deliberately asks whether quantum measure-
 2184 ment perturbation creates behaviour absent from the existing classical reference; it
 2185 does not give the classical family an additional noise or retuning budget. The model
 2186 contains no device or correlated-noise process.

2193 2194 **4.25 Statistics and reproducibility**

2196 No null-hypothesis test is used for the eight security groups, three source datasets,
 2197 or three external tasks. Every uncertainty object names its effective n , resampling
 2198 unit, and conditioning set: security intervals use five q-split clusters; external intervals
 2199 use five deterministic seed clusters per task; finite-shot intervals use 30 measurement
 2200 replicates conditional on one fixed Gram matrix. The fifteen security pipeline realiza-
 2201 tions are never treated as independent. Security intervals use the t_4 reference and are
 2202 pointwise, not simultaneous; the five underlying cluster means are published. Inclusion
 2203
 2204
 2205
 2206
 2207
 2208

of zero is never interpreted as equivalence. No multiplicity adjustment was prespec- 2209
 ified. Because multiple intervals and diagnostics are displayed, they are interpreted 2210
 descriptively rather than as a familywise error-controlled decision procedure; the BCa 2211
 analysis with five clusters is retained only as a fragile sensitivity. 2212
 2213

Sample sizes were design- and compute-constrained and fixed before outcome anal- 2214
 ysis; no power calculation was used. No completed group, task, specification, budget 2215
 scheme, caliper, or finite-shot replicate was removed based on its result. Randomized 2216
 operations use recorded seeds; external sampling, ID division, and audit tie-breaking 2217
 use deterministic SHA-256 ordering. The test suite validates source-dataset mapping, 2218
 endpoint equality across analyses, candidate-budget coverage, matching uniqueness, 2219
 seed stability, exact-reference hashes, prediction/label separation, theorem endpoints 2220
 by exhaustive enumeration, MILP consistency, monotone partial-label contraction, 2221
 and complete Monte Carlo cells. 2222
 2223
 2224
 2225
 2226
 2227
 2228
 2229

4.26 Software environment 2230

Experiments use Python 3.11/3.12, NumPy 2.3.5, pandas 2.2.3, scikit-learn 1.8.0, 2231
 Matplotlib 3.11.0, and Qiskit 2.3.1. The Windows-specific Conda snapshot and cross- 2232
 platform package specifications serve different archival purposes; neither is an exact 2233
 machine image of every full-scale run. 2234
 2235
 2236
 2237
 2238
 2239

4.27 Use of generative AI 2240

OpenAI Codex was used during manuscript revision for language editing, code 2241
 inspection, analysis implementation, and consistency checking. It did not determine 2242
 authorship, study accountability, or acceptance of any scientific claim. The authors 2243
 inspected the generated changes, verified all reported analyses and citations, and take 2244
 full responsibility for the manuscript. 2245
 2246
 2247
 2248
 2249
 2250
 2251
 2252
 2253
 2254

2255 **Declarations**

2256

2257 **Data Availability.** This study uses the public EMBER [47], UNSW-NB15 [48],
2258 ToN-IoT [49], and TableShift [46] benchmarks, which remain subject to their original
2259 access conditions. Raw source records are not redistributed. Derived split definitions,
2260 deterministic row-position hashes, shift constructions, experiment manifests, com-
2261 plete aggregated results, target-prediction matrices, physically separated audit labels,
2262 evidence frontiers, technical-failure records, and integrity manifests are archived in
2263 the immutable v1.1.6 release at <https://doi.org/10.5281/zenodo.21776862>. The all-
2264 version concept DOI is <https://doi.org/10.5281/zenodo.19147649>; the exact v1.1.5
2265 predecessor is <https://doi.org/10.5281/zenodo.21764577>, the exact v1.1.4 predeces-
2266 sor is <https://doi.org/10.5281/zenodo.21759058>, and the exact v0.8.0 predecessor is
2267 <https://doi.org/10.5281/zenodo.21717074> [58].
2268

2269 **Code Availability.** Code for constructing the splits, running the kernel experi-
2270 ments, reproducing every analysis, and regenerating every table and figure is available
2271 at https://github.com/roberto-fernandez-barrios/kernel_shift_framework and in the
2272 versioned Zenodo archive above. A Windows-specific Conda snapshot and platform-
2273 neutral package specifications document the software dependencies; neither is claimed
2274 as an exact machine image of every full-scale run. The code is released under
2275 BSD-3-Clause with no restriction on non-academic use.

2276 **Competing interests.** All authors declare no financial or non-financial competing
2277 interests.

2278 **Ethics approval and consent to participate.** This work is a secondary compu-
2279 tational analysis of public benchmark releases. The authors recruited no participants,
2280 administered no intervention, collected no new human data, and accessed no direct
2281 identifiers. The cybersecurity benchmarks contain malware or network-flow records
2282 rather than human-participant observations. For the TableShift tasks, ACS PUMS
2283

records are public-use files with Census Bureau disclosure protection [59]; BRFSS survey data and documentation are publicly downloadable from CDC [60]; the NHANES source protocols are documented by the NCHS Ethics Review Board [61]; and the diabetes-readmission records are distributed through the public UCI repository [62]. College Scorecard supplies public institution-level records [63]. Raw source records are not redistributed. This study falls outside the scope of institutional ethics review because it consists exclusively of secondary analyses of publicly available, deidentified datasets, without recruitment, interaction, intervention, access to direct identifiers, or any attempt at reidentification. Accordingly, ethics approval and additional informed consent were not required. No ethics approval or exemption is claimed.

Authors' contributions. R.F.-B.: conceptualization, methodology, software, data curation, formal analysis, visualization, investigation, writing—original draft, and writing—review and editing. I.P.-L.: supervision, validation, and writing—review and editing. A.G.-S.: validation, project administration, and writing—review and editing. P.G.B.: supervision, resources, and writing—review and editing. All authors read and approved the final manuscript.

Acknowledgements. This study received no funding.

References

- [1] Havlíček, V. *et al.* Supervised learning with quantum-enhanced feature spaces. *Nature* **567**, 209–212 (2019).
- [2] Schuld, M. & Killoran, N. Quantum machine learning in feature hilbert spaces. *Physical Review Letters* **122**, 040504 (2019).
- [3] Schuld, M. Supervised quantum machine learning models are kernel methods. *arXiv abs/2101.11020* (2021). URL <https://arxiv.org/abs/2101.11020>.

- 2347 [4] Bowles, J., Ahmed, S. & Schuld, M. Better than classical? the subtle art of bench-
2348 marking quantum machine learning models. *arXiv* **abs/2403.07059** (2024).
2349 URL <https://arxiv.org/abs/2403.07059>.
2350
2351
2352 [5] Slattery, L. *et al.* Numerical evidence against advantage with quantum fidelity
2353 kernels on classical data. *Physical Review A* **107**, 062417 (2023).
2354
2355
2356 [6] Schnabel, J. & Roth, M. Quantum kernel methods under scrutiny: A benchmark-
2357 ing study. *Quantum Machine Intelligence* **7**, 58 (2025).
2358
2359
2360 [7] Shen, J. Quantum kernels and the cross-section of stock returns: Anatomy of a
2361 vanishing advantage (2026). URL <https://arxiv.org/abs/2607.20168>. Preprint,
2362 arXiv:2607.20168.
2363
2364
2365 [8] Gulrajani, I. & Lopez-Paz, D. *In search of lost domain generalization*. *Inter-*
2366 *national Conference on Learning Representations* (International Conference
2367 on Learning Representations, 2021). URL <https://openreview.net/forum?id=lQdXeXDoWtI>.
2368
2369
2370 [9] Koh, P. W. *et al.* *Wilds: A benchmark of in-the-wild distribution shifts*. *Pro-*
2371 *ceedings of the 38th International Conference on Machine Learning*, Vol. 139
2372 of *Proceedings of Machine Learning Research*, 5637–5664. International Con-
2373 ference on Machine Learning (PMLR, Cambridge, MA, USA, 2021). URL
2374 <https://proceedings.mlr.press/v139/koh21a.html>.
2375
2376
2377 [10] Yao, H. *et al.* *Wild-time: A benchmark of in-the-wild distribution shift over*
2378 *time*. *Advances in Neural Information Processing Systems* (Neural Information
2379 Processing Systems Foundation, 2022). Datasets and Benchmarks Track.
2380
2381
2382 [11] Zhou, K., Liu, Z., Qiao, Y., Xiang, T. & Loy, C. C. Domain generalization: A
2383 survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **45**,
2384
2385
2386
2387
2388
2389
2390
2391
2392

4396–4415 (2023).	2393
	2394
[12] Huang, H.-Y. <i>et al.</i> Power of data in quantum machine learning. <i>Nature Communications</i> 12 , 2631 (2021).	2395
	2396
	2397
	2398
[13] Schreiber, F. J., Eisert, J. & Meyer, J. J. Classical surrogates for quantum learning models. <i>Physical Review Letters</i> 131 , 100803 (2023).	2399
	2400
	2401
	2402
	2403
[14] Thabet, S., Monbroussou, L., Mamon, E. Z. & Landman, J. When quantum and classical models disagree: Learning beyond minimum norm least square. <i>npj Quantum Information</i> 12 , 81 (2026).	2404
	2405
	2406
	2407
	2408
[15] Manski, C. F. <i>Partial Identification of Probability Distributions</i> Springer Series in Statistics (Springer, New York, NY, 2003).	2409
	2410
	2411
	2412
	2413
[16] Molinari, F. in <i>Microeconometrics with partial identification</i> (eds Durlauf, S. N., Hansen, L. P., Heckman, J. J. & Matzkin, R. L.) <i>Handbook of Econometrics, Volume 7A</i> 355–486 (Elsevier, 2020).	2414
	2415
	2416
	2417
	2418
[17] Polo, F. M., Maity, S., Yurochkin, M., Banerjee, M. & Sun, Y. <i>Weak supervision performance evaluation via partial identification. Advances in Neural Information Processing Systems</i> , Vol. 37 (2024).	2419
	2420
	2421
	2422
	2423
	2424
[18] Katariya, N., Iyer, A. & Sarawagi, S. <i>Active evaluation of classifiers on large datasets. 2012 IEEE 12th International Conference on Data Mining</i> , 329–338 (IEEE, 2012).	2425
	2426
	2427
	2428
	2429
	2430
[19] Kossen, J., Farquhar, S., Gal, Y. & Rainforth, T. <i>Active testing: Sample-efficient model evaluation. Proceedings of the 38th International Conference on Machine Learning</i> , Vol. 139 of <i>Proceedings of Machine Learning Research</i> , 5753–5763 (PMLR, 2021). URL https://proceedings.mlr.press/v139/kossen21a.html .	2431
	2432
	2433
	2434
	2435
	2436
	2437
	2438

- 2439 [20] Kossen, J., Farquhar, S., Gal, Y. & Rainforth, T. *Active surrogate esti-*
2440 *mators: An active learning approach to label-efficient model evaluation.*
2441 *Advances in Neural Information Processing Systems*, Vol. 35, 24557–24570
2442 (2022). URL [https://proceedings.neurips.cc/paper_files/paper/2022/hash/](https://proceedings.neurips.cc/paper_files/paper/2022/hash/9b9cfd5428153ccfbd4ba34b7e007305-Abstract-Conference.html)
2443 [9b9cfd5428153ccfbd4ba34b7e007305-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/9b9cfd5428153ccfbd4ba34b7e007305-Abstract-Conference.html).
2444
2445
2446
2447
2448 [21] Ashouritaklimi, K., Kilian, V., Huang, D., Rainforth, T. & Caron, F. Prediction-
2449 powered active testing (2026). URL <https://arxiv.org/abs/2607.08347>. Preprint,
2450 arXiv:2607.08347.
2451
2452
2453
2454 [22] Okanovic, P. *et al.* *All models are wrong, some are useful: Model selection*
2455 *with limited labels. Proceedings of The 28th International Conference on Arti-*
2456 *ficial Intelligence and Statistics*, Vol. 258 of *Proceedings of Machine Learning*
2457 *Research*, 2035–2043 (PMLR, 2025). URL [https://proceedings.mlr.press/v258/](https://proceedings.mlr.press/v258/okanovic25a.html)
2458 [okanovic25a.html](https://proceedings.mlr.press/v258/okanovic25a.html).
2459
2460
2461
2462
2463 [23] Shanmugam, D. *et al.* *Evaluating multiple models using labeled and unlabeled*
2464 *data. Advances in Neural Information Processing Systems*, Vol. 38, 30612–
2465 30648 (2025). URL [https://proceedings.neurips.cc/paper_files/paper/2025/hash/](https://proceedings.neurips.cc/paper_files/paper/2025/hash/2c0e78cc177dfeca260ef990a8c99209-Abstract-Conference.html)
2466 [2c0e78cc177dfeca260ef990a8c99209-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2025/hash/2c0e78cc177dfeca260ef990a8c99209-Abstract-Conference.html).
2467
2468
2469
2470
2471 [24] Garg, S., Balakrishnan, S., Lipton, Z. C., Neyshabur, B. & Sedghi, H. *Leveraging*
2472 *unlabeled data to predict out-of-distribution performance. International Confer-*
2473 *ence on Learning Representations* (2022). URL [https://openreview.net/forum?](https://openreview.net/forum?id=o_HsiMPYh_x)
2474 [id=o_HsiMPYh_x](https://openreview.net/forum?id=o_HsiMPYh_x).
2475
2476
2477
2478 [25] Rosenfeld, E. & Garg, S. *(almost) provable error bounds under distribution*
2479 *shift via disagreement discrepancy. Advances in Neural Information Processing*
2480 *Systems*, Vol. 36, 28761–28784 (2023).
2481
2482
2483
2484

- [26] Mishra, A. & Liu, A. *ODD: Overlap-aware estimation of model performance under distribution shift. Proceedings of the 41st Conference on Uncertainty in Artificial Intelligence*, Vol. 286 of *Proceedings of Machine Learning Research*, 3031–3047 (PMLR, 2025). URL <https://proceedings.mlr.press/v286/mishra25a.html>.
- [27] Madani, O., Pennock, D. M. & Flake, G. W. *Co-validation: Using model disagreement on unlabeled data to validate classification algorithms. Advances in Neural Information Processing Systems*, Vol. 17, 873–880 (2004). URL https://proceedings.neurips.cc/paper_files/paper/2004/hash/92f54963fc39a9d87c2253186808ea61-Abstract.html.
- [28] Musgrove, D., Radtke, A. & Haddad, T. Discordant pair analysis for sample efficient model evaluation. *Scientific Reports* **13**, 20987 (2023).
- [29] Balsubramani, A. & Freund, Y. *Optimally combining classifiers using unlabeled data. Proceedings of the 28th Conference on Learning Theory*, Vol. 40 of *Proceedings of Machine Learning Research*, 211–225 (PMLR, 2015). URL <https://proceedings.mlr.press/v40/Balsubramani15.html>.
- [30] Abreu, D., Esteve Rothenberg, C. & Abelém, A. *QML-IDS: Quantum machine learning intrusion detection system. 2024 IEEE Symposium on Computers and Communications (ISCC)*, 1–6. IEEE Computer Society (IEEE, Paris, France, 2024).
- [31] Abreu, D., Moura, D., Esteve Rothenberg, C. & Abelém, A. QuantumNetSec: Quantum machine learning for network security. *International Journal of Network Management* **35**, e70018 (2025).

- 2531 [32] Bellante, A., Fioravanti, T., Carminati, M., Zanero, S. & Luongo, A. Evaluat-
2532 ing the potential of quantum machine learning in cybersecurity: A case-study
2533 on PCA-based intrusion detection systems. *Computers & Security* **154**, 104341
2534 (2025).
2535
2536
2537
2538 [33] Bazinet, M., Zantedeschi, V. & Germain, P. Bound to disagree: Generalization
2539 bounds via certifiable surrogates (2026). URL <https://arxiv.org/abs/2602.23128>.
2540 Preprint, arXiv:2602.23128.
2541
2542
2543
2544 [34] Quiñonero-Candela, J., Sugiyama, M., Schwaighofer, A. & Lawrence, N. D. (eds)
2545 *Dataset Shift in Machine Learning* (MIT Press, Cambridge, MA, 2008).
2546
2547
2548 [35] Shaydulin, R. & Wild, S. M. Importance of kernel bandwidth in quantum machine
2549 learning. *Physical Review A* **106**, 042407 (2022).
2550
2551
2552 [36] Canatar, A., Peters, E., Pehlevan, C., Wild, S. M. & Shaydulin, R. Band-
2553 width enables generalization in quantum kernel models. *Transactions on Machine*
2554 *Learning Research* (2023).
2555
2556
2557
2558 [37] Kübler, J. M., Buchholz, S. & Schölkopf, B. *The inductive bias of quantum ker-*
2559 *nels. Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34,
2560 12661–12673 (Neural Information Processing Systems Foundation, 2021).
2561
2562
2563
2564 [38] Thanasilp, S., Wang, S., Cerezo, M. & Holmes, Z. Exponential concentration in
2565 quantum kernel methods. *Nature Communications* **15**, 5200 (2024).
2566
2567
2568 [39] Wang, X., Du, Y., Luo, Y. & Tao, D. Towards understanding the power of
2569 quantum kernels in the NISQ era. *Quantum* **5**, 531 (2021).
2570
2571
2572 [40] Gentinetta, G., Thomsen, A., Sutter, D. & Woerner, S. The complexity of
2573 quantum support vector machines. *Quantum* **8**, 1225 (2024).
2574
2575
2576

- [41] Agliardi, G. *et al.* Mitigating exponential concentration in covariant quantum kernels for subspace and real-world data. *npj Quantum Information* **12**, 12 (2026).
2577
2578
2579
2580
- [42] Liu, Y., Arunachalam, S. & Temme, K. A rigorous and robust quantum speed-up in supervised machine learning. *Nature Physics* **17**, 1013–1017 (2021).
2581
2582
2583
2584
- [43] Jerbi, S., Gyurik, C., Marshall, S. C., Molteni, R. & Dunjko, V. Shadows of quantum machine learning. *Nature Communications* **15**, 5676 (2024).
2585
2586
2587
2588
- [44] Engelen, G., Rimmer, V. & Joosen, W. *Troubleshooting an intrusion detection dataset: The CICIDS2017 case study. 2021 IEEE Security and Privacy Workshops (SPW)*, 7–12. IEEE Computer Society (IEEE, San Francisco, CA, USA, 2021).
2589
2590
2591
2592
2593
2594
2595
2596
- [45] Kakavand, S., Strohmeyer, C. & Schlotter, M. Benchmarking quantum kernel support vector machines against classical baselines on tabular data: A rigorous empirical study with hardware validation. *arXiv* **abs/2604.18837** (2026). URL <https://arxiv.org/abs/2604.18837>.
2597
2598
2599
2600
2601
2602
2603
- [46] Gardner, J., Popovic, Z. & Schmidt, L. *Benchmarking distribution shift in tabular data with TableShift. Advances in Neural Information Processing Systems*, Vol. 36 (Neural Information Processing Systems Foundation, 2023).
2604
2605
2606
2607
2608
2609
- [47] Anderson, H. S. & Roth, P. Ember: An open dataset for training static PE malware machine learning models. *arXiv* **abs/1804.04637** (2018). URL <https://arxiv.org/abs/1804.04637>.
2610
2611
2612
2613
2614
2615
- [48] Moustafa, N. & Slay, J. *UNSW-NB15: A comprehensive data set for network intrusion detection systems (UNSW-NB15 network data set),. 2015 Military Communications and Information Systems Conference (MilCIS)*, 1–6 (IEEE, 2015).
2616
2617
2618
2619
2620
2621
2622

- 2623 [49] Alsaedi, A., Moustafa, N., Tari, Z., Mahmood, A. & Anwar, A. TON_IoT teleme-
2624 try dataset: A new generation dataset of IoT and IIoT for data-driven intrusion
2625 detection systems. *IEEE Access* **8**, 165130–165150 (2020).
2626
2627
2628
2629 [50] Cortes, C. & Vapnik, V. Support-vector networks. *Machine Learning* **20**, 273–297
2630 (1995).
2631
2632
2633 [51] Rasmussen, C. E. & Williams, C. K. I. *Gaussian Processes for Machine Learning*
2634 (MIT Press, Cambridge, MA, 2006).
2635
2636
2637 [52] Javadi-Abhari, A. *et al.* Quantum computing with qiskit. *arXiv* **abs/2405.08810**
2638 (2024). URL <https://arxiv.org/abs/2405.08810>.
2639
2640
2641 [53] Ibragimov, R. & Müller, U. K. Inference with few heterogeneous clusters. *The*
2642 *Review of Economics and Statistics* **98**, 83–96 (2016).
2643
2644
2645 [54] Guo, C., Pleiss, G., Sun, Y. & Weinberger, K. Q. *On calibration of modern neural*
2646 *networks. Proceedings of the 34th International Conference on Machine Learn-*
2647 *ing (ICML)*, Vol. 70 of *PMLR*, 1321–1330 (International Conference on Machine
2648 Learning, 2017).
2649
2650
2651
2652
2653 [55] Roy, O. & Vetterli, M. *The effective rank: A measure of effective dimension-*
2654 *ality. 2007 15th European Signal Processing Conference (EUSIPCO)*, 606–610
2655 (EURASIP, 2007).
2656
2657
2658
2659 [56] Cristianini, N., Shawe-Taylor, J., Elisseeff, A. & Kandola, J. *On kernel-target*
2660 *alignment. Advances in Neural Information Processing Systems (NeurIPS)*,
2661 Vol. 14 (Neural Information Processing Systems Foundation, 2001).
2662
2663
2664 [57] Cortes, C., Mohri, M. & Rostamizadeh, A. Algorithms for learning kernels based
2665 on centered alignment. *Journal of Machine Learning Research* **13**, 795–828
2666
2667
2668

(2012). 2669

[58] Fernández-Barrios, R., Pastor-López, I., González-Santocildes, A. & 2670
García Bringas, P. Evaluation choices shape apparent quantum kernel robustness 2671
under distribution shift (2026). URL <https://doi.org/10.5281/zenodo.21717074>. 2672
2673
2674
2675
2676

[59] U.S. Census Bureau. Public use microdata sample (PUMS). <https://www.census.gov/programs-surveys/acs/microdata.html> (2026). Accessed 2 August 2026. 2677
2678
2679
2680

[60] Centers for Disease Control and Prevention. BRFSS frequently asked questions. 2681
https://www.cdc.gov/brfss/about/brfss_faq.htm (2022). Accessed 2 August 2026. 2682
2683
2684

[61] National Center for Health Statistics. Ethics review board approval: National 2685
health and nutrition examination survey. <https://www.cdc.gov/nchs/nhanes/about/erb.html> (2026). Accessed 2 August 2026. 2686
2687
2688
2689
2690

[62] UCI Machine Learning Repository. Diabetes 130-US hospitals for 2691
years 1999–2008 (2014). URL [https://archive.ics.uci.edu/dataset/296/](https://archive.ics.uci.edu/dataset/296/diabetes-130-us-hospitals-for-years-1999-2008)
diabetes-130-us-hospitals-for-years-1999-2008. 2692
2693
2694
2695
2696

[63] U.S. Department of Education. College scorecard data. <https://collegescorecard.ed.gov/data/> (2026). Accessed 2 August 2026. 2697
2698
2699
2700
2701
2702
2703
2704
2705
2706
2707
2708
2709
2710
2711
2712
2713
2714