

Quantum and Classical Kernels under Distribution

Shift: Kernel Geometry Governs

Out-of-Distribution Robustness

Roberto Fernández-Barrios^[ORCID]^{1*}, Iker Pastor-López^[ORCID]¹,
Asier González-Santocildes^[ORCID]¹, Pablo García Bringas^[ORCID]¹

^{1*}Faculty of Engineering, University of Deusto, Avda. de las
Universidades, 24., Bilbao, 48007, Spain.

*Corresponding author(s). E-mail(s): roberto.fernandez.b@deusto.es;
Contributing authors: iker.pastor@deusto.es; gonzalez.asier@deusto.es;
pablo.garcia.bringas@deusto.es;

Abstract

Whether quantum kernels improve robustness under distribution shift is usually asked as a quantum-versus-classical question. We argue, and show empirically, that it is more productively asked as a question about kernel geometry. We present a controlled kernel-swap study in which the classifier, preprocessing, split construction, and family-internal model-selection logic are held fixed while only the kernel varies, instantiated on four security benchmarks spanning two modalities and three drift mechanisms (EMBER static malware; UNSW-NB15 and ToN-IoT network flows under synthetic covariate shift and natural regime drift), two classifier families (support vector classification and a Laplace-approximation Gaussian process classifier with calibrated uncertainty), and 72 principal settings with 15 repeated runs each; all quantum kernels are computed by exact statevector simulation, and no hardware claim is made. Against the customary linear and RBF baselines, fidelity-based quantum kernels improve out-of-distribution balanced accuracy nearly uniformly (69–71 of 72 settings; Wilcoxon $p < 10^{-9}$ on network data). Against a broader classical family, a median-heuristic Laplacian kernel neutralizes the advantage on EMBER—revealing that the operative property is not quantumness but geometry—yet when both families receive the same bandwidth-tuning freedom the quantum advantage re-emerges on

EMBER and remains dominant under natural network drift (54 of 54 settings under the probabilistic classifier, $p = 1.6 \times 10^{-10}$). A mechanistic analysis explains these outcomes: within each setting, out-of-distribution accuracy tracks the survival of centered kernel-target alignment under shift, a regularity that holds in 89–99% of 360 setting-seed units across all four datasets, with high effective rank of the training kernel as its structural precondition. The classical-favorable exceptions of the original comparison are precisely the settings where the selected quantum kernel collapses to low-rank, classical-like geometry. Beyond the specific quantum-versus-classical verdict, the study contributes a reusable, fully reproducible protocol for attributing robustness effects to kernel geometry, and a design principle—prefer kernels whose label-aligned geometry survives shift—that spans classical and quantum families alike.

Keywords: quantum machine learning; quantum kernels; distribution shift; out-of-distribution robustness; kernel methods; kernel-target alignment; Gaussian processes; malware detection; network intrusion detection

1 Introduction

Quantum kernel methods are one of the clearest and most actively studied directions in quantum machine learning. By embedding classical data into quantum states and defining similarity through state overlaps, they provide a principled way to study how quantum-induced feature geometry affects supervised learning while keeping the downstream classifier fixed [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#); [Schuld \(2021\)](#). At the same time, the recent QML literature has made an equally important point: empirical conclusions about quantum models can depend strongly on the data regime, the encoding, and the evaluation protocol [Huang et al. \(2021\)](#); [Schuld et al. \(2021\)](#); [Gil-Fuster et al. \(2024\)](#); [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#). For that reason, claims about the usefulness of quantum kernels are most informative when they are established under controlled, reproducible, and carefully matched comparisons rather than through heterogeneous pipelines or loosely aligned baselines.

This need for control becomes even more important under distribution shift. In many applied settings, the relevant question is not only whether a model separates classes well under in-distribution (ID) conditions, but whether its behavior remains

reliable once the evaluation distribution changes. The broader distribution-shift and domain-generalization literature has repeatedly shown that apparent robustness gains can weaken, disappear, or become difficult to interpret when model-selection strategy, hyperparameter search, split construction, or evaluation standardization are not carefully controlled [Quiñonero-Candela et al. \(2008\)](#); [Zhou et al. \(2023\)](#); [Gulrajani and Lopez-Paz \(2021\)](#); [Koh et al. \(2021\)](#); [Yao et al. \(2022\)](#). A convincing robustness study must therefore satisfy two requirements at once: it must isolate the factor of interest tightly enough to support attribution, and it must evaluate that factor broadly enough to determine whether the observed behavior reflects a meaningful regime-level effect rather than a narrow configuration artifact.

These issues are especially acute in security applications. Drift is not incidental to the problem; it is part of the problem definition itself. Static PE malware classification has been widely studied through benchmarks such as EMBER [Anderson and Roth \(2018\)](#), while more recent work has highlighted the gap between laboratory evaluation and performance on drifted malware, as well as the difficulty of adapting with limited post-drift data [Barbero et al. \(2022\)](#); [Li et al. \(2025\)](#). The emergence of updated resources such as EMBER2024 further reflects an active push toward more realistic, broader, and more deployment-relevant malware evaluation [Joyce et al. \(2025\)](#). Network intrusion detection faces the same structural challenge: traffic distributions move as services, infrastructure, and attack campaigns evolve, and public benchmarks such as UNSW-NB15 and ToN-IoT expose exactly this kind of regime change [Moustafa and Slay \(2015\)](#); [Alsaedi et al. \(2020\)](#). Taken together, this literature suggests that a useful security-ML study should not merely report accuracy on a convenient split; it should examine how model choices behave when the evaluation explicitly targets shift, and it should do so across more than one data modality and drift mechanism.

Kernel methods are a natural setting in which to ask such questions cleanly. Because kernel-based support vector machines separate the choice of classifier from the

choice of similarity function, they offer a principled way to study how representational geometry affects ID and OOD behavior while keeping the learning machinery fixed [Cortes and Vapnik \(1995\)](#); [Schölkopf and Smola \(2002\)](#). This perspective becomes particularly interesting in quantum machine learning, where quantum kernels are defined through implicit feature maps into quantum Hilbert spaces and estimated via state overlaps or related fidelity-based similarities [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#); [Schuld \(2021\)](#). From this viewpoint, quantum kernels are not a vague alternative model class, but a concrete family of kernel constructions whose empirical value can—and should—be tested under the same standards applied to classical kernels.

At the same time, the quantum-kernel literature has matured to the point where optimistic demonstrations and critical scrutiny must now be read together. Foundational work established the feature-space interpretation of quantum kernels and motivated them as expressive alternatives to classical similarities [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#). Subsequent studies emphasized that performance in quantum machine learning can depend strongly on the data regime, the encoding, and the effective geometry induced by the model [Huang et al. \(2021\)](#); [Schuld et al. \(2021\)](#). More recent work has sharpened this message further: understanding generalization in quantum machine learning requires care beyond classical intuition, and benchmarking conclusions can be highly sensitive to protocol choices [Gil-Fuster et al. \(2024\)](#); [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#). This makes broad claims of superiority difficult to defend and reinforces the need for conservative, protocol-aware empirical studies.

The gap we address lies precisely at this intersection, and it has three parts. First, there remains limited evidence from controlled, variability-aware comparisons between classical and quantum kernel families in security-relevant settings where out-of-distribution robustness is part of the evaluation target: existing studies often differ

simultaneously in model class, preprocessing, search space, or data protocol, making it difficult to attribute observed gains to the kernel itself. Second, when such comparisons are attempted, the classical reference set is typically narrow—linear and RBF kernels under a single classifier—so an apparent quantum advantage may reflect the weakness of the baselines rather than a property of the quantum family; probabilistic classifiers with principled uncertainty quantification, and classical kernels with richer spectral structure, are usually absent. Third, and most fundamentally, empirical wins and losses are rarely accompanied by an account of *why* a kernel family should behave better or worse under shift, leaving results as observations in search of an explanation. For a question as specific as whether quantum kernels help under distribution shift, these confounds and omissions are not peripheral; they are decisive.

In this paper, we study that question through a strictly controlled *kernel-swap* protocol, instantiated broadly enough to close all three gaps at once. Within each experimental setting we hold the classifier, the preprocessing pipeline, the split definitions, and the family-internal model-selection logic fixed, and vary only the kernel. The protocol is applied to four security benchmarks spanning two data modalities and three drift mechanisms: EMBER static PE malware under two synthetic covariate-shift constructions [Anderson and Roth \(2018\)](#), and UNSW-NB15 and ToN-IoT network-flow scenarios under both an analogous synthetic construction and *natural* regime drift, in which the out-of-distribution pool consists of genuinely drifted traffic [Moustafa and Slay \(2015\)](#); [Alsaedi et al. \(2020\)](#). Two classifier families consume the same precomputed kernels: support vector classification, and a binary Laplace-approximation Gaussian process classifier [Rasmussen and Williams \(2006\)](#) that adds calibrated predictive probabilities—and hence principled uncertainty quantification under shift—without changing anything else in the pipeline. The classical family ranges from the customary linear and RBF baselines to polynomial, Laplacian, and Matérn kernels with train-only median-heuristic length scales; the quantum family

comprises fidelity kernels induced by commonly used feature-map families. Crucially, both families receive the same bandwidth-tuning freedom (a gamma sweep for RBF, an angle-scaling sweep for the quantum maps), so that neither side is compared at a fixed operating point the other is allowed to tune away from [Shaydulin and Wild \(2022\)](#); [Canatar et al. \(2023\)](#). The main evaluation comprises 72 principal settings (18 on EMBER, 54 on network flows), each repeated across 15 runs, with paired Wilcoxon signed-rank tests and Holm correction on all family-level comparisons.

Beyond measuring who wins, we ask why. For every kernel, dimension, and setting we compute geometry descriptors of the induced Gram matrices—the effective rank of the training kernel [Roy and Vetterli \(2007\)](#), centered kernel-target alignment on the ID and OOD splits [Cristianini et al. \(2001\)](#); [Cortes et al. \(2012\)](#), kernel concentration statistics, and the geometric difference between classical and quantum kernels [Huang et al. \(2021\)](#)—and test which of these quantities predict out-of-distribution behavior within and across settings. This mechanistic layer converts the benchmark from a scoreboard into an instrument for attribution.

The resulting picture is sharper than a quantum-versus-classical verdict, and more useful. Against linear and RBF baselines, fidelity-based quantum kernels improve OOD balanced accuracy nearly uniformly across all four datasets (69–71 of 72 settings under both classifiers). Against the extended classical family, the advantage on EMBER is neutralized by a median-heuristic Laplacian kernel—demonstrating that the operative property is geometric rather than quantum—but re-emerges under symmetric bandwidth tuning, and remains dominant on network flows, where under the probabilistic classifier the quantum family wins all 54 settings with gains exceeding run-to-run variability in the majority. In-distribution separability favors the quantum family in every EMBER setting regardless of the classical reference set. The mechanistic analysis explains all of these outcomes with one regularity: within a setting, OOD accuracy tracks the survival of centered kernel-target alignment under shift—positive in 89–99%

of 360 setting-seed units across all four datasets and both classifiers—with high effective rank of the training kernel as its structural precondition. The classical-favorable exceptions of the baseline comparison are exactly the settings where the selected quantum kernel collapses to a low-rank, classical-like geometry, and the Laplacian is exactly the classical kernel whose geometry most resembles the quantum maps.

Our contributions are fourfold. First, a controlled, reproducible kernel-swap protocol for robustness studies that extends across classifier families (including a probabilistic classifier with calibrated uncertainty), across a broad classical reference set with symmetric bandwidth tuning, and across datasets, modalities, and drift mechanisms. Second, a variability-aware regime map of quantum-versus-classical kernel behavior under shift on 72 settings with formal significance testing, including the honest negative result that a well-chosen classical kernel matches the quantum family on one modality at fixed bandwidth. Third, a mechanistic account—alignment survival under shift, enabled by high effective rank—that predicts within-setting OOD performance across every dataset we study, explains the structured exceptions, and yields the checkable prediction (which we confirm) that the classical kernels sharing this geometry are the competitive ones. Fourth, practical guidance and a fully public artifact: for OOD-sensitive kernel selection, prefer kernels whose label-aligned geometry survives shift; fidelity-based quantum maps exhibit this property natively, and on drifting network traffic they currently exploit it best.

2 Related Work

2.1 Distribution shift, domain generalization, and evaluation rigor

The broader literature on distribution shift has established that predictive performance can degrade substantially when training and test data are drawn from different

distributions, and that robust evaluation under shift requires more than reporting strong in-distribution performance [Quiñonero-Candela et al. \(2008\)](#); [Shimodaira \(2000\)](#); [Sugiyama et al. \(2007\)](#). Within machine learning, domain generalization (DG) has emerged as one of the main formal settings for studying out-of-distribution (OOD) performance without access to target-domain labels during training [Zhou et al. \(2023\)](#). This literature has also made clear that empirical conclusions about robustness are highly sensitive to evaluation design.

A particularly important development is the move toward benchmark-driven scrutiny of robustness claims: DomainBed showed that many reported domain-generalization improvements dissolve once datasets, architectures, hyperparameter budgets, and model-selection criteria are held consistent [Gulrajani and Lopez-Paz \(2021\)](#); WILDS demonstrated that standard training leaves large OOD gaps under naturally occurring real-world shifts [Koh et al. \(2021\)](#); and Wild-Time established temporal shift as a distinct, practically relevant regime that existing methods struggle to close [Yao et al. \(2022\)](#). These works matter here not because security detection is a standard DG benchmark, but because they establish a methodological principle we adopt wholesale: when the goal is to compare robustness under shift, fairness of protocol is part of the contribution. Our comparison therefore fixes the classifier family, preprocessing, and split definitions, so that the role of the kernel can be interpreted with minimal confounding.

2.2 Security detection under drift and evolving evaluation benchmarks

In malware detection, the relevance of distribution shift is especially acute. Malware ecosystems evolve continuously as attackers modify payloads, packers, obfuscation strategies, delivery mechanisms, and development toolchains. Static detection models

are therefore exposed to changing distributions in ways that are operationally consequential rather than merely academic. Public benchmarks have played an important role in making this problem more reproducible. EMBER remains one of the most widely used open datasets for static PE malware detection and helped standardize large-scale evaluation in this area [Anderson and Roth \(2018\)](#). More recent resources such as EMBER2024 push this effort further by broadening scope, increasing scale, and enabling more holistic evaluation across multiple malware-related tasks [Joyce et al. \(2025\)](#). Other datasets such as BODMAS and SOREL-20M likewise reflect the field’s effort to move toward larger and more realistic benchmark settings [Yang et al. \(2021\)](#); [Harang and Rudd \(2020\)](#).

At the same time, a growing body of work has shown that strong laboratory performance does not guarantee robustness to evolving malware. The Transcend line introduced conformal criteria to identify aging classifiers in deployment before their accuracy visibly degrades [Jordaney et al. \(2017\)](#), later refined by Barbero et al. [Barbero et al. \(2022\)](#); CADE detects and explains drifting samples without out-of-distribution labels [Yang et al. \(2021\)](#); and Li et al. studied adaptation to real drifted malware under limited post-drift supervision [Li et al. \(2025\)](#). On the methodological side, TESSER-ACT established that spatial and temporal experimental bias can manufacture illusory detector performance [Pendlebury et al. \(2019\)](#), and Arp et al. catalogued the recurring pitfalls of ML evaluations in computer security—sampling bias, label leakage, and inappropriate baselines among them [Arp et al. \(2022\)](#). Our protocol is, in large part, an application of those prescriptions to kernel comparison: fixed splits, leakage checks, temporal-style regime separation in the natural-drift construction, and baselines strengthened until they stop flattering the method under study. Taken together, these studies support a clear conclusion: in security detection, robustness under shift is not a niche desideratum but a primary evaluation target.

The same structural problem, with faster drift dynamics, defines network intrusion detection: benchmark corpora such as UNSW-NB15 and ToN-IoT capture traffic whose attack composition and background behavior change across collection regimes [Moustafa and Slay \(2015\)](#); [Alsaedi et al. \(2020\)](#), and adaptive NIDS pipelines are routinely retrained for exactly this reason. Using both modalities in one protocol lets us ask whether kernel-level robustness effects are properties of a dataset family or of the kernels.

Quantum machine learning has itself entered this application space. Quantum-enhanced classifiers and QSVMs have been applied to network intrusion detection since at least Gouveia and Correia [Gouveia and Correia \(2020\)](#), the area has been systematically reviewed [Nicesio et al. \(2023\)](#); [Nieto and Gil-Merino \(2026\)](#), and recent systems evaluate quantum kernels directly on the benchmarks we use, reporting quantum advantages on UNSW-NB15 and ToN-IoT [Abreu et al. \(2024, 2025\)](#); a critical evaluation on PCA-based intrusion detection found classical methods to remain competitive under realistic constraints [Bellante et al. \(2025\)](#). A common trait of the optimistic studies is decisive for the present paper: they compare against a small set of classical baselines on i.i.d. splits—precisely the two evaluation choices our results show to inflate apparent quantum advantages. Our study can be read as the missing stress test for this literature: under matched baselines, symmetric tuning, and explicit shift, part of the advantage survives and part of it does not, and the geometry explains which.

Our work differs from this line in emphasis. Rather than proposing a new drift-adaptation mechanism, drift detector, or continual-update strategy, we ask a narrower but complementary question: under a tightly controlled protocol, does the choice of kernel family itself materially affect OOD robustness? This question is well motivated by the security literature, yet comparatively underexplored in a way that isolates kernels from the many other moving parts of the pipeline.

2.3 Kernel methods, quantum kernels, and careful benchmarking in QML

Kernel methods provide a natural framework for studying robustness under controlled model comparisons because they separate the learning algorithm from the similarity function used to compare inputs. In classical machine learning, support vector machines and related kernel methods have long provided a principled route for controlling model capacity while varying the geometry of the induced feature space [Cortes and Vapnik \(1995\)](#); [Schölkopf and Smola \(2002\)](#). This makes kernels especially attractive when the scientific question concerns representation and similarity rather than wholesale changes in classifier architecture.

Quantum kernel methods arise from the observation that data can be embedded into quantum states and compared through overlaps in a quantum feature space. [Havlíček et al.](#) popularized this perspective in the context of supervised classification with quantum-enhanced feature spaces [Havlíček et al. \(2019\)](#), while [Schuld and Killoran](#) provided a broader theoretical account linking quantum models to kernel methods through feature Hilbert spaces [Schuld and Killoran \(2019\)](#). Later work sharpened this connection by arguing that supervised quantum machine learning models can often be interpreted directly as kernel methods, making the kernel viewpoint central rather than incidental [Schuld \(2021\)](#). Additional work on data encoding further showed that representational behavior depends critically on how classical inputs are mapped into quantum states, which is directly relevant when comparing quantum feature maps as alternative kernel constructions [Schuld et al. \(2021\)](#). Two spectral results frame what such kernels can and cannot do: [Kübler et al.](#) showed that the inductive bias of quantum kernels is governed by the spectrum of the induced integral operator, with flat spectra implying poor in-distribution sample efficiency [Kübler et al. \(2021\)](#), and [Thanasilp et al.](#) established that broad families of fidelity kernels concentrate exponentially toward trivial values as systems grow, erasing trainability [Thanasilp et al.](#)

(2024). Both results interact nontrivially with our findings and are taken up in the discussion.

However, the contemporary QML literature also makes clear that promising empirical behavior should not be read uncritically. Huang et al. showed that the role of data can be decisive in quantum machine learning, suggesting that observed performance gains are deeply regime-dependent rather than model-intrinsic Huang et al. (2021). Caro et al. studied generalization from few training examples Caro et al. (2022) and later proved out-of-distribution generalization guarantees for learning quantum dynamics Caro et al. (2023)—a setting whose inputs are quantum states rather than classical data, and thus complementary to the classical-input shift studied here—while Gil-Fuster et al. argued that understanding generalization in QML requires rethinking familiar classical intuitions, since standard explanations may fail to account for observed behavior in quantum models Gil-Fuster et al. (2024). In parallel, broader perspective pieces have emphasized both the promise and the unresolved challenges of QML as a field Cerezo et al. (2022).

Most importantly for the present work, recent benchmarking studies have raised the standard for empirical claims in QML. Bowles et al. argued that experimental design choices can dominate conclusions and showed, in a broad simulated benchmark, that out-of-the-box classical models often remain strong baselines against quantum alternatives Bowles et al. (2024). Schnabel and Roth extended this benchmarking perspective specifically to quantum kernel methods, conducting a large-scale study across many design choices and showing that performance depends strongly on kernel type, encoding strategy, and hyperparameter configuration rather than on a simple classical-versus-quantum dichotomy Schnabel and Roth (2025). Related controlled evidence from adjacent QML settings points in the same methodological direction: recent work in hybrid quantum reinforcement learning shows that varying multiple pipeline components at once can obscure attribution, and highlights the importance

of component-isolating, reproducible evaluation under fixed experimental conditions
[Lazaro et al. \(2025\)](#). This line of work is especially relevant to our paper because it
 supports a conservative framing: the value of quantum kernels, if any, must be demon-
 strated under controlled and transparent comparisons, and should be expected to be
 regime-dependent.

Three further strands supply the analytical tools of this study. First, work on
 kernel bandwidth in QML has shown that rescaling the data encoding is as consequen-
 tial for fidelity kernels as the length scale is for classical kernels, and can determine
 whether quantum models generalize at all [Shaydulin and Wild \(2022\)](#); [Canatar et al.](#)
[\(2023\)](#); [Egginger et al. \(2024\)](#); we operationalize this insight as a fairness requirement—
 symmetric bandwidth freedom for both families—rather than as a quantum-side
 tuning trick. Second, kernel-target alignment [Cristianini et al. \(2001\)](#); [Cortes et al.](#)
[\(2012\)](#) and effective rank [Roy and Vetterli \(2007\)](#) are classical tools for characteriz-
 ing Gram-matrix geometry, and the theory of spectral bias and task-model alignment
 shows that in-distribution generalization of kernel methods is governed precisely by
 the compatibility between the kernel’s spectrum and the target function [Bordelon](#)
[et al. \(2020\)](#); [Canatar et al. \(2021\)](#); alignment has also served as a training objective
 for quantum embedding kernels [Hubregtsen et al. \(2022\)](#), a line our negative result
 on a-priori alignment-based selection speaks to directly. To our knowledge, the joint
 use of these descriptors as *predictors of OOD behavior under distribution shift*, across
 classical and quantum kernels alike, is new: we study the shift-facing counterpart of
 task-model alignment—whether the aligned spectrum survives once the distribution
 moves. Third, Gaussian process classification [Rasmussen and Williams \(2006\)](#) provides
 the standard route to calibrated uncertainty in kernel methods; including a Laplace-
 approximation GP classifier in the kernel-swap protocol lets us evaluate uncertainty
 quality under shift without changing anything else in the pipeline.

599 Finally, we note concurrent work. A 2026 preprint studies the robustness of quan-
600 tum kernel methods under additive input corruption and proposes a spectral phase
601 encoding to mitigate it [Herrero Gómez et al. \(2026\)](#). It shares our move away from
602 clean-data peak performance, but differs in all three defining aspects: it perturbs
603 inputs with synthetic noise rather than evaluating under covariate or natural regime
604 shift, it proposes an encoding rather than auditing families under a fixed protocol,
605 and it does not use Gram-matrix geometry as the explanatory variable. Large-scale
606 quantum-kernel benchmarking without varied train–test distributions [Bowles et al.](#)
607 [\(2024\)](#); [Schnabel and Roth \(2025\)](#) remains the closest genre; the shift axis and the
608 mechanism are the present paper’s additions to it.

615 Our study is positioned at the intersection of these literatures. From the
616 distribution-shift literature, we adopt the principle that robustness claims must rest
617 on evaluation protocols that control major confounds. From the security literature,
618 we take the practical motivation that temporal and distributional change are intrinsic
619 to realistic deployment, in malware and network intrusion detection alike. From the
620 QML literature, we take both the motivation for studying quantum kernels and the
621 caution that benchmarking quality is central to interpretation. Relative to prior work,
622 our contribution is not to claim universal superiority for quantum kernels, nor to pro-
623 pose a new adaptation method for drifted security data, but to provide a controlled,
624 variability-aware, multi-dataset comparison of classical and quantum kernel families in
625 which out-of-distribution robustness is the primary evaluation target—and to accom-
626 pany the comparison with a mechanistic account of why its outcomes are what they
627 are.
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644

3 Methods

3.1 Problem setting

We study binary security classification—malware versus benign software, and attack versus benign traffic—under distribution shift. Let D_{ID} denote the in-distribution data-generating process and D_{OOD} an out-of-distribution process induced by a shift mechanism, which may be synthetically constructed or naturally occurring. Given a training set $T \sim D_{\text{ID}}$, we evaluate the learned classifier on both an ID test set S_{ID} and an OOD test set S_{OOD} . Our objective is not only to measure separability under the i.i.d. assumption, but to compare how kernel families behave when generalization is evaluated explicitly under shift.

The central methodological requirement of the study is attribution. Because the scientific question concerns the effect of the kernel, comparisons must avoid confounds from changes in classifier family, preprocessing, optimization protocol, or split design. We therefore adopt a strictly controlled *kernel-swap* protocol in which all components of the pipeline are held fixed except the kernel function itself. This design follows the same evaluation logic that motivates rigorous OOD benchmarking in the broader distribution-shift literature: when the target of inference is robustness, protocol control is not an implementation detail but part of the contribution.

3.2 Datasets and split construction

The protocol is instantiated on four benchmark scenarios spanning two data modalities. The first is EMBER, the standard public benchmark for static PE malware detection [Anderson and Roth \(2018\)](#), which is widely used, supports reproducible feature-based experimentation, and has motivated later benchmark extensions such as EMBER2024 [Joyce et al. \(2025\)](#). The remaining three are network-flow intrusion scenarios: two derived from UNSW-NB15 (with DoS and Reconnaissance, respectively, as

the drifting attack category) [Moustafa and Slay \(2015\)](#) and one from ToN-IoT (Scanning) [Alsaedi et al. \(2020\)](#). In all cases our goal is not to benchmark absolute detection performance on a standard split, but to construct controlled ID/OOD regimes that allow kernel-family comparisons under distribution shift.

For each experimental setting, we adopt a fixed three-way split design:

$$T, \quad S_{\text{ID}}, \quad S_{\text{OOD}},$$

where T is used for training, S_{ID} measures in-distribution performance, and S_{OOD} measures robustness under the induced shift. All models compared within a setting use exactly the same split definitions. This point is essential: a classical and a quantum model are never compared on different train/test partitions.

EMBER shift constructions.

On EMBER we operationalize shift through two controlled OOD-construction variants, denoted $m1$ and $m2$, both built on the histogram + byte-entropy feature blocks and both designed to preserve class balance across splits. Both are *label-conditional tail shifts*: the OOD pool is assembled from the extremes of a within-class score distribution, so the construction uses the labels to decide which samples are held out. We use this term rather than “covariate shift” deliberately. Because selection is conditioned on the class, we cannot assert that $P(Y | X)$ is preserved, which is what the covariate-shift assumption requires; class balance is preserved by construction, but that is a weaker property. The constructions are reproducible and interpretable stress tests of kernel geometry, not certified instances of a formal shift model, and none of our conclusions depend on their being so.

Variant m1: sparsity-score extremes within class.

For each sample, we compute a sparsity-related score on a predefined feature block. In the main configuration used in the experiments, this score corresponds to the number of non-zero entries within the selected block. The OOD set is formed by taking, within each class separately, the low-score and high-score extremes, i.e., the two tails of the within-class score distribution. The remaining samples form an ID pool from which the training and ID test sets are obtained by stratified sampling. Intuitively, this construction creates an OOD regime driven by atypical sparsity structure while keeping the label distribution under control.

Variant m2: distance-to-train extremes within class.

We first fix the final training set at an exact target size. Using only this training set, we compute class-conditional centroids and assign each non-training sample a score equal to its Euclidean distance to the centroid of its own class. To reduce noise and avoid leakage, this distance can be computed in a TruncatedSVD-projected space fitted on the training set only. The OOD set is then defined, again within each class separately, by the low- and high-distance extremes from the non-training pool, with exact target OOD size and proportional class allocation. The ID test set is drawn from the remainder. This produces a complementary shift regime in which OOD samples are unusual relative to the final training geometry rather than to a sparsity statistic alone.

Network-flow scenarios and shift constructions.

Each network scenario is defined by a *reference* pool (traffic from the pre-drift regime) and a *current* pool (traffic containing the drifted attack campaign), prepared with a numeric-only feature policy that removes identifiers, label-derived columns, and categorical protocol fields; UNSW-NB15 scenarios retain 39 flow features and ToN-IoT 89. We verified explicitly that no feature is label-derived: the maximum feature-label

correlation across scenarios is 0.62 (attained by well-known legitimate features such
as `sttl`), far from the near-perfect correlation that would signal leakage, and the
verification script is part of the public artifact. On these pools we instantiate two
shift mechanisms. The first, *m2-centroid*, is the direct analog of EMBER’s *m2*: within
the class-balanced reference pool, a provisional training pool is fixed, class-conditional
centroids are computed in a train-only MaxAbs+TruncatedSVD space, and the OOD
pool consists of the within-class low- and high-distance tails. The second, which we call
attack-campaign regime shift, applies no geometric selection: the training and ID pools
are drawn from the reference regime and the OOD pool from the class-balanced current
regime. The pools themselves are constructed—and we state the criterion precisely,
since it determines what the mechanism is being tested against. The reference pool
contains benign traffic plus every attack category *except* the target campaign; the
current pool contains benign traffic plus the target campaign. For UNSW-NB15 the
two pools additionally come from the two different official partitions of the testbed
capture. The OOD regime therefore combines an emerging attack campaign, absent
from training, with a capture-partition change, and it is not a temporal split in the
timestamp sense. Class priors are balanced separately in both pools, so prevalence
shift is deliberately excluded and the measured effect is attributable to the change in
traffic composition rather than to base rates. This is the most operationally realistic
of our three mechanisms, and the least synthetic, but it is not free of construction;
the full raw-to-pool pipeline is part of the public artifact.

Together, these constructions are not intended as universal models of real-world
drift. Rather, they provide reproducible shift mechanisms of three distinct kinds—
atypical-sparsity extremes, train-geometry extremes, and natural regime drift—that
are sufficiently different to support regime-level conclusions while preserving compa-
rability across kernel families. In every setting, all models are compared on exactly
the same T , S_{ID} , and S_{OOD} .




Fig. 1 The controlled kernel-swap protocol. Benchmark pools feed reproducible shift constructions; all kernels operate on a shared train-only embedding and are consumed by the same precomputed-kernel classifiers. Family-level conclusions are drawn by family-internal selection on run means with paired statistical testing, while the Gram matrices themselves feed the kernel-geometry analysis of Section 3.8.

3.3 Controlled kernel-swap protocol

Our core experimental principle is to isolate the kernel as the only variable of interest; Figure 1 summarizes the resulting pipeline. All compared models share the following components:

- **Classifier:** within each comparison, a single classifier consuming precomputed Gram matrices. Two classifier families are studied in parallel: Support Vector Classification with `kernel="precomputed"`, and a binary Gaussian process classifier under the Laplace approximation.
- **Preprocessing:** the same scaling and dimensionality-reduction pipeline for all classical and quantum models.
- **Split design:** identical T , S_{ID} , and S_{OOD} within each setting.
- **Selection logic:** the same family-internal model-selection principle applied consistently to classical and quantum candidates.

This yields a comparison between kernel families rather than between heterogeneous model pipelines. In particular, it prevents observed OOD differences from being

attributed to changes in classifier architecture, feature engineering, or split construction. The choice of precomputed-kernel classifiers is deliberate: it allows both classical and quantum similarities to be consumed by exactly the same decision rule, thereby making the comparison interpretable in kernel-method terms [Cortes and Vapnik \(1995\)](#); [Schölkopf and Smola \(2002\)](#).

The Gaussian process classifier implements the standard binary Laplace approximation with a logistic link (Algorithms 3.1 and 3.2 in [Rasmussen and Williams \(2006\)](#)), applied directly to the same unit-diagonal precomputed kernels as the SVC. Its role in the study is twofold. Methodologically, it tests whether the kernel-family conclusions are classifier-specific or transfer across learning machinery. Substantively, it produces predictive probabilities rather than scores, which allows robustness under shift to be examined through the lens of uncertainty quantification—log-loss, Brier score, expected calibration error, and predictive entropy on ID versus OOD data—rather than through accuracy alone. We stress that no post-hoc calibration procedure (temperature scaling, isotonic regression, Platt scaling) is applied: the probabilities are those the Laplace approximation produces, and we *assess* their calibration rather than assert it. Kernel-swap comparisons are always performed within a classifier, never across classifiers.

Although the classical and quantum candidate sets are not identical in cardinality, fairness in this study is defined at the protocol level: both families are evaluated under the same classifier, preprocessing, splits, dimensionality sweep, and family-internal selection criterion, so that the effect attributed to the kernel is not confounded by changes elsewhere in the pipeline.

3.4 Preprocessing and feature representation

All models use the same preprocessing pipeline. Let $x \in \mathbb{R}^p$ denote the original feature vector of the dataset at hand. We apply the following steps:

1. MaxAbs scaling. We first apply <code>MaxAbsScaler</code> to normalize feature magnitudes	921
while preserving sparsity structure.	922
	923
2. Dimensionality reduction. We project the scaled features using <code>TruncatedSVD</code>	924
to a reduced representation of dimension	925
	926
	927
	928
$d \in \{4, 6, 8, 10, 12\}.$	929
	930
	931
	932
3. Angle mapping. For the quantum pipeline, the reduced coordinates are linearly	933
mapped to an angle range $[0, \pi]$ before being passed to the quantum feature maps.	934
	935
	936
This shared preprocessing serves two purposes. First, it keeps the comparison fair	937
by ensuring that classical and quantum kernels operate on the same reduced repre-	938
sentation. Second, it makes representation size itself a controlled factor of the study,	939
enabling both aggregated setting-level comparisons and dimension-level analyses. The	940
dimensionality parameter d is therefore not a tuning artifact hidden in preprocessing,	941
but an explicit axis of the experimental grid.	942
	943
	944
	945
	946
	947
3.5 Kernel families	948
	949
<i>Baseline classical kernels.</i>	950
	951
The customary reference set consists of two standard baselines:	952
	953
	954
$K_{\text{lin}}(x, x') = x^\top x', \quad (1)$	955
	956
$K_{\text{rbf}}(x, x') = \exp\left(-\gamma\ x - x'\ _2^2\right), \quad (2)$	957
	958
	959
	960
with the RBF kernel evaluated using the standard <code>scale</code> convention for γ and the	961
linear kernel cosine-normalized to unit diagonal. These are the baselines against which	962
quantum kernels are most commonly compared, and they define our <i>classical-original</i>	963
family.	964
	965
	966

967 ***Extended classical family.***

968
969 Because an apparent quantum advantage over two baselines may say more about
970 the baselines than about the quantum family, we additionally evaluate a broader
971 *classical-extended* family: polynomial kernels of degree 2 and 3 (cosine-normalized),
972 the Laplacian kernel $K(x, x') = \exp(-\|x - x'\|_1/\ell_1)$, and Matérn kernels with
973 $\nu \in \{3/2, 5/2\}$, where all length scales ℓ are set by the median heuristic—the median
974 pairwise distance on the training set, the standard choice in kernel two-sample test-
975 ing [Gretton et al. \(2012\)](#)—so that every hyperparameter is determined from training
976 data alone, without access to ID or OOD test information. These kernels were chosen
977 deliberately to span a range of spectral behaviors: the Laplacian and low- ν Matérn
978 kernels are less smooth than the RBF and induce Gram matrices with substantially
979 heavier eigenvalue tails, which, as the mechanistic analysis will show, is precisely the
980 property at stake under shift.

981
982
983
984
985
986
987
988
989 ***Bandwidth symmetry.***

990
991 Kernel bandwidth is known to be a decisive hyperparameter for quantum kernels:
992 rescaling the input encoding plays the same role for fidelity kernels that γ plays for the
993 RBF kernel, and can qualitatively change generalization [Shaydulin and Wild \(2022\)](#);
994 [Canatar et al. \(2023\)](#). A comparison in which the classical family may tune its band-
995 width while the quantum family is pinned to a fixed encoding scale (or vice versa)
996 is therefore not protocol-fair. In the bandwidth-symmetric portion of the study, both
997 families receive the same tuning axis: the RBF kernel is evaluated at $\gamma = f \cdot \gamma_{\text{scale}}$ for
998 $f \in \{0.1, 0.3, 1, 3, 10\}$, and every quantum feature map is evaluated at angle scalings
999 $\{0.5, 1, 2\}$ of the canonical $[0, \pi]$ encoding. Family-internal selection then operates over
1000 the enlarged candidate sets exactly as before.

1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012

Quantum kernels.

Quantum kernels are defined through an implicit feature map $\phi_Q(\cdot)$ induced by a parameterized quantum circuit, with similarity measured through fidelity-like overlaps [Havlíček et al. \(2019\)](#); [Schuld and Killoran \(2019\)](#); [Schuld \(2021\)](#). For two inputs x and x' , the corresponding kernel entry can be written abstractly as

$$K_Q(x, x') = |\langle \phi_Q(x), \phi_Q(x') \rangle|^2,$$

or equivalently through the overlap of states prepared by input-dependent quantum feature maps.

In the experiments, we evaluate multiple commonly used fidelity-based feature-map families chosen to represent distinct inductive biases: ZZ-type maps, PauliXZ-type maps, and Z-map variants. Concretely, the candidate set includes two ZZ configurations with different repetition depths, one PauliXZ configuration, and one Z-map configuration. We emphasize that the study is not attempting to exhaust the quantum-kernel design space; rather, it compares a representative set of commonly used fidelity-based constructions under a tightly controlled protocol. This framing is consistent with recent work arguing that empirical conclusions in QML are often regime-dependent and highly sensitive to protocol choices [Schuld et al. \(2021\)](#); [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#).

3.6 Model selection and repeated-run protocol

A major source of ambiguity in robustness studies is the selection criterion. Strong ID performance does not necessarily imply strong OOD performance, and mixing the two objectives can make results difficult to interpret. We therefore use two complementary family-internal selection modes.

1059 ***Best-by-OOD.***

1060

1061 Within each family (classical or quantum), we select the configuration with the highest
1062 mean OOD balanced accuracy. We then compare the selected classical model against
1063 the selected quantum model. This is the robustness-oriented view and constitutes the
1064 primary comparison in the paper.

1067

1068 ***Best-by-ID.***
1069

1070 Within each family, we instead select the configuration with the highest mean ID
1071 balanced accuracy and then compare the selected classical and quantum models. This
1072 is the separability-oriented view and is reported separately to avoid conflating ID
1073 accuracy with OOD robustness.

1074 This separation is methodological, not cosmetic. It acts as a guardrail against over-
1075 interpreting ID gains as robustness gains and helps disentangle two distinct questions:
1076 which kernel family yields the strongest in-distribution class separation, and which
1077 yields the strongest out-of-distribution generalization.

1078 Selection is only one source of ambiguity in robustness studies; a second is result
1079 stability. Each configuration is therefore repeated across 15 runs, generated by com-
1080 bining five q-split seeds with three model seeds under fixed setting definitions (shift
1081 variant, master seed, and size regime). Within a setting, both classical and quan-
1082 tum candidates are evaluated on identical split definitions and summarized by mean
1083 $\pm$ standard deviation over the repeated runs, so that every reported comparison is
1084 variability-aware while preserving strict comparability between kernel families.

1095

1096 ***Statistical testing.***
1097

1098 Family-level comparisons are tested formally with paired two-sided Wilcoxon signed-
1099 rank tests over the per-setting deltas (the settings are the paired units), separately for
1100 each classifier and selection view, with Holm correction applied within each classifier
1101 across the reported comparisons. We complement these tests with the per-setting

effect-size proxy defined below, so that both the consistency of the sign pattern and the magnitude relative to run-to-run noise are visible.

3.7 Evaluation metrics

Our primary metric is balanced accuracy, reported separately on ID and OOD test sets:

$$\text{BAcc}_{\text{ID}}, \quad \text{BAcc}_{\text{OOD}}.$$

Balanced accuracy is preferred because it remains informative under moderate class imbalance and is straightforward to interpret in both settings.

The ID-to-OOD degradation

$$\text{drop} = \text{BAcc}_{\text{ID}} - \text{BAcc}_{\text{OOD}}$$

is recorded for every configuration and available in the artifact; because it conflates OOD behavior with the corresponding ID baseline, we treat OOD balanced accuracy as the primary inferential target and use drop only as a secondary descriptor.

To compare families after selection, we compute family deltas:

$$\Delta_{\text{OOD}} = \text{BAcc}_{\text{OOD}}^{(Q)} - \text{BAcc}_{\text{OOD}}^{(C)}, \quad (3)$$

$$\Delta_{\text{ID}} = \text{BAcc}_{\text{ID}}^{(Q)} - \text{BAcc}_{\text{ID}}^{(C)}. \quad (4)$$

Positive values indicate an advantage for the selected quantum model relative to the selected classical model.

1151 Because average gains can be misleading when variability is large, we also report
1152 an effect-size-to-variability proxy:

1153

1154

1155

1156

1157

1158

1159

1160

1161

1162

1163

1164

1165

1166

1167

1168

1169

1170

1171

1172

1173

1174

1175

1176

1177

1178

1179

1180

1181

1182

1183

1184

1185

1186

1187

1188

1189

1190

1191

1192

1193

1194

1195

1196

$$\text{effect} = \frac{\mu_Q - \mu_C}{\sqrt{\sigma_Q^2 + \sigma_C^2}},$$

where μ_Q, μ_C and σ_Q, σ_C are the repeated-run means and standard deviations for the metric of interest. Values greater than 1 indicate that the observed mean gain exceeds the combined variability scale, which provides a simple and interpretable robustness check on whether an apparent gain is large relative to run-to-run noise.

1167 *Probabilistic metrics.*

1169 For the Gaussian process classifier we additionally report, on both ID and OOD
1170 splits, the negative log-likelihood (log-loss), the Brier score, the expected calibration
1171 error [Guo et al. \(2017\)](#) (ECE, 15 equal-width bins), and the mean predictive entropy.
1172 These quantities characterize how the quality and honesty of predictive uncertainty
1173 change under shift: a well-behaved probabilistic model should become *more* uncertain
1174 on drifted inputs (rising entropy) while its calibration degrades as little as possible
1175 (bounded ECE increase), the desideratum established by the benchmark of Ovadia et
1176 al. [Ovadia et al. \(2019\)](#).

1184 **3.8 Kernel-geometry descriptors**

1186 To move from measuring outcomes to explaining them, we compute geometry descrip-
1187 tors of the Gram matrices that every kernel induces in every setting and dimension.
1188 All descriptors are computed from the very kernel matrices consumed by the classi-
1189 fiers (unit diagonal by construction), so the analysis describes exactly the objects the
1190 models see.

<i>Effective rank.</i>	1197
	1198
For the training Gram matrix K_{tr} with eigenvalues $\lambda_1 \geq \dots \geq \lambda_n \geq 0$ normalized	1199
to $p_i = \lambda_i / \sum_j \lambda_j$, the effective rank is $\exp(-\sum_i p_i \log p_i)$ Roy and Vetterli (2007) .	1200
	1201
It measures how many directions of feature space the kernel actually uses: a cosine-	1202
normalized linear kernel on low-dimensional embeddings is near rank one, while fidelity	1203
kernels can spread mass over many directions.	1204
	1205
	1206
	1207
<i>Centered kernel-target alignment.</i>	1208
	1209
For a kernel block K on a labeled split with labels $y \in \{-1, +1\}^n$, the centered kernel-	1210
target alignment (KTA) is $\langle K_c, yy^\top \rangle_F / (n \ K_c\ _F)$, where K_c is the doubly centered	1211
kernel Cristianini et al. (2001) ; Cortes et al. (2012) . We evaluate KTA on the training,	1212
ID, and OOD splits separately; the ID-to-OOD change in alignment measures whether	1213
the label structure remains visible to the kernel’s geometry after shift—the quantity	1214
we will call <i>alignment survival</i> .	1215
	1216
	1217
	1218
	1219
	1220
<i>Concentration and cross-split similarity.</i>	1221
	1222
We record the mean and standard deviation of off-diagonal kernel entries within the	1223
training block and of the cross blocks between each test split and the training set,	1224
together with the ratio of mean OOD-to-train and ID-to-train similarities, which mea-	1225
sures how novel the OOD inputs look under each kernel’s similarity notion; these	1226
cross-split means are the elementary statistics underlying kernel two-sample tests such	1227
as MMD Gretton et al. (2012) , used here descriptively rather than as hypothesis tests.	1228
	1229
	1230
	1231
	1232
	1233
<i>Geometric difference.</i>	1234
	1235
Between each classical training kernel K_C and each quantum training kernel K_Q	1236
we compute the regularized geometric difference $g(K_C \rightarrow K_Q) = \ \sqrt{K_Q}(K_C +$	1237
$\lambda n I)^{-1} \sqrt{K_Q}\ _\infty^{1/2}$ of Huang et al. Huang et al. (2021) , over a sweep of relative regu-	1238
larizers λ . Large values indicate that the quantum geometry can deviate substantially	1239
	1240
	1241
	1242

1243 from what the classical kernel expresses; the quantity bounds the *potential* for a
1244 performance separation but not its sign.

1246

1247 *Mechanistic analyses.*

1248

1249 These descriptors feed two complementary analyses. First, a diagnostic analysis at
1250 the setting level asks which descriptors of the Best-by-OOD-selected kernels sepa-
1251 rate classical-favorable from quantum-favorable settings (Mann–Whitney AUC over
1252 settings) and correlate with the observed OOD delta. Second, a within-setting anal-
1253 ysis asks whether geometry predicts performance across *all* kernel–dimension cells
1254 of a setting: for every setting, classifier, and q-split seed we compute the Spearman
1255 rank correlation between a descriptor (effective rank, OOD alignment) and OOD bal-
1256 anced accuracy over the 30–55 evaluated cells, and then examine the distribution of
1257 these correlations across settings and datasets. A descriptor that yields consistently
1258 positive within-setting correlations across datasets, drift mechanisms, and classifiers
1259 constitutes a dataset-general regularity rather than a benchmark artifact.

1267

1268

1269 **3.9 Experimental grid and reporting**

1270

1271 The EMBER evaluation grid comprises 18 principal settings formed by the Cartesian
1272 product of:

1273

- 1274 • two shift variants ($m1$, $m2$),
- 1275 • three master seeds,
- 1276 • and three training-size regimes.

1279

1280 The three size regimes are:

1281

1282

1283

1284 $(1000, 500, 500)$, $(2000, 1000, 1000)$, $(4000, 1800, 1800)$,

1285

1286

1287 corresponding to $(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$.

1288

Table 1 Compact summary of the controlled experimental protocol. The repeated-run design is generated by combining five split seeds with three model seeds, yielding 15 runs per evaluated configuration under fixed setting definitions.

Component	Specification
Datasets	EMBER; UNSW-NB15 (DoS, Reconnaissance); ToN-IoT (Scanning)
Shift constructions	EMBER: $m1, m2$; network: m2-centroid, natural drift
Master seeds	42, 123, 999
Split seeds	42, 123, 999, 7, 2024
Model seeds	42, 123, 999
Runs per configuration	$5 \times 3 = 15$
Size regimes ($n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}}$)	(1000, 500, 500), (2000, 1000, 1000), (4000, 1800, 1800)
Dimensions	$d \in \{4, 6, 8, 10, 12\}$
Classifiers	SVC (precomputed); Laplace-GP classifier (precomputed)
Shared preprocessing	MaxAbsScaler + TruncatedSVD + angle mapping
Classical-original family	Linear, RBF(scale)
Classical-extended family	+ poly(2,3), Laplacian, Matérn(3/2, 5/2), median heuristic
Quantum family	ZZ (1 rep), ZZ (2 reps), PauliXZ (1 rep), Z-map
Bandwidth sweep	RBF γ factors $\{0.1, 0.3, 1, 3, 10\}$; angle scales $\{0.5, 1, 2\}$
Selection views	Best-by-OOD, Best-by-ID (family-internal, on run means)
Statistical testing	Paired Wilcoxon over settings + Holm; effect proxy
Principal settings	EMBER $2 \times 3 \times 3 = 18$; network $3 \times 2 \times 3 \times 3 = 54$; total 72

The network-flow grid applies the same structure to the three scenarios (UNSW-DoS, UNSW-Reconnaissance, ToN-IoT-Scanning), each under two shift mechanisms (m2-centroid and natural drift), three master seeds, and the same three size regimes, yielding 54 principal settings. Together the study comprises 72 principal settings, each evaluated across the dimension sweep and repeated over 15 runs, for both classifiers. The bandwidth-symmetric sweep enlarges the kernel candidate sets on the full EMBER grid and on the natural-drift network settings.

For clarity and reproducibility, Table 1 summarizes the main experimental protocol used throughout the paper, including the shift constructions, controlled seed structure, size regimes, dimensionality sweep, kernel families, classifiers, and repeated-run design.

1335 Within each setting, we evaluate both kernel families across the shared prepro-
1336 cessing grid and the dimension sweep $d \in \{4, 6, 8, 10, 12\}$. This produces both a
1337 setting-level comparison, used for the main Best-by-OOD and Best-by-ID analy-
1338 ses, and a broader dimension-level view used to characterize how representation size
1339 interacts with kernel-family behavior.

1343 For reproducibility, split definitions are fixed within each setting, random seeds are
1344 controlled, and all key quantities are reported as mean $\pm$ standard deviation across
1345 repeated runs. Throughout the paper, we deliberately avoid claims of universal dom-
1346 inance or complexity-theoretic quantum advantage. The object of study is narrower
1347 and more empirical: under a fixed classifier and a fixed evaluation protocol, does
1348 changing only the kernel alter robustness under distribution shift in security-relevant
1349 tasks—and which property of the kernel is responsible?

1356 4 Results

1359 4.1 Overview

1361 We report results under the controlled kernel-swap protocol described in Section 3. All
1362 comparisons share the same classifier, preprocessing pipeline, and ID/OOD split defi-
1363 nitions within a setting; the only varying component is the kernel, and every number
1364 below is a mean over 15 repeated runs. The results are organized as an argument in four
1365 movements. First, on EMBER, the quantum family is compared against the customary
1366 linear+RBF baselines (Section 4.2). Second, the classical reference set is strength-
1367 ened to the extended family, which changes the verdict on one modality (Section 4.3).
1368 Third, both families receive the same bandwidth-tuning freedom, which changes it
1369 back (Section 4.4). Fourth, the protocol is repeated on three network-flow scenarios
1370 including natural drift (Section 4.5). Section 4.6 then explains the entire pattern with
1371 a single geometric mechanism, and Section 4.7 examines predictive uncertainty under
1372 shift.

4.2 EMBER versus the customary baselines: a near-uniform	1381
quantum advantage	1382
	1383
	1384
We begin with the comparison that dominates the literature: fidelity-based quantum	1385
kernels against linear and RBF baselines. For each setting and classifier, family-internal	1386
Best-by-OOD selection picks the configuration with the highest mean OOD balanced	1387
accuracy on the run means, and the selected representatives are compared directly.	1388
	1389
Table 2 (upper panel, first and third rows) reports the outcome.	1390
	1391
	1392
The advantage is close to uniform. Under SVC the selected quantum kernel	1393
improves OOD balanced accuracy in 17 of 18 settings (mean $\Delta_{\text{OOD}} = +0.037$, Holm-	1394
adjusted Wilcoxon $p = 7.6 \times 10^{-5}$), with the gain exceeding the combined run-to-run	1395
variability scale in 12 settings; under the GP classifier the corresponding figures	1396
are 15 of 18, $+0.028$, $p = 2.1 \times 10^{-4}$, and 11 settings. The few classical-favorable	1397
settings are marginal: all occur at the largest training size under the $m1$ construc-	1398
tion, with $ \Delta_{\text{OOD}} \leq 0.010$ and effect sizes below one, i.e., within run-to-run noise.	1399
	1400
The dimension-level view confirms that this is not a best-of-grid artifact: comparing	1401
the best quantum against the best classical candidate separately at every setting–	1402
dimension cell, the quantum family is ahead in 88 of 90 cells under SVC and 85 of 90	1403
under the GP classifier (mean $\Delta_{\text{OOD}} \approx +0.05$ in both cases).	1404
	1405
	1406
	1407
	1408
	1409
	1410
	1411
Under Best-by-ID selection (lower panel), the quantum family improves ID bal-	1412
anced accuracy in all 18 settings under both classifiers, with mean gains of $+0.078$	1413
(SVC) and $+0.065$ (GPC) and every setting exceeding the variability scale. We keep	1414
the two views separate throughout: strong ID separability is not evidence of robustness,	1415
and the following subsections show that the two properties respond very differently to	1416
stronger classical baselines. Figure 2 previews the resulting arc: the per-setting OOD	1417
deltas of this subsection (first column of each panel) collapse toward zero against	1418
the extended family (second column), partially recover under bandwidth symmetry	1419
(third), and re-emerge on network flows (fourth).	1420
	1421
	1422
	1423
	1424
	1425
	1426




Fig. 2 The four movements of the study. Per-setting Best-by-OOO deltas (quantum minus best classical, mean over 15 runs), for SVC (left) and the GP classifier (right); black bars mark medians. Against linear+RBF the quantum advantage is large and nearly uniform; the extended classical family collapses it on EMBER; symmetric bandwidth tuning partially restores it; on network flows (54 settings, including natural drift) it persists against the extended family, uniformly so under the GP classifier.

Table 2 EMBER headline results (18 settings, 15 runs each). Within each classifier (GPC = Laplace Gaussian process classifier), family-internal selection compares the best quantum candidate against the best candidate of the stated classical reference set. p -values are two-sided paired Wilcoxon over settings, Holm-corrected within classifier; the last column counts settings whose quantum-favorable gain exceeds the combined run-to-run variability scale.

Selection	Classifier	Classical ref.	Q wins	Mean Δ	Median Δ	p (Holm)	Eff. > 1
Best-by-OOO	SVC	linear+RBF	17/18	+0.0374	+0.0426	7.6×10^{-5}	12
	SVC	extended	15/18	+0.0044	+0.0064	0.03	1
	GPC	linear+RBF	15/18	+0.0283	+0.0339	2.1×10^{-4}	11
	GPC	extended	11/18	+0.0023	+0.0029	0.26	0
Best-by-ID	SVC	linear+RBF	18/18	+0.0775	+0.0794	3.1×10^{-5}	18
	SVC	extended	18/18	+0.0149	+0.0152	3.1×10^{-5}	11
	GPC	linear+RBF	18/18	+0.0649	+0.0664	3.1×10^{-5}	18
	GPC	extended	18/18	+0.0192	+0.0166	3.1×10^{-5}	11

4.3 The extended classical family neutralizes the OOD advantage on EMBER

The comparison changes character when the classical family is strengthened. Against the extended set—which adds polynomial, Laplacian, and Matérn kernels with train-only median-heuristic scales—the quantum OOD advantage on EMBER collapses by an order of magnitude (Table 2, second and fourth rows): the mean delta falls from

+0.037 to +0.004 under SVC and from +0.028 to +0.002 under the GP classifier, the gain exceeds the variability scale in at most one setting, and under the GP classifier the comparison is no longer statistically significant (Holm-adjusted $p = 0.26$). The winner of this neutralization is overwhelmingly one kernel: the median-heuristic Laplacian is the Best-by-OOD selection of the extended classical family in 13 of 18 settings under SVC and 14 of 18 under the GP classifier, almost always at the largest embedding dimension.

Two aspects of this result deserve emphasis. First, it is an honest negative result for the quantum-advantage narrative as usually framed: the OOD gains reported against linear and RBF baselines—here and, we suspect, elsewhere in the literature—measure the weakness of those baselines as much as the strength of the quantum family. Second, the *ID* advantage does not collapse: even against the extended family, the quantum kernels improve ID balanced accuracy in 18 of 18 settings under both classifiers ($p = 3.1 \times 10^{-5}$), with the gain exceeding the variability scale in 11 settings. Whatever the quantum feature maps contribute to in-distribution separability, no classical kernel in our set replicates it; what the Laplacian replicates is specifically the out-of-distribution behavior. Section 4.6 shows that this is not a coincidence: the Laplacian is precisely the classical kernel whose Gram-matrix geometry most resembles that of the fidelity maps.

4.4 Bandwidth symmetry restores the advantage

The comparison of Section 4.3 pins every quantum feature map to the canonical $[0, \pi]$ encoding while granting the classical family kernels whose length scales adapt to the data. Bandwidth, however, is as decisive for fidelity kernels as it is for the RBF kernel [Shaydulin and Wild \(2022\)](#); [Canatar et al. \(2023\)](#). When both families receive the same tuning axis—RBF gamma factors $\{0.1, 0.3, 1, 3, 10\}$ against quantum angle scales $\{0.5, 1, 2\}$ —the picture inverts again (15 runs per setting, as in the main grid). Under

SVC, the quantum family beats the bandwidth-swept extended classical family in 16 of 18 EMBER settings (mean $\Delta_{\text{OOD}} = +0.008$, Holm-adjusted $p = 1.0 \times 10^{-3}$); under the GP classifier, in 12 of 18 ($+0.006$, $p = 0.03$); the gain exceeds the run-to-run variability scale in 4 settings under each classifier. The mechanism is visible in the selections: the angle-scaled variants account for 25 of the 36 Best-by-OOD quantum choices, confirming that the fixed-bandwidth comparison of Section 4.3 had the quantum family playing with one hand tied. On the network-flow scenarios under natural drift, the same symmetric sweep (one run per setting) leaves the quantum advantage intact and significant under both classifiers (24 of 27 settings, Holm-adjusted $p = 5.2 \times 10^{-5}$ for the GP classifier; 19 of 27, $p = 9.8 \times 10^{-4}$ for SVC).

The methodological moral generalizes beyond this study: fixed-encoding quantum kernels versus tuned classical kernels is not a protocol-fair comparison, and the bias runs *against* the quantum family. Once tuning freedom is symmetric, the extended classical family no longer neutralizes the quantum OOD advantage.

4.5 Network flows: the advantage survives the extended family and peaks under natural drift

The 54 network-flow settings test whether the EMBER conclusions are properties of one dataset family or of the kernels. They are neither, exactly: the network results are *stronger* for the quantum family than EMBER’s (Table 3). Against linear+RBF, the selected quantum kernel improves OOD balanced accuracy in 54 of 54 settings under both classifiers (Holm-adjusted $p = 6.5 \times 10^{-10}$), with the gain exceeding the variability scale in 35 (SVC) and 44 (GPC) settings. More importantly, against the *extended* classical family—the set that neutralized the advantage on EMBER—the quantum family remains ahead in 54 of 54 settings under the GP classifier (mean $\Delta_{\text{OOD}} = +0.020$, $p = 6.5 \times 10^{-10}$, effect > 1 in 33) and in 41 of 54 under SVC

Table 3 Network-flow headline results (54 settings across UNSW-NB15 DoS/Reconnaissance and ToN-IoT Scanning, two shift mechanisms, 15 runs each). Columns as in Table 2.

Selection	Classifier	Classical ref.	Q wins	Mean Δ	Median Δ	p (Holm)	Eff. > 1
Best-by-OOD	SVC	linear+RBF	54/54	+0.0256	+0.0264	6.5×10^{-10}	35
	SVC	extended	41/54	+0.0066	+0.0028	3.2×10^{-5}	7
	GPC	linear+RBF	54/54	+0.0344	+0.0338	6.5×10^{-10}	44
	GPC	extended	54/54	+0.0204	+0.0192	6.5×10^{-10}	33
Best-by-ID	SVC	linear+RBF	54/54	+0.0210	+0.0104	6.5×10^{-10}	29
	SVC	extended	39/54	+0.0036	+0.0028	1.5×10^{-4}	11
	GPC	linear+RBF	53/54	+0.0188	+0.0100	6.5×10^{-10}	24
	GPC	extended	40/54	+0.0057	+0.0044	7.9×10^{-7}	14

($p = 3.2 \times 10^{-5}$). The Laplacian is again the dominant classical challenger, and again it is not enough.

The breakdown by shift mechanism sharpens the picture. Under SVC, the quantum-versus-extended comparison is essentially tied on the synthetic centroid construction (15 of 27 settings) but nearly uniform under natural drift (26 of 27): the advantage is largest precisely on the drift mechanism that was not constructed by us and most resembles deployment conditions. Under the GP classifier the advantage is uniform on both mechanisms (27 of 27 each). Read jointly with Section 4.3, the regime map is: on EMBER’s synthetic covariate extremes a heavy-tailed classical kernel can match the quantum family at fixed bandwidth; on drifting network traffic, under a probabilistic classifier, nothing in our classical set does.

4.6 One geometric mechanism explains the pattern

The results so far pose a puzzle with three faces: why do quantum kernels beat linear and RBF baselines so uniformly, why does the Laplacian—alone among our classical kernels—catch them on EMBER, and why does natural network drift favor them again? The geometry descriptors of Section 3.8 answer all three at once.

Table 4 Within-setting Spearman correlation between kernel-geometry descriptors and OOD balanced accuracy, summarized per dataset and classifier over all setting-seed units (five q-split seeds per setting; 30–55 kernel–dimension cells per unit). “> 0” is the fraction of units with a positive correlation.

Dataset	Classifier	Units	Effective rank		KTA on OOD	
			median ρ	> 0	median ρ	> 0
EMBER	GP classifier	90	0.78	99%	0.68	89%
EMBER	SVC	90	0.83	98%	0.65	89%
ToN-IoT Scanning	GP classifier	90	0.84	100%	0.54	99%
ToN-IoT Scanning	SVC	90	0.75	100%	0.59	99%
UNSW-NB15 DoS	GP classifier	90	0.55	98%	0.51	96%
UNSW-NB15 DoS	SVC	90	0.55	90%	0.42	96%
UNSW-NB15 Recon.	GP classifier	90	0.26	98%	0.57	94%
UNSW-NB15 Recon.	SVC	90	0.19	83%	0.46	99%

A *cross-dataset regularity.*

Within every setting we correlate, across all evaluated kernel–dimension cells, each geometry descriptor with OOD balanced accuracy (Table 4 and Figure 3). The alignment-survival quantity—centered kernel-target alignment evaluated on the OOD split—predicts OOD accuracy within settings essentially everywhere: the within-setting Spearman correlation is positive in 89–99% of the 360 setting-seed units, with median ρ between 0.42 and 0.74, across all four datasets, both classifiers, and all three drift mechanisms. Effective rank of the training kernel—a train-only quantity—is an equally strong predictor on EMBER and ToN-IoT (median $\rho \approx 0.75$ –0.84), moderate on UNSW-DoS, and weak (though still predominantly positive) on UNSW-Reconnaissance. The natural reading is that high effective rank is the structural precondition—a kernel whose geometry uses only one or two directions has nothing with which to track label structure into the tails—while alignment survival is the proximate predictor of OOD performance.




Fig. 3 The mechanism is a cross-dataset regularity. Distribution of within-setting Spearman correlations between each geometry descriptor and OOD balanced accuracy, over all setting-seed units, per dataset and classifier. Alignment survival (left) is positively correlated with OOD accuracy essentially everywhere; effective rank (right) is strong on EMBER and ToN-IoT and weaker, though still predominantly positive, on the UNSW scenarios.

A dose-response gradient within the classical family.

The mechanism makes a prediction that can be tested without any quantum kernel at all: classical kernels with higher effective rank should retain label alignment better under shift. Table 5 and Figure 4 confirm a clean monotone gradient across the seven classical kernels (18 EMBER settings $\times$ 5 dimensions each): as mean effective rank rises from 1.19 (linear) through 2.76 (RBF) and 6–8 (Matérn) to 17.7 (Laplacian), the mean ID-to-OOD change in alignment rises monotonically from -0.017 to $+0.019$, and the fraction of cells where alignment *improves* under shift doubles from 32% to 66%. This explains Section 4.3 precisely: the Laplacian is the classical kernel that shares the quantum maps’ geometric signature, and it is exactly the kernel that neutralizes their OOD advantage. Conversely, the fidelity maps’ near-uniform advantage over linear and RBF baselines is an advantage over the two lowest-rank geometries in the study.

What the geometry does not do.

Two negative results complete the mechanistic picture, and both are informative. First, the geometric difference of Huang et al. [Huang et al. \(2021\)](#)—the quantity a reader versed in the quantum-kernel literature would reach for first—does not predict which family wins a setting; at moderate regularization it even correlates *negatively* with the




Fig. 4 Effective rank versus alignment survival across both kernel families. Markers are means over the 18 EMBER settings and 5 dimensions; bars span interquartile ranges. The relationship is monotone across both families: the classical kernels that retain alignment under shift are exactly the heavy-tailed ones, and the fidelity maps occupy the high-rank end of the same continuum.

Table 5 Dose-response gradient in the classical family (18 EMBER settings $\times$ 5 dimensions, q-split seed 42). Higher effective rank of the training Gram matrix is associated with better survival—indeed improvement—of kernel-target alignment from ID to OOD.

Kernel	Eff. rank	ΔKTA (OOD-ID)	Improves
Linear	1.19	-0.017	32%
Polynomial (2)	1.29	-0.017	32%
Polynomial (3)	1.41	-0.015	38%
RBF (scale)	2.76	-0.011	38%
Matérn 5/2	6.41	+0.000	46%
Matérn 3/2	7.71	+0.004	47%
Laplacian	17.69	+0.019	66%

observed OOD delta. This is consistent with its definition: g bounds the potential for a separation between the families, not its sign, so a quantum kernel can differ greatly from the classical reference and be worse. Second, we found no reliable fully a-priori selection rule: choosing the quantum configuration by training-set alignment alone




Fig. 5 Predictive uncertainty under shift. Per-setting ID-to-OOD change in mean predictive entropy (left) and expected calibration error (right) for the Best-by-OOD configuration of each family under the GP classifier; black bars mark medians.

does not identify the settings where the family will win. The geometry explains and diagnoses; it does not yet provide a deployment-time oracle, and we flag this honestly as an open problem.

4.7 Predictive uncertainty under shift

The GP classifier’s calibrated probabilities add a complementary robustness lens (Figure 5). Averaged over Best-by-OOD selections, predictive entropy rises from ID to OOD for every family and dataset group—the qualitative behavior one wants from a probabilistic detector facing drift—and rises most for the quantum kernels (e.g., +0.029 versus +0.014–+0.018 for the classical families on EMBER; +0.072 versus +0.041–+0.061 on network flows). Calibration degrades moderately for all families (ECE increases of +0.02–+0.05), with the quantum family’s ECE increase slightly larger than the classical families’. In other words, the quantum kernels’ OOD accuracy advantage does not come with hidden overconfidence—their uncertainty grows under shift at least as fast as the classical kernels’—but neither do they calibrate better; the advantage is in the decision geometry, not in the probability calibration.

1795 5 Discussion

1796

1797 5.1 What is, and is not, a quantum advantage here

1798

1799
1800 The study’s central interpretive move is to replace the question “do quantum ker-
1801 nels beat classical kernels under shift?” with “which property of a kernel governs its
1802 behavior under shift, and who has it?”. Read through that lens, the results decompose
1803 cleanly. The near-uniform advantage over linear and RBF baselines is real, replicable
1804 across four datasets and two classifiers, and statistically unambiguous—but it is an
1805 advantage of *high-effective-rank, alignment-preserving geometry* over the two lowest-
1806 rank geometries in the study, a property the fidelity maps possess natively rather than
1807 exclusively. The Laplacian result makes this precise: one classical kernel with the same
1808 geometric signature closes the OOD gap on EMBER, so the property, not the physics,
1809 is doing the work there. What remains after every control we could devise—extended
1810 baselines, symmetric bandwidth tuning, a probabilistic classifier, natural drift—is a
1811 genuine residual: the ID-separability advantage survives everything on EMBER, and
1812 the OOD advantage survives everything on drifting network traffic. Those two residuals
1813 are, on present evidence, the defensible quantum contribution.

1814
1815 From the perspective of the distribution-shift literature, this reinforces a general
1816 lesson: conclusions about robustness are only as credible as the protocol used to obtain
1817 them [Quiñonero-Candela et al. \(2008\)](#); [Gulrajani and Lopez-Paz \(2021\)](#); [Koh et al. \(2021\)](#); [Yao et al. \(2022\)](#). Our results add a QML-specific corollary: the composition
1828 of the classical reference set and the symmetry of tuning freedom are not implemen-
1829 tation details but the difference between a significant advantage, a null result, and a
1830 significant advantage again (Sections 4.2–4.4). Any single one of those three verdicts,
1831 reported alone, would have been misleading.

1832

1833

1834

1835

1836

1837

1838

1839

1840

5.2 The ID–OOD dissociation is informative

The contrast between the two selection views is sharper in this study than in its predecessors, and it now has an interpretation. The quantum ID-separability advantage is universal (18 of 18 EMBER settings against every classical reference, with effects exceeding run-to-run noise in most) and is not replicated by any classical kernel we evaluate, including the Laplacian. The OOD advantage, by contrast, is contingent: it depends on the reference family, the bandwidth protocol, and the drift mechanism. This dissociation is exactly what the mechanism predicts. In-distribution separability rewards expressive geometry *per se*; robustness rewards geometry whose label alignment *survives* the shift, which is a joint property of the kernel and the shift mechanism. A kernel family can therefore dominate the first criterion while merely competing on the second—and mistaking one for the other is precisely the error the two-view protocol is designed to prevent [Gulrajani and Lopez-Paz \(2021\)](#). The dissociation also speaks to the “accuracy on the line” literature, which documents a strong empirical coupling between ID and OOD accuracy across models and the settings where it breaks [Miller et al. \(2021\)](#): at the level of kernel families under our shift constructions, the coupling breaks systematically, and the geometry identifies the axis along which it does.

5.3 The mechanism in context

The alignment-survival regularity connects several strands of literature that have so far run in parallel. The bandwidth results extend the findings of Shaydulin and Wild [Shaydulin and Wild \(2022\)](#) and Canatar et al. [Canatar et al. \(2023\)](#) from in-distribution generalization to robustness under shift, and give them a comparative twist: bandwidth is not merely important for quantum kernels, it is important *symmetrically*, and asymmetric tuning protocols systematically bias quantum-versus-classical comparisons. The failure of the geometric difference [Huang et al. \(2021\)](#) to predict OOD

outcomes sharpens its interpretation rather than contradicting it: g was derived as a bound on the potential for in-distribution separations, and our results show empirically that potential-for-difference and benefit-under-shift are distinct quantities—large g is compatible with large losses. Finally, the dose-response gradient within the classical family (Table 5) grounds the mechanism independently of anything quantum: kernels with heavier spectral tails hold their label alignment into the tails of the data distribution, a statement about kernels generally, for which the quantum maps are simply the most extreme instances in our study.

The high-rank finding must also be reconciled with two spectral results that superficially point the other way. Kübler et al. argue that quantum kernels generalize in-distribution only when the relevant spectral mass is concentrated in few components Kübler et al. (2021), and Thanasilp et al. show that fidelity kernels concentrate exponentially—toward effectively rank-one, uninformative Gram matrices—as qubit counts grow Thanasilp et al. (2024). There is no contradiction, but there is a genuine tension that our results locate precisely. Both results concern the asymptotic or i.i.d. regime; our effective ranks (roughly 10–150 on training sets of 10^3 – 4×10^3 samples at 4–12 qubits) sit far below the concentration regime, and the studied maps demonstrably retain in-distribution power (the 18-of-18 ID advantage). Within that benign band, the quantity that governs *robustness* is how much of the spectrum survives with its label alignment intact under shift—and there, more rank helps. The bandwidth sweep makes the trade-off operational: angle scaling moves kernels along the rank axis, and the selected scales are those that raise rank without approaching the concentration pathology. Alignment-preserving rank is therefore best understood as bounded from both sides—enough spectral richness to track label structure into the tails, not so much that the kernel concentrates into noise.

We emphasize what the mechanism does not yet deliver. It is explanatory and diagnostic—given the Gram matrices of a setting, it accounts for which kernels will

generalize under the shift—but our attempts to convert it into an a-priori selection rule from training data alone failed (Section 4.6). Alignment on OOD inputs requires OOD labels; effective rank is train-only but is a precondition, not a guarantee. Turning alignment survival into a deployable model-selection signal, perhaps via unlabeled OOD data or distributionally robust surrogates, is in our view the most interesting open problem this study raises.

5.4 Implications for security machine learning practice

From a practical perspective, the paper suggests three takeaways for security ML. The first is substantive: for OOD-sensitive kernel selection, prefer kernels whose label-aligned geometry survives shift. Concretely, if quantum kernels are unavailable or too costly, a median-heuristic Laplacian kernel is a far stronger robustness baseline than the RBF kernel that pipelines default to—on our EMBER settings it recovers most of what the fidelity maps offer. Where quantum kernels are on the table, they warrant serious consideration for drifting-traffic workloads, where their advantage was largest and most uniform in our study.

The second is methodological. Independently of the preferred kernel family, the protocol used here provides a reusable blueprint for robust evaluation in security settings:

1. keep the classifier and preprocessing fixed, and swap only the kernel;
2. define explicit OOD splits tied to interpretable shift mechanisms, including at least one natural drift mechanism where the data permit;
3. give every family the same hyperparameter freedom, bandwidth included;
4. repeat runs, report variability, and test family-level claims formally;
5. and separate model-selection criteria according to the actual objective (robustness vs. separability).

Security ML frequently suffers from comparisons in which model class, tuning effort, and evaluation conditions change simultaneously; the three-verdict sequence of Sections 4.2–4.4 is a concrete demonstration of how much the conclusion depends on getting this right.

The third concerns diagnostics: the geometry descriptors are cheap to compute from Gram matrices that the pipeline already builds, and the alignment-survival analysis can be run post hoc on any deployed kernel system with a labeled drift sample. Practitioners should nonetheless resist reading these results as a blanket recommendation to replace classical kernels; the responsible recommendation is to test the candidate families under the expected shift regime, using exactly the kind of protocol described here.

5.5 Limitations and threats to validity

Several limitations define the scope of our conclusions.

Shift operationalization.

Our OOD splits cover three reproducible mechanisms—within-class sparsity extremes, train-geometry extremes, and natural regime drift on network traffic—but not the full space of real-world drift processes. Explicit temporal splits on timestamped malware, family-based drift, and adversarially induced shift remain outside the study’s scope, and the alignment-survival regularity should be re-tested under them rather than assumed.

Kernel and encoding coverage.

The classical family now spans seven kernels plus a bandwidth sweep, and both families receive symmetric tuning freedom, but neither candidate set is exhaustive: the quantum side is restricted to four commonly used fidelity feature maps at low qubit counts, and trained or task-adapted encodings, projected kernels, and deeper maps are

not evaluated. The mechanism gives a concrete prediction for such extensions—what should matter is their Gram-matrix geometry, not their provenance—which future work can falsify. At substantially larger qubit counts, exponential kernel concentration [Thanasilp et al. \(2024\)](#) would erode the alignment-preserving rank the mechanism rewards, so our conclusions should not be extrapolated beyond the low-qubit regime studied here.

Selection over candidate sets.

Best-by-OOD and Best-by-ID are selection procedures over finite candidate sets, applied identically to both families and evaluated on run means with formal paired testing. Residual sensitivity to the candidate space and search granularity nonetheless remains, and nested or cross-validated selection protocols would strengthen the analysis further.

Simulation, not hardware.

All quantum kernels are computed by exact statevector simulation. The study therefore characterizes the geometry of fidelity kernels, not the behavior of noisy devices: finite sampling of kernel entries, hardware noise, and error mitigation would all perturb the Gram matrices [Wang et al. \(2021\)](#), and exponential concentration implies that with polynomially many shots the very spectral richness our mechanism rewards becomes expensive to resolve at larger scales [Thanasilp et al. \(2024\)](#); the effect of either on alignment survival is unknown and worth measuring. We make no claim about hardware feasibility, resource advantage, or quantum speedup—the only rigorously established quantum-classical learning separation concerns a specifically constructed problem, not natural data [Liu et al. \(2021\)](#). At the dimensions studied, the fidelity kernels are also classically simulable [Schreiber et al. \(2023\)](#); [Jerbi et al. \(2024\)](#), which is precisely what makes them available as practical kernel constructions today.

2071 *Dataset scope and feature caveats.*

2072

2073 Four benchmark scenarios across two modalities is a substantial widening relative
2074 to single-dataset studies, but security is broader still: dynamic malware analysis,
2075 encrypted traffic, and non-security modalities remain untested. Within the network
2076 scenarios, the feature sets inherited from the public benchmarks retain port num-
2077 bers and protocol indicators, which are legitimate but known shortcut-prone features
2078 in NIDS benchmarks [Engelen et al. \(2021\)](#); because every kernel receives identical
2079 features, this affects absolute accuracies, not the family comparisons. Finally, the
2080 mechanism’s effective-rank component is itself regime-dependent (strong on EMBER
2081 and ToN-IoT, weak on UNSW-Reconnaissance), a gradation we report rather than
2082 smooth over.

2083

2084

2085

2086

2087 **5.6 Conservative interpretation**

2088

2089 In light of these limitations and the broader QML benchmarking literature, the
2090 appropriate interpretation of the paper is deliberately conservative. We do not claim
2091 universal superiority, complexity-theoretic quantum advantage, or hardware relevance.
2092 We claim, with formal statistical support, three narrower things: that fidelity-based
2093 quantum kernels beat the customary baselines under shift essentially wherever we
2094 looked; that this fact is explained by a geometric property a well-chosen classical
2095 kernel can partially share; and that under symmetric protocols the quantum family
2096 retains a significant OOD advantage on drifting network traffic and a universal ID-
2097 separability advantage on EMBER. This regime-aware view is consistent with recent
2098 calls for careful benchmarking in QML [Bowles et al. \(2024\)](#); [Schnabel and Roth \(2025\)](#)
2099 and, we believe, is scientifically stronger than either an unqualified advantage claim or
2100 an unqualified debunking: it says what the advantage is made of, where it lives, and
2101 how to check for it.

2102

2103

2104

2105

6 Conclusion

This paper asked why kernel families behave differently under distribution shift, using the quantum-versus-classical comparison as both motivation and stress test. A controlled kernel-swap protocol—classifier, preprocessing, splits, selection logic, and tuning freedom held fixed while only the kernel varies—was instantiated on four security benchmarks spanning two modalities and three drift mechanisms, two classifier families including a probabilistic one, and 72 principal settings with 15 repeated runs and formal paired testing.

The empirical arc has four movements. Fidelity-based quantum kernels beat the customary linear and RBF baselines under shift nearly everywhere. A stronger classical family neutralizes that advantage on EMBER, revealing that the operative property is geometric rather than quantum. Symmetric bandwidth tuning restores it, revealing that the neutralization itself depended on an asymmetric protocol. And on drifting network traffic the advantage is uniform under the probabilistic classifier, while the in-distribution separability advantage survives every control on EMBER. A single mechanism accounts for the whole pattern: kernels whose training geometry has high effective rank retain their kernel-target alignment under shift, and alignment survival predicts out-of-distribution accuracy within settings in 89–99% of 360 setting-seed units across all four datasets. The mechanism also explains the exceptions—the competitive classical kernels are exactly those sharing the geometric signature—and makes falsifiable predictions for kernel constructions we did not test.

The contribution is therefore threefold. Empirically, the study provides the broadest controlled, variability-aware regime map of quantum-versus-classical kernel behavior under shift that we are aware of, including the honest negative results (the Laplacian neutralization, the non-predictiveness of the geometric difference, the absence of an a-priori selection rule). Mechanistically, it converts an observation into an explanation with cross-dataset support. Methodologically, it leaves behind a fully

reproducible protocol—and a cautionary demonstration that the quantum-versus-classical verdict flips twice as the protocol is made progressively fairer, which we offer as a concrete standard for future QML benchmarking.

The most important open problems are two. Practically, turning alignment survival into a deployment-time selection signal without OOD labels would convert the diagnosis into a tool; the security community’s label-free drift detectors [Yang et al. \(2021\)](#) suggest this is not hopeless. Theoretically, the dose-response relationship between spectral tails and alignment survival calls for a formal account, and the ingredients exist: task-model alignment theory explains in-distribution kernel generalization through spectrum–target compatibility [Bordelon et al. \(2020\)](#); [Canatar et al. \(2021\)](#), and recent minimax analyses characterize RKHS regression under covariate shift with bounded density ratios [Ma et al. \(2023\)](#). Extending that line to the survival of aligned spectral mass under the shift models studied here—asking under which conditions high effective rank provably protects kernel-target alignment—is a concrete research program rather than a gap. Answering either question would benefit classical and quantum kernel methods alike—which is, in the end, the point: the productive question is not whether one family dominates, but which similarity geometry aligns best with the shift regime one actually faces.

Appendix A Additional Results

This appendix reports the full setting-level Best-by-OOD comparisons between the selected quantum kernel and the best kernel of the *extended* classical family (the strongest comparison in the study), for both classifiers and all 72 principal settings. Every value is a mean $\pm$ standard deviation over 15 repeated runs; Effect denotes $\Delta_{\text{OOD}}/\sqrt{\sigma_C^2 + \sigma_Q^2}$. Size labels denote $(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$: S = (1k, 500, 500), M = (2k, 1k, 1k), L = (4k, 1.8k, 1.8k). Shift labels for the network scenarios abbreviate scenario/mechanism: UD = UNSW-NB15 DoS, UR = UNSW-NB15

Reconnaissance, TS = ToN-IoT Scanning; /m2 is the centroid construction and /nat is natural drift. Quantum kernel names carry the angle scale after @ where it differs from 1. All tables are generated directly from the public artifact by `scripts/reporting/make_v2_tables.py`.

Declarations

Availability of data and materials. This study is based on the public EMBER Anderson and Roth (2018), UNSW-NB15 Moustafa and Slay (2015), and ToN-IoT Alsaedi et al. (2020) benchmark datasets. The datasets supporting the conclusions of this article—derived split definitions, shift constructions, experiment manifests, aggregated results, and summary tables—are available in the Zenodo repository, <https://doi.org/10.5281/zenodo.19147649>. All code used to generate the splits, run the classical and quantum kernel experiments, and reproduce every table and figure in this article is part of the same artifact:

- Project name: Controlled Kernel Evaluation under Distribution Shift
- Project home page: https://github.com/roberto-fernandez-barrios/kernel_shift_framework
- Archived version: <https://doi.org/10.5281/zenodo.19147649>
- Operating system(s): platform independent
- Programming language: Python (3.12)
- Other requirements: pinned environment provided (`environment.lock.yml`); qiskit 2.3.1
- License: BSD-3-Clause
- Any restrictions to use by non-academics: none

Competing interests. The authors declare that they have no competing interests.

Funding. The authors received no specific funding for the submitted work.

2255

2256

2257 **Table A1** EMBER, SVC: full Best-by-ODD comparison against the extended classical family.

2258

	Shift	Seed	Size	Classical	Quantum	OOD C	OOD Q	Δ_{OOD}	Eff.
2259	m1	42	S	Laplacian (d4)	PauliXZ (d12)	0.756 ± 0.007	0.758 ± 0.020	+0.002	+0.09
2260	m1	42	M	Laplacian (d12)	ZZ-r1 (d12)	0.753 ± 0.019	0.761 ± 0.007	+0.009	+0.42
2261	m1	42	L	Laplacian (d12)	ZZ-r1 (d12)	0.750 ± 0.014	0.757 ± 0.008	+0.006	+0.38
2262	m1	123	S	Laplacian (d12)	PauliXZ (d12)	0.762 ± 0.019	0.750 ± 0.020	-0.012	-0.43
2263	m1	123	M	Laplacian (d12)	ZZ-r1 (d10)	0.760 ± 0.027	0.756 ± 0.029	-0.003	-0.08
2264	m1	123	L	Laplacian (d12)	ZZ-r1 (d10)	0.751 ± 0.019	0.753 ± 0.014	+0.002	+0.09
2265	m1	999	S	Laplacian (d12)	PauliXZ (d12)	0.767 ± 0.028	0.768 ± 0.022	+0.000	+0.00
2266	m1	999	M	Matérn-5/2 (d12)	ZZ-r1 (d12)	0.755 ± 0.015	0.762 ± 0.013	+0.007	+0.34
2267	m1	999	L	RBF (d12)	ZZ-r1 (d10)	0.763 ± 0.024	0.753 ± 0.005	-0.010	-0.42
2268	m2	42	S	Laplacian (d10)	ZZ-r1 (d10)	0.762 ± 0.012	0.763 ± 0.018	+0.000	+0.02
2269	m2	42	M	Laplacian (d12)	ZZ-r2 (d12)	0.767 ± 0.012	0.784 ± 0.009	+0.017	+1.10
2270	m2	42	L	Laplacian (d12)	ZZ-r2 (d12)	0.783 ± 0.008	0.795 ± 0.011	+0.012	+0.85
2271	m2	123	S	Laplacian (d12)	ZZ-r1 (d12)	0.762 ± 0.023	0.771 ± 0.019	+0.009	+0.30
2272	m2	123	M	Laplacian (d12)	ZZ-r1 (d12)	0.775 ± 0.017	0.783 ± 0.013	+0.007	+0.35
2273	m2	123	L	Laplacian (d12)	ZZ-r2 (d12)	0.785 ± 0.011	0.794 ± 0.007	+0.009	+0.71
2274	m2	999	S	Laplacian (d12)	ZZ-r1 (d12)	0.769 ± 0.014	0.775 ± 0.022	+0.006	+0.25
2275	m2	999	M	Laplacian (d10)	ZZ-r2 (d12)	0.776 ± 0.016	0.780 ± 0.020	+0.004	+0.14
2276	m2	999	L	Laplacian (d12)	ZZ-r2 (d12)	0.787 ± 0.013	0.803 ± 0.009	+0.015	+0.97

2275

2276

2277

2278

2279

2280 **Table A2** EMBER, GPC: full Best-by-ODD comparison against the extended classical family.

2281

	Shift	Seed	Size	Classical	Quantum	OOD C	OOD Q	Δ_{OOD}	Eff.
2282	m1	42	S	Laplacian (d12)	PauliXZ (d12)	0.767 ± 0.018	0.760 ± 0.017	-0.007	-0.29
2283	m1	42	M	Laplacian (d12)	ZZ-r1 (d12)	0.765 ± 0.021	0.766 ± 0.011	+0.002	+0.07
2284	m1	42	L	RBF (d10)	PauliXZ (d12)	0.764 ± 0.015	0.761 ± 0.011	-0.003	-0.16
2285	m1	123	S	Laplacian (d12)	PauliXZ (d12)	0.765 ± 0.016	0.753 ± 0.025	-0.012	-0.41
2286	m1	123	M	Laplacian (d12)	PauliXZ (d12)	0.775 ± 0.035	0.772 ± 0.023	-0.004	-0.09
2287	m1	123	L	Matérn-5/2 (d10)	ZZ-r2 (d12)	0.758 ± 0.017	0.754 ± 0.014	-0.004	-0.18
2288	m1	999	S	Laplacian (d12)	PauliXZ (d12)	0.766 ± 0.028	0.762 ± 0.027	-0.004	-0.11
2289	m1	999	M	Matérn-5/2 (d12)	Z-map (d12)	0.768 ± 0.016	0.776 ± 0.009	+0.008	+0.43
2290	m1	999	L	RBF (d12)	ZZ-r1 (d8)	0.763 ± 0.009	0.754 ± 0.008	-0.008	-0.74
2291	m2	42	S	Laplacian (d12)	ZZ-r1 (d12)	0.762 ± 0.024	0.764 ± 0.015	+0.002	+0.06
2292	m2	42	M	Laplacian (d12)	ZZ-r2 (d12)	0.767 ± 0.013	0.778 ± 0.008	+0.010	+0.68
2293	m2	42	L	Laplacian (d12)	ZZ-r2 (d12)	0.783 ± 0.008	0.793 ± 0.009	+0.010	+0.85
2294	m2	123	S	Laplacian (d12)	ZZ-r2 (d12)	0.767 ± 0.025	0.772 ± 0.015	+0.005	+0.19
2295	m2	123	M	Laplacian (d12)	ZZ-r1 (d12)	0.772 ± 0.020	0.783 ± 0.015	+0.011	+0.44
2296	m2	123	L	Laplacian (d12)	ZZ-r2 (d12)	0.780 ± 0.010	0.791 ± 0.010	+0.012	+0.83
2297	m2	999	S	Laplacian (d12)	ZZ-r1 (d10)	0.767 ± 0.021	0.775 ± 0.021	+0.008	+0.27
2298	m2	999	M	Laplacian (d12)	ZZ-r2 (d12)	0.776 ± 0.016	0.780 ± 0.025	+0.004	+0.13
2299	m2	999	L	Laplacian (d12)	ZZ-r2 (d12)	0.786 ± 0.013	0.798 ± 0.011	+0.012	+0.70

2298

2299

2300

Table A3 Network flows, SVC: full Best-by-OOD comparison against the extended classical family.

Shift	Seed	Size	Classical	Quantum	OOD C	OOD Q	Δ_{OOD}	Eff.
TS/m2	42	S	Laplacian (d12)	ZZ-r1 (d6)	0.941 ± 0.013	0.954 ± 0.009	+0.014	+0.90
TS/m2	42	M	Laplacian (d12)	ZZ-r1 (d10)	0.941 ± 0.014	0.958 ± 0.005	+0.017	+1.15
TS/m2	42	L	Laplacian (d12)	ZZ-r1 (d8)	0.959 ± 0.009	0.964 ± 0.006	+0.005	+0.44
TS/m2	123	S	Laplacian (d12)	ZZ-r1 (d8)	0.943 ± 0.017	0.959 ± 0.011	+0.015	+0.75
TS/m2	123	M	Laplacian (d12)	ZZ-r1 (d6)	0.954 ± 0.007	0.958 ± 0.006	+0.003	+0.35
TS/m2	123	L	Laplacian (d10)	ZZ-r1 (d8)	0.955 ± 0.003	0.959 ± 0.003	+0.003	+0.72
TS/m2	999	S	Laplacian (d10)	ZZ-r1 (d10)	0.936 ± 0.011	0.957 ± 0.003	+0.021	+1.80
TS/m2	999	M	Laplacian (d12)	ZZ-r1 (d8)	0.954 ± 0.009	0.962 ± 0.007	+0.008	+0.72
TS/m2	999	L	Laplacian (d12)	ZZ-r1 (d10)	0.965 ± 0.004	0.969 ± 0.003	+0.003	+0.71
TS/nat	42	S	Laplacian (d8)	ZZ-r1 (d10)	0.987 ± 0.003	0.988 ± 0.005	+0.001	+0.10
TS/nat	42	M	Laplacian (d10)	ZZ-r1 (d12)	0.989 ± 0.004	0.990 ± 0.002	+0.001	+0.24
TS/nat	42	L	Laplacian (d12)	Z-map (d12)	0.989 ± 0.005	0.991 ± 0.003	+0.002	+0.31
TS/nat	123	S	Laplacian (d8)	ZZ-r1 (d12)	0.990 ± 0.005	0.992 ± 0.005	+0.002	+0.36
TS/nat	123	M	Laplacian (d12)	Z-map (d10)	0.989 ± 0.003	0.992 ± 0.002	+0.003	+0.61
TS/nat	123	L	Laplacian (d10)	Z-map (d12)	0.990 ± 0.002	0.991 ± 0.003	+0.002	+0.40
TS/nat	999	S	Laplacian (d12)	Z-map (d12)	0.979 ± 0.006	0.985 ± 0.002	+0.006	+1.05
TS/nat	999	M	Laplacian (d10)	Z-map (d10)	0.990 ± 0.002	0.991 ± 0.004	+0.002	+0.38
TS/nat	999	L	Laplacian (d10)	Z-map (d12)	0.992 ± 0.002	0.992 ± 0.003	+0.000	+0.12
UD/m2	42	S	Laplacian (d12)	PauliXZ (d12)	0.920 ± 0.016	0.919 ± 0.014	-0.001	-0.06
UD/m2	42	M	Laplacian (d12)	PauliXZ (d12)	0.934 ± 0.017	0.937 ± 0.007	+0.003	+0.16
UD/m2	42	L	Laplacian (d12)	PauliXZ (d10)	0.942 ± 0.008	0.940 ± 0.006	-0.002	-0.18
UD/m2	123	S	Laplacian (d12)	PauliXZ (d12)	0.917 ± 0.020	0.926 ± 0.013	+0.008	+0.33
UD/m2	123	M	Laplacian (d12)	PauliXZ (d12)	0.943 ± 0.007	0.940 ± 0.010	-0.004	-0.28
UD/m2	123	L	Laplacian (d12)	Z-map (d12)	0.948 ± 0.003	0.947 ± 0.005	-0.001	-0.20
UD/m2	999	S	Laplacian (d12)	PauliXZ (d12)	0.923 ± 0.023	0.927 ± 0.009	+0.004	+0.16
UD/m2	999	M	Laplacian (d12)	PauliXZ (d12)	0.935 ± 0.013	0.939 ± 0.007	+0.004	+0.26
UD/m2	999	L	Laplacian (d12)	Z-map (d12)	0.947 ± 0.004	0.946 ± 0.004	-0.001	-0.22
UD/nat	42	S	Linear (d4)	ZZ-r2 (d10)	0.781 ± 0.040	0.809 ± 0.017	+0.028	+0.65
UD/nat	42	M	Laplacian (d12)	Z-map (d12)	0.829 ± 0.019	0.838 ± 0.014	+0.009	+0.38
UD/nat	42	L	Laplacian (d12)	Z-map (d12)	0.829 ± 0.011	0.829 ± 0.008	+0.000	+0.01
UD/nat	123	S	Laplacian (d12)	ZZ-r2 (d6)	0.775 ± 0.026	0.818 ± 0.018	+0.043	+1.37
UD/nat	123	M	Laplacian (d12)	ZZ-r1 (d10)	0.790 ± 0.020	0.812 ± 0.007	+0.022	+1.04
UD/nat	123	L	Laplacian (d12)	ZZ-r1 (d12)	0.809 ± 0.012	0.825 ± 0.011	+0.016	+0.97
UD/nat	999	S	Linear (d4)	ZZ-r2 (d6)	0.759 ± 0.051	0.800 ± 0.018	+0.042	+0.77
UD/nat	999	M	Laplacian (d12)	ZZ-r2 (d10)	0.788 ± 0.016	0.809 ± 0.014	+0.021	+0.95
UD/nat	999	L	Laplacian (d12)	ZZ-r1 (d12)	0.804 ± 0.006	0.826 ± 0.013	+0.023	+1.63
UR/m2	42	S	Laplacian (d12)	Z-map (d12)	0.909 ± 0.025	0.909 ± 0.015	-0.000	-0.01
UR/m2	42	M	Laplacian (d12)	PauliXZ (d12)	0.931 ± 0.019	0.935 ± 0.008	+0.004	+0.18
UR/m2	42	L	Laplacian (d12)	PauliXZ (d12)	0.944 ± 0.005	0.943 ± 0.006	-0.001	-0.12
UR/m2	123	S	Laplacian (d12)	PauliXZ (d12)	0.946 ± 0.014	0.942 ± 0.007	-0.004	-0.24
UR/m2	123	M	Laplacian (d12)	PauliXZ (d12)	0.951 ± 0.011	0.948 ± 0.008	-0.004	-0.26
UR/m2	123	L	Laplacian (d12)	Z-map (d12)	0.954 ± 0.007	0.954 ± 0.006	-0.000	-0.03
UR/m2	999	S	Laplacian (d12)	Z-map (d12)	0.915 ± 0.016	0.918 ± 0.014	+0.003	+0.12
UR/m2	999	M	Laplacian (d12)	Z-map (d12)	0.947 ± 0.006	0.944 ± 0.008	-0.003	-0.33
UR/m2	999	L	Laplacian (d12)	Z-map (d12)	0.956 ± 0.003	0.948 ± 0.007	-0.008	-1.03
UR/nat	42	S	Linear (d12)	ZZ-r2 (d6)	0.770 ± 0.019	0.778 ± 0.013	+0.008	+0.36
UR/nat	42	M	Laplacian (d10)	ZZ-r1 (d4)	0.809 ± 0.035	0.815 ± 0.010	+0.007	+0.18
UR/nat	42	L	Laplacian (d10)	ZZ-r1 (d4)	0.812 ± 0.020	0.815 ± 0.008	+0.003	+0.12
UR/nat	123	S	Linear (d12)	ZZ-r1 (d6)	0.795 ± 0.012	0.797 ± 0.018	+0.002	+0.11
UR/nat	123	M	Linear (d10)	ZZ-r1 (d4)	0.800 ± 0.018	0.802 ± 0.024	+0.003	+0.09
UR/nat	123	L	Laplacian (d12)	ZZ-r1 (d4)	0.818 ± 0.018	0.811 ± 0.009	-0.007	-0.34
UR/nat	999	S	Linear (d12)	ZZ-r1 (d6)	0.780 ± 0.020	0.781 ± 0.008	+0.000	+0.02
UR/nat	999	M	Linear (d12)	ZZ-r2 (d6)	0.783 ± 0.014	0.809 ± 0.015	+0.026	+1.26
UR/nat	999	L	Laplacian (d8)	ZZ-r2 (d6)	0.804 ± 0.015	0.814 ± 0.008	+0.010	+0.59

Table A4 Network flows, GPC: full Best-by-OOD comparison against the extended classical family.

Shift	Seed	Size	Classical	Quantum	OOD C	OOD Q	Δ_{OOD}	Eff.
TS/m2	42	S	Laplacian (d12)	ZZ-r1 (d8)	0.906 ± 0.019	0.954 ± 0.011	+0.049	+2.27
TS/m2	42	M	Laplacian (d8)	ZZ-r1 (d12)	0.921 ± 0.004	0.957 ± 0.005	+0.037	+5.70
TS/m2	42	L	Laplacian (d8)	ZZ-r1 (d8)	0.931 ± 0.007	0.964 ± 0.006	+0.034	+3.58
TS/m2	123	S	Laplacian (d8)	ZZ-r1 (d10)	0.910 ± 0.056	0.958 ± 0.006	+0.048	+0.86
TS/m2	123	M	Laplacian (d4)	ZZ-r1 (d12)	0.932 ± 0.007	0.958 ± 0.007	+0.025	+2.59
TS/m2	123	L	Laplacian (d4)	ZZ-r1 (d10)	0.934 ± 0.004	0.960 ± 0.003	+0.026	+4.84
TS/m2	999	S	Laplacian (d6)	ZZ-r1 (d10)	0.918 ± 0.021	0.954 ± 0.009	+0.036	+1.58
TS/m2	999	M	Laplacian (d4)	ZZ-r1 (d10)	0.931 ± 0.002	0.961 ± 0.006	+0.030	+4.49
TS/m2	999	L	Laplacian (d8)	ZZ-r1 (d10)	0.941 ± 0.003	0.968 ± 0.003	+0.027	+6.65
TS/nat	42	S	Laplacian (d8)	ZZ-r1 (d8)	0.981 ± 0.005	0.990 ± 0.002	+0.009	+1.67
TS/nat	42	M	Laplacian (d8)	ZZ-r1 (d10)	0.987 ± 0.002	0.991 ± 0.003	+0.004	+1.08
TS/nat	42	L	Laplacian (d8)	ZZ-r1 (d12)	0.983 ± 0.003	0.992 ± 0.002	+0.009	+2.30
TS/nat	123	S	Laplacian (d8)	ZZ-r1 (d8)	0.981 ± 0.017	0.992 ± 0.003	+0.011	+0.65
TS/nat	123	M	Laplacian (d8)	ZZ-r1 (d10)	0.985 ± 0.004	0.993 ± 0.002	+0.008	+1.72
TS/nat	123	L	Laplacian (d8)	ZZ-r1 (d12)	0.985 ± 0.004	0.992 ± 0.002	+0.007	+1.64
TS/nat	999	S	Laplacian (d10)	ZZ-r1 (d10)	0.965 ± 0.027	0.980 ± 0.015	+0.015	+0.48
TS/nat	999	M	Laplacian (d8)	ZZ-r1 (d10)	0.984 ± 0.005	0.991 ± 0.003	+0.007	+1.20
TS/nat	999	L	Laplacian (d8)	ZZ-r1 (d10)	0.988 ± 0.002	0.992 ± 0.003	+0.004	+1.41
UD/m2	42	S	Laplacian (d12)	PauliXZ (d12)	0.894 ± 0.028	0.904 ± 0.016	+0.009	+0.29
UD/m2	42	M	Laplacian (d12)	PauliXZ (d12)	0.911 ± 0.007	0.934 ± 0.009	+0.023	+1.96
UD/m2	42	L	Laplacian (d12)	Z-map (d12)	0.927 ± 0.008	0.939 ± 0.010	+0.013	+1.02
UD/m2	123	S	Laplacian (d12)	Z-map (d12)	0.889 ± 0.019	0.917 ± 0.019	+0.028	+1.04
UD/m2	123	M	Laplacian (d12)	PauliXZ (d12)	0.916 ± 0.016	0.937 ± 0.007	+0.022	+1.24
UD/m2	123	L	Laplacian (d12)	Z-map (d12)	0.932 ± 0.014	0.945 ± 0.007	+0.013	+0.86
UD/m2	999	S	Laplacian (d12)	Z-map (d12)	0.878 ± 0.022	0.911 ± 0.031	+0.033	+0.87
UD/m2	999	M	Laplacian (d12)	PauliXZ (d12)	0.914 ± 0.011	0.931 ± 0.009	+0.017	+1.19
UD/m2	999	L	Laplacian (d12)	PauliXZ (d12)	0.936 ± 0.009	0.944 ± 0.003	+0.008	+0.82
UD/nat	42	S	Laplacian (d12)	ZZ-r2 (d6)	0.764 ± 0.023	0.801 ± 0.023	+0.037	+1.13
UD/nat	42	M	Laplacian (d12)	PauliXZ (d12)	0.813 ± 0.010	0.829 ± 0.008	+0.016	+1.20
UD/nat	42	L	Laplacian (d12)	Z-map (d12)	0.818 ± 0.010	0.831 ± 0.009	+0.013	+0.97
UD/nat	123	S	Laplacian (d12)	ZZ-r2 (d10)	0.775 ± 0.027	0.813 ± 0.017	+0.038	+1.20
UD/nat	123	M	Laplacian (d12)	ZZ-r2 (d6)	0.781 ± 0.020	0.805 ± 0.016	+0.024	+0.93
UD/nat	123	L	Laplacian (d12)	ZZ-r1 (d12)	0.800 ± 0.005	0.816 ± 0.006	+0.016	+2.09
UD/nat	999	S	Linear (d4)	ZZ-r2 (d6)	0.766 ± 0.053	0.797 ± 0.021	+0.032	+0.55
UD/nat	999	M	Laplacian (d12)	ZZ-r2 (d12)	0.784 ± 0.018	0.806 ± 0.013	+0.022	+1.02
UD/nat	999	L	Laplacian (d12)	ZZ-r2 (d12)	0.805 ± 0.006	0.822 ± 0.012	+0.017	+1.31
UR/m2	42	S	Laplacian (d12)	Z-map (d12)	0.880 ± 0.021	0.902 ± 0.022	+0.022	+0.73
UR/m2	42	M	Laplacian (d12)	Z-map (d12)	0.901 ± 0.008	0.925 ± 0.009	+0.024	+1.99
UR/m2	42	L	Laplacian (d12)	PauliXZ (d12)	0.929 ± 0.012	0.942 ± 0.003	+0.013	+1.05
UR/m2	123	S	Laplacian (d12)	Z-map (d12)	0.905 ± 0.012	0.926 ± 0.023	+0.021	+0.81
UR/m2	123	M	Laplacian (d12)	Z-map (d12)	0.914 ± 0.010	0.946 ± 0.010	+0.032	+2.29
UR/m2	123	L	Laplacian (d12)	PauliXZ (d12)	0.946 ± 0.004	0.953 ± 0.006	+0.007	+0.91
UR/m2	999	S	Laplacian (d10)	PauliXZ (d12)	0.883 ± 0.023	0.910 ± 0.025	+0.027	+0.81
UR/m2	999	M	Laplacian (d12)	Z-map (d12)	0.912 ± 0.004	0.941 ± 0.008	+0.029	+3.13
UR/m2	999	L	Laplacian (d10)	PauliXZ (d12)	0.937 ± 0.014	0.949 ± 0.003	+0.012	+0.84
UR/nat	42	S	Poly-2 (d4)	PauliXZ (d6)	0.760 ± 0.020	0.784 ± 0.024	+0.025	+0.79
UR/nat	42	M	Laplacian (d10)	ZZ-r1 (d4)	0.795 ± 0.040	0.817 ± 0.013	+0.022	+0.51
UR/nat	42	L	Laplacian (d8)	ZZ-r1 (d4)	0.796 ± 0.007	0.813 ± 0.005	+0.017	+1.95
UR/nat	123	S	Poly-2 (d12)	ZZ-r2 (d6)	0.780 ± 0.024	0.797 ± 0.016	+0.017	+0.59
UR/nat	123	M	Linear (d12)	ZZ-r1 (d4)	0.794 ± 0.014	0.803 ± 0.020	+0.008	+0.34
UR/nat	123	L	Laplacian (d12)	ZZ-r2 (d12)	0.810 ± 0.023	0.812 ± 0.019	+0.002	+0.06
UR/nat	999	S	Poly-2 (d12)	ZZ-r2 (d6)	0.777 ± 0.016	0.782 ± 0.005	+0.005	+0.30
UR/nat	999	M	Poly-2 (d12)	ZZ-r2 (d6)	0.780 ± 0.017	0.814 ± 0.016	+0.034	+1.51
UR/nat	999	L	Laplacian (d8)	ZZ-r1 (d4)	0.790 ± 0.010	0.807 ± 0.011	+0.017	+1.10

Authors' contributions.	2393
RFB: conceptualization, methodology, software, data	2394
curation, formal analysis, visualization, investigation, writing—original draft,	2395
writing—review and editing. IPL: supervision, validation, writing—review and editing.	2396
AGS: validation, project administration, writing—review and editing. PGB: supervi-	2397
sion, resources, writing—review and editing. All authors read and approved the final	2398
manuscript.	2399
	2400
	2401
	2402
Acknowledgements.	2403
Not applicable.	2404
	2405
	2406
References	2407
	2408
Abreu, D., Esteve Rothenberg, C., Abelém, A.: QML-IDS: Quantum machine learning	2409
intrusion detection system. In: 2024 IEEE Symposium on Computers and Commu-	2410
nications (ISCC), pp. 1–6. IEEE, ??? (2024). https://doi.org/10.1109/ISCC61673.	2411
2024.10733655	2412
	2413
	2414
	2415
	2416
Abreu, D., Moura, D., Esteve Rothenberg, C., Abelém, A.: QuantumNetSec: Quan-	2417
tum machine learning for network security. International Journal of Network	2418
Management 35 (4), 70018 (2025) https://doi.org/10.1002/nem.70018	2419
	2420
	2421
	2422
Alsaedi, A., Moustafa, N., Tari, Z., Mahmood, A., Anwar, A.: TON_IoT teleme-	2423
try dataset: A new generation dataset of IoT and IIoT for data-driven intrusion	2424
detection systems. IEEE Access 8 , 165130–165150 (2020) https://doi.org/10.1109/	2425
ACCESS.2020.3022862	2426
	2427
	2428
	2429
Arp, D., Quiring, E., Pendlebury, F., Warnecke, A., Pierazzi, F., Wressnegger, C.,	2430
Cavallaro, L., Rieck, K.: Dos and don'ts of machine learning in computer security.	2431
In: 31st USENIX Security Symposium, pp. 3971–3988 (2022)	2432
	2433
	2434
	2435
Anderson, H.S., Roth, P.: Ember: An open dataset for training static PE malware	2436
machine learning models. arXiv abs/1804.04637 (2018)	2437
	2438

2439 Bowles, J., Ahmed, S., Schuld, M.: Better than classical? the subtle art of bench-
 2440 marking quantum machine learning models. arXiv **abs/2403.07059** (2024) <https://doi.org/10.48550/arXiv.2403.07059>
 2441
 2442
 2443
 2444 Bordelon, B., Canatar, A., Pehlevan, C.: Spectrum dependent learning curves in ker-
 2445 nel regression and wide neural networks. In: Proceedings of the 37th International
 2446 Conference on Machine Learning (ICML). PMLR, vol. 119, pp. 1024–1034 (2020)
 2447
 2448
 2449 Bellante, A., Fioravanti, T., Carminati, M., Zanero, S., Luongo, A.: Evaluating the
 2450 potential of quantum machine learning in cybersecurity: A case-study on PCA-
 2451 based intrusion detection systems. Computers & Security **154**, 104341 (2025) <https://doi.org/10.1016/j.cose.2025.104341>
 2452
 2453
 2454
 2455
 2456
 2457 Barbero, F., Pendlebury, F., Pierazzi, F., Cavallaro, L.: Transcending TRANSCEND:
 2458 Revisiting malware classification in the presence of concept drift. In: 2022 IEEE
 2459 Symposium on Security and Privacy (SP), pp. 805–823 (2022). <https://doi.org/10.1109/SP46214.2022.9833659>
 2460
 2461
 2462
 2463
 2464 Canatar, A., Bordelon, B., Pehlevan, C.: Spectral bias and task-model align-
 2465 ment explain generalization in kernel regression and infinitely wide neural
 2466 networks. Nature Communications **12**, 2914 (2021) <https://doi.org/10.1038/s41467-021-23103-1>
 2467
 2468
 2469
 2470
 2471
 2472 Caro, M.C., Huang, H.-Y., Cerezo, M., Sharma, K., Sornborger, A., Cincio, L., Coles,
 2473 P.J.: Generalization in quantum machine learning from few training data. Nature
 2474 Communications **13**, 4919 (2022) <https://doi.org/10.1038/s41467-022-32550-3>
 2475
 2476
 2477
 2478 Caro, M.C., Huang, H.-Y., Ezzell, N., Gibbs, J., Sornborger, A.T., Cincio, L.,
 2479 Coles, P.J., Holmes, Z.: Out-of-distribution generalization for learning quan-
 2480 tum dynamics. Nature Communications **14**, 3751 (2023) <https://doi.org/10.1038/s41467-023-4644-4>
 2481
 2482
 2483
 2484

s41467-023-39381-w	2485
	2486
Cortes, C., Mohri, M., Rostamizadeh, A.: Algorithms for learning kernels based on	2487
centered alignment. <i>Journal of Machine Learning Research</i> 13 (28), 795–828 (2012)	2488
	2489
	2490
Canatar, A., Peters, E., Pehlevan, C., Wild, S.M., Shaydulin, R.: Bandwidth enables	2491
generalization in quantum kernel models. <i>Transactions on Machine Learning</i>	2492
<i>Research</i> (2023)	2493
	2494
	2495
	2496
Cristianini, N., Shawe-Taylor, J., Elisseeff, A., Kandola, J.: On kernel-target align-	2497
ment. In: <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , vol. 14	2498
(2001)	2499
	2500
	2501
	2502
Cortes, C., Vapnik, V.: Support-vector networks. <i>Machine Learning</i> 20 (3), 273–297	2503
(1995) https://doi.org/10.1007/BF00994018	2504
	2505
	2506
	2507
Cerezo, M., Verdon, G., Huang, H.-Y., Cincio, L., Coles, P.J.: Challenges and oppor-	2508
tunities in quantum machine learning. <i>Nature Computational Science</i> 2 , 567–576	2509
(2022) https://doi.org/10.1038/s43588-022-00311-3	2510
	2511
	2512
Engelen, G., Rimmer, V., Joosen, W.: Troubleshooting an intrusion detection dataset:	2513
The CICIDS2017 case study. In: <i>2021 IEEE Security and Privacy Workshops (SPW)</i> ,	2514
pp. 7–12. IEEE, ??? (2021). https://doi.org/10.1109/SPW53761.2021.00009	2515
	2516
	2517
	2518
Egginger, S., Sakhnenko, A., Lorenz, J.M.: A hyperparameter study for quantum	2519
kernel methods. <i>Quantum Machine Intelligence</i> 6 , 44 (2024) https://doi.org/10.	2520
1007/s42484-024-00172-1	2521
	2522
	2523
	2524
Gretton, A., Borgwardt, K.M., Rasch, M.J., Schölkopf, B., Smola, A.: A kernel two-	2525
sample test. <i>Journal of Machine Learning Research</i> 13 (25), 723–773 (2012)	2526
	2527
	2528
Gouveia, A., Correia, M.: Towards quantum-enhanced machine learning for network	2529
	2530

2531 intrusion detection. In: 2020 IEEE 19th International Symposium on Network Com-
 2532 puting and Applications (NCA) (2020). [https://doi.org/10.1109/NCA51143.2020.](https://doi.org/10.1109/NCA51143.2020.9306691)
 2533 [9306691](https://doi.org/10.1109/NCA51143.2020.9306691)
 2534
 2535
 2536
 2537 Gil-Fuster, E., Eisert, J., Bravo-Prieto, C.: Understanding quantum machine learning
 2538 also requires rethinking generalization. *Nature Communications* **15**, 2277 (2024)
 2539 <https://doi.org/10.1038/s41467-024-45882-z>
 2540
 2541
 2542
 2543 Gulrajani, I., Lopez-Paz, D.: In search of lost domain generaliza-
 2544 tion. In: *International Conference on Learning Representations* (2021).
 2545 <https://openreview.net/forum?id=IQdXeXDoWtl>
 2546
 2547
 2548
 2549 Guo, C., Pleiss, G., Sun, Y., Weinberger, K.Q.: On calibration of modern neural net-
 2550 works. In: *Proceedings of the 34th International Conference on Machine Learning*
 2551 (ICML). PMLR, vol. 70, pp. 1321–1330 (2017)
 2552
 2553
 2554 Huang, H.-Y., Broughton, M., Mohseni, M., Babbush, R., Boixo, S., Neven, H.,
 2555 McClean, J.R.: Power of data in quantum machine learning. *Nature Communica-*
 2556 *tions* **12**, 2631 (2021) <https://doi.org/10.1038/s41467-021-22539-9>
 2557
 2558
 2559
 2560 Havlíček, V., Córcoles, A.D., Temme, K., Harrow, A.W., Kandala, A., Chow,
 2561 J.M., Gambetta, J.M.: Supervised learning with quantum-enhanced feature spaces.
 2562 *Nature* **567**, 209–212 (2019) <https://doi.org/10.1038/s41586-019-0980-2>
 2563
 2564
 2565
 2566 Herrero Gómez, P., Jimeno Morenilla, A., Muñoz-Hernández, D., Mora Mora, H.:
 2567 Spectral Phase Encoding for Quantum Kernel Methods. *arXiv:2602.19644* (2026)
 2568
 2569
 2570 Harang, R., Rudd, E.M.: Sorel-20m: A large scale benchmark dataset for malicious
 2571 PE detection. *arXiv abs/2012.07634* (2020)
 2572
 2573
 2574 Hubregtsen, T., Wierichs, D., Gil-Fuster, E., Derks, P.-J.H.S., Faehrmann, P.K.,
 2575
 2576

Meyer, J.J.: Training quantum embedding kernels on near-term quantum computers. <i>Physical Review A</i> 106 , 042431 (2022) https://doi.org/10.1103/PhysRevA.106.042431	2577 2578 2579 2580 2581 2582
Jerbi, S., Gyurik, C., Marshall, S.C., Molteni, R., Dunjko, V.: Shadows of quantum machine learning. <i>Nature Communications</i> 15 , 5676 (2024) https://doi.org/10.1038/s41467-024-49877-8	2583 2584 2585 2586 2587 2588
Joyce, R.J., Miller, G., Roth, P., Zak, R., Zaresky-Williams, E., Anderson, H., Raff, E., Holt, J.: Ember2024 – a benchmark dataset for holistic evaluation of malware classifiers. <i>arXiv</i> abs/2506.05074 (2025)	2589 2590 2591 2592 2593
Jordaney, R., Sharad, K., Dash, S.K., Wang, Z., Papini, D., Nouretdinov, I., Cavallaro, L.: Transcend: Detecting concept drift in malware classification models. In: 26th USENIX Security Symposium, pp. 625–642 (2017)	2594 2595 2596 2597 2598 2599
Kübler, J.M., Buchholz, S., Schölkopf, B.: The inductive bias of quantum kernels. In: <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , vol. 34, pp. 12661–12673 (2021)	2600 2601 2602 2603 2604 2605
Koh, P.W., Sagawa, S., Marklund, H., Xie, S.M., Zhang, M., Balsubramani, A., Hu, W., Yasunaga, M., Phillips, R.L., Gao, I., Lee, T., David, E., Stavness, I., Guo, W., Earnshaw, B., Haque, I., Beery, S.M., Leskovec, J., Kundaje, A., Pierson, E., Levine, S., Finn, C., Liang, P.: Wilds: A benchmark of in-the-wild distribution shifts. In: <i>Proceedings of the 38th International Conference on Machine Learning. Proceedings of Machine Learning Research</i> , vol. 139, pp. 5637–5664. PMLR, ??? (2021). https://proceedings.mlr.press/v139/koh21a.html	2606 2607 2608 2609 2610 2611 2612 2613 2614 2615 2616 2617
Liu, Y., Arunachalam, S., Temme, K.: A rigorous and robust quantum speed-up in supervised machine learning. <i>Nature Physics</i> 17 , 1013–1017 (2021) https://doi.org/10.1038/s41567-021-01017-1	2618 2619 2620 2621 2622

2623 [10.1038/s41567-021-01287-z](https://doi.org/10.1038/s41567-021-01287-z)
2624
2625 Li, A.S., Iyengar, A., Kundu, A., Bertino, E.: Revisiting concept drift in windows
2626 malware detection: Adaptation to real drifted malware with minimal samples. In:
2627 Network and Distributed System Security Symposium (NDSS) (2025). [https://doi.](https://doi.org/10.14722/ndss.2025.240830)
2628 [org/10.14722/ndss.2025.240830](https://doi.org/10.14722/ndss.2025.240830)
2629
2630
2631
2632
2633 Lazaro, J., Vazquez, J.-I., Garcia-Bringas, P.: Dissecting quantum reinforcement learn-
2634 ing: A systematic evaluation of key components. arXiv **abs/2511.17112** (2025)
2635 <https://doi.org/10.48550/arXiv.2511.17112>
2636
2637
2638
2639 Ma, C., Pathak, R., Wainwright, M.J.: Optimally tackling covariate shift in RKHS-
2640 based nonparametric regression. The Annals of Statistics **51**(2), 738–761 (2023)
2641 <https://doi.org/10.1214/23-AOS2268>
2642
2643
2644
2645 Moustafa, N., Slay, J.: UNSW-NB15: A comprehensive data set for network intrusion
2646 detection systems (UNSW-NB15 network data set),. In: 2015 Military Commu-
2647 nications and Information Systems Conference (MilCIS), pp. 1–6 (2015). [https:](https://doi.org/10.1109/MilCIS.2015.7348942)
2648 [//doi.org/10.1109/MilCIS.2015.7348942](https://doi.org/10.1109/MilCIS.2015.7348942)
2649
2650
2651
2652 Miller, J.P., Taori, R., Raghunathan, A., Sagawa, S., Koh, P.W., Shankar, V., Liang,
2653 P., Carmon, Y., Schmidt, L.: Accuracy on the line: On the strong correlation between
2654 out-of-distribution and in-distribution generalization. In: Proceedings of the 38th
2655 International Conference on Machine Learning (ICML). PMLR, vol. 139, pp. 7721–
2656 7735 (2021)
2657
2658
2659
2660
2661 Nieto, A., Gil-Merino, R.: Survey of quantum kernel methods for detection systems
2662 and security analytics. In: Distributed Computing and Artificial Intelligence, Spe-
2663 cial Sessions II, 22nd International Conference. Lecture Notes in Networks and
2664 Systems, vol. 1630, pp. 281–290. Springer, Cham (2026). <https://doi.org/10.1007/>
2665
2666
2667
2668

978-3-032-05748-8_23	2669
	2670
Nicesio, O.K., Leal, A.G., Gava, V.L.: Quantum machine learning for network intrusion	2671
detection systems, a systematic literature review. In: 2023 IEEE 2nd Interna-	2672
tional Conference on AI in Cybersecurity (ICAIC) (2023). https://doi.org/10.1109/	2673
ICAIC57335.2023.10044125	2674
	2675
	2676
	2677
	2678
Ovadia, Y., Fertig, E., Ren, J., Nado, Z., Sculley, D., Nowozin, S., Dillon, J.V.,	2679
Lakshminarayanan, B., Snoek, J.: Can you trust your model's uncertainty? Evalu-	2680
ating predictive uncertainty under dataset shift. In: Advances in Neural Information	2681
Processing Systems (NeurIPS), vol. 32 (2019)	2682
	2683
	2684
	2685
Pendlebury, F., Pierazzi, F., Jordaney, R., Kinder, J., Cavallaro, L.: TESSERACT:	2686
Eliminating experimental bias in malware classification across space and time. In:	2687
28th USENIX Security Symposium, pp. 729–746 (2019)	2688
	2689
	2690
	2691
Quiñonero-Candela, J., Sugiyama, M., Schwaighofer, A., Lawrence, N.D. (eds.):	2692
Dataset Shift in Machine Learning. MIT Press, Cambridge, MA (2008). https://doi.org/10.7551/mitpress/9780262170055.001.0001	2693
	2694
	2695
	2696
	2697
Roy, O., Vetterli, M.: The effective rank: A measure of effective dimensionality. In:	2698
2007 15th European Signal Processing Conference (EUSIPCO), pp. 606–610 (2007)	2699
	2700
	2701
Rasmussen, C.E., Williams, C.K.I.: Gaussian Processes for Machine Learning. MIT	2702
Press, Cambridge, MA (2006). https://doi.org/10.7551/mitpress/3206.001.0001	2703
	2704
	2705
Schuld, M.: Supervised quantum machine learning models are kernel methods. arXiv	2706
abs/2101.11020 (2021) https://doi.org/10.48550/arXiv.2101.11020	2707
	2708
	2709
	2710
Schreiber, F.J., Eisert, J., Meyer, J.J.: Classical surrogates for quantum learn-	2711
ing models. Physical Review Letters 131 , 100803 (2023) https://doi.org/10.1103/	2712
	2713
	2714

2715 [PhysRevLett.131.100803](#)
 2716
 2717 Shimodaira, H.: Improving predictive inference under covariate shift by weighting the
 2718 log-likelihood function. *Journal of Statistical Planning and Inference* **90**(2), 227–244
 2719 (2000) [https://doi.org/10.1016/S0378-3758\(00\)00115-4](https://doi.org/10.1016/S0378-3758(00)00115-4)
 2720
 2721
 2722
 2723 Schuld, M., Killoran, N.: Quantum machine learning in feature hilbert spaces. *Physical*
 2724 *Review Letters* **122**(4), 040504 (2019) [https://doi.org/10.1103/PhysRevLett.122.](https://doi.org/10.1103/PhysRevLett.122.040504)
 2725 [040504](https://doi.org/10.1103/PhysRevLett.122.040504)
 2726
 2727
 2728
 2729 Sugiyama, M., Krauledat, M., Müller, K.-R.: Covariate shift adaptation by importance
 2730 weighted cross validation. *Journal of Machine Learning Research* **8**(35), 985–1005
 2731 (2007)
 2732
 2733
 2734
 2735 Schnabel, J., Roth, M.: Quantum kernel methods under scrutiny: A benchmark-
 2736 ing study. *Quantum Machine Intelligence* **7**, 58 (2025) [https://doi.org/10.1007/](https://doi.org/10.1007/s42484-025-00273-5)
 2737 [s42484-025-00273-5](https://doi.org/10.1007/s42484-025-00273-5)
 2738
 2739
 2740
 2741 Schölkopf, B., Smola, A.J.: *Learning with Kernels: Support Vector Machines, Regu-*
 2742 *larization, Optimization, and Beyond*. MIT Press, Cambridge, MA (2002)
 2743
 2744
 2745 Schuld, M., Sweke, R., Meyer, J.J.: Effect of data encoding on the expressive power of
 2746 variational quantum-machine-learning models. *Physical Review A* **103**(3), 032430
 2747 (2021) <https://doi.org/10.1103/PhysRevA.103.032430>
 2748
 2749
 2750
 2751 Shaydulin, R., Wild, S.M.: Importance of kernel bandwidth in quantum machine learn-
 2752 ing. *Physical Review A* **106**(4), 042407 (2022) [https://doi.org/10.1103/PhysRevA.](https://doi.org/10.1103/PhysRevA.106.042407)
 2753 [106.042407](https://doi.org/10.1103/PhysRevA.106.042407)
 2754
 2755
 2756
 2757 Thanasilp, S., Wang, S., Cerezo, M., Holmes, Z.: Exponential concentration in
 2758 quantum kernel methods. *Nature Communications* **15**, 5200 (2024)
 2759
 2760

Wang, X., Du, Y., Luo, Y., Tao, D.: Towards understanding the power of quantum kernels in the NISQ era. <i>Quantum</i> 5 , 531 (2021) https://doi.org/10.22331/q-2021-08-30-531	2761 2762 2763 2764 2765 2766
Yao, H., Choi, C., Cao, B., Lee, Y., Koh, P.W., Finn, C.: Wild-time: A benchmark of in-the-wild distribution shift over time. In: <i>Advances in Neural Information Processing Systems</i> (2022). Datasets and Benchmarks Track	2767 2768 2769 2770 2771 2772
Yang, L., Ciptadi, A., Laziuk, I., Ahmadzadeh, A., Wang, G.: Bodmas: An open dataset for learning based temporal analysis of PE malware. In: <i>2021 IEEE Security and Privacy Workshops (SPW)</i> , pp. 78–84 (2021). https://doi.org/10.1109/SPW53761.2021.00020	2773 2774 2775 2776 2777 2778 2779
Yang, L., Guo, W., Hao, Q., Ciptadi, A., Ahmadzadeh, A., Xing, X., Wang, G.: CADE: Detecting and explaining concept drift samples for security applications. In: <i>30th USENIX Security Symposium</i> , pp. 2327–2344 (2021)	2780 2781 2782 2783 2784 2785
Zhou, K., Liu, Z., Qiao, Y., Xiang, T., Loy, C.C.: Domain generalization: A survey. <i>IEEE Transactions on Pattern Analysis and Machine Intelligence</i> 45 (4), 4396–4415 (2023) https://doi.org/10.1109/TPAMI.2022.3195549	2786 2787 2788 2789 2790 2791 2792 2793 2794 2795 2796 2797 2798 2799 2800 2801 2802 2803 2804 2805 2806