

Evaluation Choices Shape Apparent Quantum Kernel Robustness under Distribution Shift

Roberto Fernández-Barrios^{[ORCID]1*}, Iker Pastor-López^{[ORCID]1},
Asier González-Santocildes^{[ORCID]1}, Pablo García Bringas^{[ORCID]1}

^{1*}Faculty of Engineering, University of Deusto, Avda. de las
Universidades, 24., Bilbao, 48007, Spain.

*Corresponding author(s). E-mail(s): roberto.fernandez.b@deusto.es;
Contributing authors: iker.pastor@deusto.es; gonzalez.asier@deusto.es;
pablo.garcia.bringas@deusto.es;

Abstract

Claims of quantum-kernel robustness can depend on how candidates are tuned, selected, and compared. We map this dependence with a controlled kernel-swap benchmark under distribution shift, a ten-specification curve, and prospectively frozen corroboration on three external domain shifts. In eight fixed security scenario-groups, same-test target-label selection, fixed regularization, and weak classical references produce apparent gains up to **+0.037**. Train-only regularization, target-label-free selection, and structurally matched candidate budgets instead give source-dataset-equal quantum-minus-classical effects of **−0.0056** for support-vector classification and **−0.0009** for Gaussian-process classification. Across the external tasks, the controlled support-vector effect is **−0.0119** (conditional 95% interval [**−0.0205**, **−0.0033**]) and protocol contraction is **+0.0208**. Effective rank and alignment associate with accuracy but occur in both families. Repeated finite-shot perturbations reveal estimator-dependent deviations without simulating hardware. The apparent robustness verdict is therefore evaluation-dependent in these low-qubit regimes.

Keywords: quantum machine learning; quantum kernels; distribution shift;
out-of-distribution robustness; kernel methods; kernel-target alignment; Gaussian
processes; malware detection; network intrusion detection

047 1 Introduction

048
049 Quantum kernel methods embed classical data into quantum states and use state over-
050 laps as similarities for otherwise classical learners [1–3]. This separation makes them
051 a clean setting for asking whether quantum-induced feature geometry provides useful
052 inductive bias. It also makes comparison unusually sensitive to evaluation design: the
053 downstream learner can be held fixed, but conclusions still depend on the encoding,
054 bandwidth, regularization, classical reference set, and rule used to select a candidate
055 [4–7].

056 Recent benchmarks have accordingly shifted the question from whether an isolated
057 quantum model can perform well to whether a claimed gain survives strong and sym-
058 metric controls. Broad QML studies find competitive classical baselines once pipelines
059 are aligned [6]; quantum-kernel benchmarks show substantial sensitivity to feature
060 map and hyperparameter choice [7]; and a concurrent tabular study using nested cross-
061 validation, spectral analysis, and hardware validation finds no significant quantum
062 win across its tested pairs [8]. Out-of-distribution generalization in QML is itself an
063 established topic: prior work analyzes learning quantum dynamics outside the training
064 distribution and relates data-dependent generalization to quantum Fisher information
065 [9, 10]. Our narrower deployment question is whether a family-level quantum-kernel
066 robustness claim on classical data survives target-label-free selection and symmetric
067 tuning, reference, and search-budget controls.

068 Distribution shift makes model selection part of the estimand. Work on domain
069 generalization and real-world shift has shown that robustness gains can disappear
070 when search budgets, hyperparameters, and selection criteria are standardized [11–14].
071 A selector that inspects the shifted test labels estimates an oracle upper bound, not
072 deployable performance; a family with more candidates has more opportunities to win
073 that selection. A controlled quantum–classical evaluation under shift must therefore

074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092

hold fixed not only the classifier and data split, but also regularization, candidate budget, and label access.

Security detection provides a useful stress test because shift is intrinsic to the application. Malware evolves over time, and experimental bias can overstate detector performance [15–17]; network traffic changes with infrastructure and attack campaigns [18, 19]. Quantum classifiers have already been evaluated on intrusion-detection benchmarks, sometimes with reported gains over small classical reference sets on i.i.d. splits [20–22]. We use these datasets not to propose a security system, but to expose kernel families to reproducible constructed shifts and to a held-out attack-campaign shift.

The closest kernel-robustness-specific concurrent work studies additive input corruption and proposes a spectral phase encoding [23]. Our question is complementary: rather than optimize a new encoding or claim the first QML study under shift, we audit how evaluation choices change a family-level conclusion under constructed and held-out-campaign shifts. The distinguishing contribution is a controlled classical-data kernel swap combining no-OOD-label selection, equal candidate budgets, per-configuration train-only regularization, a complete fixed-case specification curve, and prospectively frozen external corroboration.

We implement a controlled *kernel-swap* benchmark in which preprocessing, splits, classifier, and family-internal selection logic are fixed while only the kernel changes. Four security scenarios from EMBER, UNSW-NB15, and ToN-IoT yield eight scenario-groups across constructed and held-out-campaign shifts. Support-vector and Laplace Gaussian-process classifiers consume the same precomputed Gram matrices. The quantum family contains four fidelity feature-map families with angle-scale sweeps; the classical family extends linear and RBF references with polynomial, Laplacian, and Matérn kernels and symmetric bandwidth freedom. Each configuration’s regularization is tuned from training data alone. The primary deployed candidate is selected on a disjoint ID-validation split, and the two extended families receive equal

139 candidate budgets. Fifteen pipeline realizations per setting quantify conditional vari-
140 ability; the benchmarks are treated as fixed case studies rather than as random draws
141 from an undefined population.
142

143 The comparison is coupled to a geometry audit. We measure effective rank, cen-
144 tered kernel-target alignment, concentration, and geometric difference [4, 24–26]; use
145 partial association, leave-one-source-dataset-out prediction, and effective-rank match-
146 ing to restrict interpretation; and repeat finite-shot perturbations under unprojected,
147 independent-square PSD, and coherent train-eigenspace PSD estimators.
148

149 The evaluation choices change the conclusion. With same-test oracle selection,
150 fixed regularization, and customary linear+RBF baselines, fidelity kernels appear to
151 improve OOD balanced accuracy by as much as +0.037. With train-only regular-
152 ization, no-OOD-label selection, and structurally matched candidate budgets, the
153 source-dataset-equal quantum-minus-classical delta is -0.0056 for SVC and -0.0009
154 for the GP classifier. The primary SVC comparison is classical-favoured or condition-
155 ally tied in seven of eight scenario-groups. Effective rank and OOD alignment remain
156 associated with accuracy, but the association is regime-dependent and shared by both
157 families. Nearest-rank pairing within the predefined overlap yields negative median
158 differences in all eight groups, including under one-to-one and caliper sensitivities.
159 Repeated finite-shot estimates inflate exact-statevector effective rank and produce
160 condition-dependent OOD deviations.
161

162 The contribution is therefore the combined evaluation design: distribution shift,
163 target-label-free selection, train-only regularization, structurally matched search bud-
164 gets, a prespecified specification curve, and prospectively frozen external corrobor-
165 ation. Together these controls show how the reported quantum-kernel robustness result
166 changes across evaluation bundles and that the observed geometry does not uniquely
167 identify quantum provenance. The scope is deliberately limited to the low-qubit fidelity
168 kernels, datasets, and shift constructions studied here.
169

170
171
172
173
174
175
176
177
178
179
180
181
182
183
184




Fig. 1 The controlled kernel-swap protocol. Benchmark pools feed reproducible shift constructions; all kernels operate on a shared train-only embedding, receive per-configuration regularization tuned on the training split, and are consumed by the same precomputed-kernel classifiers. Family-level conclusions are drawn by target-label-free, equal-budget selection on an in-distribution validation split (P1'). Each scenario-group interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the fixed benchmark pools and design. Gram matrices feed the associational geometry audit and repeated finite-shot sensitivity.

2 Results

2.1 Overview

We report results under the controlled kernel-swap protocol described in Section 4. All comparisons share the same classifier, preprocessing pipeline, and ID/OOD split definitions within a scenario-group; the only varying component is the kernel. Each setting contains fifteen pipeline realizations arranged in five q-split seed clusters, not fifteen independent observations. The argument proceeds in three movements. First, under test-peeking oracle selection and fixed regularization, the fidelity kernels show an apparent OOD advantage over linear and RBF baselines (Section 2.2). Second, we tune each configuration on training data alone, select on an in-distribution validation split disjoint from both test splits, and match both the count and block structure of the extended-family search budgets. Under this controlled comparison no robust family-level advantage remains; the primary SVC comparison is classical-favoured or tied in seven of eight scenario-groups (Section 2.4). Third, a geometry

analysis shows that effective rank and OOD alignment are associated with—but do not causally determine—performance, that the association is regime-dependent, and that it favours neither family (Section 2.6); repeated finite-shot perturbations then quantify how sampling and positive-semidefinite projection alter exact-statevector endpoints (Section 2.7). Throughout, the eight scenario-groups are fixed case studies; conditional intervals describe pipeline variation given the benchmark pools and design, not uncertainty over a population of datasets.

2.2 The apparent advantage under oracle selection

We begin with a deliberately optimistic reference comparison: fidelity-based quantum kernels against linear and RBF baselines, with each family’s representative chosen by Best-by-OOD selection—the configuration with the highest balanced accuracy on the OOD test split itself—and the classical kernels left at the default regularization $C = 1$. Related intrusion-detection studies motivate the small-reference context [20–22], but we do not claim that this exact protocol dominates QML practice. Under these conventions, on EMBER, the quantum family shows an OOD advantage of +0.037 under SVC and +0.028 under the GP classifier against linear+RBF, shrinking to +0.004/+0.002 once the classical family is enlarged with heavier-tailed kernels. Supplementary Table 5 gives the complete oracle results.

Two features favour this apparent quantum result. Best-by-OOD is an *oracle*: it uses the labels it then reports. Fixing $C = 1$ also leaves classical candidates at a potentially unsuitable regularization point. These values are therefore descriptive upper bounds, not deployable effects.

2.3 A specification curve locates the reversal

To avoid privileging a preferred endpoint, we enumerated all ten logically valid, already-computed specifications fixed before this extension (Figure 2). They trace a descriptive path over target-label access, SVC regularization, classical reference

strength, and candidate budget. S1–S4 use the legacy generation with fixed $C = 1$, whereas S5–S10 use the corrected v4 generation with train-cross-validated C ; the S4–S5 transition therefore does not isolate regularization from generation. Every specification retains all eight security scenario-groups.

The family-level conclusion spans both sides of zero. For SVC, the source-dataset-equal quantum-minus-classical mean ranges from $+0.047$ under legacy ID selection with fixed $C = 1$ and the customary reference (S2) to -0.006 under ID-validation selection, train-only regularization, the extended reference, and equal budgets (S10). For the GP classifier, for which the regularization axis is not applicable, the corresponding range is $+0.042$ to -0.003 . Expanding the classical reference contracts the positive effect before target-label-free selection or budget matching is imposed; train-only SVC regularization and ID-validation selection then move the controlled estimates to zero or slightly negative. The path is not monotone in every intermediate specification, which is precisely why reporting only an oracle and a final controlled endpoint would be insufficient. The curve shows that the apparent advantage is conditional on a bundle of evaluation decisions rather than a stable property of quantum provenance.

2.4 The controlled comparison removes the apparent advantage

A deployable comparison must tune each configuration’s regularization without looking at the test data, select the deployed configuration without OOD labels, and compare families at the same candidate budget. We implement all three (Section 4). For every kernel candidate, the SVC regularization C is chosen by five-fold cross-validation on the training Gram matrix alone. The deployed configuration is then the one with the best balanced accuracy on an in-distribution *validation* split, held out from the ID test split (protocol P1’); no OOD label is ever consulted in selection. The primary comparison uses 60 candidates per extended family in every scenario-group and classifier, with the full 115-candidate classical pool retained as a sensitivity.




Fig. 2 A frozen specification curve shows how the apparent family effect changes across evaluation bundles. Grey points are the eight fixed security scenario-groups; diamonds average groups within EMBER, UNSW-NB15, and ToN-IoT before giving the three source datasets equal weight. Specifications are shown in their prespecified S1–S10 order, not sorted by effect. S1–S4 use the legacy generation with fixed $C = 1$; S5–S10 use the corrected v4 generation with train-cross-validated C , so their boundary is generation-confounded and does not isolate a regularization effect. Positive values favour the fidelity-kernel family. GPC has no C -regularization axis. These are descriptive fixed-case summaries, not per-axis causal effects; no specification-level hypothesis test or population interval is reported.

We treat the eight scenario-groups as fixed case studies and report, per group, the quantum-minus-classical OOD delta with a conditional interval over the five split realizations, and no population p -value (Section 4).

Table 1 and Figure 3 give the result: no robust advantage remains against either reference. Against the equal-budget extended classical family, the P1' delta is classical-favoured or within the conditional uncertainty of zero in seven of the eight scenario-groups under the primary SVC comparison. Averaging groups within each source dataset and then giving EMBER, UNSW-NB15, and ToN-IoT equal weight gives -0.00565 (leave-one-source-dataset-out range $[-0.00702, -0.00450]$); the GP-classifier mean is -0.00094 ($[-0.00190, +0.00094]$). Against the customary

Table 1 Target-label-free selection results for eight fixed scenario-groups. Values are quantum-minus-classical OOD balanced-accuracy differences under $P1'$ (select on ID-validation, tune per-configuration C on train, evaluate on OOD test). The frozen extended-family analysis uses 60 candidates per family in every group-classifier cell. Each scenario interval is a 95% conditional cluster- t interval over five q-split seed clusters ($n = 5$), conditional on the benchmark pools and design; the fifteen pipeline realizations per setting are nested within these clusters. An interval containing zero is not an equivalence result. Source-dataset-equal rows average groups within three source datasets and report leave-one-source-dataset-out ranges, not confidence intervals. The customary linear+RBF comparison is contextual rather than equal-budget.

Scenario	Clf.	vs lin./RBF	vs ext. (equal B)	conditional interval
EMBER $m1$	SVC	−0.0095	−0.0058	[−0.0340, +0.0223]
	GPC	−0.0050	−0.0009	[−0.0101, +0.0083]
EMBER $m2$	SVC	+0.0139	−0.0064	[−0.0137, +0.0009]
	GPC	+0.0149	+0.0027	[−0.0027, +0.0081]
UNSW-DoS (campaign)	SVC	−0.0043	−0.0033	[−0.0139, +0.0073]
	GPC	+0.0041	−0.0057	[−0.0120, +0.0005]
UNSW-DoS (constructed)	SVC	−0.0085	−0.0044	[−0.0119, +0.0031]
	GPC	−0.0056	−0.0108	[−0.0177, −0.0039]
UNSW-Recon (campaign)	SVC	+0.0004	+0.0050	[−0.0058, +0.0157]
	GPC	+0.0074	+0.0052	[−0.0019, +0.0123]
UNSW-Recon (constructed)	SVC	−0.0076	−0.0089	[−0.0146, −0.0031]
	GPC	−0.0083	−0.0075	[−0.0208, +0.0059]
ToN-IoT (campaign)	SVC	−0.0172	−0.0148	[−0.0268, −0.0029]
	GPC	−0.0046	−0.0090	[−0.0190, +0.0010]
ToN-IoT (constructed)	SVC	−0.0016	−0.0011	[−0.0049, +0.0028]
	GPC	+0.0313	+0.0109	[+0.0017, +0.0202]
Source-dataset-equal mean, SVC		−0.0041	−0.0056	[−0.0070, −0.0045]
Source-dataset-equal mean, GPC		+0.0059	−0.0009	[−0.0019, +0.0009]

linear+RBF baseline—the comparison that produced the +0.037 oracle result of Section 2.2—the target-label-free delta is −0.00407 under SVC and +0.00591 under the GP classifier, with heterogeneous signs across scenario-groups. The single regime that retains a small quantum lean under both equal-budget comparisons is UNSW-NB15 reconnaissance under held-out attack-campaign shift (+0.005 for both classifiers); both conditional intervals include zero. The full-pool extended-family sensitivity is similar in direction and scale (source-dataset-equal means −0.00378 for SVC and −0.00293 for GPC), so the conclusion does not depend on giving the classical family its larger pool.




Fig. 3 No robust quantum advantage under target-label-free selection. Per-scenario-group P1' OOD differences (quantum minus classical) are shown for SVC (left) and the GP classifier (right). Filled circles compare against the equal-budget extended classical family; their error bars are 95% conditional cluster- t intervals over five q-split seed clusters ($n = 5$), conditional on the fixed benchmark pools and design. Open diamonds compare against the customary linear+RBF reference. All eight fixed scenario-groups are shown; no population interval is implied.

The candidate-budget result is insensitive to the sampling scheme.

The primary `kernel_blocked` scheme samples 12 whole classical kernel-shape blocks crossed with all five dimensions, matching both the quantum count and its feature-map-by-dimension structure. Uniform and kernel-stratified samples match the count but not this block structure. The source-dataset-equal resampling estimates remain close under all three schemes: -0.00576 , -0.00555 , and -0.00548 for SVC, and -0.00185 , -0.00232 , and -0.00253 for GPC, respectively. The largest scheme range in any group-classifier cell is 0.00353. Supplementary Table 2 reports all 48 cells. Candidate-count equality is therefore distinct from search-space isomorphism, but this distinction does not drive the conclusion here.

Rank matching does not reveal a residual quantum advantage.

We pair each quantum configuration with the classical configuration of nearest training effective rank within the same run and dimension. This observational matching reuses




Fig. 4 Nearest-rank pairing yields negative median quantum-minus-classical differences in all eight fixed scenario-groups. Each quantum configuration is paired, with replacement, to the nearest-effective-rank classical configuration within run and dimension, subject to rank ratio ≤ 1.25 . The filter retains 48,708/64,800 pairs (75.2%; 69.3–83.3% by group). Bars show per-group median OOD balanced-accuracy differences; annotations give the fraction of retained pairs in which the quantum member is more accurate. This is observational matching; classical candidates may be reused and no causal family effect is implied.

classical candidates and applies the prespecified caliper $\max(r_Q, r_C) / \min(r_Q, r_C) \leq 1.25$. It retains 48,708 of 64,800 pairs (75.2%; 69.3–83.3% by group). Figure 4 reports negative median quantum-minus-classical differences in all eight groups (–0.0013 to –0.0094); the quantum member is more accurate in 22–37% of retained pairs. The pooled median is –0.0056. Minimum-cost one-to-one matching retains 24,675 pairs at the same caliper and yields a pooled median of –0.0053, again negative in every group. Results remain negative over calipers 1.10, 1.50, 2.00, and without a caliper (Supplementary Table 3). These pairings describe the predefined region of rank overlap; they do not identify a causal effect of kernel family.

2.5 Protocol sensitivity is prospectively corroborated on external domain shifts

We next applied the prospectively frozen protocol to College Scorecard, diabetes readmission, and ACS income, preserving the published TableShift domain splits and withholding all OOD labels from the controlled selector. The complete external grid contains 30 task-size-seed units and 37,800 split-level results; every prespecified unit passed the independent audit. Figure 5 shows all task-size effects and conditional seed-cluster intervals.

The primary SVC result satisfies the machine-readable outcome category frozen as `replicated_protocol_sensitivity`; narratively, protocol sensitivity is prospectively corroborated across three fixed external tasks. The task-equal controlled effect is -0.0119 (95% conditional seed-cluster interval $[-0.0205, -0.0033]$, $n = 5$ deterministic subsampling seeds per task), whereas the weak-reference, fixed- C , same-test-oracle effect is $+0.0089$ ($[+0.0047, +0.0137]$). Their contraction is $+0.0208$ ($[+0.0113, +0.0307]$), with a leave-one-task-out range of $[+0.0116, +0.0260]$. No task has a positive controlled SVC effect whose interval excludes zero in both sizes. Size-equal controlled effects are -0.0160 for College Scorecard, -0.0066 for diabetes readmission, and -0.0131 for ACS income.

The secondary GP classifier shows a larger contraction of $+0.0324$ ($[+0.0223, +0.0427]$), but its task-equal controlled effect is $+0.0077$ with an interval spanning zero ($[-0.0015, +0.0183]$). College Scorecard leans quantum under this classifier, diabetes leans classical, and ACS is slightly positive. This classifier dependence limits the conclusion appropriately: the external study corroborates the sensitivity of the family verdict to evaluation design, not universal classical dominance.




Fig. 5 External domain shifts separate optimistic and deployable conclusions. Effects are quantum minus classical OOD balanced accuracy for the weak-reference fixed-regularization same-test oracle, the equal-budget train-tuned same-test oracle, and the primary equal-budget train-tuned validation selector. Rows show both nested sizes for all three fixed tasks plus the task-equal mean. Error bars are 95% conditional seed-cluster intervals over five deterministic subsampling seeds per task ($n = 5$), preserving the nested-size relation. They quantify pipeline variation conditional on the fixed tasks and design, not uncertainty over a population of tasks.

2.6 Geometry is associated with, but does not determine, robustness

The results are consistent with both families reaching similar informative geometries. Geometry carries partial information about behaviour under shift, but neither the analysis nor the matching identifies it as a causal mediator: the association is regime-dependent, not the single governing law that optimistic accounts (our own earlier report included) proposed.

The geometry–OOD link is a regime-dependent association, not a law.

Two descriptors of the training and OOD Gram matrices carry information about which kernels generalize: the effective rank of the training kernel and the centered kernel-target alignment on the OOD split. But their association with OOD accuracy is regime-dependent, not a single mechanism. Within scenario-groups (Figure 6 and Supplementary Table 4), the median Spearman correlation between effective rank and




Fig. 6 The geometry-OOO association is regime-dependent. Points are median within-unit Spearman correlations between each geometry descriptor and OOD balanced accuracy, per fixed scenario-group for SVC; each unit is a run-dimension cell. Effective rank orders accuracy strongly in some regimes (EMBER *m2*, ToN-IoT *m2*) and weakly or negatively in others (the UNSW constructed-shift scenarios); OOD alignment is more consistently positive but never uniform. These descriptive associations have no population interval and do not identify a causal mechanism.

OOD accuracy ranges from +0.80 (EMBER *m2*) down to -0.15 (UNSW-NB15 reconnaissance under the constructed shift); partial correlations that control for family label and length scale track the raw values closely. OOD alignment is more consistently positive (+0.22 to +0.64) but likewise varies by regime. Because OOD KTA uses the labels of the same OOD split on which accuracy is measured, it is a post hoc, label-reusing diagnostic rather than an independently validated mechanism. A geometry model trained on two source datasets and used to rank kernels in the third achieves only Spearman 0.22–0.40 (EMBER 0.28, UNSW-NB15 0.22, ToN-IoT 0.40). We therefore describe a genuine but partial association, not a causal mechanism or deployable predictor.

2.7 Finite-shot estimation: the geometry is a statevector idealization

All quantum kernels above are computed by exact statevector simulation, which isolates kernel geometry but is not what a device would return. We therefore froze a post-confirmatory Monte Carlo sensitivity before computing its outcomes, then froze

a reviewer-motivated out-of-sample correction before computing that extension. The subset contains one $q1000$ run per scenario-group, all at master, q -split, and model seed 42; within each run, the exact $P1'$ -selected quantum SVC configuration is fixed. For each of four shot counts, 30 independent measurement replicates replace every fidelity entry by a binomial estimate and reselect C from training data under three conditions: no PSD correction, independent projection of square train and OOD matrices with rectangular blocks unchanged, and a coherent train-eigenspace Nyström extension (Methods). Stable SHA-256 seeds are distinct by run, configuration, shot count, replicate, and matrix block.

Table 2 summarizes the eight conditional Monte Carlo distributions. Sampling does not induce a uniform accuracy direction. Under the original independent-square heuristic, the median across the eight fixed scenario-group medians is -0.0005 at 128 shots (group range $[-0.0252, +0.0131]$), -0.0051 at 512, -0.0052 at 2,048, and 0.0000 at 8,192 shots. The coherent Nyström values are -0.0010 , -0.0026 , 0.0000 , and 0.0000 ; the unprojected values are -0.0044 , -0.0009 , -0.0005 , and 0.0000 . PSD handling is therefore part of the estimand rather than an innocuous cleanup step, but no condition produces a uniform accuracy direction. Central 95% Monte Carlo intervals for all runs are shown in Supplementary Figure 2.

Effective rank is measurement-sensitive: the median perturbed-to-exact ratio falls from 1.93 at 128 shots to 1.13 at 8,192. Median absolute OOD-accuracy deviations are 0.0065 – 0.0069 under the Nyström extension and 0.0070 – 0.0084 under the independent-square heuristic; non-monotonicity is expected because each perturbed training matrix reselects a discrete C . OOD-KTA changes under Nyström are not directly comparable to those from the independently sampled OOD square because Nyström derives that square from OOD–train features; we report them as estimator-dependent diagnostics in the artifact. These distributions quantify measurement-sampling variability

Table 2 Repeated finite-shot fidelity-estimation sensitivity for eight fixed exact Gram matrices (SVC). Each row summarizes, across scenario-groups, the within-run median over 30 measurement replicates. Δ_{OOD} is perturbed minus exact OOD balanced accuracy; brackets give the range of the eight medians, not a confidence interval. “Independent PSD” projects train and OOD square matrices separately while leaving evaluation–train blocks unchanged; it is an algorithmic heuristic, not one coherent out-of-sample kernel. “Nyström PSD” derives all corrected blocks from the sampled training matrix’s positive eigenspace. Rank ratio is perturbed/exact effective rank after clipping negative spectral mass and is shared across conditions. Per-run central 95% Monte Carlo intervals are in Supplementary Figure 2.

Shots	Condition	median Δ_{OOD} [group range]	median $ \Delta_{\text{OOD}} $	rank ratio
128	Unprojected	−0.0044 [−0.0563, +0.0137]	0.0118	1.93
128	Independent PSD	−0.0005 [−0.0252, +0.0131]	0.0070	1.93
128	Nyström PSD	−0.0010 [−0.0227, +0.0140]	0.0068	1.93
512	Unprojected	−0.0009 [−0.0226, +0.0192]	0.0068	1.48
512	Independent PSD	−0.0051 [−0.0281, +0.0041]	0.0084	1.48
512	Nyström PSD	−0.0026 [−0.0239, +0.0048]	0.0069	1.48
2,048	Unprojected	−0.0005 [−0.0222, +0.0055]	0.0080	1.24
2,048	Independent PSD	−0.0052 [−0.0238, +0.0041]	0.0072	1.24
2,048	Nyström PSD	0.0000 [−0.0225, +0.0093]	0.0065	1.24
8,192	Unprojected	0.0000 [−0.0158, +0.0114]	0.0101	1.13
8,192	Independent PSD	0.0000 [−0.0149, +0.0143]	0.0081	1.13
8,192	Nyström PSD	0.0000 [−0.0158, +0.0162]	0.0070	1.13

conditional on eight fixed exact Gram matrices. They do not model a device, circuit compilation, gate or readout noise, mitigation, or hardware feasibility.

2.8 Predictive uncertainty under shift

The Laplace GP probabilities provide a secondary uncertainty check (Supplementary Figure 1). For the target-label-free selected candidate, predictive entropy rises from ID to OOD for both families in every scenario-group, while expected calibration error changes moderately and comparably. Because neither family has an accuracy advantage, the fidelity kernels show no calibration benefit unavailable to classical kernels. We do not interpret rising entropy alone as evidence of calibration.

3 Discussion

The study’s central interpretive move is to replace the question “do quantum kernels beat classical kernels under shift?” with “which properties of a kernel are associated with its behaviour under shift, and does either family have exclusive access to them?”. Read through that lens, the results support a narrower conclusion: the apparent advantage is evaluation-sensitive. It does not survive the combined control of target-label access, train-only regularization, and candidate search, does not reappear against the equal-budget extended family, and nearest-rank pairing yields negative median differences throughout the predefined overlap. The geometry descriptors carry information as regime-dependent associations rather than a governing law, and they are shared: well-tuned heavy-tailed classical kernels occupy observed geometric regimes also reached by fidelity maps. On present evidence there is no robust out-of-distribution advantage for fidelity kernels in the low-qubit settings studied. A modest in-distribution separability edge under the probabilistic classifier remains, but is small, classifier-dependent, and shrinks under target-label-free selection.

From the perspective of the distribution-shift literature, this reinforces a general lesson: conclusions about robustness are conditional on the protocol used to obtain them [11–13, 27]. Our results add a QML-specific corollary. The strength of the classical baseline, whether regularization is tuned, whether model selection may consult target labels, and the structure of each family’s candidate search can each move the quantum-versus-classical estimate. Reporting these choices is therefore necessary for interpreting an apparent $+0.037$ advantage or comparing it with a controlled endpoint near zero.

The external analysis provides prospective corroboration beyond the constructed security shifts. Under a protocol frozen before result inspection, three education, health, and socioeconomic tasks corroborate a positive protocol contraction and a

negative controlled SVC effect. The result is stable to omitting any one task, yet heterogeneous by task and classifier: the GP endpoint retains a small positive task-equal lean whose interval crosses zero. The strongest defensible conclusion is therefore about evaluation sensitivity, not the universal ordering of kernel families.

In-distribution separability is the one axis on which the fidelity kernels retain an edge, and even it must be stated carefully. Under the oracle that selects and evaluates on the same in-distribution split, the quantum family separates the classes slightly better than linear+RBF in every EMBER setting; but under target-label-free selection—choosing on a validation split and reporting on a disjoint test split—the edge falls to +0.003 under SVC (mixed in sign across groups) and +0.012 under the probabilistic classifier, and it does not translate into an observed out-of-distribution benefit. This modest ID–OOD dissociation does not by itself establish operational value for shift-sensitive deployment. The pattern is consistent with the “accuracy on the line” literature [28], in which ID and OOD accuracy usually move together; here they decouple weakly, while descriptors shared by both families mark part of the axis along which they do.

The geometry association connects several strands of literature. That length scale is decisive for kernel behaviour is known for both families [29, 30]; our comparative result is that, once length scale is tuned symmetrically and regularization is tuned, the two families reach overlapping geometric regimes, so asymmetric tuning can change the family comparison. The geometric difference of Huang et al. [4] does not predict which family wins a scenario, consistent with its definition as a bound on the *potential* for a separation rather than its sign; a large geometric difference is compatible with equal or worse OOD accuracy, as we observe. Effective rank and OOD alignment are therefore descriptors of kernels in general, not unique signatures of quantum provenance.

The high-rank finding must also be reconciled with two spectral results that superficially point the other way. Kübler et al. argue that quantum kernels generalize

in-distribution only when the relevant spectral mass is concentrated in few components [31], and Thanasilp et al. show that fidelity kernels concentrate exponentially—toward effectively rank-one, uninformative Gram matrices—as qubit counts grow [32]. There is no contradiction, but there is a genuine tension that our results locate precisely. Both results concern the asymptotic or i.i.d. regime; our effective ranks (roughly 10–150 on training sets of 10^3 – 4×10^3 samples at 4–12 qubits) sit far below the concentration regime. The repeated finite-shot sensitivity adds an empirical caveat these theoretical results anticipate: effective rank from sampled fidelities is inflated by a median factor of 1.93 at 128 shots and 1.13 at 8,192 shots (Section 2.7). Exact effective rank is therefore a statevector idealization even at these modest qubit counts.

We emphasize what the geometry analysis does not deliver. It is a post-hoc association, not a deployable selection rule: choosing a configuration by training-set geometry alone does not reliably identify the better kernel, and OOD alignment requires OOD labels. Because the controlled comparisons do not yield a consistent quantum advantage under target-label-free selection, a more useful open problem is whether any label-free, train-only signal can select a robust kernel at all—a question that concerns classical and quantum kernels equally.

From a practical perspective, the paper suggests three takeaways for security ML. The first is substantive: on the evidence here, fidelity-based quantum kernels do not offer a robustness advantage under distribution shift over well-tuned classical kernels, while their pairwise overlap estimation introduces circuit and shot costs absent from direct evaluation of the classical baselines. A median-heuristic Laplacian kernel, especially one with a shortened length scale and its regularization tuned, is a strong and inexpensive robustness baseline, and it matches the fidelity maps under no-OOD-label selection on both modalities. Practitioners weighing a quantum-kernel pipeline for OOD-sensitive security detection should, on this evidence, prefer a tuned classical kernel unless a specific regime demonstrably favours the quantum construction after the same controls.

875 The second is methodological. Independently of the preferred kernel family, the
876
877 protocol used here provides a reusable blueprint for robust evaluation in security
878 settings:
879

- 880 1. keep the classifier and preprocessing fixed, and swap only the kernel;
- 881 2. define explicit OOD splits tied to interpretable shift mechanisms, including at least
882 one non-constructed shift mechanism where the data permit;
- 883 3. give every family the same freedom—the same candidate budget, symmetric length-
884 scale tuning, and per-configuration regularization tuned on training data alone;
- 885 4. select the deployed configuration without out-of-distribution labels, and keep an
886 in-distribution test split disjoint from the validation split used for selection;
- 887 5. repeat the pipeline, report variability as conditional intervals, and treat correlated
888 benchmarks as fixed case studies rather than pooling them into a single population
889 *p*-value.

890
891
892 Security ML frequently suffers from comparisons in which model class, tuning effort,
893 and evaluation conditions change simultaneously; the way the estimate of Sections 2.2–
894 2.4 moves from a large apparent advantage to a near-zero or negative endpoint
895 as controls are added demonstrates how strongly the conclusion depends on the
896 evaluation design.
897

898 The third concerns diagnostics: the geometry descriptors are cheap to compute
899 from Gram matrices that the pipeline already builds, and the alignment-survival anal-
900 ysis can be run post hoc on any deployed kernel system with a labeled drift sample.
901
902 Practitioners should nonetheless resist reading these results as a blanket recommen-
903 dation to replace classical kernels; the responsible recommendation is to test the
904 candidate families under the expected shift regime, using exactly the kind of protocol
905 described here.
906

907 Several limitations define the scope of our conclusions.
908
909
910

Our shift operationalization covers three reproducible mechanisms—within-class sparsity-tail extremes, train-geometry-tail extremes, and a held-out attack-campaign shift on network traffic (an unseen attack campaign plus a capture-partition change, not a timestamped temporal split)—but not the full space of real-world drift processes. Explicit temporal splits on timestamped malware, family-based drift, and adversarially induced shift remain outside the study’s scope, and our conclusions should be re-tested under them rather than assumed.

Kernel and encoding coverage also remains finite. The classical family spans seven kernels plus a bandwidth sweep, and both families receive symmetric tuning freedom, but neither candidate set is exhaustive: the quantum side is restricted to four commonly used fidelity feature maps at low qubit counts, and trained or task-adapted encodings, projected kernels, and deeper maps are not evaluated. The observed association suggests that future extensions should report their Gram-matrix geometry as well as provenance, but it does not predict their outcome. At substantially larger qubit counts, exponential kernel concentration [32] creates an additional regime not represented here, so our conclusions should not be extrapolated beyond the low-qubit study.

Selection is necessarily over finite candidate sets. P1’ (ID-validation selection), P2 (cross-realization OOD-supervised selection), and P3 (same-test OOD oracle) are selection procedures over finite candidate sets. The primary comparison matches candidate budgets, and the full-pool and budget curves provide sensitivities, but residual dependence on candidate definitions and search granularity remains. Fully nested selection over additional independent benchmark pools would strengthen the analysis further.

The quantum kernels are simulated rather than executed on hardware. The study therefore characterizes fidelity-kernel geometry, not noisy-device behaviour. Repeated

967 binomial sampling quantifies one component of measurement uncertainty, but hard-
968 ware noise, transpilation, readout, mitigation, and circuit-dependent correlations
969 remain unmodelled [33]; exponential concentration may also make informative spectral
970 structure expensive to resolve at larger scales [32]. We make no claim about hardware
971 feasibility, resource advantage, or quantum speedup—the only rigorously established
972 quantum-classical learning separation concerns a specifically constructed problem, not
973 natural data [34]. At the dimensions studied, the fidelity kernels are also classically
974 simulable [35, 36], which is precisely what makes them available as practical kernel
975 constructions today.

981
982 Dataset scope and feature caveats provide a final boundary. Four benchmark
983 scenarios across two modalities is a substantial widening relative to single-dataset
984 studies, but security is broader still: dynamic malware analysis and encrypted traffic
985 remain untested. Within the network scenarios, the feature sets inherited from the
986 public benchmarks retain port numbers and protocol indicators, which are legitimate
987 but known shortcut-prone features in NIDS benchmarks [37]; because every kernel
988 receives identical features, this affects absolute accuracies, not the family comparisons.
989 Finally, the mechanism’s effective-rank component is itself regime-dependent (strong
990 on EMBER and ToN-IoT, weak on UNSW-Reconnaissance), a gradation we report
991 rather than smooth over.

992
993 In light of these limitations, the appropriate interpretation of the paper is delib-
994 erately conservative. We do not claim universal superiority, complexity-theoretic
1000 quantum advantage, or hardware relevance—nor do we claim to have refuted quan-
1001 tum kernels in general. We claim three narrower things, supported by conditional
1002 effect estimates rather than a population p -value: that the apparent out-of-distribution
1003 advantage of fidelity kernels over customary baselines disappears when same-test
1004 oracle selection and untuned regularization are removed; that under equal candi-
1005 date budgets and no-OOD-label selection no robust family-level advantage remains
1006
1007
1008
1009
1010
1011
1012

against either the customary or the extended classical reference, across eight fixed
scenario-groups and two classifiers; and that the geometric descriptors associated with
out-of-distribution behaviour occur in both families and do not identify kernel prove-
nance as the source of robustness. This view is consistent with recent calls for careful
benchmarking in QML [6–8] and is more informative than either an unqualified advan-
tage claim or an unqualified dismissal: it identifies the evaluation choices that generate
the optimistic result and shows how the conclusion changes when they are controlled.

This paper asked why kernel-family verdicts change under distribution shift,
using the quantum-versus-classical comparison as both motivation and stress test. A
controlled kernel-swap protocol—classifier, preprocessing, splits, selection logic, and
tuning freedom held fixed while only the kernel varies—was instantiated on four
security scenarios spanning two modalities and three shift mechanisms, two classifier
families including a probabilistic one, and 72 principal settings. Each setting contains
fifteen pipeline realizations nested in five effective q-split clusters. Effects are fixed-case
summaries with conditional cluster intervals, not population significance tests.

The empirical arc has three movements. Under same-test oracle selection and fixed
regularization, fidelity-based quantum kernels appear to beat the customary linear and
RBF baselines under shift. That result does not survive the controlled comparison:
once each configuration is tuned on training data, selection uses no OOD labels, and
both extended families search a structurally matched candidate budget, the source-
dataset-equal OOD difference is -0.0056 under the primary SVC analysis and -0.0009
under the GP classifier. The primary comparison is classical-favoured or tied in seven
of eight scenario-groups. Nearest-rank pairing yields negative median differences in
every group under with-replacement, one-to-one, and caliper sensitivities. Effective
rank and OOD alignment associate with performance in a regime-dependent way,
occur in both families, and change under repeated finite-shot estimation.

1059 The central contribution is the combined protocol: distribution shift, target-
 1060 label-free selection, train-only regularization, structurally matched search budgets,
 1061 a prespecified specification curve, and prospectively frozen external corroboration.
 1062 Empirically, no robust family-level advantage survives these controls in the studied
 1063 regimes. Methodologically, the study maps the evaluation bundles associated with the
 1064 optimistic endpoint and provides a reproducible protocol for controlling their com-
 1065 ponents; it does not assign a causal effect to every axis of the specification curve.
 1066 Analytically, it shows that geometric quantities invoked to explain such advantages
 1067 describe overlapping regions of both families and change under finite-shot estimation.
 1068 The most useful open problem this leaves is not “which family wins” but whether
 1069 any label-free, train-only signal can select a robust kernel under shift at all; the security
 1070 community’s label-free drift detectors [38] and the theory of task–model alignment
 1071 [39–41] are natural starting points. Answering it would benefit classical and quantum
 1072 methods alike: the operative question is which similarity geometry aligns with the shift
 1073 regime one actually faces, and how to identify it without target-label access.
 1074
 1075
 1076
 1077
 1078
 1079
 1080
 1081
 1082
 1083
 1084

1085 4 Methods

1088 4.1 Problem setting and estimands

1089 We study binary security classification—malware versus benign software and attack
 1090 versus benign traffic—under distribution shift. Let D_{ID} denote the in-distribution
 1091 process and D_{OOD} a constructed stress-test or held-out-campaign process. Each run
 1092 contains a training set T , an ID holdout that is divided into validation V_{ID} and test
 1093 S_{ID} , and an OOD test set S_{OOD} . Preprocessing, samples, classifier implementation,
 1094 tuning rule, and family-selection rule are shared; only the kernel changes. Comparisons
 1095 are always within a classifier.
 1096
 1097
 1098
 1099
 1100
 1101
 1102
 1103
 1104

The primary endpoint is OOD balanced accuracy after P1' deployment selection.	1105
A family effect is	1106
$\Delta_{\text{OOD}} = \text{BAcc}_{\text{OOD}}^{(Q)} - \text{BAcc}_{\text{OOD}}^{(C)},$	1107
	1108
so positive values favour the quantum family. ID balanced accuracy and ID-to-OOD	1109
degradation are secondary. P2 selects from OOD information in other realizations,	1110
and P3 selects and evaluates on the same OOD test labels; both are explicitly non-	1111
deployable descriptive oracles.	1112
	1113
	1114
	1115
	1116
	1117
4.2 Security datasets and shared split sizes	1118
	1119
The source datasets are EMBER [15], UNSW-NB15 [18], and ToN-IoT [19]. Four	1120
binary scenarios—EMBER malware, UNSW-NB15 DoS, UNSW-NB15 Reconnaissance,	1121
and ToN-IoT Scanning—produce two shift mechanisms each and hence eight	1122
scenario-groups. Classical and quantum kernels always receive identical row indices	1123
and class balance.	1124
	1125
	1126
	1127
	1128
The 72 principal settings cross eight groups with master seeds {42, 123, 999} and	1129
$(n_{\text{train}}, n_{\text{ID}}, n_{\text{OOD}})$ equal to (1000, 500, 500), (2000, 1000, 1000), or (4000, 1800, 1800).	1130
	1131
Each setting is instantiated with q-split seeds {42, 123, 999, 7, 2024} and model seeds	1132
{42, 123, 999}, yielding fifteen pipeline realizations arranged in five q-split clusters.	1133
	1134
Indices are stored explicitly; automated gates verify split-role disjointness, required	1135
cardinalities, both classes, identical family inputs, and absence of label-derived	1136
features.	1137
	1138
	1139
	1140
	1141
4.3 EMBER shift construction	1142
	1143
EMBER uses the histogram and byte-entropy feature blocks and two label-conditional	1144
tail stress tests. They are not described as certified covariate shifts because condition-	1145
ing the split on class does not establish invariance of $P(Y X)$. For $m1$, the sparsity	1146
score is the number of non-zero entries in the selected feature block. Within each class,	1147
	1148
	1149
	1150

low- and high-score extremes supply the exact-size OOD set; training and ID samples are stratified draws from the remainder.

For $m2$, the final training set is fixed first. A MaxAbs scaler and TruncatedSVD representation are fitted on training data only, and a centroid is computed for each class. Every non-training sample is scored by Euclidean distance to its own-class centroid. Within-class low- and high-distance extremes form S_{OOD} , and the ID holdout is sampled from the remaining pool. This construction probes unusual geometry relative to the fixed training set without using test information to fit the representation.

4.4 Network-flow shift construction

Numeric preparation removes identifiers, explicit label derivatives, and categorical protocol fields, leaving 39 UNSW-NB15 and 89 ToN-IoT features. A leakage audit records every exclusion and feature-label association; the largest observed absolute correlation is 0.62 and belongs to a legitimate traffic feature. Each scenario has reference and current pools and uses two mechanisms.

The *m2-centroid* mechanism is the network analogue of EMBER *m2*: train-only MaxAbs+TruncatedSVD class centroids define within-class distance tails, which supply the OOD set. The *attack-campaign* mechanism performs no geometric selection. Training and ID pools contain benign traffic and all non-target attack categories; OOD contains benign traffic and the held-out target campaign. UNSW-NB15 reference and OOD pools additionally use different official capture partitions. This is an unseen-campaign plus capture-partition shift, not a timestamped temporal split.

4.5 Shared preprocessing and representations

For every run and dimension $d \in \{4, 6, 8, 10, 12\}$, a MaxAbs scaler and TruncatedSVD are fitted on T ; the reduced coordinates are standardized and mapped linearly with training-fitted parameters so that the *training* coordinates lie in $[0, \pi]$. Held-out coordinates receive the same affine map and may extrapolate outside that interval under

shift. The same embedded samples are supplied to both families, making the comparison a kernel swap rather than a representation swap. Dimension is a candidate axis. No validation, ID-test, or OOD-test row fits a transformation. The SVD random state equals the model seed.

4.6 Classifiers and train-only regularization

The primary learner is an SVC with a precomputed Gram matrix and balanced class weights [42]. For every kernel geometry, $C \in \{0.01, 0.1, 1, 10, 100\}$ is selected by five-fold stratified cross-validation on the training Gram matrix; ties favour the smaller C . When $n_{\text{train}} > 2000$, this selection step uses a class-stratified random-state-20260717 subsample capped at 2,000 rows, while the final model is always refitted on the complete training block. Because C is resolved before family selection, its grid points are not counted as separate family candidates.

The secondary learner is a binary Gaussian-process classifier with a logistic link and Laplace approximation [43]. It consumes the same precomputed kernel, uses diagonal jitter 10^{-8} , at most 100 Newton iterations, and convergence tolerance 10^{-6} . The logistic-probit approximation supplies predictive probabilities; no post-hoc calibration is applied.

4.7 Classical kernel family

The customary reference contains a cosine-normalized linear kernel and RBF kernel. The extended family adds cosine-normalized degree-2 and degree-3 polynomial kernels with $\gamma = 1/d$ and intercept 1, the Laplacian kernel $\exp(-\|x - x'\|_1/\ell_1)$, and Matérn-3/2 and Matérn-5/2 kernels. The base ℓ_1 and ℓ_2 scales are median pairwise Manhattan and Euclidean distances, respectively, calculated from at most 2,000 training rows with the model seed; the RBF base scale is fitted from training variance. RBF, Laplacian, and both Matérn kernels use factors $\{0.1, 0.3, 1, 3, 10\}$ around these training-only bases.

1243 Deduplication yields 23 kernel-shape/scale blocks crossed with five dimensions, or 115
1244 extended-classical candidates in full-coverage groups.

1246

1247 4.8 Quantum fidelity-kernel family

1249

1250 The quantum family uses exact statevector fidelities

1251

1252

$$1253 K_Q(x, x') = |\langle \phi_Q(x) | \phi_Q(x') \rangle|^2$$

1254

1255

1256 [1, 2]. Four feature-map blocks are evaluated: full-entanglement ZZ maps with one and
1257 two repetitions, a full-entanglement Pauli map with one repetition and Pauli strings X
1258 and Z, and a Z map with two repetitions. Each map uses angle-scale factors $\{0.5, 1, 2\}$
1259 after the shared $[0, \pi]$ mapping and all five dimensions, giving 60 candidates. The
1260 dimensions correspond to 4–12 simulated qubits. Classical bandwidth and quantum
1261 angle-scale freedom serve the symmetric purpose of controlling locality [29, 30]; no
1262 circuit is trained.

1268

1269 4.9 Target-label-free family selection

1270

1271 The original ID holdout is divided deterministically and class-stratifiably into V_{ID}
1272 and S_{ID} . Within each class, global row indices are ordered by SHA-256 of the
1273 salt `ksf-v4-idsplit::` and index, then assigned alternately to validation and test.
1274 The halves are disjoint and approximately class-balanced. P1' deploys, within each
1275 family, the candidate with highest V_{ID} balanced accuracy and reports it once on
1276 S_{ID} and S_{OOD} . Ties follow lexicographic candidate order. OOD labels never enter
1277 regularization, preprocessing, or P1' selection.

1283

1284

1285

1286

1287

1288

4.10 Candidate-budget matching	1289
	1290
Best-of-family selection rewards a larger pool. The confirmatory coverage audit found	1291
115 extended-classical and 60 quantum candidates in every one of the eight scenario-	1292
groups and both classifiers. The primary comparison therefore draws 60 candidates	1293
from each family in all 16 group-classifier cells. For each cell, 5,000 extended-classical	1294
subsamples are generated with seed 20260715. The primary <code>kernel_blocked</code> scheme	1295
samples 12 of the 23 complete kernel-shape/scale blocks, each crossed with all five	1296
dimensions, to mirror the 12 quantum feature-map/angle-scale blocks crossed with the	1297
same dimensions. This matches both candidate count and block structure, although it	1298
does not make the two search spaces identical. <code>uniform</code> samples individual candidates,	1299
while <code>kernel_stratified</code> distributes the draw across kernel-shape/scale strata. Each	1300
subsample selects its P1' winner; the classical endpoint is the resampling expectation.	1301
The full 115-candidate pool is secondary. Resampling percentiles describe sensitivity	1302
to the sampled candidate set, not population uncertainty.	1303
	1304
	1305
	1306
	1307
	1308
	1309
	1310
	1311
	1312
	1313
4.11 Security fixed-case aggregation and conditional intervals	1314
	1315
The effective replication axis within a scenario-group is q-split seed. Master seed is a	1316
design stratum rather than a replicate, some EMBER OOD pools recur across master	1317
seeds, and training sizes are nested. Audited OOD overlap is at most 1.2% for EMBER	1318
and 7.8% for network groups. Within each q-split cluster, model seeds are averaged	1319
within master seed, master seeds within size, and the three sizes with equal weight. The	1320
five resulting cluster means T_j give the group effect $\bar{T}$ and 95% conditional cluster- t	1321
interval	1322
	1323
	1324
	1325
	1326
	1327
	1328
	1329
This few-cluster t construction follows the cluster-mean approach of Ibragimov and	1330
Müller [44]. The interval ($n = 5$ q-split clusters) is conditional on the fixed benchmark	1331
pools and experimental design; inclusion of zero is not evidence of equivalence. A	1332
	1333
	1334

1335 9,999-draw BCa bootstrap over cluster means and leave-one-q-split calculations are
1336
1337 secondary diagnostics, reported for all 16 primary cells in Supplementary Table 6.

1338 Across groups, EMBER $m1/m2$ are averaged within source dataset, the four
1339
1340 UNSW-NB15 groups within UNSW-NB15, and the two ToN-IoT groups within
1341
1342 ToN-IoT. These three source-dataset means receive equal weight. Leave-one-source-
1343
1344 dataset-out ranges omit one of the three at a time. The eight scenario-groups and
1345
1346 three datasets were not sampled from defined populations; therefore no population
1347
1348 p -value or dataset-population confidence interval is reported.

1348

1349 4.12 Specification curve

1350
1351 The specification curve is a descriptive multiverse over every logically valid end-
1352
1353 point with complete paired legacy and v4 data. Its ten prespecified rows cross, where
1354
1355 available, legacy or corrected generation; same-test OOD-oracle, legacy ID-based, or
1356
1357 disjoint ID-validation selection; fixed $C = 1$ or train-only cross-validation for SVC;
1358
1359 customary linear/RBF or extended references; and native or equal budgets. S1–S4
1360
1361 come from the legacy generation with fixed $C = 1$ and S5–S10 from v4 with train-only
1362
1363 cross-validation, so S4–S5 is generation-confounded and cannot identify the isolated
1364
1365 effect of regularization. Selection, reference, and budget contrasts within v4 are cleaner
1366
1367 within-generation comparisons, but the curve as a whole remains associational. S1–S10
1368
1369 were ordered before aggregation and are never sorted by effect. Realizations are aver-
1370
1371 aged within setting and scenario-group, then groups within the three source datasets,
1372
1373 and finally source datasets with equal weight. All eight group effects remain visible
1374
1375 and no specification-level hypothesis test is performed.

1373

1374 4.13 Prospectively frozen external corroboration

1375
1376 Before acquiring TableShift data or inspecting any external result, we committed
1377
1378 the tasks, sampling, preprocessing, candidate pools, estimands, uncertainty proce-
1379
1380 dure, and outcome-category rule. The source implementation is pinned to commit

fca9429814703a07e3902d005d46563a207b7f0a [45]. Three public tasks span distinct domains and published shift variables: College Scorecard (institution type), diabetes readmission (admission source), and ACS income (geographic region). TableShift’s `train`, `validation`, `id_test`, and `ood_test` roles are preserved; `ood_validation` is never loaded. College Scorecard is read from its public CC0 archive because the pinned downloader invokes authenticated Kaggle access; the pinned parser, schema, and domain split are unchanged. Source-file SHA-256 hashes and the container digest are recorded.

Each task has nested strata $(n_{\text{train}}, n_{\text{val}}, n_{\text{ID}}, n_{\text{OOD}}) = (1000, 250, 250, 500)$ and $(2000, 500, 500, 1000)$. For seeds $\{42, 123, 999, 7, 2024\}$, source rows are ordered by SHA-256 of task, split, seed, and original position, giving deterministic label-blind samples. Gates verify source-role disjointness, size nesting, both classes, source-schema preservation, and at least 12 non-constant encoded features. All 30 task-size-seed units passed.

All transformations fit `train` only. Constant columns are removed; numeric variables receive median imputation; Boolean, categorical, and string variables receive most-frequent imputation and one-hot encoding with unknown categories ignored. MaxAbs scaling, TruncatedSVD at the five study dimensions, standardization, and a training-fitted affine map that places training coordinates in $[0, \pi]$ follow; held-out coordinates may extrapolate beyond that range. The SVD random state is 20260729. The security kernel and classifier pools are then reused.

For each task, size, seed, and classifier, the validation-selected 60-candidate quantum endpoint is compared with 5,000 `kernel_blocked` 60-candidate subsamples of the 115-candidate classical pool using seed 20260729. The primary effect is quantum minus expected classical OOD balanced accuracy. The equal-budget OOD oracle selects on `ood_test`; the weak-fixed oracle uses linear/RBF, fixes SVC at $C = 1$, and also selects on the reported OOD labels. Protocol contraction is weak-fixed-oracle effect minus controlled P1’ effect.

Seeds are averaged within task and size, nested sizes receive equal weight within task, and the three fixed tasks receive equal weight. Conditional 95% intervals use 9,999 resamples of the five seed clusters within task while preserving nested sizes. Leave-one-task-out ranges are descriptive. The intervals quantify deterministic-subsampling variability conditional on three fixed tasks; they do not represent a population of tasks.

4.14 Performance and uncertainty metrics

Balanced accuracy on S_{OOD} is primary. Accuracy, macro and positive-class F1, ROC AUC, and precision–recall AUC are recorded on every split. The GP analysis additionally records log loss, Brier score, expected calibration error with 15 equal-width probability bins [46], and mean predictive entropy. No threshold or calibration parameter is fitted on OOD labels.

4.15 Geometry descriptors and association analyses

For positive-semidefinite K with eigenvalues λ_i and $p_i = \lambda_i / \sum_j \lambda_j$, effective rank is

$$r_{\text{eff}} = \exp\left(-\sum_i p_i \log p_i\right)$$

[24]. For $y \in \{-1, +1\}^n$ and doubly centered K_c , centered kernel-target alignment is

$$\text{KTA}(K, y) = \frac{\langle K_c, yy^\top \rangle_F}{n \|K_c\|_F}$$

[25, 26]. KTA is computed separately on training, ID, and OOD square blocks. We also record off-diagonal mean and standard deviation, cross-split similarities, and regularized geometric difference

$$g(K_C \rightarrow K_Q) = \left\| \sqrt{K_Q} (K_C + \lambda n I)^{-1} \sqrt{K_Q} \right\|_\infty^{1/2}$$

[4].	1473
Within each run–dimension cell, Spearman correlations relate effective rank, OOD	1474
KTA, and KTA survival (OOD minus ID) to SVC OOD balanced accuracy for all,	1475
classical-only, and quantum-only candidates. Partial Spearman association residualizes	1476
ranked variables against family and log length/angle scale. For portability, a descrip-	1477
tive model bins log effective rank into ten training-dataset quantiles, learns mean	1478
OOD accuracy per bin from two source datasets, and ranks cells in the third; held-out	1479
Spearman correlation is reported once per source dataset.	1480
Rank matching is observational. Within each run and dimension, every quantum	1481
candidate is first paired to the classical candidate with smallest absolute effective-rank	1482
difference; classical candidates may be reused. Rank ratio $\max(r_Q, r_C)/\min(r_Q, r_C)$ is	1483
filtered at prespecified calipers 1.10, 1.25, 1.50, and 2.00, plus unfiltered reporting. A	1484
sensitivity uses minimum-cost one-to-one assignment with absolute log-rank difference,	1485
leaving excess candidates unmatched. Counts, retained fractions, rank discrepancies,	1486
mean and median OOD differences, and the fraction favouring quantum are reported	1487
for every caliper. None of these analyses identifies a causal geometry or family effect.	1488
	1489
	1490
	1491
	1492
	1493
	1494
	1495
	1496
	1497
	1498
	1499
	1500
4.16 Repeated finite-shot perturbation	1501
The original v0.6.0 finite-shot analysis and the reviewer-motivated v0.7.0 out-of-	1502
sample PSD extension were each frozen before their outcomes were computed. They	1503
fix one $q1000$, master-seed-42, q-split-42, model-seed-42 run per security group and,	1504
within each, the exact-statevector quantum SVC candidate selected by $P1'$ on V_{ID} .	1505
Lexicographic candidate order resolves ties. Input <code>summary_v4.csv</code> SHA-256 hashes,	1506
chosen kernels, dimensions, and exact C values are stored in the frozen specifications.	1507
For shots $s \in \{128, 512, 2048, 8192\}$ and measurement replicates $r = 0, \dots, 29$,	1508
each exact fidelity p is sampled independently as $\text{Binomial}(s, p)/s$. A separate NumPy	1509
generator is used for train, ID-validation, ID-test, OOD-test, and OOD-square blocks.	1510
	1511
	1512
	1513
	1514
	1515
	1516
	1517
	1518

1519 Its unsigned 64-bit seed is the first eight SHA-256 bytes of
 1520
 1521
 1522 `ksf-v6-shots::<run>::<kernel>::<dim>::<shots>::<replicate>::<block>`.
 1523
 1524
 1525 Python’s process-randomized hash is not used.
 1526
 1527 Square train and OOD blocks are sampled on the upper triangle, mirrored, clipped
 1528
 1529 to $[0, 1]$, and assigned unit diagonal; rectangular blocks receive binomial sampling and
 1530
 1531 clipping. Three prespecified conditions use the same sampled blocks. The unprojected
 1532
 1533 condition makes no PSD correction. The independent-square heuristic clips negative
 1534
 1535 eigenvalues of the train and OOD square matrices separately while leaving every
 1536
 1537 rectangular evaluation–train block unchanged; it is an algorithmic intervention for a
 1538
 1539 precomputed SVC, not a coherent correction of one out-of-sample kernel.
 1540
 1541 The coherent Nyström condition eigendecomposes the sampled training matrix as
 1542
 1543 $K_T = U \text{diag}(\lambda) U^\top$ and retains $\lambda_i > \tau$, where $\tau = \max\{10^{-12}, \epsilon n \max(1, \max_i |\lambda_i|)\}$.
 1544
 1545 For retained U_+ and λ_+ ,
 1546
 1547
$$K_T^+ = U_+ \text{diag}(\lambda_+) U_+^\top, \quad \Phi_E = K_{E,T} U_+ \text{diag}(\lambda_+)^{-1/2}.$$

 1548
 1549 The corrected evaluation–train and evaluation-square blocks are $\Phi_E \Phi_T^\top$ and $\Phi_E \Phi_E^\top$,
 1550
 1551 respectively, with $\Phi_T = U_+ \text{diag}(\lambda_+)^{1/2}$. Thus the OOD square used for KTA is
 1552
 1553 derived from the sampled OOD–train block rather than independently projected. In
 1554
 1555 all three conditions, C is reselected by the same five-fold train-only procedure and
 1556
 1557 the final SVC is refitted. Outputs include signed and absolute OOD change from the
 1558
 1559 fixed exact endpoint, effective-rank ratio, OOD-KTA change, retained rank, eigen-
 1560
 1561 value tolerance, negative-eigenvalue fraction, and Frobenius changes for every block.
 1562
 1563
 1564

Within each fixed Gram matrix, medians and central 2.5–97.5% Monte Carlo intervals summarize 30 replicates. Across the eight fixed matrices, medians and ranges are descriptive. This model contains no device or correlated-noise process.

4.17 Statistics and reproducibility

No null-hypothesis test is used for the eight security groups, three source datasets, or three external tasks. Every reported uncertainty object names its effective n , resampling unit, and conditioning set: security group intervals use five q-split clusters; external intervals use five deterministic seed clusters per task; finite-shot intervals use 30 measurement replicates conditional on one fixed Gram matrix. The fifteen security pipeline realizations are never treated as independent. Small- n security intervals use the t_4 reference and descriptive ranges accompany cross-dataset summaries. An interval including zero is never interpreted as equivalence. No multiplicity adjustment is required because no family of p -values is produced.

Sample sizes were design- and compute-constrained and fixed before outcome analysis; no power calculation was used. No completed group, task, specification, budget scheme, caliper, or finite-shot replicate was removed based on its result. Randomized operations use recorded seeds; external sampling and ID division use deterministic SHA-256 ordering. The test suite validates source-dataset mapping, endpoint equality across analyses, candidate-budget coverage, matching uniqueness, seed stability, exact-reference hashes, and complete Monte Carlo cells.

4.18 Code availability and software

Experiments use Python 3.11/3.12, NumPy 2.3.5, pandas 2.2.3, scikit-learn 1.8.0, Matplotlib 3.11.0, and Qiskit 2.3.1. The exact cross-platform environment is pinned. Code, manifests, tests, aggregated results, and figure/table generators are available at https://github.com/roberto-fernandez-barrios/kernel_shift_framework. The exact

v0.6.0 submission release is identified by the reserved immutable DOI <https://doi.org/10.5281/zenodo.21672470>; the concept DOI resolves all versions.

4.19 Use of generative AI

OpenAI Codex was used during manuscript revision for language editing, code inspection, analysis implementation, and consistency checking. It did not determine authorship, study accountability, or acceptance of any scientific claim. The authors inspected the generated changes, verified all reported analyses and citations, and take full responsibility for the manuscript.

Declarations

Data Availability. This study uses the public EMBER [15], UNSW-NB15 [18], ToN-IoT [19], and TableShift [45] benchmarks, which remain subject to their original access conditions. Raw source records are not redistributed. Derived split definitions, deterministic row-position hashes, shift constructions, experiment manifests, complete aggregated results, and summary tables are available through the project Zenodo concept record at <https://doi.org/10.5281/zenodo.19147649>. The exact v0.6.0 submission artifact is identified by the reserved immutable version DOI <https://doi.org/10.5281/zenodo.21672470>.

Code Availability. Code for constructing the splits, running the kernel experiments, reproducing every analysis, and regenerating every table and figure is available at https://github.com/roberto-fernandez-barrios/kernel_shift_framework and in the versioned Zenodo archive above. The platform-independent environment is pinned in `environment.lock.yml`; the code is released under BSD-3-Clause with no restriction on non-academic use.

Competing interests. All authors declare no financial or non-financial competing interests.

Authors' contributions. R.F.-B.: conceptualization, methodology, software, data curation, formal analysis, visualization, investigation, writing—original draft, and writing—review and editing. I.P.-L.: supervision, validation, and writing—review and editing. A.G.-S.: validation, project administration, and writing—review and editing. P.G.B.: supervision, resources, and writing—review and editing. All authors read and approved the final manuscript.

Acknowledgements. This study received no funding.

References

- [1] Havlíček, V. *et al.* Supervised learning with quantum-enhanced feature spaces. *Nature* **567**, 209–212 (2019).
- [2] Schuld, M. & Killoran, N. Quantum machine learning in feature hilbert spaces. *Physical Review Letters* **122**, 040504 (2019).
- [3] Schuld, M. Supervised quantum machine learning models are kernel methods. *arXiv abs/2101.11020* (2021). URL <https://arxiv.org/abs/2101.11020>.
- [4] Huang, H.-Y. *et al.* Power of data in quantum machine learning. *Nature Communications* **12**, 2631 (2021).
- [5] Schuld, M., Sweke, R. & Meyer, J. J. Effect of data encoding on the expressive power of variational quantum-machine-learning models. *Physical Review A* **103**, 032430 (2021).
- [6] Bowles, J., Ahmed, S. & Schuld, M. Better than classical? the subtle art of benchmarking quantum machine learning models. *arXiv abs/2403.07059* (2024). URL <https://arxiv.org/abs/2403.07059>.

- [7] Schnabel, J. & Roth, M. Quantum kernel methods under scrutiny: A benchmarking study. *Quantum Machine Intelligence* **7**, 58 (2025).
- [8] Kakavand, S., Strohmeyer, C. & Schlotter, M. Benchmarking quantum kernel support vector machines against classical baselines on tabular data: A rigorous empirical study with hardware validation. *arXiv* **abs/2604.18837** (2026). URL <https://arxiv.org/abs/2604.18837>.
- [9] Caro, M. C. *et al.* Out-of-distribution generalization for learning quantum dynamics. *Nature Communications* **14**, 3751 (2023).
- [10] Haug, T. & Kim, M. S. Generalization of quantum machine learning models using quantum fisher information metric. *Physical Review Letters* **133**, 050603 (2024).
- [11] Gulrajani, I. & Lopez-Paz, D. *In search of lost domain generalization*. *International Conference on Learning Representations* (International Conference on Learning Representations, 2021). URL <https://openreview.net/forum?id=lQdXeXDoWtI>.
- [12] Koh, P. W. *et al.* *Wilds: A benchmark of in-the-wild distribution shifts*. *Proceedings of the 38th International Conference on Machine Learning*, Vol. 139 of *Proceedings of Machine Learning Research*, 5637–5664. International Conference on Machine Learning (PMLR, Cambridge, MA, USA, 2021). URL <https://proceedings.mlr.press/v139/koh21a.html>.
- [13] Yao, H. *et al.* *Wild-time: A benchmark of in-the-wild distribution shift over time*. *Advances in Neural Information Processing Systems* (Neural Information Processing Systems Foundation, 2022). Datasets and Benchmarks Track.
- [14] Zhou, K., Liu, Z., Qiao, Y., Xiang, T. & Loy, C. C. Domain generalization: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **45**,

4396–4415 (2023).	1749
	1750
[15] Anderson, H. S. & Roth, P. Ember: An open dataset for training static PE	1751
malware machine learning models. <i>arXiv abs/1804.04637</i> (2018). URL https://arxiv.org/abs/1804.04637 .	1752
	1753
	1754
	1755
	1756
[16] Pendlebury, F., Pierazzi, F., Jordaney, R., Kinder, J. & Cavallaro, L. <i>TESSER-</i>	1757
<i>ACT: Eliminating experimental bias in malware classification across space and</i>	1758
<i>time. 28th USENIX Security Symposium</i> , 729–746 (USENIX Association, 2019).	1759
	1760
	1761
	1762
[17] Arp, D. <i>et al. Dos and don'ts of machine learning in computer security. 31st</i>	1763
<i>USENIX Security Symposium</i> , 3971–3988 (USENIX Association, 2022).	1764
	1765
	1766
[18] Moustafa, N. & Slay, J. <i>UNSW-NB15: A comprehensive data set for network</i>	1767
<i>intrusion detection systems (UNSW-NB15 network data set),. 2015 Military</i>	1768
<i>Communications and Information Systems Conference (MilCIS)</i> , 1–6 (IEEE,	1769
2015).	1770
	1771
	1772
	1773
	1774
[19] Alsaedi, A., Moustafa, N., Tari, Z., Mahmood, A. & Anwar, A. TON-IoT teleme-	1775
try dataset: A new generation dataset of IoT and IIoT for data-driven intrusion	1776
detection systems. <i>IEEE Access</i> 8 , 165130–165150 (2020).	1777
	1778
	1779
	1780
[20] Abreu, D., Esteve Rothenberg, C. & Abelém, A. <i>QML-IDS: Quantum machine</i>	1781
<i>learning intrusion detection system. 2024 IEEE Symposium on Computers and</i>	1782
<i>Communications (ISCC)</i> , 1–6. IEEE Computer Society (IEEE, Paris, France,	1783
2024).	1784
	1785
	1786
	1787
[21] Abreu, D., Moura, D., Esteve Rothenberg, C. & Abelém, A. QuantumNetSec:	1788
Quantum machine learning for network security. <i>International Journal of Network</i>	1789
<i>Management</i> 35 , e70018 (2025).	1790
	1791
	1792
	1793
	1794

- 1795 [22] Bellante, A., Fioravanti, T., Carminati, M., Zanero, S. & Luongo, A. Evaluat-
1796 ing the potential of quantum machine learning in cybersecurity: A case-study
1797 on PCA-based intrusion detection systems. *Computers & Security* **154**, 104341
1798 (2025).
1800
- 1801 [23] Herrero Gómez, P., Jimeno Morenilla, A., Muñoz-Hernández, D. & Mora Mora, H.
1802 Spectral phase encoding for quantum kernel methods. *EPJ Quantum Technology*
1803 **13**, 66 (2026).
1804
- 1805 [24] Roy, O. & Vetterli, M. *The effective rank: A measure of effective dimension-*
1806 *ality. 2007 15th European Signal Processing Conference (EUSIPCO)*, 606–610
1807 (EURASIP, 2007).
1808
- 1809 [25] Cristianini, N., Shawe-Taylor, J., Elisseeff, A. & Kandola, J. *On kernel-target*
1810 *alignment. Advances in Neural Information Processing Systems (NeurIPS)*,
1811 Vol. 14 (Neural Information Processing Systems Foundation, 2001).
1812
- 1813 [26] Cortes, C., Mohri, M. & Rostamizadeh, A. Algorithms for learning kernels based
1814 on centered alignment. *Journal of Machine Learning Research* **13**, 795–828
1815 (2012).
1816
- 1817 [27] Quiñonero-Candela, J., Sugiyama, M., Schwaighofer, A. & Lawrence, N. D. (eds)
1818 *Dataset Shift in Machine Learning* (MIT Press, Cambridge, MA, 2008).
1819
- 1820 [28] Miller, J. P. *et al. Accuracy on the line: On the strong correlation between*
1821 *out-of-distribution and in-distribution generalization. Proceedings of the 38th*
1822 *International Conference on Machine Learning (ICML)*, Vol. 139 of *PMLR*,
1823 7721–7735 (International Conference on Machine Learning, 2021).
1824
- 1825 [29] Shaydulin, R. & Wild, S. M. Importance of kernel bandwidth in quantum machine
1826 learning. *Physical Review A* **106**, 042407 (2022).
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840

- [30] Canatar, A., Peters, E., Pehlevan, C., Wild, S. M. & Shaydulin, R. Bandwidth enables generalization in quantum kernel models. *Transactions on Machine Learning Research* (2023). 1841–1845
- [31] Kübler, J. M., Buchholz, S. & Schölkopf, B. *The inductive bias of quantum kernels*. *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34, 12661–12673 (Neural Information Processing Systems Foundation, 2021). 1846–1850
- [32] Thanasilp, S., Wang, S., Cerezo, M. & Holmes, Z. Exponential concentration in quantum kernel methods. *Nature Communications* **15**, 5200 (2024). 1851–1852
- [33] Wang, X., Du, Y., Luo, Y. & Tao, D. Towards understanding the power of quantum kernels in the NISQ era. *Quantum* **5**, 531 (2021). 1853–1856
- [34] Liu, Y., Arunachalam, S. & Temme, K. A rigorous and robust quantum speed-up in supervised machine learning. *Nature Physics* **17**, 1013–1017 (2021). 1857–1860
- [35] Schreiber, F. J., Eisert, J. & Meyer, J. J. Classical surrogates for quantum learning models. *Physical Review Letters* **131**, 100803 (2023). 1861–1863
- [36] Jerbi, S., Gyurik, C., Marshall, S. C., Molteni, R. & Dunjko, V. Shadows of quantum machine learning. *Nature Communications* **15**, 5676 (2024). 1864–1868
- [37] Engelen, G., Rimmer, V. & Joosen, W. *Troubleshooting an intrusion detection dataset: The CICIDS2017 case study*. *2021 IEEE Security and Privacy Workshops (SPW)*, 7–12. IEEE Computer Society (IEEE, San Francisco, CA, USA, 2021). 1869–1874
- [38] Yang, L. *et al.* *CADE: Detecting and explaining concept drift samples for security applications*. *30th USENIX Security Symposium*, 2327–2344 (USENIX Association, 2021). 1875–1880

1887 [39] Bordelon, B., Canatar, A. & Pehlevan, C. *Spectrum dependent learning curves*
1888 *in kernel regression and wide neural networks. Proceedings of the 37th Interna-*
1889 *tional Conference on Machine Learning (ICML)*, Vol. 119 of *PMLR*, 1024–1034
1890
1891 (International Conference on Machine Learning, 2020).
1892
1893
1894 [40] Canatar, A., Bordelon, B. & Pehlevan, C. Spectral bias and task-model alignment
1895 explain generalization in kernel regression and infinitely wide neural networks.
1896 *Nature Communications* **12**, 2914 (2021).
1897
1898
1899
1900 [41] Ma, C., Pathak, R. & Wainwright, M. J. Optimally tackling covariate shift in
1901 RKHS-based nonparametric regression. *The Annals of Statistics* **51**, 738–761
1902 (2023).
1903
1904
1905
1906 [42] Cortes, C. & Vapnik, V. Support-vector networks. *Machine Learning* **20**, 273–297
1907 (1995).
1908
1909
1910 [43] Rasmussen, C. E. & Williams, C. K. I. *Gaussian Processes for Machine Learning*
1911 (MIT Press, Cambridge, MA, 2006).
1912
1913
1914 [44] Ibragimov, R. & Müller, U. K. Inference with few heterogeneous clusters. *The*
1915 *Review of Economics and Statistics* **98**, 83–96 (2016).
1916
1917
1918 [45] Gardner, J., Popovic, Z. & Schmidt, L. *Benchmarking distribution shift in tabular*
1919 *data with TableShift. Advances in Neural Information Processing Systems*, Vol. 36
1920 (Neural Information Processing Systems Foundation, 2023).
1921
1922
1923
1924 [46] Guo, C., Pleiss, G., Sun, Y. & Weinberger, K. Q. *On calibration of modern neural*
1925 *networks. Proceedings of the 34th International Conference on Machine Learn-*
1926 *ing (ICML)*, Vol. 70 of *PMLR*, 1321–1330 (International Conference on Machine
1927 Learning, 2017).
1928
1929
1930
1931
1932